

Smoldyn Code Documentation

Version 2.75

Steve Andrews

©July, 2025

Contents

1	Programmer's introduction	9
1.1	What is Smoldyn?	9
1.2	Smoldyn design philosophy	9
2	Smoldyn code and build system	11
2.1	Repositories	11
2.1.1	Primary site	11
2.1.2	Source code	11
2.1.3	Documentation	11
2.1.4	Python	11
2.2	Source code	11
2.2.1	Source directories	11
2.2.2	Dependencies	13
2.3	CMake build system	16
2.3.1	Basic building	16
2.3.2	Build errors	18
2.3.3	CMake build system code	19
2.3.4	CMake subdirectories	23
2.3.5	CMake file locations	23
2.3.6	Python module after building	25
2.3.7	Cross-compiling for Windows with MinGW	25
2.3.8	Building for Windows on Windows with MinGW	26
2.3.9	Building for Windows on Windows with Visual C	26
2.4	Unit and regression testing	27
2.5	GPU Smoldyn	27
2.6	Code techniques and tricks	28
2.6.1	New releases	28
2.6.2	Formatting	28
2.6.3	Code merging	29
2.7	Documentation	29
2.8	Continuous Integration	30
2.9	Interfacing with Biosimulators, Simularium, and Vivarium	31
2.9.1	Biosimulators	31
2.9.2	Simularium	31
2.9.3	Vivarium	31
3	Python binding code	33
3.1	module.cpp	33
3.2	util.h	33
3.3	Simulation.h	34
3.4	Callbackfunc.cpp and Callbackfunc.h	35

Contents

3.5	Python scripts in smoldyn directory	36
3.6	Python calling sequence	37
4	C code: files, macros, variables, etc.	39
4.1	Smoldyn source files	39
4.2	Constants and global variables	39
4.3	Macros	40
4.4	Local variables	41
5	C Code: structures and functions	45
5.1	Header files	45
5.2	Molecules (functions in smolmolec.c)	45
5.3	Walls (functions in smolwall.c)	62
5.4	Reactions (functions in smolrxn.c)	64
5.4.1	enumerated types	65
5.4.2	reaction structure	65
5.4.3	reaction superstructure	67
5.4.4	packed species identities	68
5.4.5	reaction functions	68
5.5	Rules (functions in smolrule.c)	80
5.5.1	enumerated types	80
5.5.2	rule structure	80
5.5.3	rule superstructure	81
5.5.4	rule functions	81
5.6	Surfaces (functions in smolsurf.c)	84
5.6.1	Surface geometry	84
5.6.2	Molecule-surface interactions	86
5.6.3	Surface data structures	87
5.6.4	Function interdependence	91
5.6.5	Surface functions	92
5.7	Boxes (functions in smolboxes.c)	112
5.8	Compartments (functions in smolcompart.c)	117
5.8.1	Data structures	117
5.9	Ports (functions in smolport.c)	121
5.9.1	Data structures	122
5.10	Lattices (functions in smollattice.c)	124
5.10.1	Data structures	125
5.10.2	Functions	126
5.11	Filaments (functions in smofilament.c)	130
5.11.1	Data structures declared in smoldyn.h	130
5.11.2	Filament math - 3D conformations	135
5.11.3	Filament math - quaternions	139
5.11.4	Filament math - segment 0 orientation	141
5.11.5	Filament math - filament stretching forces	141
5.11.6	Filament math - filament bending forces	142
5.11.7	Filament math - filament thermal forces	147
5.11.8	Filament math - filament dynamics	148
5.11.9	Filament simulation approach	149

5.11.10	Filament math - numbers	150
5.11.11	Filament status	150
5.11.12	Functions.	151
5.12	BioNetGen (functions in smolbng.c)	163
5.13	Complexes (not written yet)	171
5.14	Graphics (functions in smolgraphics.c)	172
5.15	Simulation structure (functions in smolsim.c)	177
5.15.1	Data structures	177
5.15.2	Functions	179
5.16	Commands (functions in smolcmd.c)	185
5.16.1	Writing commands	185
5.16.2	Externally accessible function	188
5.16.3	Individual command functions	188
5.17	Top-level code (functions in smoldyn.c)	198
5.18	Libsmoldyn (functions in libsmoldyn.cpp)	198
5.18.1	Error handling	198
6	Code design	201
6.1	Memory management	201
6.2	Data structure preparation and updating	203
6.3	User input	205
6.4	Simulation algorithm sequence	205
6.5	Wildcards, species groups, and patterns	207
6.5.1	Patterns	207
6.5.2	Control flow and functions - species names	208
6.5.3	Reaction rule storage	208
6.5.4	Control flow and functions - reactions	209
6.5.5	Control and flow - network generation	210
6.6	Units	210
6.6.1	Unit formatting	211
6.6.2	Unit parsing control flow	211
7	Smoldyn modifications	213
7.0.1	Modifications for version 1.5 (released 7/03)	213
7.0.2	Modifications for version 1.51 (released 9/5/03)	214
7.0.3	Modifications for version 1.52 (released 10/24/03)	214
7.0.4	Modifications for version 1.53 (released 2/9/04)	215
7.0.5	Modifications for version 1.54 (released 3/3/04)	216
7.0.6	Modifications for version 1.55 (released 8/20/04)	216
7.0.7	Modifications for version 1.56 (released 1/14/05)	216
7.0.8	Modifications for version 1.57 (released 2/17/05)	216
7.0.9	Modifications for version 1.58 (released 7/22/05)	216
7.0.10	Modifications for version 1.59 (released 8/26/05)	217
7.0.11	Modifications for version 1.60 (not released, but given to Karen 9/30/05)	217
7.0.12	Modifications for version 1.70 (released 5/17/06)	217
7.0.13	Modifications for version 1.71 (released 12/8/06)	217
7.0.14	Modifications for version 1.72 (released 2/26/07)	218
7.0.15	Modifications for version 1.73 (released 9/25/07)	218

Contents

7.0.16	Modifications for version 1.74 (released 10/22/07)	219
7.0.17	Modifications for version 1.75 (released 11/6/07)	219
7.0.18	Modifications for version 1.76 (released 11/7/07)	220
7.0.19	Modifications for version 1.77 (released 11/18/07)	220
7.0.20	Modifications for version 1.78 (released 11/29/07)	220
7.0.21	Modifications for version 1.79 (released 12/6/07)	221
7.0.22	Modifications for version 1.80 (released 12/22/07)	222
7.0.23	Modifications for version 1.81 (released 1/22/08)	222
7.0.24	Modifications for version 1.82 (released 2/28/08)	223
7.0.25	Modifications for version 1.83 (released 3/14/08)	223
7.0.26	Modifications for version 1.84 (released 4/11/08)	223
7.0.27	Modifications for version 1.85 (released 6/3/08)	224
7.0.28	Modifications for version 1.86 (released 11/17/08)	225
7.0.29	Modifications for version 1.87 (released 12/7/08)	225
7.0.30	Modifications for version 1.88 (released 1/16/09)	225
7.0.31	Modifications for version 1.89 (released 2/11/09)	226
7.0.32	Modifications for version 2.00 (released 2/17/09)	226
7.0.33	Modifications for version 2.01 (released 3/3/09)	226
7.0.34	Modifications for version 2.02 (released 5/5/09)	227
7.0.35	Modifications for version 2.03 (released 5/22/09)	227
7.0.36	Modifications for version 2.04 (released 6/27/09)	227
7.0.37	Modifications for version 2.05 (released 7/23/09)	228
7.0.38	Modifications for version 2.06 (released 11/6/09)	228
7.0.39	Modifications for version 2.07 (released 11/17/09)	228
7.0.40	Modifications for version 2.08 (unofficially released 11/20/09)	229
7.0.41	Modifications for version 2.09 (released 1/6/10)	229
7.0.42	Modifications for version 2.10 (released 3/24/10)	229
7.0.43	Modifications for version 2.11 (released 5/4/10)	229
7.0.44	Modifications for version 2.12 (released 6/10/10)	230
7.0.45	Modifications for version 2.13 (released 7/15/10)	230
7.0.46	Modifications for version 2.14 (released 7/18/10)	230
7.0.47	Modifications for version 2.15 (released 7/20/10)	230
7.0.48	Modifications for version 2.16 (released 9/24/10)	230
7.0.49	Modifications for version 2.17 (released 11/19/10)	231
7.0.50	Modifications for version 2.18 (released 1/6/11)	231
7.0.51	Modifications for version 2.19 (released 2/11/11)	231
7.0.52	Modifications for version 2.20 (released 3/4/11)	232
7.0.53	Modifications for version 2.21 (released 3/11/11)	232
7.0.54	Modifications for version 2.22 (released 3/22/11)	232
7.0.55	Modifications for version 2.23 (released 6/24/11)	232
7.0.56	Modifications for version 2.24 (released 7/27/11)	234
7.0.57	Modifications for version 2.25 (released 9/26/11)	235
7.0.58	Modifications for version 2.26 (released 3/2/12)	235
7.0.59	Modifications for version 2.27 (released 7/26/12)	235
7.0.60	Modifications for version 2.28 (released 8/28/12)	237
7.0.61	Modifications for version 2.29 (released 4/10/13)	237
7.0.62	Modifications for version 2.30 (released 8/21/13)	238
7.0.63	Modifications for version 2.31 (released 9/9/13)	238

7.0.64	Modifications for version 2.32 (released 8/29/14)	238
7.0.65	Modifications for version 2.33 (released 10/9/14)	239
7.0.66	Modifications for version 2.34 (released 1/8/15)	239
7.0.67	Modifications for version 2.35 (released 4/23/15)	240
7.0.68	Modifications for version 2.36 (released 6/9/15)	240
7.0.69	Modifications for version 2.37 (released 10/7/15)	240
7.0.70	Modifications for version 2.38 (released 10/22/15)	240
7.0.71	Modifications for version 2.39 (released 1/15/16)	242
7.0.72	Modifications for version 2.40 (released 3/22/16)	242
7.0.73	Modifications for version 2.41 (released 4/8/16)	243
7.0.74	Modifications for version 2.42 (released 4/29/16)	243
7.0.75	Modifications for version 2.43 (released 5/18/16)	243
7.0.76	Modifications for version 2.44 (released 6/27/16)	243
7.0.77	Modifications for version 2.45 (released 7/15/16)	244
7.0.78	Modifications for version 2.46 (released 7/30/16)	244
7.0.79	Modifications for version 2.47 (released 8/30/16)	244
7.0.80	Modifications for version 2.48 (released 11/16/16)	245
7.0.81	Modifications for version 2.49 (released 2/17/17)	245
7.0.82	Modifications for version 2.50 (released 2/27/17)	246
7.0.83	Modifications for version 2.51 (released 3/15/17)	246
7.0.84	Modifications for version 2.52 (released 5/16/17)	246
7.0.85	Modifications for version 2.53 (released 5/25/17)	246
7.0.86	Modifications for version 2.54 (released 8/27/17)	247
7.0.87	Modifications for version 2.55 (released 7/16/18)	247
7.0.88	Modifications for version 2.56 (released 9/18/18)	248
7.0.89	Modifications for version 2.57 (released 3/21/19)	248
7.0.90	Modifications for version 2.58 (released 3/29/19)	249
7.0.91	Modifications for version 2.59 (released 5/15/19)	249
7.0.92	Modifications for version 2.60 (released 6/7/19)	249
7.0.93	Modifications for version 2.61 (released 7/22/19)	250
7.0.94	Modifications for version 2.62 (released 10/14/20)	250
7.0.95	Modifications for version 2.63 (released 12/15/20)	250
7.0.96	Modifications for version 2.64 (released 3/25/21)	251
7.0.97	Modifications for version 2.65 (released 5/18/21)	252
7.0.98	Modifications for version 2.66 (released 6/18/21)	252
7.0.99	Modifications for version 2.67 (released 11/2/21)	253
7.0.100	Modifications for version 2.68 (released 7/18/22)	253
7.0.101	Modifications for version 2.69 (released 10/17/22)	253
7.0.102	Modifications for version 2.70 (released 10/24/22)	253
7.0.103	Modifications for version 2.71 (released 2/5/23)	253
7.0.104	Modifications for version 2.72 (released 8/17/23)	254
7.0.105	Modifications for version 2.73 (released 2/27/24)	255
7.0.106	Modifications for version 2.74 (released 1/22/25)	256
7.0.107	Modifications for version 2.75 (released 7/30/25)	257
7.0.108	Modifications for version 2.76 (not released yet)	258

8 The wish/ to do list 259

8.1	Bugs and issues to fix	259
-----	------------------------	-----

Contents

8.2	Desired features	260
8.2.1	Libsmoldyn	260
8.2.2	Code acceleration	260
8.2.3	Electrostatics	260
8.2.4	BioNetGen	260
8.2.5	Math	261
8.2.6	Commands	261
8.2.7	Distribution	261
8.2.8	Core Smoldyn	262
8.2.9	Graphics and I/O	262
8.2.10	Code improvements	263
8.2.11	Major additions	263

1 Programmer's introduction

1.1 What is Smoldyn?

Smoldyn is a Brownian dynamics simulator. It represents space as a 1-, 2-, or 3-dimensional continuum, as opposed to a lattice, and it steps through time using finite length time steps. Smoldyn represents molecules as individual point-like particles and membranes as infinitely thin surfaces. Smoldyn simulates molecular diffusion, chemical reactions between individual molecules, and a wide variety of molecule-surface interactions. So far, Smoldyn has been used primarily for either detailed biophysics research problems, such as on diffusion-influenced reaction dynamics, or for investigating the effects of spatial organization on simple biological systems, such as the *Escherichia coli* chemotaxis system.

Smoldyn is also a community development project. I wrote nearly all of the core code, Jim Schaff and his colleagues at UCHC added some code for integrating Smoldyn into VCell, Nathan Addy wrote a rule-based modeling module called Libmoleculizer that almost worked, Lorenzo Dematte and Denis Gladkov independently parallelized Smoldyn to run on graphics processing units (GPUs), Martin Robinson added adjacent-space hybrid simulation capabilities to Smoldyn, and Dilawar Singh added Python bindings to Libsmoldyn.

In order to maintain and enhance Smoldyn's value for computational biologists, as opposed to letting it become a hodgepodge of mismatched code, it is helpful to carefully define what Smoldyn is, and what it isn't.

1.2 Smoldyn design philosophy

Central to Smoldyn's design philosophy is the concept of two distinct levels of approximation between physical reality and numerical simulation.

In the first level, physical reality is approximated to a perfectly defined *model system*. Smoldyn's model system builds on the one that Smoluchowski developed in 1917. It uses continuous space and continuous time. Depending on the context, molecules are represented as either point-like particles or as perfect spheres. In particular, they are (usually) point-like particles without excluded volume when considering diffusion and they are perfect spheres when considering chemical reactions. Model molecules do not have orientations, momenta, or kinetic energies because these aspects equilibrate in solution on much shorter time scales than the dynamics that the Smoluchowski model focuses on. As an example of the approximation from physical reality to the model system, molecular diffusion in reality is driven by a very high rate of intermolecular collisions and relies on van der Waals and steric molecular interactions. However, it is approximated in the model system as mathematically perfect Brownian motion, where each molecule behaves independently of every other molecule and moves with an infinitely detailed trajectory.

In the second level of approximation, the dynamics of the model system are approximated using numerical algorithms to yield the *simulated system*. In particular, Smoldyn approximates the model system through the use of finite simulation time steps. For example, Smoldyn simulates Brownian motion using Gaussian-distributed displacements at each time step. Because Smoldyn uses finite

1 Programmer's introduction

size time steps, it is nonsensical to ask about the state of the simulated system during time steps. Instead, the simulation produces what can be seen as simulation system snapshots at the end of each time step. When the results in these snapshots are completely indistinguishable from those found with the ideal model system, to within computer round-off error, then the simulation algorithms are called “exact”. Note that exactness only refers to agreement between the simulated system and the model system; correspondence between the model system and physical reality is a completely separate issue, and one which depends very much on the specific dynamics that the modeler wishes to investigate.

The Smoldyn software represents a balance between algorithms that are exact for the Smoluchowski model system and those that are computationally efficient. Often, these goals are actually complementary because highly accurate algorithms enable the use of long simulation time steps and hence enable fast simulations. However, there are also often tradeoffs, where better accuracy leads to slower simulations. The challenge with seeking a balanced approach between accuracy and computational efficiency is that the software users (i.e. modelers) generally aren't comfortable trusting software that is known to be inaccurate. For this reason, every algorithm in Smoldyn is exact in two ways. First, the simulated *rates* of all isolated algorithms are exact for any length time step. For example, in a simulation of the irreversible bimolecular reaction $A + B \rightarrow C$, Smoldyn always gets the macroscopic reaction rate exactly correct, although the exact positions of the molecules are not necessarily in perfect agreement with those found for the model system. This focus on rates is important for Smoldyn to yield accurate equilibrium constants. Secondly, all simulated dynamics, regardless of how many algorithms are used in a simulation, approach exactness as simulation time steps are reduced towards zero. In the process, of course, simulation run times approach infinity, so exactness isn't actually achievable. However, having results that can approach exactness is important because it enables modelers to understand and quantify their simulation errors.

2 Smoldyn code and build system

2.1 Repositories

2.1.1 Primary site

Smoldyn's primary website is <http://www.smoldyn.org>. This website has information about the software and offers downloads of the current release, whether as pre-compiled versions or tar/zip/tgz files with the source code and other parts of the package.

2.1.2 Source code

Smoldyn's code repository is github at <https://github.com/ssandrews/Smoldyn>. This repository was named Smoldyn_Official from its first github upload with version 2.57, up to version 2.62.dev, and was then renamed to just Smoldyn. Prior code repositories used Subversion servers that are now discontinued. I also started a SourceForge account at some point for Smoldyn, which I have not maintained. There are several other github repositories that have Smoldyn in their names, all of which are for the same overall software, but are not maintained by myself and don't have the official releases.

2.1.3 Documentation

The most current documentation is included in the download package as LaTeX and pdf files for the User's Manual and code documentation (this file).

Also, Dilawar Singh started a readthedocs website for Smoldyn at <https://readthedocs.org/projects/smoldyn/>. I'll try to maintain that site but it may not be as current as the documentation that's packaged with the downloads. In addition, note that the documentation is stored on the same github site as the code, so if you really want the latest un-released documentation, then you can get it at the github site.

2.1.4 Python

Dilawar started a PyPI site for Smoldyn at <https://pypi.org/project/smoldyn/>. It is updated from github through nightly builds, so it should always be up to date.

2.2 Source code

2.2.1 Source directories

All of Smoldyn's code is in subdirectories within the "source" directory. These subdirectories and their contents are listed below.

There is also one file in the source directory, which is `smoldynconfigure.h.in`. This file is used as a template by the CMake build system for creating a header file that describes the actual build configuration.

2 Smoldyn code and build system

BioNetGen I think this is the complete BioNetGen download, from about 2016. Smoldyn uses the BNG2.pl perl file for rule-based modeling when using the BNG language. This code is not included in Smoldyn during building, but a link is made to this file. Also, this entire BioNetGen directory gets copied to an install location (typically /usr/local/bin for Mac or Linux) when Smoldyn is installed. (This copying occurs if installing with the install script, but it doesn't appear to happen if installing using the "make install" target in the CMake build process.)

libSteve This is a library of code that I wrote, which is used throughout the Smoldyn software. All of it is plain C, but can be compiled as C++ as well. Essentially none of the code here is specifically for Smoldyn or, more importantly, essentially none of the code here calls functions that are part of the core Smoldyn program (e.g. smolsim.c, smolreact.c, etc.). I think that the sole exception is SimCommand.c, which can call smolcmd.c but can also be run without calling it. In addition to the code that I wrote, the SFMT directory contains code for the Mersenne Twister random number generator, which I did not write.

MSVCLibs These are headers and libraries needed for compiling Smoldyn with the MSVC compiler, on Windows. All of the code here was downloaded on a Windows computer, compiled with MSVC, and then copied over to this directory. This code is only used when compiling with MSVC. Not all of the code here is actually used.

NextSubVolume This is code that Martin Robinson wrote for hybrid Smoldyn, for simulating adjacent spatial regions with differing levels of detail. It is only used when building with the "OPTION_NSV" option turned on.

pybind11 This is an external library on which Smoldyn's Python bindings are built. It has its own CMakeLists.txt file. I have not modified anything in here, and I think it is just the pybind11 download with no modifications.

python This directory has all of Dilawar's code for Smoldyn's Python bindings. It has its own CMakeLists.txt file, which I have not tried to understand.

SmolCrowd This is not compiled into Smoldyn, but a separate text-based utility. It generates random crowders for importing into Smoldyn configuration files.

Smoldyn This directory includes the core of the Smoldyn program. Essentially all of the code here has a ".c" suffix, because it was originally written in C and is still mostly just C. However, some files have slight C++ additions, so I now compile it with a C++ compiler. The file smoldynlib.py was written by Sean Garbett in 2013 and appears to be a Python interface for libSmoldyn. I haven't used it.

SmolEmulate This is a separate utility, which is not compiled into Smoldyn. It emulates Smoldyn algorithms, for calibrating reaction rates. This code evolved from versions that I wrote to derive 3D reaction parameters for the Andrews and Bray, 2004 paper, and then to derive 1D surface interaction parameters for the Andrews, 2009 paper. It also supports 2D reactions, but I haven't finished that work yet (as of 2020).

vcell This directory contains code used for linking Smoldyn into VCell. It was written by the VCell team, and I just keep it here for future use by them.

vtk This is a wrapper for an external library for the VTK visualization toolkit. It's used by the NSV code. I'm not sure if VTK needs to be installed on the host computer for Smoldyn to offer VTK functionality.

wrl2smol This is another utility that is not compiled into Smoldyn. It converts VRML language triangle data into Smoldyn style triangle input.

2.2.2 Dependencies

As much as possible, Smoldyn is compiled statically so that its dependencies are automatically included in the binary distributions. However, if you compile Smoldyn yourself, then you may need to get them.

Smoldyn's code dependencies have changed some over time, but are mostly shown below in a tree structure, such that each file depends on the files that are indented below it. This tree might not be fully accurate for the current version. As much as possible, Smoldyn builds and runs without its dependencies, but offers fewer features if they aren't available. In particular, I've struggled some with getting some of them to work on Windows.

```
Smoldyn
  OpenGL
  OpenGL-glut
  libTiff
  zlib
  libiconv
  NSV
    VTK
  pybind11
```

Getting the code for these dependencies has ranged from easy to impossible, so my approaches to solve these issues are listed below.

python, pip, wheel . These aren't code dependencies as the others are, but are still required software tools that need to be installed. Note that wheel and pip need to be affiliated with the same copy of python that is being run; if it's not, then you get the error `invalid command 'bdist_wheel'`. This error arises from wheel not being installed; that can be solved by configuring Smoldyn with CMake and then running `make install.deps`. Note that the CMake output shows which copy of Python is being run. You can also enter `which python` to see what Python copy is being used (if it's an alias, then `ls -l path/python` will show what the alias points to). For example, my CMake shows `Found Python3: /opt/local/bin/python3.10 (found version "3.10.12")`; if I enter `which python` and then resolve the alias, I get `/opt/local/bin/python -> /opt/local/bin/python3.10`, which is the same file; also, if I enter `pip install wheel` it says that the requirement is already satisfied and it's at `/Users/sandrews/Library/Python/3.10/lib/python/site-packages`.

glut is complicated. It's part of the standard OpenGL libraries on Macs, but doesn't seem to be standard on Windows. Also, glut itself is obsolete, having been replaced by, which is a superset of glut; nevertheless, the glut.h header file is supposed to be used, unless one wants functions that are only available in freeglut.

2 Smoldyn code and build system

For Windows with MSVC, I downloaded freglut 3.2.1 from <http://freeglut.sourceforge.net> and built it on a Windows computer using the Visual Studio compiler. To build, I created and then moved to a build directory, entered “cmake ..”, and entered “msbuild freglut.sln”. This created the files lib/Debug/freglutd.lib, lib/Debug/freglut_staticd.lib, and bin/Debug/freglutd.dll. The “d” characters show that these are the debug versions of the library. These work in Smoldyn but aren’t ideal. I then found that I can create the release versions of the library with the same CMake step and then “msbuild freglut.sln /p:Configuration=Release”. I put the compiled libraries in the source/MSVCLibs directory, along with a header file from the same freglut download.

For Windows with MinGW, I found that the MinGW cross-compiler download that I use (see below) has freglut included in its collection of libraries. Thus, I link to that. Everything seems to work and Smoldyn builds well, but the resulting smoldyn.exe doesn’t actually run.

To address this, I also tried to compile the freglut 3.2.1 download on a Windows computer but with the MinGW compiler. This didn’t work. Thus, instead, I downloaded the freglut 3.2.1 source code to my Mac and cross-compiled it for Windows using “cmake .. -DCMAKE_TOOLCHAIN_FILE=../Toolchain-mingw32.cmake”. This actually worked and produced libfreeglut_static.a, libfreeglut.dll.a, libfreeglut.dll, and many more files. I was able to build Smoldyn while linking to these files, but the result was the same as before in that I got smoldyn.exe just fine, but it didn’t run.

For Mac, I’ve been using regular glut, not freglut, which works easily but has the drawback that control never returns to Smoldyn after it is handed over to `glutmainloop`. To address this (12/14/20), I installed freglut using MacPorts, which was trivial. However, entering `port contents freglut` showed the files that were installed, and none were a static library. I tried for a while to link to it anyhow using the CMakeLists.txt file, without luck. I also tried to reinstall the MacPorts version (3/29/21), with edits in the Portfile so that static libraries would be built. This seemed easy, but took a long time and didn’t work out in the end. The Portfile says that static libraries should be built, but they aren’t. In another attempt, I downloaded freglut from <http://freeglut.sourceforge.net>, as I did for Windows, and tried to build using CMake. This failed at the linking stage. There was a prior report of it not working for Macs (see <https://sourceforge.net/projects/freeglut/>), with the same error and no solution given, so this seems to be a building bug and not easily fixed. The issue seems to be that Apple deprecated the GLX functions that libfreeglut depends on, so libfreeglut can no longer be built statically. Later on (1/11/22), it turned out that CMake had been modified so that it found my freglut preferentially over the built-in glut, which then caused it to not build. The solution was to remove freglut from my computer.

libtiff is optional. It is used for saving graphics images as TIFF files.

For Mac, I’ve installed libtiff many times. MacPorts has worked, but the latest version has lots of dependencies that I didn’t want. My latest approach (March, 2021 and then again in February, 2022) was to download the latest version (4.3.0) as source code from <http://download.osgeo.org/libtiff/>. Then, I created the subdirectory “mybuild”, moved to there, and entered “cmake .. -Dzlib=OFF -Dlibdeflate=OFF -Dpixmaplog=OFF -Djpeg=OFF -Dold-jpeg=OFF -Djbig=OFF -Dlzma=OFF -Dzstd=OFF -Dwebp=OFF -Djpeg12=OFF”. This turned off all external codecs, which reduced dependencies, hopefully without reducing the functionality that Smoldyn needs. Afterward, I entered “make” and “sudo make install”, which installed to /usr/local/include and /usr/local/lib. Then, I copied the /usr/local/include/*tiff* and /usr/local/lib/*tiff* files to the Mac/libtiff subdirectory of

Smoldyn. After that, I repeated the configure, build, and install process, using the same flags but also “-DBUILD_SHARED_LIBS=OFF”. This built the static libraries. I copied those to the Smoldyn/Mac/libtiff directory as well.

For Windows with MSVC, I downloaded version 3.8.0 from <http://download.osgeo.org/libtiff/>, but it required Autotools to compile, which didn’t work for me on Windows. So, I downloaded version 4.1.0, which uses CMake. It already has a “build” directory, so I worked in a new directory called “mybuild”. I configured using the option “cmake .. -DBUILD_SHARED_LIBS=FALSE” and built with “msbuild tiff.sln /p:Configuration=Release”. This created the libraries, libtiff/Release/tiff.lib and libtiff/Release/tiffxx.lib. I copied these libraries over to source/MSVCLibs, and also included several header files (tiff.h, tiffio.h, tiffvers.h, tiffconf.h). The result works.

libiconv is a library that converts between different character encodings. I don’t know why it’s necessary, or even if it is necessary (the default Smoldyn build option settings have the iconv option turned off), but I cared about it at some point. I downloaded it from <http://www.gnu.org/software/libiconv/>. I got version 1.14. I configured with “./configure --enable-static” so that I’d get the static library.

libXML++ is a wrapper for the libxml2 XML parser. Again, I don’t know why it’s necessary or even if it is necessary, but I cared about it at some point. It might have been a dependency for libMoleculizer, which was a disastrous Smoldyn module that I eventually got rid of.

Anyhow, I downloaded libXML++ from <http://libxmlplusplus.sourceforge.net/>. While more recent versions are available, this is likely to be a major mistake because they have loads of dependencies, and the dependencies have dependencies, and so on. Instead, get libXML++ version 1.0.5. This is fully sufficient, and it works well. After downloading, extract the archive, change to the libxml++-1.0.5 directory, enter “./configure”, “make”, and “sudo make install”. This was straightforward for me.

vtk is the Visualization Toolkit, useful for graphical visualization. This functionality was added to Smoldyn with the NSV addition, but I’m not sure if it only applies to the NSV portion of Smoldyn, or to all graphics.

I downloaded my Mac version from <http://www.vtk.org/VTK/resources/software.html>, getting version 5.10.1. To build it, I created subdirectory called “build” within the vtk download directory, changed to the build directory, and entered “cmake ..” followed by “make”. This is a very large package which took about a half hour to build. Lots of warnings were emitted, but the build completed successfully. Then, “sudo make install” installed the result. Quite a lot of files were installed. As part of including VTK in the project, there is a “source/vtk” directory within Smoldyn, which has just a few files.

For Windows, I downloaded VTK from <https://vtk.org/download/>, getting the VTK-9.0.1.tar.gz file. I then extracted at the command line with “tar -xvzf VTK-9.0.1.tar.gz”, created a build subdirectory, moved into it, and ran “cmake ..” and “msbuild VTK.sln”. This required almost 3 hours to build on my very slow Windows computer. I copied the build directory over to “source/VTKlibs”, and planned to link it in, following advice at <https://vtk.org/Wiki/VTK/Tutorials/CMakeListsFile>. However, I discovered that this VTK build directory is truly enormous (3.8 Gb) and it offers very little additional capability to Smoldyn, so I decided to leave it out for the Windows build.

2.3 CMake build system

2.3.1 Basic building

Smoldyn builds with CMake, which can be downloaded from <http://www.cmake.org>. I use version 3.16, but any version above 3.4 should work (lower versions don't work due to requirements from pybind11). CMake installs trivially (at least on Mac), with a standard installer and no building required.

CMake can be run from either a command line interface (my preference) or with a GUI. At a command line interface, change directories to cmake. Every time you change CMake settings, you'll probably want to do a clean build. To do so, enter "rm -r *", while in the cmake directory (verify that you're in this directory!), to remove any prior build results. If you're asked about whether manifest.txt should be removed, say yes; this file shows the directories where Smoldyn was installed previously, thus providing information for you to remove it. For a default build, enter "cmake ..". A few test results will be printed out, and then configuring will be complete. See below for custom builds. The other option is to use the CMake GUI. It can be started by entering "cmake-gui" at a command line. Either way, when CMake is done, it will have written a lot of stuff to the cmake directory. Important files are "Makefile", which is the standard Makefile for the code and also smoldynconfigure.h, which is a C header file that the Smoldyn code uses for knowing what the build parameters are. On Windows with the MSVC compiler, there's no makefile, but smoldyn.sln instead; this is the "solution" file for compiling.

Once configuring is complete, enter "make" (on Mac or Linux, or "msbuild smoldyn.sln /p:Configuration=Release" on Windows with MSVC). Hopefully, Smoldyn will build, again with build files being put into the cmake directory. For me, "make" often doesn't work, because I get the error "ld: can't write output file: ../../../../source/python/smoldyn/_smoldyn.cpython-310-darwin.so for architecture x86_64". The solution is to run "sudo make" and enter a password because, for some reason, I don't have the correct ownership permissions for my build directory. To help debug issues with the "make" step, try "make VERBOSE=1" to see what commands CMake is sending to the system.

Once the "make" step works, Smoldyn can be run. If you want to install, enter "sudo make install" and enter your password, to install Smoldyn to the usual place (/usr/local/bin on Linux and Mac systems).

The Python wheel installs itself automatically, but often doesn't update if it was already there. To force an update, run `pip install smoldyn-2.72-...whl` (look for this filename near the bottom of the `sudo make install` output). This installs the wheel to the correct location. For me, that location is /opt/local/Library/Frameworks/Python.framework/Versions/3.10/lib/python3.10/site-packages/smoldyn/.

Here is the summary of the building process:

```
> cd Smoldyn
> mkdir build
> cd build          # empty the directory if necessary with: rm -r *
> cmake ..          # add any options here
> make
> sudo make install
```

For custom builds, you need to set various options to non-default settings. This is straightforward in the CMake GUI. There, you just check or uncheck boxes, as desired. Alternatively, from a command line interface, you can start CMake with "cmake .. -i" for interactive mode, and then CMake will ask you about each option. For each, you can just press return to select the default,

2.3 CMake build system

or enter in values of your choice. Finally, you can also list each non-default option directly on the command line (preceded with a ‘D’, presumably for define).

Following are all of the build options in the Smoldyn CMake files, plus some of the more helpful standard ones.

Smoldyn option	default	effect when ON
-DSMOLDYN_VERSION	2.62.dev	Smoldyn version number
-DOPTION_TARGET_SMOLDYN	ON	Build stand-alone Smoldyn program
-DOPTION_TARGET_LIBSMOLDYN	ON	Build LibSmoldyn library
-DOPTION_VCELL	OFF	Build for inclusion within VCell
-DOPTION_NSV	ON	Build with Next Subvolume support
-DOPTION_VTK	ON	Build with support for VTK visualization
-DOPTION_STATIC	OFF	Build using static libraries
-DOPTION_USE_OPENGL	ON	Build with graphics support
-DOPTION_USE_LIBTIFF	ON	Build with libtiff support
-DOPTION_USE_ICONV	OFF	Build with Libiconv support (use default)
-DOPTION_USE_ZLIB	OFF	Build with zlib support (use default)
-DOPTION_PYTHON	ON	Build Python module
-DOPTION_EXAMPLES	OFF	Run Libsmoldyn tests
-DOPTION_STRICT_BUILD	OFF	Treat warnings as errors and enable address sanitizer
-DOPTION_DOCS	OFF	Convert documentation formats and run Doxygen
-DPYTHON_VERSION	Auto	Choose Python version (I use 3.10)
CMake option	default	function
-DCMAKE_BUILD_TYPE	Release	Choose CMake build type options are: None, Debug, Release, RelWithDebInfo, and MinSizeRel
-DCMAKE_CXX_COMPILER:FILEPATH	clang	Compile with specific compiler for example: /usr/bin/g++
-DCMAKE_TOOLCHAIN_FILE	None	Cross-compiling toolchain file for example: ../Toolchain-mingw32.cmake

CMake has many different targets, which are entered after **make**. The ones listed in the upper section of the following table are likely to be useful, whereas I think the ones in the lower section are probably less useful and/or are run automatically.

Target	CMake file	Description
none listed	many	Compiles and links code
install	many	Installs executables and other code to system locations
docs	docs	Generate documentation with Sphinx and doxygen
install_deps	python	Install Python dependencies
wheel	python	Generate the Python wheel
examples	examples	Runs test code
TeX2RST	docs	Dummy target to convert all TeX to RST files
Doxygen	docs	Run doxygen
Sphinx	docs	Run Sphinx
doc_livehtml	docs	Generate live html
copy_python_tree	python	Copies python source tree to binary directory
pyinstall	python	Used by install script (useful in packaging)
pyinstall_venv	python	Runs “python3 setup.py install”

2 Smoldyn code and build system

<code>pyuninstall</code>	<code>python</code>	Runs “pip uninstall smoldyn”
<code>pydevel</code>	<code>python</code>	Runs “python3 setup.py develop –user”
<code>lint</code>	<code>python</code>	Static type checking using mypy
<code>lint_strict</code>	<code>python</code>	Static type checking using mypy

2.3.2 Build errors

```
[ 97%] Linking CXX shared module .../_smoldyn.cpython-313-darwin.so
ld: can't write output file: .../_smoldyn.cpython-313-darwin.so
```

This arose due to the incorrect Python version. My computer is currently setup to run Python 3.10, but CMake chose version 3.13. I fixed this by adding the configuration flag `-DPYTHON_VERSION=3.10`.

object file was built for newer ‘macOS’ version than being linked

Look at the CMake output. You should see that the `CMAKE_CXX_FLAGS` is set with `-mmacosx-version-min` equal to a lower version number than the current OS. Fix this in the main CMake.txt file in a line that looks like

```
set(CMAKE_CXX_FLAGS "${CMAKE_CXX_FLAGS} -Wall -Wno-deprecated -mmacosx
-version-min=14")
```

Update the line to give the current OS.

`ModuleNotFoundError: No module named ‘setuptools’`

This means that Python setuptools isn’t installed. I installed it with MacPorts with `sudo port install py310-setuptools`. This solved the problem.

`dyld[93366]: Library not loaded: @rpath/libtiff.5.dylib`

This error arose when building Smoldyn 2.72 with a new Mac OS upgrade to version 14.1. I solved it by adding the following line to the CMakeLists.txt file:

```
if(RUN_LIBS)
    set_target_properties(smoldyn PROPERTIES INSTALL_RPATH ${RUN_LIBS
    })
endif()
```

This line is in the “smoldyn” target portion of the Targets part of the file (it’s lines 513-515). I also added the `RUN_LIBS` variable to the LibTiff portion of the CMakeLists.txt code. I don’t know why this wasn’t needed before, and isn’t needed for other libraries, but is for LibTiff, but it appears to work.

`error: invalid command ‘bdist_wheel’`

This arises for me when two different Python versions are being used. Entering `python --version` at a command line shows which version is being used. Also, entering `which python` shows which copy of Python is being run. The CMake output has a line that says `Found Python3:....`. If its output doesn’t match with the copy of Python being used, then that’s a problem. Fix it with the CMake option `-DPYTHON_VERSION`.

2.3.3 CMake build system code

The CMake code for Smoldyn, in CMakeLists.txt in the top-most directory, is relatively straightforward, but has still caused me vastly more grief than it should have. For that reason, this section gives a detailed look at the cmake code, focusing particularly on the variable definitions.

Some aspects of this file use somewhat deprecated approaches. For example, it uses the “use_directories” statement, but one is supposed to use “target_use_directories” instead. The advantage to the latter approach is that it’s better about keeping variables within their scope, but this file is simple enough that I’m not bothering to change things.

Basic setup

Defines the project name, lists the CMake minimum version number, and sets the Smoldyn version number. Also tells CMake which version of C++ to use.

Variable	Description
SMOLDYN_VERSION	Smoldyn version number string
CMAKE_CXX_STANDARD	C++ version number

Targets to build Whether to build Smoldyn, Libsmoldyn, or both.

Variable	Description
OPTION_TARGET_SMOLDYN	Create stand-alone Smoldyn program
OPTION_TARGET_LIBSMOLDYN	Create LibSmoldyn library

Compiling options Most of the compiling options are listed here. They determine which components to include and what the code is being built for. The default values listed in the code sometimes get overridden later in this same section, depending on other settings. Hopefully, the code outputs a warning if a user-requested value gets overridden.

Variable	Description
OPTION_VCELL	Compile Smoldyn for VCell
OPTION_NSV	Compile Smoldyn with NextSubvolume functionality
OPTION_VTK	Compile Smoldyn with VTK functionality
OPTION_STATIC	Compile Smoldyn with static libraries
OPTION_USE_OPENGL	Build with OpenGL support
OPTION_USE_ZLIB	Build with Zlib support
OPTION_USE_LIBTIFF	Whether to include the LibTIFF library
OPTION_USE_ICONV	Whether to include the Libiconv library
OPTION_PYTHON	Build Python module
OPTION_EXAMPLES	Run Libsmoldyn tests
OPTION_LATTICE	Whether to build lattice code
HAVE_ZLIB	Already have Zlib library
HAVE_ICONV	Already have Libiconv library

Core code information This section lists all of the Smoldyn header files and source files, plus a few extra header files from some of the optional features (e.g. vtkwrapper.h). Some of the source files are only included for some of the options. All of the source and main files properties are

2 Smoldyn code and build system

set to the C++ language, which is necessary because most of them have a “.c” ending (note that this setting only applies to the files listed so far, and does not apply to any subsequently listed files). The include directories are set here to all of the directories with header files.

Variable	Description
HEADER_FILES	List of all header files
SRC_FILES	List of all source files
MAIN_FILES	Name of the main file

Compiler flags This section sets most of the compiler flags for the build. It starts by determining what type of build this is, where the options are ‘Release’, ‘Debug’, ‘None’, ‘RelWithDebInfo’, and ‘MinSizeRel’; it sets the value to ‘Release’ as a default. This is a built-in CMake variable, so CMake defines several compiler flags automatically based on this build type, without them needing to be set here. This section also addresses the strict-build option, setting the build type to “Debug” and setting `CMAKE_CXX_FLAGS_DEBUG` and `CMAKE_LINKER_FLAGS_DEBUG` to address sanitizing options. However, neither variable is used anywhere else in this file, so I’m questioning if it actually does anything useful. Next, this section creates a list of possible build platforms and goes through the list of possible platforms, adding to the compiler flags as needed as it goes along. As far as this is concerned, “MinGW” means cross-compile from Mac to Windows using the MinGW compiler, and “Windows” means compile on Windows using the Visual Studio compiler. This section also sets the path to the BNG2 perl script.

Variable	Description
CMAKE_BUILD_TYPE	Build type (default: Release)
CMAKE_C_FLAGS	C compiler flags string
CMAKE_CXX_FLAGS	C++ compiler flags string
DEP_LIBS	Library dependencies for linking
VCELL_BUILD	Build for VCell (ON or OFF)
APPLE_BUILD	Build for Apple (ON or OFF)
NIX_BUILD	Build for Unix or Linux (ON or OFF)
WINDOWS_BUILD	Build for Windows (ON or OFF)
BNG2_PATH	Path to BNG2 perl script

OpenGL (gl and glu, not glut) OpenGL is a bit of a challenge to include correctly, and platform dependent. The gl and glu libraries are usually built in, while the glut library is not necessarily, so this section just focuses on the gl and glu libraries. OpenGL is becoming deprecated, but I still use it, so this code starts by stating that it uses legacy preferences. The `include(FindOpenGL)` statement is used where possible; it automatically defines the variables `OPENGL_FOUND`, `OPENGL_INCLUDE_DIR`, and `OPENGL_LIBRARIES`. However, the code overrides the include directories for Windows because the automatic system doesn’t seem to find the built-in gl.h header file, and a different version of the header file from the Smoldyn directory seems to work well.

Variable	Description	Purpose
HAVE_OPEN_GL	TRUE if available, FALSE if not	smoldynconfig.h
HAVE_GL_H	TRUE for header file gl.h	smoldynconfig.h
HAVE_GL_GL_H	TRUE for header file GL/gl.h	smoldynconfig.h

HAVE_OPENGL_GL_H	TRUE for header file OpenGL/gl.h	smoldynconfig.h
HAVE_GLU_H	TRUE for header file glu.h	smoldynconfig.h
HAVE_GL_GLU_H	TRUE for header file GL/glu.h	smoldynconfig.h
HAVE_OPENGL_GLU_H	TRUE for header file OpenGL/glu.h	smoldynconfig.h
OPENGL_FOUND	TRUE if OpenGL is found	local only
OPENGL_INCLUDE_DIR	Include directory	local only
OPENGL_LIBRARIES	OpenGL Library directory	local only
GLU32_LIBRARIES	glu library directory	local only
DEP_LIBS	Library dependencies for linking	Appended

OpenGL glut The OpenGL glut libraries are found and added next. They are standard libraries on Mac and Linux, but aren't included on Windows. As a result, this uses the freeglut library for Windows, using a version provided with the Smoldyn package. From the freeglut documentation, one is supposed to use glut.h if only glut functionality is wanted, and freeglut.h if additional functionality is used that is only supported by freeglut; thus, this should only use glut.h, although parts of it incorrectly use freeglut.h. The `include(FindGLUT)` statement is used where possible; it automatically defines the variables `GLUT_FOUND`, `GLUT_INCLUDE_DIR`, and `GLUT_LIBRARIES`.

Variable	Description	Purpose
HAVE_GLU_H	TRUE for header file glut.h	smoldynconfig.h
HAVE_GL_GLU_H	TRUE for header file GL/glut.h	smoldynconfig.h
HAVE_GLU_GLU_H	TRUE for header file GLUT/glut.h	smoldynconfig.h
HAVE_GL_FREEGLUT_H	TRUE for header file GL/freeglut.h	smoldynconfig.h
GLUT_FOUND	TRUE if glut is found	local only
GLUT_INCLUDE_DIR	Include directory	local only
GLUT_LIBRARIES	glut library directory	local only
DEP_LIBS	Library dependencies for linking	Appended

LibX11 This section used to add the LibX11 library, which was apparently only needed for a static build on a Mac. However, the section is fully commented out at present, so it presumably wasn't needed.

LibTiff This section adds the LibTiff library, to enable the user to save TIFF format images of the simulation. I haven't been able to get LibTiff to work correctly on Windows yet, so this section only applies to non-Windows systems.

Variable	Description
HAVE_LIBTIFF	TRUE if LibTiff is found, FALSE if not
TIFF_INCLUDE_DIR	Include directory
TIFF_LIBRARY	Library directory
DEP_LIBS	Appended: library dependencies for linking

Zlib This section adds the Zlib library.

Variable	Description
HAVE_ZLIB	TRUE if Zlib is found, FALSE if not
ZLIB_INCLUDE_DIRS	Include directory

2 Smoldyn code and build system

ZLIB_LIBRARIES	Library directory
DEP_LIBS	Appended: library dependencies for linking

Libiconv This section adds the Libiconv library.

Variable	Description
HAVE_INCONV	TRUE if Libiconv is found, FALSE if not
ICONV_INCLUDE_DIRS	Include directory
ICONV_LIBRARIES	Library directory
DEP_LIBS	Appended: library dependencies for linking

VTK This section adds the VTK library, used for saving graphical output to VTK format files. The VTK source code is in the Smoldyn source/vtk subdirectory.

Variable	Description
VTK_FOUND	TRUE if VTK is found, FALSE if not
VTK_INCLUDE_DIRS	Include directory
DEP_LIBS	Appended: library dependencies for linking

NextSubvolume This section adds the code for hybrid simulation using lattices, which is the nsv code written by Martin Robinson. All of the source code is in the Smoldyn subdirectory source/NextSubVolume. This section used to call the CMakeLists file in source/NextSubVolume with the add_subdirectory statement, but then I simplified the build by moving all of that source into this CMakeLists.txt file. The result is a fairly simple build. This section adds code for signal.h if MinGW is used, and adds a boost include directory. It also appends the lists of source files and header files with the NSV code.

Variable	Description
SIGNAL_H_DIR	Directory to signal.h, only for MinGW
Boost_INCLUDE_DIR	Include directory for Boost, which is in the NSV source
SRC_FILES	Appended: Smoldyn source files
HEADER_FILES	Appended: Smoldyn header files

Targets Here, the target building is set up. The two targets are Smoldyn and Libsmoldyn. Smoldyn has an executable, with source files, main files, and header files. Libsmoldyn comes in both shared and static versions, both of which have source files and header files. No variables are defined here.

Python module This section simply adds two CMake subdirectories, for pybind11 and the python code that Dilawar Singh wrote.

Install This section describes how to install the software. Just the Smoldyn executable is installed if that is the only target. If the Libsmoldyn target is built as well, then this also installs: the shared and static Libsmoldyn library files and the header files: libsmoldyn.h, smoldyn.h, smoldynconfigure.h.

Package This was an effort to create a nice package with CPack. I haven't finished this topic yet.

Testing Here are a few tests of Libsmoldyn. This section adds the examples subdirectory.

2.3.4 CMake subdirectories

The CMakeLists file lists a few subdirectories.

(1) The subdirectory “source/pybind11” includes the pybind11 library, which offers “Seamless operability between C++11 and Python... pybind11 is a lightweight header-only library that exposes C++ types in Python and vice versa, mainly to create Python bindings of existing C++ code”. It is widely used and can be downloaded from github. It does not require pre-compiling like the other dependencies because it is compiled with the rest of Smoldyn. There is no custom code here, but only code from the pybind11 project.

(2) The subdirectory “source/python” includes Python binding code and was written, by Dilawar, specifically for Smoldyn.

(3) The subdirectory “examples” is for automatic testing, which has been partially set up but not completed. To run these examples, enter “make examples”.

(4) The subdirectory “docs” is for generating documentation, meaning converting the manual and code documentation to web-friendly formats and running Doxygen on all of the code. Run this by using the `-DOPTION_DOCS=ON` configure option.

Several different targets can be built. The core software has the Smoldyn and LibSmoldyn targets, which are selected or unselected using the CMake options `OPTION_TARGET_SMOLDYN` and `OPTION_TARGET_LIBSMOLDYN`, described above. I’m not sure if `smoldyn_static` and `smoldyn_shared` qualify as targets or not, but they are both built as part of LibSmoldyn. When building, these targets create a Smoldyn executable, called just smoldyn on a Mac or smoldyn.exe on Windows, and static and shared libraries called libsmoldyn_static.a and libsmoldyn_shared.dylib on a Mac. Also, one can write `make install`, although this isn’t really a different target. In addition, `make examples` does some testing.

Within the python subdirectory, additional targets are: `wheel`, `pyinstall`, `pyinstall_venv` (virtual environment), `pyuninstall`, `pydevel`, and `doc_api_html`. I don’t think that most of these targets need to be called specifically, but are instead called automatically when running just “make”. These targets are fairly simple, just running a line of Python code in most cases (it runs the setup.py program, which is somewhat involved, but carries out the installation). The `pyuninstall` line runs pip, and removes the Smoldyn version that pip installed, but not the Smoldyn version that `pyinstall` installs.

2.3.5 CMake file locations

Running CMake from the build directory with `cmake ..` creates a lot of files in the build directory. I don’t believe that it creates files anywhere else. Some notable ones are:

- `cmake_install.cmake`. A CMake file that includes installation instructions. It’s not very readable.
- `CMakeCache.txt`. List of variables that were defined in the CMakeLists.txt file and their values.
- `CMakeFiles` directory. Lots of files are in here, plus subdirectories that will be used for compiled object code.
- `Makefile`. The primary makefile for building.
- `smoldynconfigure.h`. The C configuration file that was written by CMake based on `source/smoldynconfigure.h.in`.

2 Smoldyn code and build system

Running `make` then builds Smoldyn. As it goes along, it writes object files to `build/CMakeFiles/smoldyn_static.dif/source/....`, or to other subdirectories within `build/CMakeFiles`. It puts the compiled results at the top-level of the build directory, where these include `libsmoldyn_shared.dylib`, `libsmoldyn_static.a`, and the smoldyn executable.

When it gets to the Python code, it first builds the C++ target `lib_pysmoldyn.a`, which appears to have exactly the same object files as the other targets, without any Python-specific files. This target gets put into the `build/source/python` directory. Next, `CallbakeFunc.cpp` and `module.cpp` get compiled. I think they get linked to `lib_pysmoldyn.a` and then turned into `_smoldyn.cpython-310-darwin.so`, which gets saved in `source/python/smoldyn` (this is *not* within the build directory). The `.so` suffix is for a shared library, using the standard Linux suffix. No other files appear to get written to the source directory.

Next, the Python source files (the source tree) get copied from the source directory to the appropriate places in the build directory. Then it builds the Python wheel; I'm not sure if this is a sub-portion of the copying step or the same thing. It checks for a pre-existing wheel in `/Users/sandrews/Library/Python/3.10/lib/python/site-packages`; it says that one is found, but that seems odd to me because I can't find any Smoldyn files in there at all. Maybe it's looking for some other wheel. In any case, it next copies the Smoldyn Python files (`smoldyn.py`, `__init__.py`, `types.py`, etc.), and also the newly created `_smoldyn.cpython-310-darwin.so` over from `source/python/smoldyn` to the `build/lib/smoldyn` directory. Then, it creates a directory for the wheel and copies the same files over into there. Then, it does a bunch of stuff with egg information, although I was under the impression that eggs and wheels are alternate approaches for similar results, and weren't used together, but I guess not. After that, it creates an actual wheel file, and adds the Python files, the `_smoldyn.cpython-310-darwin.so` file, and some more files with wheel metadata and similar stuff. Finally, the wheel directory is removed.

By the end of the `make` step, the important files are: `libsmoldyn_shared.dylib`, `libsmoldyn_static.a`, the smoldyn executable, and the wheel file in the top-level build directory. One file was added to the source directory, and there are vast numbers of other files within subdirectories of the build directory, but nothing else.

Next, running `sudo make install` does lots of file copying. This starts by repeating the `make` functionality, although there's no need to recompile the code since that's already done. However, this does repeat checking for the wheel, finding that it's in `/Users/sandrews/Library/Python/3.10/lib/python/site-packages`, even though it isn't. Then, this repeats all steps of the Python copying, wheel directory creation, egg creation, wheel creation, copying into the wheel, and removing the wheel directory.

Finally, this copies files into system locations. Here is the `install_manifest.txt` file, which lists all of the install files and locations:

```
/usr/local/bin/smoldyn
/usr/local/lib/libsmoldyn_shared.dylib
/usr/local/lib/libsmoldyn_static.a
/usr/local/include/libsmoldyn.h
/usr/local/include/smoldyn.h
/usr/local/include/smoldynconfigure.h
```

There's no evidence that this copies the wheel anywhere.

Running `sudo make pyinstall` installs the Python code, *probably to the wrong place*. This make target is listed in the `source/python/CMakeLists.txt` file, where it simply says to run `setup.py` and install in the `${CMAKE_INSTALL_PREFIX}` directory, which is `/usr/local` for me. It puts the files

in `/usr/local/lib/python3.10/site-packages`. However, this isn't actually useful for me because my Python doesn't look there for files.

Instead, run `pip install smoldyn-2.64-...whl`. This installs the wheel to the correct location. For me, that is `/opt/local/Library/Frameworks/Python.framework/Versions/3.10/lib/python3.10/site-packages/smoldyn/`.

2.3.6 Python module after building

Rather than installing the Python module, it's also possible to import the build version into Python, directly from the `build/py` directory. Add this directory to your `PYTHONPATH` temporarily with `export PYTHONPATH=/path/to/Smoldyn/build/py:$PYTHONPATH`. With this, the module can be accessed from any directory. Also, somehow, it's possible to add the module directory to `sys.path`, which is then a permanent solution. However, better is to actually install it with `pip`.

Note that it's possible to see where the library is imported from by checking `smoldyn.__file__` while in Python. For example,

```
>>> import smoldyn
>>> print(smoldyn.__file__)
/home/dilawars/Work/FORKES/Smoldyn/build/py/smoldyn/__init__.py
```

2.3.7 Cross-compiling for Windows with MinGW

Smoldyn was cross-compiled for Windows from Mac using MinGW up to version 2.58, although it didn't work well for the last several of those versions. Versions 2.59 to 2.61 used a different build system in which I copied source code files from Mac to Windows and then built on a Windows computer using my own compile script and the MinGW compiler. For version 2.62, I downloaded the MSVC compiler and returned to using CMake, but now running on the Windows computer. I'd like to get back to cross-compiling, but it's not working currently. Following are cross-compiling notes.

In one failed attempt, I got MinGW from MacPorts using `port install mingw-w64`. This installed the meta-package `mingw-w64`, along with two sets of other packages: `i686-w64-mingw32-...` and `x86_64-w64-mingw32-...`, where the `...` refers to the following 5 endings: `binutils`, `crt`, `gcc`, `headers`, and `winpthread`. The `i686-w64-mingw32-gcc` compiler sort of works, but not really. I wrote the standard `hello.c`, which compiled nicely on my Mac with `gcc -Wall hello.c -o hello`. To cross-compile, I tried to compile with `i686-w64-mingw32-gcc -Wall hello.c -o hello.exe` but this returned the error message `"i686-w64-mingw32-gcc: error trying to exec 'cc1': execvp: No such file or directory"`. However, this did work when I prefaced the instruction with `"sudo"` and then entered my password. This clearly implied that the issue had to do with permissions, but I was never able to figure out what was wrong and my query to Stack Overflow about this was useless.

I had success with MinGW when I switched to compiling with the `x86_64-w64-mingw32-gcc` compiler, which I also downloaded with MacPorts (this didn't work in 2018 and 2019, but magically did again in 2020). It installed files with the same root but ending in `-binutils`, `-ctr`, `-headers`, and `-winpthread`. Compiling hello world using the line `"x86_64-w64-mingw32-gcc -Wall hello.c -o hello.exe"` worked without permission issues. This also successfully built the `wrl2smol` utility with simply `"x86_64-w64-mingw32-gcc -Wall ../source/wrl2smol/wrl2smol.c -o wrl2smol.exe"`. It works excellently with non-graphics applications, including `wrl2smol`, `SmolCrowd`, and `Smoldyn` without OpenGL. Adding in OpenGL still builds properly, but the result won't run on a Windows computer, returning an error that it won't run. Debugging this with a tool called "Dependency walker"

2 Smoldyn code and build system

indicated that lots of Windows libraries were missing, such as “api-ms-win-appmodel-runtime”. I wasn’t able to solve the problem.

Part of the MacPorts MinGW download is a lot of library code that’s pre-compiled for MinGW. The useful portions seem to be in `/opt/local/x86_64-w64-mingw32/include/` for the header files and `/opt/local/x86_64-w64-mingw32/lib/` for the source files. These include most of the libraries that Smoldyn uses. It didn’t used to include glut, but the 2020 version does. Back when it didn’t include glut, I downloaded the “freeglut 3.0.0 MinGW Package” from <https://www.transmissionzero.co.uk/software/freeglut-devel/>. The download was in the Smoldyn source directory, from which I copied the include and lib directory contents to the `/opt/local/x86_64-w64-mingw32` directories.

For Libtiff, I had several failed attempts. I copied the i386 tiff*.h files from MinGW directory to `/opt/local/x86...include` directory. They seem ok, but they’re clearly for a different architecture. I also tried the i386 libtiff.a static libtiff library, but Smoldyn building complained that it wasn’t compatible. I also tried downloading libtiff.a from many different websites, but got the same result every time, that they weren’t compatible. I gave up on offering tiff support for Windows, but I now think that I might be able to download the Libtiff source code and cross-compile myself. I successfully compiled the Libtiff source code on Windows with MSVC, but haven’t tried yet for MinGW.

2.3.8 Building for Windows on Windows with MinGW

I think I have two different versions of MinGW on my Windows computer. For the second one, I downloaded MinGW using the mingw-get-setup.exe installer, from osdn.net, which was remarkably hard to figure out. I next ran “mingw-get.exe install mingw32-gcc” and it seemed to install the compiler.

For versions 2.59 to 2.61, I built Windows versions on a Windows computer with MinGW. I hadn’t figured out CMake there yet, so I copied the smoldynconfigure.h file and all of the source files to the Windows computer, compiled, copied the .exe file back to the Mac, and continued with the build process. I distributed old version of SmolCrowd and wrl2smol in the Windows package (they only change very rarely, so the old ones were also the current ones).

2.3.9 Building for Windows on Windows with Visual C

For version 2.62 and maybe future versions, I’m compiling the Windows version on a Windows computer using the Visual C compiler (MSVC). After lots of work, this now functions using the CMake build system. The method starts with creating a build directory, as usual (I haven’t figured out how to remove all of the contents of the build directory, so I clean it out by moving up to the Smoldyn directory, removing build with “rmdir /s build”, and then recreating it with “mkdir build”, “cd build”). Then, in the build directory, “cmake ..” as usual, with any options, and “msbuild smoldyn.sln /p:Configuration=Release”. The final option is essential because Visual Studio ignores the CMake directive to make this a release build, creating a debug build instead, which then needs to link to debug libraries that most people don’t have. This option tells Visual Studio that it should actually be a release build.

I haven’t figured out how to build alternate targets yet. The CMake documentation for `build.command` (and an email from Ciaran) suggests that this might be possible with “cmake .. -target wheel”, but this doesn’t work for me. Also, the Microsoft documentation at <https://docs.microsoft.com/en-us/visualstudio/msbuild/msbuild-command-line-reference?view=vs-2019> say that

I can build with “msbuild smoldyn.sln -target:wheel”. This doesn’t seem totally wrong, but it says that the target doesn’t exist. Yet another approach is to change directories to build/source/python and enter “msbuild wheel.vcxproj”. This seemed promising for me but led to the error “invalid command ‘bdist_wheel’”, which a Google search seemed to suggest arose from not using pip. Also, I can’t seem to figure out how to get pip to work on that computer, so perhaps that’s the problem.

Here is some more advice, from Dilawar on 11/26/20.

- `cmake .. -DOPTION_PYTHON=ON -DOPTION_EXAMPLES=ON`
- `cmake --build . --config Release`
- `ctest -C Release`

The build step should install the wheel.

2.4 Unit and regression testing

Smoldyn does not include unit tests in their purest form. Instead, I tested Smoldyn’s algorithms, both for qualitative and quantitative performance, using simple Smoldyn configuration files, which function as unit tests. Many of these test files are in the examples directory. Individual files focus on specific algorithms, such as diffusion, unimolecular reactions, absorbing surface interactions, etc. Simply watching the simulation graphics is typically adequate for assessing qualitative performance, while data analysis and comparison against analytical theory is generally required for assessing quantitative performance.

The examples/S95_regression directory includes files used for regression testing. This directory includes many of the original unit tests, although typically modified for longer time steps, shorter total run times, and a fixed random number seed. In addition, I removed the original output for these unit tests and instead instructed them to output all molecule positions at the end of the simulation. The idea is that this is a very sensitive way of detecting whether all interactions during the simulation were the same between two runs, or not. This directory also includes a Python script `regression.py`. This script runs all unit tests, outputs their results to a subdirectory, and compares the results to those that were previously generated. All differences are reported, enabling detection of potential bugs.

To add a new unit test to the regression testing suite, simply make sure that the unit test runs relatively quickly (a few seconds) and has no text output. Then, copy and paste the top few lines from any of the current unit tests (e.g. `bounce2.txt`), so that the new unit test will output molecule positions or other relevant output. Finally, list the unit test name in the Python script.

2.5 GPU Smoldyn

Two versions of GPU Smoldyn have been written, one by Lorenzo Dematté, and one by Denis Gladkov. I never managed to get either of them to work. However, here are some dependencies that I downloaded while trying to get Denis’s version to work

For GPU Smoldyn, you will need some other things too. First is the CUDA library, which is from NVIDIA at <https://developer.nvidia.com/cuda-downloads>. This downloaded and installed itself. Next is the GLEW library (OpenGL Extension Wrangler library), which is from <http://glew.sourceforge.net>. This builds with simply “make” and “make install” (no “./configure” required). Next, the CUDPP library is from <http://code.google.com/p/cudpp/>. It builds with CMake (make a build directory, change to that directory, enter “cmake ..”, “make”,

2 Smoldyn code and build system

and “sudo make install”). For some reason, my build did not install the `cuda_config.h` file, so I had to do so by hand. From the CUDPP build directory, I entered “sudo cp ../include/cudpp_config.h /usr/local/include/” and that fixed problems.

The Boost library is from <http://www.boost.org/users/download/>. This doesn’t get installed with an installer, but instead the whole directory gets copied over. It didn’t work when I put it in a system location, but did work when I copied the “boost” subdirectory into the GPU code directory (Smoldyn/trunk/GPU/Gladkov/smoldyn-gpu-dg/).

2.6 Code techniques and tricks

2.6.1 New releases

To release a new version, there are many things to do.

- Go through any libSteve documentation that has changed and print to pdf all of the new files.
- Describe the changes in the history section of this code documentation.
- Update the version number in: (1) the top-level CMakeLists.txt file, (2) the Code documentation, (3) the User’s manual, (4) the release shell script, (5) the install batch file for Windows, (6) CHANGELOG.md.
- Run the regression tests using the latest version number and confirm that all results are correct.
- Run my release shell script (it’s not on git and is customized for my computer system, but I’m happy to share it with others if they want).
- Commit changes to github. Tag the commit. Example: `git tag -a v2.65 -m "release Version 2.65"; git push origin v2.65`.
- Update the www.smoldyn.org webpage. Upload the new files, delete the old files, and update the `index.html` and `download.html` web pages.

2.6.2 Formatting

I like a very compact C source code style, where the indentation shows the code depth rather than the braces. It just was that way for the most part, but then the VCell team added code in their style, and then Dilawar reformatted the code to yet another style. I figured out how to partially revert to my style, and also standardize others’ additions. I use the “uncrustify” software, which I downloaded from SourceForge at <https://sourceforge.net/projects/uncrustify/>. The package currently resides in my Smoldyn top level directory. I wrote a configuration file called `ssa.cfg`, which I use to format code; it is in the uncrustify directory.

As an example, to reformat `smolreact.c`, enter: `uncrustify -c ../../uncrustify/ssa.cfg smolreact.c`. This creates a new file, `smolreact.c.uncrustify`. However, this still leaves lots of things that I like to fix. First, remove closing braces on their own lines. To do so, do a find and replace, first replace all “`}`” with “`}`”, repeating as many times as necessary. Then replace all “`\n}`” with “`}`”. Another thing to look for is multi-line for statements. It would be nice if these could be removed by uncrustify, but it doesn’t.

For the Python code, Dilawar uses the black software, available at <https://pypi.org/project/black/>. He says that if I change the Python code, just run `black filename.py` to auto-format the code, thus keeping the format consistent.

Similarly, he uses clang-format for C++ files, but is less worried about format consistency there.

2.6.3 Code merging

Here are some options for merging code.

Minimalist text based merging can be done as follows:

```
mydir="../../../gccCode/Library"
    set variable
```

```
diff Geometry.c $mydir
    compares local version with version in $mydir, and prints out differences.
```

```
grep -n 'Geo_Sphere_Normal' ../Smoldyn/source/*.c
    finds all lines of Smoldyn source code that call Geo_Sphere_Normal.
```

Alternatively, XCode offers FileMerge.app, which is at `/Developer/Applications/Utilities/FileMerge.app`. It's very easy to use. Yet another option is Eclipse, at www.eclipse.org, which works well.

2.7 Documentation

The documentation directory is “docs”.

This directory contains original hand-written documentation files for Smoldyn, it's utility programs, and the libraries that Smoldyn calls. Essentially all of these files were written by me (Steve). An exception is that the BioNetGen documentation is a published paper on BioNetGen by Faeder, Blinov, and Hlavacek. The documentation specifically for Smoldyn is within the Smoldyn subdirectory. Here are the User's Manual, figures for it, and this programmer's documentation, along with a few especially relevant published papers about Smoldyn. Unfortunately, LaTeX requires far too many files for a single document, but that's just how it is.

The “docs” directory also contains things for automatically converting documentation from the hand-written files to rst and other file formats. These procedures are run through the “docs/CMakeLists.txt” file. It can be used as a submodule to the top-level cmake file (enabled by `-DOPTION_DOCS=ON`) or run as a standalone cmake project i.e., `cd docs; cmake .; make docs`. The CMakeLists.txt file does the following things:

(1) It converts the `SmoldynManual.tex` and `SmoldynCodeDoc.tex` to rst format, which is the restructured text format. This format is used for Python documentation. This conversion is done by calling the pandoc utility, which is a “universal document converter”, used for converting between markup languages (pandoc is available for Macs through MacPorts and elsewhere). This way one can update the LaTeX file in the documentation folder and have the rst docs in sync with them. Pandoc is a soft dependency; if it is not found, the rst files are not updated and old ones are used.

(2) The CMake file installs any software dependencies that are needed, which are listed in the `docs/requirements.txt` file, using the Python pip utility, which seems a bit concerning because it changes the system without warning the user. These dependencies are (plus a few others that are components of the ones listed here):

2 Smoldyn code and build system

Pweave — “A scientific report generator and a literate programming tool for Python. Pweave can capture the results and plots from data analysis and works well with NumPy, SciPy and matplotlib. It is able to run python code from source document and include the results and capture matplotlib plots in the output.”

sphinx — “Sphinx converts reStructuredText files into HTML websites and other formats including PDF, Epub, Texinfo and man.”

breathe — “Breathe provides a bridge between the Sphinx and Doxygen documentation systems. It is an easy way to include Doxygen information in a set of documentation generated by Sphinx. The aim is to produce an autodoc like support for people who enjoy using Sphinx but work with languages other than Python. The system relies on the Doxygen’s xml output.”

rst_include — “Since you can not include files into RST files on github or PyPi, you can resolve such imports with this software. That means you can locally write your RST documents (for instance with pycharm) and use there the `.. include:` option to include other RST Files or code snippets into Your Document. Afterwards, you can run this software to create a monolithic README.rst that can be viewed on Github or Pypi.”

commonmark — “A strongly defined, highly compatible specification of Markdown. Markdown is a plain text format for writing structured documents, based on formatting conventions from email and usenet.”

recommonmark — “A docutils-compatibility bridge to CommonMark, enabling you to write CommonMark inside of Docutils and Sphinx projects.”

smoldyn — This installs Smoldyn using pip.

(3) It runs Doxygen on all of the source code files to generate source code documentation as rst files. This step is configured by the Doxyfile.in file, which is just a long list of configuration switches.

(4) It runs Sphinx and Breathe to convert the rst documents to html documents. It then runs Sphinx and Doxygen to create documentation for the code.

(5) It uses Sphinx-autobuild to create live HTML. I’m not sure what live HTML is versus just HTML.

Dilawar wrote: The readthedocs service runs docs/conf.py file which has been tweaked to run doxygen and sphinx on RTD server. The RTD website should have the same text as the documentation folder now, with some formatting issues.

I ran this documentation build process, and it required lots of software installations first. I installed most of them using MacPorts, but some weren’t available there. For those, I installed using pip. The combination worked well. The final result was a huge number of warnings from the documentation tools, but essentially no errors. Actually, there was one error, which was that sphinx complained about a character in a .tex file that was not UTF-8. The solution was to search for it using “`grep -axv ‘.*’ SmoldynCodeDoc.tex`”, which successfully located the offending character.

Running “`sudo make install`” writes the documentation to `/usr/local/share/doc/smoldyn-doc/sphinx`.

2.8 Continuous Integration

Another new directory is called “scripts”. It includes development related scripts for building Smoldyn on online platforms like Travis, Open Build Service, OSX, and docker. I don’t believe

that these are complete yet. The `build_wheels_osx.sh` script has stuff about homebrew, which I avoid, so this is going to take some work before I'll want to run it.

2.9 Interfacing with Biosimulators, Simularium, and Vivarium

Smoldyn is interfaced with several other tools.

2.9.1 Biosimulators

Biosimulators is an online access point to a wide variety of simulators, including Smoldyn. It was headed up initially by Jonathan Karr and his student Bilal Shaikh. Jonathan emailed me about joining the project on May 17, 2021. Afterward, he and Bilal, along with Dilawar and myself, integrated Smoldyn into the project. The work was published in a paper titled “BioSimulators: a central registry of simulation engines and services for recommending specific tools” in *Nucleic Acids Research* 50:W108-W114, 2022. Jonathan and Bilal both left Biosimulators around the time the paper was published, and the work was taken over by others.

In May, 2023, I started working with Eran Agmon and Logan Drescher (both at UCHC) to improve Smoldyn integration into Biosimulators. The goal at this point is to get the E. coli Min simulation running, but there are still some bugs to fix at their end. Ideally, one would go to the Biosimulators website at <https://biosimulations.org>, click on “Run simulations”, click on “Utilities/Create a project”, upload a Smoldyn file, check that the website fields auto-filled correctly, and click on the “Simulate” button. So far, this isn't working.

For Biosimulators to display graphical output from Smoldyn, configuration files need to have the following commands, or something equivalent.

```
output_files FILER00Tout.txt
cmd i 0 TIME_STOP 1 executiontime FILER00Tout.txt
cmd i 0 TIME_STOP 1 listmols FILER00Tout.txt
```

Then, the Smoldyn output is converted to graphics using the Simularium conversion tool. It is at https://github.com/simularium/simulariumio/blob/main/examples/Tutorial_smoldyn.ipynb.

2.9.2 Simularium

Simularium is a project of the Allen Institute for Cell Biology and is run by Blair Lyons. I was invited to join it on Dec. 2, 2020. It is a simulation viewer for previously written and submitted models. They wrote the converter that converts Smoldyn molecule position output to Simularium input.

2.9.3 Vivarium

Vivarium is another online tool for multiple simulators. I was invited to join it around May, 2021. It was being run by Eran Agmon (Stanford University, Covert lab) and Ryan Spangler (Allen Institute). Eran and Ryan added Smoldyn to Vivarium using Smoldyn's Python interface, which sort of worked, but I don't think that we fully completed the task. As of May, 2023, Eran says that Vivarium is still going, but it sounds like Biosimulators is of greater interest and will likely replace Vivarium.

3 Python binding code

This section focuses on the code in the source/python directory, essentially all of which was written by Dilawar Singh. Most of this code is in C++ but it's here to support the Python bindings.

3.1 module.cpp

```
int init_and_run(const string &filepath, const string &flags, bool wflag, bool
    quit_at_end = false)]
```

This function is an alternate point of entry to `main`, from my `smoldyn.c` code. It is called by `PYBIND11_MODULE` with the `loadModel` low-level Python function.

Send in the `filepath`, including both the path and the filename in `filepath` and the Smoldyn flags in `flags`. In addition, `wflag` is another copy of the suppress warnings flag, perhaps only for whether output files should be overwritten, and `quit_at_end` is a flag for whether Smoldyn should just quit when the simulation is done. This function passes these things on to `simInitAndLoad` and `simUpdateAndDisplay` (both in `smolsim.c`). Here, `cursim_` is the current simulation structure. This also opens output files and runs the simulation with `smolsimulate` or `smolsimulategl`, as appropriate. At the end, it frees the simulation structure and returns. Returns 0 for success or an error code for failure.

```
PYBIND11_MODULE(_smoldyn, m)
```

This is a macro that creates a function for the interface between C code and Python. From the website <https://pybind11.readthedocs.io/en/stable/basics.html>, with some edits to make it apply to the code here: “The `PYBIND11_MODULE()` macro creates a function that will be called when an `import` statement is issued from within Python. The module name (`_smoldyn`) is given as the first macro argument. The second argument (`m`) defines a variable of type `py::module_` which is the main interface for creating bindings. The method `module_::def()` generates binding code that exposes the Smoldyn functions to Python.”

This code exposes the enumerated types, the members of the simulation structure, and the C API to the Python interface. Most of this function is simply a wrapper of the C API. The python functions created at the end of this function are not in the C API, but probably should be. Most of their code is in `module.h`.

Also, `module.cpp` has simple utilities, which really should be in the C API but weren't added to it. Most of these functions are called by the `PYBIND11_MODULE` function. Many of these functions could also be usefully merged into the `PYBIND11_MODULE` code to improve overall simplicity.

The `module.h` header file includes global variables and declarations for all of these functions.

3.2 util.h

`util.h` includes a few inline utility functions, including `splitPath` for splitting a filename path into separate root and filename strings (used in `module.cpp` `init_and_run`) and `color2RGBA` for

3 Python binding code

converting a color string to RGBA numbers (this calls my `graphicsreadcolor` function).

3.3 Simulation.h

```
bool connect(const py::function& func, const py::object& target, const size_t
step, const py::list& args)
```

Creates a python callback for the Smoldyn main simulation loop. `func` is the function to be called, `target` is some sort of target for it (I don't know what this is), `step` is the number of time steps between callbacks, and `args` are arguments for the function that is called.

```
bool addToSimpPtrVec(simpPtr ptr)
```

Adds `ptr`, which is a simulation pointer, to the list of simulation pointers called `simpPtrs_`, which is declared as a global variable in this file. Doesn't do anything if the pointer was already in the list. Returns true if it was added and false if it was already there.

```
bool deleteSimpPtr(simpPtr ptr)
```

Searches for `ptr` in the list of simulation pointers `simpPtrs_`. If found, it is deleted. Also, the memory is cleared. Returns true if found and deleted and false if not found.

```
size_t getDim()
```

Returns the simulation dimensionality, using the value in the global `dim_` variable. This should probably be a C API function as well.

```
void setDim(size_t dim)
```

Sets the dimensionality of the current simulation structure, as well as the global `dim_` variable. This function should only be called when a simulation structure is first being set up. Otherwise, it is a dangerous function because lots of memory allocation in the simulation structure depends on the dimensionality, so really shouldn't be modified without completely erasing and rebuilding the simulation structure.

```
void printSimpPtrNotInitWarning(const char* funcname)
```

Simply prints a warning to say that the function named `funcname` is complaining about the `simpPtr` not being initialized.

```
void setRandomSeed(size_t seed)
```

Sets the random number seed for the current simulation to `seed` if this value is greater than or equal to 0, or to a randomly chosen value otherwise. This could be but isn't a wrapper for the C API function of the same name.

```
size_t getRandomSeed(void)
```

Returns the current random number seed. This should be in C API.

```
bool initialize()
```

This creates a new simulation structure, putting it into the global variable `curSim_`, using information stored in other global variables `lowbounds`, `highbounds`, and `dim` (this last one is retrieved through `getDim`). Returns true for success or false for failure.

```
void runUntil(const double breaktime, const double dt, bool display)
```

Runs the simulation until time `breaktime` using time step `dt`. Output is sent to the display if `display` is true. This is a wrapper for `smolRunSimUntil`.

3.4 *Callbackfunc.cpp and Callbackfunc.h*

`bool run(double stoptime, double dt, bool display)`

Runs the simulation up to time `stoptime`. This is a wrapper for `smolRunSim`.

`void setBoundaries(const vector<pair<double, double>>& bounds)`

`void setBoundaries(const vector<double>& lowbounds, const vector<double>& highbounds)`

Both functions set the boundaries for a simulation by writing them to global variables and then calling `initialize` to add them to the simulation structure.

`pair<vector<double>, vector<double>> getBoundaries()`

Gets and returns the boundaries.

`ErrorCode setDt(double dt)`

Sets the simulation time step.

`ErrorCode getDt(double dt)`

Returns the simulation time step.

`bool setModelpath(const string& modelpath)`

This takes in the path for the model, including the model filename, and splits it into the path and filename separately. It then assigns the results to the current simulation filepath and filename. The function always returns true.

3.4 **Callbackfunc.cpp and Callbackfunc.h**

These files include the code for the Python callback functions. Use things like `setFuncName`, `setStep`, `setTarget`, `setFunc`, and `setArgs` to register and set up callback functions.

`CallbackFunc::CallbackFunc()`

Constructor.

`CallbackFunc::~CallbackFunc()`

Destructor.

`bool CallbackFunc::evalAndUpdate(double t)`

Evaluate the callback function at time `t`. This also calls the target function if there is one, and otherwise uses a “magic string”, which I didn’t understand.

`bool CallbackFunc::isValid() const`

Returns whether the callback function has a size greater than zero, meaning that it’s valid.

`void CallbackFunc::setFunc(const py::function& func)`

Sets the callback function to `func`.

`py::function CallbackFunc::getFunc() const`

Gets and returns the callback function.

`void CallbackFunc::setFuncName(const string& fname)`

Sets the callback function name to `name`.

`string CallbackFunc::getFuncName() const`

Gets and returns the callback function name.

3 Python binding code

```
void CallbackFunc::setStep(size_t step)
    Sets the step size for the callback function to step.

size_t CallbackFunc::getStep() const
    Gets and returns the stepsize for the callback function.

void CallbackFunc::setTarget(const py::handle& target)
    Sets the target function for the callback function to target.

py::handle CallbackFunc::getTarget() const
    Gets and returns the target for the callback function.

void CallbackFunc::setArgs(const py::list& args)
    Sets the arguments for the callback function to the list args.

py::list CallbackFunc::getArgs() const
    Gets the arguments for the callback function.
```

3.5 Python scripts in smoldyn directory

I tried to determine which scripts get imported and in what order by adding print statements to the beginning of them. If I start a Python session and then enter “import smoldyn”, it runs the following scripts, in order: `__init__.py`, `smoldyn.py`, `types.py`, `utils.py`. My understanding is that each of these imports the next in sequence, so `__init__.py` imports `smoldyn.py`, which then imports `types.py`, which then imports `utils.py`. Running a Python Smoldyn simulation goes through the same files when the “import smoldyn” statement is reached, and they are not imported again, at least from the top. It doesn’t appear that `__main__.py` is ever imported at all. The files are listed below in the same order that they actually imported.

`__init__.py`

This file only has a few lines since it’s only role is to import `smoldyn.py`. When I’m in a Python session and enter “`smoldyn.__file__`”, the result is a path to the `__init__.py` file, again showing that this is the top-level file.

`smoldyn.py`

This file creates the “user API” for the Python bindings. In the process, it defines all of the user API classes, and the setting and getting functions.

`types.py`

This declares a color type as a union of a string name, and a RGBA vector. It also defines the class `Color`, which can convert name to RGBA.

`utils.py`

This defines a `load_model` function with a path and arguments, which then sends the information onto the compiled Smoldyn code. It also converts colors from string to RGBA, using matplotlib if available or Smoldyn color conversion if not.

__main__.py

This file parses the command line arguments using the `argparse` standard library and then runs the model using `smoldyn.load_model()`. This file includes the Python command-line options, including “overwrite”, “quit-at-end”, and “args”. This code doesn’t necessarily run. According to Dilawar, if Smoldyn is started with “python3 -m smoldyn template.txt”, then this will execute the `__main__.py` file, with command line arguments passed to it. Here, the Python package is being used as a proxy for the Smoldyn executable.

3.6 Python calling sequence

The Python interface can be better understood by considering what happens when it is called.

Example 1. The user’s Python script includes `s.addReaction("rxn", subs=[A, B], prds=[C], rate=-1, binding_radius=0.05, reaction_probability=0.1)`. This sends control to the `addReaction` method in `smoldyn.py` (line 2130). That creates a Python reaction object, which is defined in `smoldyn.py` line 1272 and instantiated in the `__init__` method in line 1286. This method adds the reaction to the simulation with `self.simulation.addReaction`, which connects to a `module.cpp` line that reads `.def("addReaction", ...` (line 1129) and then on to the `libsmoldyn.cpp` function `smolAddReaction` (line 1775). This line, finally, connects to the Smoldyn core code, by calling `RxnAddReaction`, which is in `smolreact.c`.

Control then moves back up the stack to the reaction instantiation in `smoldyn.py`, where the rate needs to get set next. The instantiation runs the method `self.setRate`, which is also in `smoldyn.py`, in line 1396 and reads `def setRate(self, rate, reaction_probability=0.0, binding_radius=0.0):..` This has the line `self.simulation.setReactionRate`, which connects to `module.cpp` `.def("setReactionRate", ...`, goes to `libsmoldyn.cpp` function `smolSetReactionRate`, and then on to the Smoldyn core with `RxnSetValue`.

The following table summarizes these connections.

File	Line	Code
User code		<code>s.addReaction("rxn", subs=[A, B], prds=[C], rate=1)</code>
<code>smoldyn.py</code>	2130	<code>def addReaction(...)</code>
<code>smoldyn.py</code>	1272	<code>class Reaction(object): ... def __init__(...)</code>
<code>module.cpp</code>	1129	<code>.def("addReaction", ...)</code>
<code>libsmoldyn.cpp</code>	1775	<code>enum ErrorCode smolAddReaction(simptr sim, ...)</code>
<code>smolreact.c</code>	2044	<code>rxnptr RxnAddReaction(simptr sim, ...)</code>
<code>smoldyn.py</code>	1396	<code>def setRate(self, rate, ...)</code>
<code>module.cpp</code>	1190	<code>.def("setReactionRate", ...)</code>
<code>libsmoldyn.cpp</code>	1887	<code>enum ErrorCode smolSetReactionRate(simptr sim, ...)</code>
<code>smolreact.c</code>	1721	<code>int RxnSetValue(simptr sim, ...)</code>

Example 2. For running the simulation, the user enters `s.run(...)`. This goes to `smoldyn.py` line 1946 where `run` is defined. This method calls `runSim` in line 1994. This might go to several possible locations. In `module.c`, `runSim` is defined in line 407, which then calls `Simulation::runSim`. Also, `module.c` defines `runSim` in line 1381, which calls `libsmoldyn.cpp` `smolRunSim`. Additionally, `simulation.cpp` has a function called `Simulation::runSim` in line 152. This last one calls `libsmoldyn.cpp` `smolRunSim`. So, everything appears to converge on `smolRunSim` before too long.

4 C code: files, macros, variables, etc.

The Smoldyn code is separated into several sets of files. (1) Library files, such as `math2.c`, are general-purpose C library files, nearly all of which I wrote. Smoldyn uses some of the functions in them, but far from all. (2) Each of these library files has a header, such as `math2.h`, that declares the structures and functions within that library file. These library files and headers are documented in separate documents, such as the file `math2_doc.pdf`. (3) The core Smoldyn source code is in files that begin with “smol”, such as `smolmolec.c`. Smoldyn uses all of these functions. The main entry point to the program is in the file `smoldyn.c`, in the main function. This file also includes some high level functions for running the simulation. The other files take care of different portions of the simulation, such as molecules, virtual boxes, or surfaces. I have tried to encapsulate the code so that functions in any file are allowed to read directly from any structure, but only the functions in the file that corresponds to a structure is allowed to write to it. (4) The Smoldyn core source code files share a single header file, which is called `smoldyn.h`. It declares all data structures and function declarations. This header and the core Smoldyn files are documented here and in part I of the documentation.

4.1 Smoldyn source files

file	function
<code>smolboxes.c</code>	virtual boxes
<code>smolcmd.c</code>	runtime interpreter commands
<code>smolcomparts.c</code>	compartments
<code>smoldyn.c</code>	top level functions
<code>smoldyn.h</code>	data structures and function declarations
<code>smolgraphics.c</code>	OpenGL graphics
<code>smolmolec.c</code>	molecules
<code>smolport.c</code>	ports
<code>smolreact.c</code>	reactions
<code>smolsim.c</code>	simulation structure and high level functions
<code>smolsurface.c</code>	surfaces
<code>smolwall.c</code>	walls

4.2 Constants and global variables

smoldyn.h

```
#define SMOLDYN_VERSION 2.16 // current Smoldyn version number
```

This is the current version number of Smoldyn.

```
#define DIMMAX 3
```

This is the maximum dimensionality permitted.

4 C code: files, macros, variables, etc.

```
#define VERYCLOSE 1.0e-12
```

Distance that is certain to be safe from round-off error during calculations.

```
enum StructCond {SCinit, SCok, SCparams, SClists};
```

This is used in multiple structures to report the structure condition. SCinit is for just initialized, or initial initialization; SCok is for fully updated and ready for use; SCparams is for complete except that internal parameters need computation; and SClists is for structure lists and maybe also parameters need computation.

smoldyn.c

```
simptr Sim;
```

Sim is a global variable for the current simulation structure. This is only used when graphics are being shown using OpenGL, because OpenGL does not allow variables to be passed in the normal way between functions.

4.3 Macros

Macros are defined near the top of smoldynfuncs.h. The global strings that they use are defined at the top of smolsim.cpp.

```
#define CHECK(A) if(!(A)) ErrorType=1;goto failure; else (void)0
```

This is a useful macro for several routines in which any of several problems may occur, but all problems result in freeing structures and leaving. Program flow goes to the label failure if A is false. Many people would consider both the use of a macro function and the use of a goto statement to be bad programming practice, and especially bad when used together. However, in this case it significantly improves code readability. As usual, partially defined structures should always be kept traversable and in good order so they can be freed at any time. The “else (void)0” termination of the macro is used so that if CHECK(...) is followed by an else, that else will refer to the prior if, and not to the CHECK. Because of the trailing else, compilers may complain if the CHECK macro follows an if and is not surrounded by braces.

```
#define CHECKMEM(A) if(!(A)) ErrorType=3;snprintf(ErrorString,STRCHARLONG,"Cannot  
allocate memory"); goto failure; else (void)0
```

Check that memory allocation worked.

```
#define CHECKS(A,...) if(!(A)) ErrorType=2;  
snprintf(ErrorString,STRCHARLONG,__VA_ARGS__); goto failure; else (void)0
```

This is identical to CHECK, except that it also copies the included string and other arguments to the variable ErrorString if a failure occurs.

```
#define CHECKM(A,...) if(!(A) || strmatherror(ErrorMath,1)) ErrorType=2;  
snprintf(strcat(ErrorString,ErrorMath),STRCHARLONG,__VA_ARGS__); goto failure;  
else (void)0
```

This is used after a call to `strmathsscanf`, which is used for reading numerical values. It works the same as the CHECKS macro, but it also checks for math errors using the `strmatherror` function and appends those to the end of the error string.

```
#define CHECKBUG(A,...) if(!(A)) ErrorType=4;
    snprintf(ErrorString,STRCHARLONG,__VA_ARGS__); goto failure; else (void)0
```

This is identical to CHECKS, except for creating a higher level of error type.

4.4 Local variables

It has proven useful to use consistent names for local variables for code readability. In places, there are exceptions, but the following table lists the typical uses for most local variables. This table is also quite out of date.

variable	type	use
a	double	binding radius for bimolecular reaction
b, b2	int	box address
blist	boxptr*	list of boxes, index is [b]
boxs	boxssptr	pointer to box superstructure
bptr	boxptr	pointer to box
bval	double	unbinding radius for bimolecular reaction
c	int	index of compartment
ch	char	generic character
cmd	cmdptr	pointer to a command
cmds	cmdssptr	pointer to the command superstructure
cmpt	compartptr	pointer to a compartment
cmptss	compartssptr	pointer to a compartment superstructure
d	int	dimension number
dc1, dc2	double	diffusion coefficients for molecules
dead	moleculeptr*	list of dead molecules, index [m]
difc	double*	list of diffusion coefficients, index [i]
dsum	double	sum of diffusion coefficients
dim	int	dimensionality of space
dt	double	time step
er	int	error code
erstr	char*	error string
face	enum PanelFace	face of a panel
flt1, flt2	double	generic double variable
fptr	FILE*	file stream
got	int[]	flag for if parameter is known yet
i	int	molecule identity, reactant number, or generic integer
i1, i2, ...	int	molecule identities
indx	int*	dim dimensional index of box position
itct	int	count of number of items read from a string
j	int	number of reaction for certain i
k	int	index of points within a compartment
lctr	int	line number counter for reading text file
line	char[]	complete line of text
line2	char*	pointer to unparsed portion of string
live	moleculeptr**	list of live molecules, index [ll][m]
ll	int	index of live list
m, m2, m3	int	index of molecule in list

4 C code: files, macros, variables, etc.

m1, m2, m3	int*	scratch space matrices of size dimxdim
mlist	moleculeptr*	list of molecules, index is [m]
mols	molssptr	pointer to molecule superstructure
mptr	moleculeptr	pointer to molecule
mptr1, mptr2	moleculeptr	pointer to more molecules
ms	enum MolecState	molecule state
name	char**	names of molecules, index is [i]
nbox	int	number of boxes
ni2o	int	value of nidentorder
nident	int	number of molecule identities
nl	int*	number of live molecules in a live list, index [ll]
nm, nm2	char[]	name of molecule, reaction, or surface
nmol	int	number of molecules in list
npnl	int	number of panels for a surface
npts	int	number of points for a surface panel
nprod	int*, int	number of products for reaction [r]
nrxn	int*	number of reactions for [i]
nsrf	int	number of surfaces
o2	int	order of second reaction
optr	int*	pointer to the order of reaction
order	int	order of reaction
p	int	reaction product number
p	int	panel number for surfaces
pfp	ParseFilePtr	configuration file pointer and information
pgemptr	double*	pointer to probability of geminate recombination
point	double**	list of points that define a panel [p][pt]
pos	double*	a position
prod	moleculeptr**	list of products for reaction [r], index is [p]
pnl	panelptr	pointer to a panel
pnls	panelptr*	list of panels
ps	enum PanelShape	panel shape
pt	int	index for points, for surfaces
r	int	reaction number
r2	int	reaction number for second reaction
rate	double*	requested rate of reaction [r]
rate2	double*	internal rate parameter of reaction [r]
rate3	double	actual rate of reaction
rev	int	code for reaction reversibility; see findreverserxn
rname	char**	names of reactions, index is [r]
rpar	double, double*	reversible parameter, index is [r]
rpart	char, char*	reversible parameter type, index is [r]
rprr	int*	pointer to reaction number
rxn	rxnptr	pointer to a reaction structure
rxn2	rxnptr	pointer to a second reaction structure
s	int	surface number
side	int*	number of boxes on each side of space, index [d]
sim	simprr	pointer to simulation structure
smpr	simprr*	pointer to pointer to simulation structure

4.4 Local variables

srf	surfaceptr	pointer to a surface structure
srfs	surfacessptr	pointer to a surface superstructure
step	double	rms step length of molecule or molecules
str1	char []	generic string
table	int**	table of reaction numbers for [i][j]
topd	int	top of empty molecules in dead list
total	int	total number of reactions in list
v1, v2, v3	int*	scratch space vectors of size dim
w	int	index of wall
wlist	wallptr*	list of walls, index is [w]
word	char []	first word of a line of text
wptr	wallptr	pointer to wall

5 C Code: structures and functions

Smoldyn is written in C, with a C style. The proper maintenance of structures, which are described below, is a central aspect of the program. In general, the basic objects are molecules, walls, surfaces, and virtual boxes, each of which has its own structure. In many cases, these items are grouped together into superstructures, which are basically just a list of fundamental elements, along with some more information that pertains to the whole list. Reactions, compartments, and ports aren't really objects, but are also among the core data structures. Finally, a simulation structure is a high level structure which contains all the parameters and the current state of the simulation.

An aspect of structures that is important to note, especially if changes are made, is which structures own what elements. For example, a molecule owns its position vector, meaning that that piece of memory was allocated with the molecule and will be freed with the molecule. On the other hand, a molecule does not own a virtual box, but merely points to the one that it is in.

All allocation routines return either a pointer to the structure that was allocated, or NULL if memory wasn't available. Assuming that they succeed, all structure members are initialized, typically to 0 or NULL depending on the member type. All the memory freeing routines are robust in that they don't mind NULL inputs or NULL internal pointers. However, this is only useful and robust if allocation is done in an order that always keeps the structure traversable and keeps pointers set to NULL until they are ready to be initialized.

Both the code and the description below are sorted into categories: molecules, walls, reactions, boxes, surfaces, compartments, ports, and the simulation structure. In many cases, functions within each category work with only their respective object. However, the core program is highly integrated so that functions in one category may use objects in another category. Functions in one category (defined by the file that they are listed in) are not supposed to write to objects in other categories, although some exceptions may exist.

5.1 Header files

Smoldyn has several header files. They are: (i) `smoldyn.h`, which lists all of the structure declarations, (ii) `smoldynfuncs.h`, which lists all of the basic Smoldyn function declarations, (iii) `libsmoldyn.h`, which lists all declarations for Libsmoldyn, and (iv) `smoldyn_config.h`, which is automatically generated during the configuration process and which lists the compilation configure options.

5.2 Molecules (functions in `smolmolec.c`)

Each individual molecule is stored with a `moleculestruct` structure, pointed to by a `moleculeptr`. This contains information about the molecule's position, identity, and other characteristics that are specific to each individual molecule. These molecules are organized using a molecule superstructure, which contains lists of the active molecules, a list of unused molecule storage space called dead molecules, and other information about the molecules in general.

```
#define MSMAX 5
```

5 C Code: structures and functions

```
#define MSMAX1 6
enum MolecState {MSsoln, MSfront, MSback, MSup, MSdown, MSbsoln, MSall,
    MSnone};
enum MolListType { , MLTport, MLTnone};
#define PDMAX 6
enum PatternData {PDalloc, PDnresults, PDnspecies, PDmatch, PDsubst,
    PDrule};
```

`MolecState` enumerates the physical states that a molecule can be in, which are respectively: in solution (i.e. not bound to a surface), bound to the front of a surface, bound to the back of a surface, transmembrane in the up direction, and transmembrane in the down configuration. `MSMAX` is the number of enumerated elements. While not a state that molecules are allowed to be in, `MSbsoln` is useful for reactions or surface actions to differentiate between in solution on the front of a surface versus on the back of a surface; `MSsoln` and `MSbsoln` are used for these, respectively. `MSMAX1` accounts for the additional destination state. Also, `MSall` and `MSnone` are useful enumerations for some functions to use as inputs or outputs.

`MolListType` enumerates the types of lists of molecules. `MLTsystem` is for molecules that are in the simulation system and `MLTport` is for molecule buffers for porting. `MLTnone` is not a type but is the absence of any type.

The `PatternData` enumeration, of which there are `PDMAX` options, are for the first `PDMAX` elements of the `mols->patindex` array, called the header. This array, as explained below, lists the species indices that match to text patterns. Each listing requires some information to describe how long the listing is, what's stored in it, and what it corresponds to, which is included in the header. These header elements should be retrieved using the enumerated values. See below for their description and see the Wildcards section of the Code Design chapter.

```
typedef struct moleculestruct {
    unsigned long long serno;           // serial number
    int list;                           // destination list number (ll)
    double *pos;                        // dim dimensional vector for position [d]
    double *posx;                       // dim dimensional vector for old position[d]
    ]
    double *via;                        // location of last surface interaction [d]
    double *posoffset;                  // position offset arising from jumps [d]
    int ident;                          // species of molecule; 0 is empty (i)
    enum MolecState mstate;             // physical state of molecule (ms)
    struct boxstruct *box;              // pointer to box which molecule is in
    struct panelstruct *pnl;           // panel that molecule is bound to if any
    struct panelstruct *pnlx;          // old panel that molecule was bound to if
    any
} *moleculeptr;
```

`moleculestruct` is a structure used for each molecule.

`serno` is the unique serial number that each live molecule is given; it is assigned by the utility function `getnextmol`, which should be used to add new live molecules to the system. If a molecule is being imported from elsewhere, it is legitimate to overwrite the serial number with the previous value. The serial number should normally never exceed a long int, except when using single molecule tracking when two serial numbers are concatenated into one, which then requires an unsigned long long int.

`list` is the master list number that the molecule should be listed in (-1 for dead, other values for live lists); this is modified with `molkill`, `molchange`, or one of the `addmol` functions.

`pos` and `posx`, both of which are owned by the structure, are always valid positions, although

5.2 Molecules (functions in smolmolec.c)

not necessarily within the system volume. `posx` is the position from the previous time step, used to determine if a molecule crossed a wall or surface. `posoffset` is the cumulative position offset that should be added to the position to correct for jumps and periodic boundaries (it replaced the `wrap` element). `via` is the location of the most recent surface interaction. It is work space and so should not be assumed to represent anything, unless it was just set to that value.

`ident` should always be between 0 and `nident-1`, inclusive. A molecule type of 0 is an empty molecule for transfer to the dead list (and should also have `list` equal to -1).

Except during set up, `box` should always point to a valid box.

`mstate` is the physical state of the molecule, which might be solvated or any of several surface-bound positions.

If this molecule is bound to a surface, `pnl` points to that surface panel. Also, `pnlx` points to the panel that `posx` was on.

```
typedef struct molsuperstruct {
    enum StructCond condition;    // structure condition
    struct simstruct *sim;        // simulation structure
    int maxspecies;               // maximum number of species
    int nspecies;                 // number of species, including empty mols.
    char **spname;                // names of molecular species
    int maxpattern;               // maximum number of patterns
    int npattern;                 // actual number of patterns
    char **patlist;               // list of patterns [pat]
    int **patindex;               // species indices for patterns [pat][j]
    char **patrname;              // pattern reaction name if any [pat]
    double **difc;                // diffusion constants [i][ms]
    double **difstep;             // rms diffusion step [i][ms]
    double ***difm;               // diffusion matrix [i][ms][d]
    double ***drift;              // drift vector [i][ms][d]
    double *****surfdrift;      // surface drift [i][ms][s][ps][d]
    double **display;             // display size of molecule [i][ms]
    double ***color;              // RGB color vector [i][ms]
    int **exist;                  // flag for if molecule could exist [i][ms]
    moleculeptr *dead;            // list of dead molecules [m]
    int maxdlimit;                // maximum allowed size of dead list
    int maxd;                     // size of dead molecule list
    int nd;                       // total number of molecules in dead list
    int topd;                     // index for dead list; above are resurrected
    int maxlist;                  // allocated number of live lists
    int nlist;                    // number of live lists
    int **listlookup;             // lookup table for live lists [i][ms]
    char **listname;              // names of molecule lists [ll]
    enum MolListType *listtype;   // types of molecule lists [ll]
    moleculeptr **live;           // live molecule lists [ll][m]
    int *maxl;                    // size of molecule lists [ll]
    int *nl;                      // number of molecules in live lists [ll]
    int *topl;                    // live list index; above are reborn [ll]
    int *sortl;                   // live list index; above need sorting [ll]
    int *diffuselist;             // 1 if any listed moles diffuse [ll]
    unsigned long serno;          // serial number for next resurrected
    molec.
    int ngausstbl;                // number of elements in gausstbl
    double *gausstbl;            // random numbers for diffusion
    int *expand;                  // whether species expand with libmzr [i]
```

5 C Code: structures and functions

```
long int touch;           // counter for molecule modification
} *molsspnr;
```

molsuperstruct contains and owns information about molecular properties and it also contains and owns lists of molecules. **condition** is the current condition of the superstructure and **sim** is a pointer to the simulation structure that owns this superstructure. **maxspecies** is the number of molecular species for which the arrays are allocated, **nspecies** is the actual number of defined species, and **spname** is the list of names for those species. Other superstructures that have **maxspecies** elements, such as surfaces, are intended to be copies of the one here for internal use, while this one remains the master.

The definitive version of the pattern stuff is described in the Wildcards section of the Code Design chapter. However, a little is described here, too. Patterns represent one or more species names. Each string that is used gets recorded here. This list is only updated when it is used. **maxpattern** slots are allocated for patterns, of which **npattern** are actually used. The actual list of patterns is called **patlist**; these are sorted by alphabetical order. Sorted in parallel are the lists **patindex** which is a list of lists, and **patrname**, which is a list of reaction names. **patrname** is only used if the pattern represents a reaction, and is used to differentiate between multiple different reactions that have identical patterns. In **patindex**, the outer list corresponds to the patterns. In the sublists, the first several elements, called the header, give important information about the list. The header occupies the first **PDMAX** list elements.

Diffusion is described with **difc**, a list of diffusion constants; **difstep** is a vector of the rms displacements on each coordinate during one time step if diffusion is isotropic; and **difm** is a list where each element is either a **NULL** value if diffusion is isotropic or a **dimxdim** size diffusion matrix (actually the square root of the matrix). **drift** is the vector for molecular drift, relative to system coordinates. **surfdrift** is for molecular drift relative to the local surface panel coordinates; this memory is only allocated as required. All of these are arrays on the molecule identity, followed by arrays on the molecule state (size **MSMAX**). **display** is simply the size of molecules for graphical output (which scales differently for different output styles) and **color** is the 3-dimensional color vector for each molecule; again size of state list is **MSMAX**.

exist is 1 for each identity and state that could be a part of the system and 0 for those that are not part of the system. This is set in **molupdate**, were any molecules and states that exist then are recorded, as are all reaction products. If commands create molecules, they should also set the exist flag with **molsetexist**.

expand is a flag for on-the-fly rule-based modeling. It is initialized to 0 and stays that way so long as there have never been any molecules of this species, it is increased to 1 if at least one molecule of this species has been created but it has not yet been used for rule expansion, and is set to either 2 or 3 if it has been used for expansion.

touch is a counter that counts the number of times that the list of molecules has been modified. No meaning is ascribed to any particular value. Instead, it can be used to determine if the molecule state has changed between one call of a function and another call of a function, used to prevent recomputing things if it hasn't changed. The **touch** value should be incremented by any function that directly changes molecules, whether it creates new ones, kills existing ones, or moves them. Functions that call other functions for these purposes (e.g. that call **addmol**, **molkill**, or **molchangeident**) do not increment **touch**. Molecules are not considered to be changed if they are merely re-sorted between molecule lists or re-assigned to boxes.

The molecule lists are separated into two parts. The first set is the live list, which are those molecules that are actually in the system or that are being stored for transfer elsewhere (i.e. buffers for ports are also live lists); the others are in the dead list, are empty molecules, and have

5.2 Molecules (functions in *smolmolec.c*)

no influence on the system. If more molecules are needed in the system than the total number allocated, the program sends an error message and ends; in the future, it may be possible to dynamically create larger lists. Upon initialization, all molecules are created as empty molecules in the dead list, whereas during program execution, all lists are typically partially full. After sorting, each live list, `ll`, has active molecules from element 0 to element `nl[ll]-1`, inclusive, and has undefined contents from `nl[ll]` to `maxl[ll]-1`. Similarly, the dead list is filled with empty molecules from 0 to `nd-1`, and has undefined contents from `nd` to `maxd-1`; in this case, `topd` equals `nd`.

Functions other than `mol sort`, such as chemical reactions, are allowed to kill live molecules (with `mol kill`) or resurrect dead ones (with `getnextmol`, `addmol`, or `addsurfmol`) but they should not move molecules or change the list indices. With new molecules that are gotten with `getnextmol`, set the molecule identity, state, list (with `mol s->listlookup`), position, old position, panel if appropriate, and box. Set the `box` element of the molecule to point to the proper box, but do not add the molecule to that box's molecule list. It is now in the resurrected list, which is the top of the dead list between `topd` and `nd-1`, inclusive. Routines should be written so that these mis-sorted molecules do not cause problems. They are sorted with `mol sort`, which moves the empty molecules in the live lists to the dead list, moves the resurrected ones to the top of the proper live list, compacts the live lists (molecule order is not maintained), and identifies the newly reborn molecules in the live lists by setting `topl[ll]`; the reborn molecules extend from `topl[ll]` to `nl[ll]`.

Example of the lists:

index	live[0]	live[1]	dead
8	?	?	maxd ?
7	maxl[0] ?	?	-
6	-	?	-
5	-	maxl[1] ?	-
4	-	-	-
3	nl[0] -	-	nd -
2	2 topl[1], nl[1] -	-	topd 1
1	topl[0] 1	0	0
0	0	3	0

Here, each list has `max=8`, and so is indexed with `m` from 0 to 7. A '?' is memory that is not part of that which was allocated, a '-' is a NULL value, a '0' is an empty molecule, and other numbers are other identities ('1' and '2' are mobile, whereas '3' is immobile). The '0's in the the two live lists are to be transferred to the dead list during the next sort, while the '1' in the dead list has been resurrected and is to be moved to mobile live list. Based on the `topl[0]` index, it can be seen that the '1' and '2' in the mobile live list were just put there during the last sorting, and so are reborn molecules.

There is one dead list and there are `nlist` live lists. The dead list has total size `maxd` and is filled to level `nd`. The index `topd`, which is between 0 and `nd`, separates the dead molecules (list element is -1) with indices from 0 to `topd-1` from the resurrected molecules (list elements 0) that have indices between `topd` and `nd-1`. The dead list is automatically expanded up to size `maxdlimit` if this value is positive, and is automatically expanded without limit if `maxdlimit` is negative (the default). Live list `ll` has total size `maxl[ll]` and is filled to level `nl[ll]`. The index `topl[ll]`, which is between 0 and `nl[ll]` separates the molecules that were there on the prior time step that have smaller indices from those that were created in the last time step, called the reborn molecules,

5 C Code: structures and functions

which have higher indices.

enumerated type functions

`enum MolecState molstring2ms(char *string)`

Returns the enumerated molecule state value, given a string input. Permitted input strings are “solution”, “aq” (aqueous), “front”, “back”, “up”, “down”, “fsoln”, “bsoln”, and “all”. Returns `MSnone` if input is “none” or is not recognized.

`char *molms2string(enum MolecState ms, char *string)`

Returns the string that corresponds to the enumerated molecule state `ms` in `string`, which must be pre-allocated. Also, the address of `string` is returned to allow for function nesting.

`enum MolListType molstring2mlt(char *string)`

Returns the enumerated molecule list type, given a string input. Permitted input strings are “system” and “port”. Returns `MLTnone` for all other input.

`char *molmlt2string(enum MolListType mlt, char *string)`

Returns the string that corresponds to the enumerated molecule list type `mlt`. The string needs to be pre-allocated; it is returned to allow function nesting.

low level utilities

`char *molserno2string(unsigned long long serno, char *string)`

Writes the molecule serial number to a pre-allocated string, which needs to be of size `STRCHAR`, returning the pointer to that string. If the serial number is not a concatenated number, then just writes that number to the string. If the serial number has been concatenated, then this writes the left portion, a period, and then the right portion. Note that this uses the defined string `LLUFORMAT` because Windows doesn’t use the standard `%llu` format specifier for some reason.

`unsigned long long molstring2serno(char *string)`

Converts a string with a serial number value to a serial number. If the string is simply a number, then this reads the number and returns it. If the string has two numbers separated by a period, then this takes the left number, shifts it left 32 places, and adds the right number, to create a concatenated serial number. Returns 0 on failure (serial numbers equal to 0 are not allowed).

`unsigned long long molfindserno(simptr sim, unsigned long long def, long int pserno, unsigned long long r1serno, unsigned long long r2serno, unsigned long long *sernolist)`

Computes a serial number for a new molecule, based upon coded instructions in `pserno`, a default value in `def`, reactant serial numbers in `r1serno` and `r2serno`, and previously assigned product serial numbers in `sernolist`. This does not use `sim` at all, except if the coded instructions are for a serial number with the format “new.new”, in which case this uses the default value as the first new value and then gets the next one off of `sim->mols->serno` and also increments the value. This function does no checking to make sure that inputs are valid.

A positive value of `pserno` indicates that it is an actual serial number and a negative value indicates that it is a bit-inverted code, labeled the bitcode here. The bitcode is 4 hex numbers,

5.2 Molecules (functions in *smolmolec.c*)

with the first two for the left half of the serial number and the second two for the right half of the serial number. The other functions that deal with bitcodes are `rxnstring2sernocode` and `rxnsernocode2string`.

bits	mask	use
1, 2, 3	0xE0	not used
4	0x10	0 for right side, 1 for left side
5	0x08	1 for p, 0 for r or new
6	0x04	1 for r, 0 for p or new
7, 8	0x03	which p or r, or 1 for new

`int molismobile(simpstr sim, int species, enum MolecState ms)`

Returns 1 if molecules of species `species` and state `ms` are mobile at all and 0 if they are not. Mobility includes isotropic and anisotropic diffusion, drift, and surface drift. `MSbsoln` is an allowed input, which always returns the same result as a `MSsoln` input.

`int molstring2pattern(const char *str, enum MolecState *msptr, char *pat, int mode)`

This assembles a species pattern string from molecule species strings. For example, if the user enters a reaction as `A*(front) + B|C(soln) -> A*{B|C}(front)`, then the function that reads this string (which are `molstring2index1` and `rxnparsereaction`) will send the `A*(front)`, `B|C(soln)`, and `A*{B|C}(front)` strings to this function for them to be assembled in the pattern `A*B|C\nA*{B|C}`, and for the states to be returned to the calling function. The pattern is then generally sent off to `molpatternindex`.

Enter `str` as a string of which the first word is the species name or pattern to be processed. The state, if listed at the end of the name string, is returned in `msptr`. Enter `mode` as 0 if this is a new pattern, 1 if the current text being added to the pattern is part of the “match” side of the pattern (the first two words in the above example), or 2 if the current text being added to the pattern is part of the “substitute” side of the pattern (the last word in the above example). The `pat` string needs to be allocated beforehand to size `STRCHAR`.

Returns 0 for success, -1 if `str` or `pat` are missing or if there is no first word in `str`, -2 if the parentheses in `str` don’t match up, -3 if the molecule state could not be read, or -4 if the pattern length exceeds the maximum number of allowed characters (which is `STRCHAR`).

`int molreversepattern(const char *pattern, char *patternrev)`

Takes in a reaction pattern in `pattern` and reverses it for the reverse reaction, returning the result in the string `patternrev`. This only works for patterns that represent reactions. This simply writes the reverse pattern as the products of the original and then the reactants. Returns 0 for success or -1 if `pattern` is not a reaction.

`int molpatternindex(simpstr sim, const char *pattern, const char *rname, int isrule, int update, int **indexptr)`

This function takes in a pattern string, in `pattern`, and returns the list of species or species combinations that correspond to this pattern. To make this function efficient, it records its answers in the molecule superstructure `patlist` and `patindex` lists, so that the list does not need to be recomputed if it is asked for multiple times. The returned list of species is the proper `patindex` list from the molecule superstructure, including its header data. The data returned by this function is a pointer to the original data, not a copy of it, so it should

5 C Code: structures and functions

generally not be modified. See the pattern discussion under the molecule superstructure description for details about the data that are returned. Enter **isrule** with 1 if this is a rule, in which case new species names that arise from expanding the pattern get added to the simulation (but not their reactions), or with 0 if this is not a rule, meaning that new species names are errors. Enter **update** with 0 if the index list should not be updated at all, with 1 if the index list should be created if it doesn't exist already, and with 2 if the index list should be updated to the current list of species (or to the recently created species for reactions and on-the-fly generation).

This function does not support species groups that are defined without wildcards.

Returns 0 for success, -1 for inability to allocate memory, -2 if no wildcards were entered and one or more of the match species is unknown, -3 if a substitute species is unknown and the substitute species had no wildcards in it, -4 if the match string included more words than allowed by this function (which is 4 currently), -5 if a trial match string was too long to fit in **STRCHAR** characters (even if this wasn't actually a match), -6 if species generation failed, -11 for inability to allocate memory, -12 for missing ' ' operand, -13 for missing & operand, -15 for mismatched braces, or -20 for a destination pattern that is incompatible with the matching pattern (i.e. it has to have either 1 destination or the same number of destination options as pattern options). If **update** is set to 0, then no errors are possible. In this case, if the pattern is not in the list, then the function does not add it to the list, but simply returns a value of 0 and **indexptr** pointing to NULL. If **update** is set to 1, then the only error possible is -1, for inability to allocate memory.

First, this function looks to see if **pattern** is already in the pattern list. If the pattern is not found, this function adds it, while maintaining alphabetical order. This may require expanding the list of patterns and pattern indices if the list is full.

Next, for both new and old patterns, this function determines if the indices are up to date, meaning that the **PDnspecies** value is negative, which means that it never needs updating, or it is equal to the current number of species in the simulation, which means that it is already up to date. If the indices are up to date, then the function is done; it simply returns the correct index list. If not, then the function figures out the header to the **index** list, if necessary, and prepares several variables for later use. These variables are: **matchstr**, **substr**, and **newline**, which refer to the pattern string; **istart** and **jstart**, which are the species number to start updating from and the starting result in the index list to start updating to; **matchwords**, **subwords**, and **totalwords**, which are the numbers of each type of word in the pattern; **haswildcard** and **hasspeciesgroup**, which are whether the pattern has wildcard characters and/or a species group; and **nspecies** which is the current number of species in the simulation.

Finally, the function goes through a number of pattern types, from simplest to most general. The simple ones are simply special cases of the most general one, but they are included because their code will run much faster when applicable, and because the simple cases help to make sense of the general case.

(1) If the pattern is "all", then the list should include all species names, but there's nothing else to worry about.

(2) If the pattern has no newline, no wildcards, and 1 matchword, then it must be just the name of a species. It could also be the name of a species group, but it shouldn't be because those are set for updating not required. Assuming it's a single species, this makes space for it in the index, finds the identity value of the species, sticks it in the index, and sets the header

5.2 Molecules (functions in *smolmolec.c*)

values. If the function didn't find the species name (or if the pattern is a species group name), then this returns error code -2 to indicate an unknown species name.

(3) If the pattern has no newline character and one matchword, then it must be a single species name with a wildcard character. If that's the case, then the function goes through all species that haven't been considered before, sees if each one matches to the match word, and adds them to the `index` list if so.

(4) Any other patterns without newline characters, meaning those that have multiple matchwords, are not allowed, so they result in an error.

(5) Next, the function considers patterns with a newline but no wildcards and no species groups, and one matchword. In this case, it is for a reaction with exactly one reactant. It will also have exactly one entry in the `index` list. In this case, the function reads through the `matchwords` to get each reactant name and puts those in `ispecies`. Then, typically, it keeps on going, putting the reactant identities in the `index` variable. Then, it reads through `subwords` to get the product names, which it converts to identities, and puts those in `index`. The `PDnspecies` value is set to -1 because this `index` should never need updating again. This function can also handle rules, possibly on-the-fly. For example, a reaction rule could be $A + B \rightarrow C$. If this is generated on-the-fly, then it simply says that species C should not be created until there is a molecule of A or B.

(6) Next, the same thing but 2 matchwords.

(7) Next, with newline, species group or wildcards, and 1 matchword.

(8) Finally, with newline, species group or wildcards, and 2 matchwords.

```
int molstring2index1(simptr sim, char *str, enum MolecState *msptr, int
**indexptr)
```

This reads the first word of the string in `str`, parses it to find the state listed, if any, and determines what species name or names it refers to. The state is returned in `msptr` and the species index list is pointed to by `indexptr`. The function simply calls `molstring2pattern` and then `molpatternindex` sequentially. On success, if there is exactly one result in the list and exactly one match item in the pattern, this returns the index of the result. If there are multiple results, it returns 0, which is still a successful result.

It can also return the following error codes: -1 if `str` is missing or has no first word, -2 if the parentheses in `str` don't match up, -3 if the molecule state could not be read, -4 if there are no wildcards in the string and the species name is unknown (or if the number of characters in `str` is more than the maximum allowed in a pattern (256), which shouldn't ever happen), -5 if the species is "all", -6 if the logic expansion failed due to missing & operand or mismatched braces, or -7 if memory could not be allocated.

```
int moladdspeciesgroup(simptr sim, char *group, char *species, int imol)
```

Creates a species group named `group` and adds the species (or species group, including species names with wildcards) named `species` to that group. If the group already exists, then this just adds the species to the existing group. The `species` value is allowed to be NULL for creating an empty group and it is also allowed to be a group as well, including a group specified using wildcard characters. This function uses the same pattern infrastructure as for wildcard characters, storing the data in the same `patstring` and `patindex` data structures. Also, `species` can be NULL and a single species can be added instead using `imol`, where this is the identity of a single species.

5 C Code: structures and functions

Returns 0 for success, -1 if the group name `name` is missing, -2 if there are parentheses mismatches (which shouldn't be there anyhow), -3 if a molecule state could not be read (which shouldn't be there), -4 if the species name does not correspond to any species, -5 if the group name is "all", -6 if logic expansion failed due to missing & operand or mismatched braces, -7 if memory could not be allocated, -8 if a molecule state is given and isn't "all", or -9 if the group name is the same as an existing molecule name.

`char *molpos2string(simptr sim, moleculeptr mptr, char *string)`

Writes molecule position in `mptr->pos` to `string` using "%g" formatting code for `sprintf`. Each coordinate value, including the first one, is preceded by a space. If the simulation includes surfaces, this function ensures that the written position, including round-off errors, is both in the same box and on the same side of all surface panels (not including the panel that the molecule is bound to, if any) as the actual position. If this function cannot achieve these criteria after 50 attempts, it prints a warning, and returns the string.

set structure values

`int molsetgausstable(simptr sim, int size)`

Sets the size of the Gaussian look-up table to `size` and also allocates the table, if needed. Setting `size` to 0 or a negative number keeps the current size if it has already been allocated, or creates a table with the default size (4096) if not. Otherwise, `size` is required to be an integer power of two. This will replace an existing table if the new size is different from the previous one. Returns 0 for success, 1 for insufficient memory, or 3 if the size is not an integer power of two.

`void molsetdifc(simptr sim, int ident, int *index, enum MolecState ms, double difc)`

Sets the diffusion coefficient for molecule `ident` and state `ms` to `difc`. For multiple identities, enter them in `index` using the pattern index header. If `ms` is `MSall`, this sets the diffusion coefficient for all states. This does not update rms step sizes or reaction rates.

`int molsetdifm(simptr sim, int ident, int *index, enum MolecState ms, double *difm)`

Sets the diffusion matrix for molecule `ident` and state `ms` to `difm`. Any required matrices that were not allocated previously are allocated here. For multiple species, enter them in `index` using the pattern index header. If `ms` is `MSall`, this sets the diffusion matrix for all states. This returns 0 for successful operation and 1 for failure to allocate memory. This updates the isotropic diffusion coefficient but does not update rms step sizes or reaction rates.

`int molsetdrift(simptr sim, int ident, int *index, enum MolecState ms, double *drift)`

Sets the drift vector for molecule `ident` and state `ms` to `drift`. Any required vectors that were not allocated previously are allocated here. For multiple species, enter them in `index` using the pattern index header. If `ms` is `MSall`, this sets the drift vector for all states. This returns 0 for successful operation and 1 for failure to allocate memory.

`int molsetsurfdrift(simptr sim, int ident, int *index, enum MolecState ms, int surface, enum PanelShape ps, double *drift)`

Sets the surface drift vector for molecule `ident`, state `ms`, surface `surface`, and panel shape `ps` to `drift`. Any required memory that was not allocated previously is allocated here. For

5.2 Molecules (functions in *smolmolec.c*)

multiple species, enter them in `index` using the pattern header. If `ms` is `MSall`, this sets the surface drift vector for all surface-bound states. In addition, `surface` can be -1 to indicate all surfaces and `ps` can be `PSall` to indicate all panel shapes. Any combination of “all” conditions is permitted. This returns 0 for successful operation and 1 for failure to allocate memory.

```
void molsetdisplaysize(simptr sim, int ident, int *index, enum MolecState ms,
double dsize)
```

Sets the display size for molecule `ident` and state `ms` to `dsize`. For multiple species, enter them in `index` using the pattern header. If `ms` is `MSall`, this sets the display size for all states.

```
void molsetcolor(simptr sim, int ident, int *index, enum MolecState ms, double
*color)
```

Sets the color for molecule `ident` and state `ms` to the 3-dimensional RGB vector `color`. For multiple species, enter them in `index` using the pattern header. If `ms` is `MSall`, this sets the color for all states.

```
void molsetlistlookup(simptr sim, int ident, int *index, enum MolecState ms, int
ll)
```

Sets the list lookup table value to live list number `ll` for molecule `ident` and state `ms`. For multiple species, enter them in `index` using the pattern header. Special codes are also possible in the `ident` input: `ident=-7` implies all diffusing molecules, and `ident=-8` implies all non-diffusing molecules. Using `ms=MSall` implies all states. Note that the `listlookup` element is defined for both `MSsoln` and `MSbsoln`, and they are always set to the same values.

```
void molsetexist(simptr sim, int ident, enum MolecState ms, int exist)
```

Sets the `exist` element of the molecule superstructure for identity `ident` and state `ms` to `exist`; “all” inputs are not permitted.

```
int molcount(simptr sim, int i, enum MolecState ms, int max)
```

Counts the number of molecules of type `i` and state `ms` currently in the simulation. If `max` is -1 it is ignored, and otherwise the counting stops as soon as `max` is reached. Either or both of `i` and `ms` can be set to “all”; enter `i` as a negative number and enter `ms` as `MSall`. All molecule lists and the dead list are checked; porting lists are included. This function returns correct molecule counts whether molecule lists have been sorted since recent changes or not. It runs fastest if molecule lists have been sorted. If `i` is less than zero, this implies all species; if `i` is greater than zero, this implies that specific species; and if `i` equals zero, this implies that the `index` entry should be used instead. In this last case, enter lists of species using `index`, using the standard index pattern header. This function returns 0 if the molecule superstructure still has an “SCinit” condition, regardless of whether there are molecules in it or not.

This function is essentially the same exact thing several times in a row, for the different input cases and with slightly different outer loops. This could be substantially shortened, but is done this way for better speed. This function requires that the `index` list be sorted.

```
void molscancmd(simptr sim, int i, int *index, enum MolecState ms, cmdptr cmd,
enum CMDcode(*fn)(simptr, cmdptr, char*))
```

Scans over all molecules that meet the criteria listed and calls the function `fn` for each one. This function is similar to `molcount`, except that it calls a function for each molecule rather than just counting them. As with `molcount`, this checks all molecule lists and the dead list; it also works whether molecules have been sorted or not. Send in `i` as the species number, which

5 C Code: structures and functions

should be less than 0 for all species, 0 to indicate that `i` should be ignored and `index` should be used instead, or as a positive number for that specific species. Enter `index` as a sorted list of species, using the standard index pattern header and format; it is only used if `i` equals zero. Enter `ms` as `MSall` for all states and as a specific state otherwise. The function entry is designed for a Smoldyn command function and so has the same format, despite the fact that this format isn't really ideal in this case. The function gets passed the simulation pointer in the first argument, `cmd` in the second argument, and a pointer to the current molecule in the third parameter (cast to a `char*` due to the non-ideal formatting). The function should return `CMDok` to indicate that the scan should continue or `CMDstop` to indicate that the scan should stop. See the commands section of the manual to see how this function can be used.

```
void molscanfn(simptr sim, int i, int *index, enum MolecState ms, char *erstr,
double(*fn)(void*, char*, char*))
```

This is identical to `molscancmd` but has a slightly more versatile function declaration that's not designed just for commands. The `erstr` string is for returning errors. The call-back function, `fn` should take the simulation structure as its first argument, the error string as its second argument, and the molecule pointer, cast as a `char*`, as its third argument.

```
int molismatch(moleculeptr mptr, int i, int *index, enum MolecState ms)
```

Tests to see if the molecule in `mptr` matches the conditions given in species number `i`, index list `index`, and state `ms`. These latter three elements should be the values returned by `molstring2index1`, so they can be for single species, species groups, etc.

```
int MolCalcDifcSum(simptr sim, int i1, enum MolecState ms1, int i2, enum
MolecState ms2)
```

Calculate and returns diffusion coefficient sums. This allows `ms1` and/or `ms2` to be the `MSbsoln` state. Also, enter `i1` and/or `i2` as 0 to not include it in the sum.

memory management

```
moleculeptr molalloc(int dim)
```

`molalloc` allocates and initializes a new `moleculestruct`. The serial number is set to 0, the `list` to -1 (dead list), positional vectors to the origin, the identity to the empty molecule (0), the state to `MSsoln`, and `box` and `pn1` to NULL. The molecule is returned unless memory could not be allocated, in which case NULL is returned.

```
void molfree(moleculeptr mptr)
```

`molfree` frees the space allocated for a `moleculestruct`, as well as its position vectors. The contents of `box` and `pn1` are not freed because they are references, not owned by the molecule structure.

```
molexpandsurfdrift(simptr sim, int oldmaxspec, int oldmaxsrf)
```

Expands the surface drift data structure, when the species list and/or the surface list is expanded. Enter `oldmaxspec` and `oldmaxsrf` with the maximum number of species and surfaces before expansion (if only one needs to be expanded, then both still need to be listed, but one will match the current maximum). This function simply calls `molsetsurfdrift` with all of the data in the current data structure, which re-builds the data structure in a larger format. This function is called by `surfacessalloc` and `molssalloc`.

```
void molfreesurfdrift(double *****surfdrift, int maxspec, int maxsrf)
```

Frees the space allocated for all surface drift data, which is stored in a molecule superstructure.

```
int molpatternindexalloc(int **indexptr, int n)
```

Allocates space for a single index list of the species pattern string lookup table, or expands it. If space is needed for a new index list, send in `indexptr` pointing to NULL; if an existing index list needs expansion, send in `indexptr` pointing to the start of the list that needs expansion. Send in `n` as the number of total spaces (including the overhead spaces) that are desired, or send in `n` as -1 for automatic allocation and expansion. This allocates memory and sends the result back pointed to by `indexptr`. Returns 0 for success or 1 for inability to allocate memory. This sets element 0 of the result to equal the allocated length of the array. It sets all other newly created elements to 0.

```
int molpatternalloc(simpstr sim, int maxpattern)
```

Allocates space for species patterns and their indices. Send in `sim` as the simulation structure and `maxpattern` for the desired total allocated number of pattern spaces. This allocates space, copies over any existing data, and initializes the new spaces to empty values. This takes care of the `patlist` and `patindex` lists.

```
molsspstr molssalloc(molsspstr mols, int maxspecies)
```

`molssalloc` allocates and initializes a molecule superstructure. This function may be called multiple times, in order to increase the maximum number of species. The Gaussian table is left empty; it is filled in in `molupdate`. Returns NULL if there is insufficient memory. Enter `maxspecies` with your desired number of simulated species. One more than this will actually be allocated because this assigns species number 0 to the “empty” species.

```
int mollistalloc(molsspstr mols, int maxlist, enum MolListType mlt)
```

Allocates `maxlist` new live lists of list type `mlt` for the already existing molecule superstructure `mols`. This works whether there were already live lists or not. Returns the index of the first live list that was just added for success or a negative code for failure: -1 for out of memory, -2 for a negative `maxlist` input value, or -3 for a NULL `mols` input. The `maxlist` element of the superstructure is updated. The `nlist` element of the superstructure is unchanged.

This does all of the allocation separately from the molecule superstructure. At the end, if all goes well, it frees the current memory and replaces it with the new memory.

```
int molexpandlist(molsspstr mols, int dim, int ll, int nspaces, int nmolecs)
```

Expands molecule list, where `mols` is the molecule superstructure and `dim` is the system dimensionality. This both creates new lists or expands existing lists, as required. If `ll` is negative, the dead list is expanded and otherwise live list number `ll` is expanded. If `nspaces` is negative, the list size is doubled and otherwise `nspaces` spaces are added to the list. The first `nmolecs` of these spaces are filled with new dead molecules (`mptr->list` element set to -1). Because this shouldn't normally be called with `ll ≥ 0` and `nmolecs > 0`, error code 2 is returned if this happens. This returns 0 for success, 1 for out of memory during list expansion, 2 for illegal inputs, 3 for more molecules are being created than will fit in the list even after expansion, and 4 for out of memory during molecule allocation.

```
void molssfree(molsspstr mols, int maxident, int maxsrf)
```

`molssfree` frees both a superstructure of molecules and all the molecules in all its lists.

data structure output

5 C Code: structures and functions

void molssoutput(simptr sim)

molssoutput prints all the parameters in a molecule superstructure including: molecule diffusion constants, rms step lengths, colors, and display sizes; and dead list and live list sizes and indices.

void writemols(simptr sim, FILE *fptr)

Writes all information about the molecule superstructure to the file **fptr** using a format that can be read by Smoldyn. Does not write information about individual molecules. This allows a simulation state to be saved.

void writemolecules(simptr sim, FILE *fptr)

Writes information about all individual molecules to the file **fptr** using a format that can be read by Smoldyn. This allows a simulation state to be saved.

int checkmolparams(simptr sim, int *warnptr)

Checks some parameters in a molecule superstructure and substructures to make sure that they are legitimate and reasonable. Prints error and warning messages to the display. Returns the total number of errors and, if **warnptr** is not NULL, the number of warnings in **warnptr**.

structure setup

int molenablemols(simptr sim, int maxspecies)

Enables molecules. This function can be called multiple times. Enter **maxspecies** as -1 for default species allocation, or to a positive number for the number of species that should be allocated. In the default, the number of species is set to 5 for the initial call, and is either left unchanged if there is spare space or doubled if there isn't space for subsequent calls. Returns 0 for success, 1 if memory could not be allocated, or 2 if **maxspecies** is less than the currently allocated number of species.

void molsetcondition(molssptr mols, enum StructCond cond, int upgrade)

Sets the molecule superstructure condition to **cond**, if appropriate. Set **upgrade** to 1 if this is an upgrade, to 0 if this is a downgrade, or to 2 to set the condition independent of its current value. If the condition is downgraded, this also downgrades the simulation structure condition.

int addmollist(simptr sim, char *nm, enum MolListType mlt)

Adds a molecule list named **nm** and of type **mlt** to the molecule superstructure, allocating it if needed. Returns the index of the list for success, -1 if memory could not be allocated, -2 if the list name has already been used, or -3 for illegal inputs (**mols** or **nm** was NULL).

int molsetmaxspecies(simptr sim, int max)

Sets the maximum number of molecular species to **max+1**, where the additional species represents empty molecules. This function is only supplied for backward compatibility, as it is now (version 2.23) completely identical to **molenablemols**, which should be called instead. Returns 0 for success, 1 for insufficient memory, or 2 if **maxspecies** is smaller than the prior allocated number of species.

int molsetmaxmol(simptr sim, int max)

Sets the maximum number of molecules that the simulation is allowed to use to **max**. Enter **max** as -1 to specify that molecules should be allocated as needed without bound, which is the default behavior. This does not allocate any molecules or molecule lists. This function

5.2 Molecules (functions in *smolmolec.c*)

does not need to be called at all. This works during initial setup, or later on. Returns 0 for success, 1 if memory could not be allocated, or 5 if the requested `max` value is less than the current number of allocated molecules.

`int moladdspecies(simptr sim, char *nm)`

Adds species named `nm` to the list of species that is in the molecule superstructure. This enables molecule support if it hasn't been enabled already. Returns a positive value corresponding to the index of a successfully adds species for success, -1 for failure to allocate memory, -4 if if trying to add a species named "empty", -5 if the species already exists, or -6 if the species name includes wildcards (which are forbidden).

`int molgeneratespecies(simptr sim, const char *name, int nparents, int parent1, int parent2)`

Generates a new molecular species which has name `name`. If this molecule is generated as the product of a reaction, then enter the number of reactants in `nparents` and the reactant species numbers in `parent1` and `parent2`. Only `nparents` of these values are considered by the function. Using the parent values, this function generates diffusion coefficient and display parameters for the new species. If there is one parent, the new values are the same as those for the parent. If there are two parents, the new diffusion coefficient is $D_{new} = (D_1^{-3} + D_2^{-3})^{-1/3}$, the new display size is $r_{new} = (r_1^3 + r_2^3)^{1/3}$, and the new color is $rgb_{new} = (r_1rgb_1 + r_2rgb_2)/(r_1 + r_2)$, where D is a diffusion coefficient, r is a display size, and rgb is one of the red-green-blue color values. This is the same scheme used in *smolbng.c*.

Returns the new species index for success, -1 for failure to allocate memory, -4 if trying to add a species named "empty", -5 if the species already exists, or -6 if the species name includes wildcards (which is forbidden here).

`void molsupdateparams(molssptr mols, double dt)`

Calculates the `difstep` parameter of the molecule superstructure and also sets the `diffuselist` set of flags in the molecule superstructure. `dt` is the simulation time step. This function should be called during initial setup (this is called from `molsupdate`), if any diffusion coefficient changes (performed with `molsetdifc`), or if any diffusion matrix changes (performed with `molsetdifm`, which also updates the diffusion coefficient).

`void molsupdatelists(simptr sim)`

Updates molecule superstructure from the level of `SClists` to `SCparams`. Can be run multiple times.

`int molsupdate(simptr sim)`

This sets up or updates the molecule superstructure. It may be called at program startup or at any later time. This sets up, or updates all molecule superstructure parameters, and works in all situations. It sets up the Gaussian table, live lists, live list lookup numbers, and diffusion step lengths. It does not process individual molecules (i.e. sorting and boxes). Returns 0 for success, or 1 for insufficient memory.

adding and removing molecules

`void molkill(simptr sim, moleculeptr mptr, int ll, int m)`

Kills a molecule from one of the live lists. `mptr` is a pointer to the molecule and `ll` is the list that it is currently listed in (probably equal to `mptr->list`, but not necessarily). If it is known, enter the index of the molecule in the master list (i.e. not a box list) in `m`; if it's

5 C Code: structures and functions

unknown set `m` to -1. If the molecule should be killed without triggering list sorting (a rare occurrence), then send in `ll` as -1. This function resets most parameters of the molecule structure, but leaves it in the master list and in a box for later sorting by `molssort`. The appropriate `sortl` index is updated.

`moleculeptr getnextmol(molssptr mols)`

Returns a pointer to the next molecule on the dead list so that its data can be filled in and it can be added to the system. The molecule serial number is assigned. In the process, this increments the `serno` element of the molecule superstructure, which is an unsigned long int and wraps around when it reaches all 1 values. This updates the `topd` element of the molecule superstructure. Returns NULL if there are no more available molecules. The intention is that this function should be called anytime that molecules are to be added to the system.

`moleculeptr newestmol(molssptr mols)`

Returns a pointer to the molecule that was most recently added to the system, assuming that `molssort` has not been called in the meantime. For example, if 1 molecule is successfully added with `addmol`, `addsurfmol`, or `addcompartmol`, this will return a pointer to that molecule.

`int addmol(simpstr sim, int nmol, int ident, double *poslo, double *poshi, int sort)`

Adds `nmol` molecules of type `ident` and state `MSsoln` to the system. These molecules are not added to surfaces. Their positions are chosen randomly within the rectanguloid that is defined by its corners `poslo` and `poshi`. Set these vectors equal to each other for all molecules at the same point. Set `sort` to 1 for complete sorting immediately after molecules are added and 0 for not. Returns 0 for success, 1 for out of memory, or 3 for more molecules being added than permitted with `mols->maxdlimit`.

`int addsurfmol(simpstr sim, int nmol, int ident, enum MolecState ms, double *pos, panelptr pnl, int surface, enum PanelShape ps, char *pname)`

Adds `nmol` surface-bound molecules, all of type `ident` and state `ms`, to the system. They can be added to a specific panel by specifying the panel in either of two ways: send in its pointer in `pnl`, or specify the panel shape in `ps` and the panel name in `pname`. To add to all panels on the surface, send in `pnl` equal to NULL and/or set `ps` to `PSall`. To add the molecules to a certain point, send it in with `pos`, and otherwise set `pos` to NULL for random positions (there is no check that `pos` is actually on or near the panel). The function returns 0 for successful operation, 1 for inability to allocate temporary memory space, 2 for no panels match the criteria listed, or 3 for insufficient permitted molecules. See the `surfacearea` description for more information about the parameter input scheme.

For multiple panels, this function creates tables that list the cumulative areas of the included panels and the panel pointer for each included panel.

`int addcompartmol(simpstr sim, int nmol, int ident, compartptr cmpt)`

Adds `nmol` molecules of type `ident` and state `MSsoln` to the system with random locations that are within compartment `cmpt`. Returns 0 for success, 2 if a random point cannot be found, or 3 if there aren't enough available molecules.

molecule manipulations

`void molchangeident(simpstr sim, moleculeptr mptr, int ll, int m, int i, enum MolecState ms, panelptr pnl, double *crsspt)`

5.2 Molecules (functions in *smolmolec.c*)

Changes the identity or state of a molecule that is currently in the system to species *i* and state *ms*. It is permissible for *i* to equal 0 for the molecule to be killed, which is equivalent to calling `molkill`. `mptr` is a pointer to the molecule and `ll` is the list that it is currently listed in (probably equal to `mptr->list`, but not necessarily). If the molecule's list information should be updated to reflect the identity change, which is nearly always, then `ll` should be the molecule's list; if not though, then send in both `ll` and `m` as -1. If `m` is known, then enter the index of the molecule in the master list (i.e. not a box list) in `m`; if it's unknown set `m` to -1. If the molecule is to be bound to a panel (independent of whether it was bound to a panel before or not), enter the panel in `pnl`. Or, if it is to be in solution but adjacent to a panel, enter this panel in `pnl`.

If this identity change occurred mid-time step at some location, such as due to a surface collision, then send in that location in `crsspt`. This function will then rescale the molecule's motion over the course of the time step to address any change in diffusion coefficient that arises from the species and/or state change. Send in `crsspt` as `NULL` to not use this correction.

This function sets some parameters of the molecule structure, fixes the location as needed, and, if appropriate, updates `sortl` to indicate to `molsort` that sorting is needed. This also increments the molecule `touch` value to show that molecules have been touched.

```
int molmovemol(simptr sim, moleculeptr mptr, const double *delta)
```

Moves molecule `mptr` by distance `delta`. This translates the molecule, including keeping it on a surface if it's surface-bound, and dealing with surface and/or wall collisions.

core simulation functions

```
int molsort(simptr sim, int onlydead2live)
```

Sorts molecules between live and dead lists, and between live lists. This also takes care of the live lists within boxes, as well as all list indices. Sorting is based solely on the list element of the molecule structure. Molecule ordering in lists is not preserved. If a molecule is in the system (in a master live list of type `MLTsystem`), its box element must point to a box, and those boxes' molecule lists must list the respective molecules. Resurrected molecules need to have the proper box listed in the molecule structure, but should not be listed in the box list; this listing is taken care of here. The routine returns 0 for normal operation and 1 if memory could not be allocated.

Under normal operation, the `onlydead2live` option is set to 0. In this case, the function first sets the `topl` indices to the ends of the live lists, so that future functions can see what's new in those lists. Then, it checks molecules in live lists, starting with the `sortl` molecule, which is the bottom one that is known to need sorting, and checks up to the `topl-1` molecule, which is the end of the list, as it was when this function was called. Any missorted molecules get moved. Next, the function processes all molecules in the dead list, starting from `topd` and going to `nd`, which is called the resurrected portion of the dead list. All of these molecules are moved to the appropriate live list (unless they are empty molecules, in which case they get put on the top of the dead list). Afterwards, `nd` is decreased to equal `topd`, meaning that the resurrected list has zero length and that there's nothing more to do there. Finally, the function sets the `sortl` indices to equal the `n1` indices, to show that all of the molecules in the live lists have been sorted. Thus, at the end, $\text{topl} \leq \text{sortl} = \text{n1}$, where the sublist from `topl` and `n1` is the reborn list. Also, `topd = nd`, showing that there is no resurrected list.

On occasion, it's helpful to set the `onlydead2live` option to 1. In this case, the function only sorts resurrected molecules in the dead list into the live lists. It does not change the

5 C Code: structures and functions

`sortl` or `topl` indices. One result is that any molecules that were considered to be reborn before are still considered to be reborn (and, in fact, there are now more reborn molecules due to their addition from the resurrected list). This option is helpful when molecules that have been added to the system need to be made available for other functions, such as in some commands that add molecules and for molecules added from the lattice code.

int moldosurfdrift(simp_{ptr} sim, molecule_{ptr} mptr, double dt)

Performs surface drift on molecule `mptr` over time step `dt`. This function should only be called if it is known that this molecule is surface-bound and that the surface drift data structure has been allocated at least down to the level of `surfdrift[i][ms]`. It should also be called before other drift or diffusion functions, because the molecule's position on the surface may affect its surface drift vector.

int diffuse(simp_{ptr} sim)

`diffuse` does the diffusion for all molecules over one time step using single-threaded operation. Collisions with walls and surfaces are ignored and molecules are not reassigned to the boxes. If there is a diffusion matrix, it is used for anisotropic diffusion; otherwise isotropic diffusion is done, using the `difstep` parameter. The `posx` element is updated to the prior position and `pos` is updated to the new position. Surface-bound molecules are diffused as well, and they are returned to their surface. Returns 0 for success and 1 for failure (which is impossible for this function).

5.3 Walls (functions in smolwall.c)

The simulation volume is defined by its bounding walls. If no other surfaces are defined, these walls can be reflecting, periodic, absorbing, or transparent. Because walls can be transparent, molecules can leave the simulation volume. However, this can be a bad idea because the virtual boxes are defined to exactly fill the volume within the walls, so molecules or surfaces outside of the simulation volume can lead to very slow simulations. Also, the graphics are designed for the simulation volume within the walls. If surfaces are defined, then walls, regardless of how they are set up, are simulated as though they are transparent.

Walls are quite simple, defined with only a simple structure and no superstructure. A simulation always has $2 \cdot \text{dim}$ walls.

```
typedef struct wallstruct {
    int wdim;           // dimension number of perpendicular to wall
    int side;           // low side of space (0) or high side (1)
    double pos;         // position of wall along dim axis
    char type;          // properties of wall
    struct wallstruct *opp; // pointer to opposite wall
} *wallptr;
```

`wallstruct` (declared in `smolib.h`) is a structure used for each wall. The type may be one of four characters, representing the four possible boundary conditions.

type	boundary
r	reflecting
p	periodic
a	absorbing
t	transparent

5.3 Walls (functions in *smolwall.c*)

Pointers to the opposite walls are used for wrap-around diffusion, but are simply references. There is no superstructure of walls, but, instead a list of walls is used. Walls need to be in a particular order: walls numbered 0 and 1 are the low and high position walls for the 0 coordinate, the next pair are for the 1 coordinate, and so on up to the $2*\text{dim}-1$ wall. These walls are designed to be bounds of simulated space, and are not configured well to act as membranes. Wall behaviors are completely ignored if any membranes are declared.

low level utilities

`void systemrandpos(simptr sim, double *pos)`

Returns a random point within the system volume, chosen with a uniform distribution.

`double systemvolume(simptr sim)`

Returns the total volume of the system.

`void systemcorners(simptr sim, double *poslo, double *poshi)`

Returns the low and high corners of the system volume in `poslo` and `poshi`, respectively. Both results are optional; enter NULL if a point is unwanted.

`void systemcenter(simptr sim, double *center)`

Returns the center of the system.

`double systemdiagonal(simptr sim)`

Returns the diagonal length of the system, or just the length if it is 1-D.

`int posinsystem(simptr sim, double *pos)`

Returns 1 if `pos` is within the system boundaries (equal to the edges counts as inside) and 0 if it is outside.

`double wallcalcdist2(simptr sim, double *pos1, double *pos2, int wpcode, double *vect)`

Calculates squared distance between point `pos1` and point `pos2`, while accounting for periodic boundaries. These are accounted for using `wpcode`, which is the wrapping code for the box that `pos1` is in. This code needs to be entered. If it's not known, then find the box pointer for the two positions with `bptr1=pos2box(sim, pos1)` and similarly for `pos2`, then set `b2` to be the index of `bptr2` within the list `bptr1->neigh`, and finally use `bptr1->wpneigh[b2]` as `wpcode`. Also, `vect` needs to be entered as a `dim`-dimensional vector of doubles. It is returned as the vector from `pos1` to `pos2`, while accounting for wrapping.

memory management

`wallptr wallalloc(void)`

`wallalloc` allocates and initializes a new wall. The pointer to the opposite wall needs to be set.

`void wallfree(wallptr wptr)` `wallfree` frees a wall.

`wallptr *wallsalloc(int dim)`

`wallsalloc` allocates an array of pointers to $2*\text{dim}$ walls, allocates each of the walls, and sets them to default conditions (reflecting walls at 0 and 1 on each coordinate) with correct pointers in each opp member.

5 C Code: structures and functions

`void wallsfree(wallptr *wlist, int dim)`
wallsfree frees an array of 2*dim walls, including the walls.

data structure output

`void walloutput(simptr sim)`
walloutput prints the wall structure information, including wall dimensions, positions, and types, as well as the total simulation volume.

`void writewalls(simptr sim, FILE *fptr)`
Writes all information about the walls to the file `fptr` using a format that can be read by Smoldyn. This allows a simulation state to be saved.

`int checkwallparams(simptr sim, int *warnptr)`
Checks some parameters of simulation walls to make sure that they are reasonable. Prints warning messages to the display. Returns the total number of errors and, if `warnptr` is not NULL, the number of warnings in `warnptr`.

structure setup

`int walladd(simptr sim, int d, int highside, double pos, char type)`
Adds a wall to the system. If no walls have been added yet, this allocates the necessary memory. `d` is the dimension that the wall bounds, `highside` is 0 if the wall is on the low side of the system and 1 if it is on the high side of the system, `pos` is the location of the wall in the `d` dimension, and `type` describes the boundary condition (if there aren't any surfaces). Returns 0 for success, 1 for unable to allocate memory, or 2 if the simulation structure `dim` element hasn't been set up yet.

`int wallsettype(simptr sim, int d, int highside, char type)`
Sets the type of an existing wall for dimension `d` to `type`. Set `highside` to 0 if the wall is on the low side of the system and 1 if it is on the high side of the system. Enter `d` and/or `highside` with a negative number to indicate "all" dimensions and/or system sides.

core simulation functions

`void checkwalls1mol(simptr sim, moleculeptr mptr)`
Checks the walls for a single molecule, reflecting, wrapping around, or absorbing as needed. This is identical to `checkwalls`, but only for one molecule.

`void checkwalls(simptr sim, int ll, int reborn, boxptr bptr)`
`checkwalls` does the reflection, wrap-around, or absorption of molecules at walls by checking the current position, relative to the wall positions (as well as a past position for absorbing walls). Only molecules in live list `ll` are checked. If `reborn` is 1, only the newly added molecules are checked; if it's 0, the full list is checked. It does not reassign the molecules to boxes or sort the live and dead ones. It does not matter if molecules are assigned to the proper boxes or not. If `bptr` is NULL, all diffusing molecules are checked, otherwise only those in box `bptr` are checked.

5.4 Reactions (functions in smolrxn.c)

Reactions were overhauled for Smoldyn version 1.82, so the following text describes the current version. Reactions are stored with several structures. There is a reaction superstructure for each

reaction order (which may be 0, 1, or 2). Within each superstructure, there is a separate structure for each reaction.

5.4.1 enumerated types

Following are the enumerated types and the structures.

```
#define MAXORDER 3
#define MAXPRODUCT 16
enum RevParam {RPhone, RPirrev, RPconfspread, RPbounce, RPpgem, RPpgemmax,
               RPpgemmaxw, RPratio, RPunbindrad, RPpgem2, RPpgemmax2, RPratio2,
               RPOffset, RPfixed};
```

The constant `MAXORDER` is one more than the maximum reaction order that is permitted. For now, order 3 and higher reactions are not supported, although much of the code should function with any reaction order. High-order reactions may be supported in future versions. `MAXPRODUCT` is the maximum number of products that a reaction can have, which is only used at present in loading reactions from a configuration file. The enumerated type `RevParam` lists the possible “reversible parameter types” that are allowed.

5.4.2 reaction structure

```
typedef struct rxnstruct {
    struct rxnsuperstruct *rxnss; // pointer to superstructure
    char *rname;                  // pointer to name of reaction
    int *rctident;                 // list of reactant identities [rct]
    enum MolecState *rctstate;     // list of reactant states [rct]
    int *permit;                   // permissions for reactant states [ms]
    int nprod;                     // number of products
    int *prdidnt;                  // list of product identities [prd]
    enum MolecState *prdstates;    // list of product states [prd]
    long int *prdserno;            // list of product serno rules [prd]
    int *pr dintersurf;            // list of product intersurface rules [prd]
    listptrli logserno;            // list of serial nums for logging reaction
    char *logfile;                 // filename for logging reaction
    double rate;                   // requested reaction rate
    double multiplicity;           // rate multiplier
    double bindrad2;               // squared binding radius, if appropriate
    double prob;                   // reaction probability
    double chi;                    // diffusion-limited fraction
    double tau;                    // characteristic reaction time
    enum RevParam rparamt;         // type of parameter in rpar
    double rparam;                 // parameter for reaction of products
    double unbindrad;              // unbinding radius, if appropriate
    double **pr dpos;              // product position vectors [prd][d]
    int disable;                   // 1 if reaction is disabled
    struct compartstruct *cmpt;    // compartment reaction occurs in, or NULL
    struct surfacestruct *srf;     // surface reaction on, or NULL
} *rxnptr;
```

Each individual reaction, of any order, is stored in a reaction structure, `rxnstruct`. `rname` is a pointer to the reaction name that is stored in, and owned by, the reaction superstructure. `rctident` is a list of the reactant identities for the reaction, listed in the same order in which

5 C Code: structures and functions

they were listed in the configuration file. Other than the order of the reactants, which is not stored elsewhere, the list of reactants that is stored here is redundant with the **table** element of the reaction superstructure. **rctstate** is a list of the allowed reactant states, again listed in the same order in which the reactants were listed in the configuration file. Each item of the **rctstate** list may be a single state, **MSnone**, **MSall**, or **MSsome**. These are fairly self-explanatory; **MSsome** means that more than one reactant state is allowed to react, but not all states. **permit**, which is largely redundant with **rctstate**, is a list of flags for which reactant states, or state combinations, are allowed to react in this reaction. State combinations can be created or interpreted with the functions **rxnpackstate** and **rxnunpackstate**. For order 2 and above, **permit** is not necessarily symmetric: for example, if solution and front are permitted, this does not imply that front and solution are permitted (however, **permit** is symmetric if multiple reactant identities are the same).

nprod is the number of products for the listed reaction, which may be any non-negative number. **prdident** and **prdstate**, which are arrays that are indexed from 0 to **rxn->nprod-1**, list the product identities and states; in this case, only single states are allowed (i.e. not **MSall** or **MSnone**, although **MSbsoln** is allowed).

prdserno is a list of serial number rules for products. The default is that this list is **NULL**, meaning not used. If it is used, then it lists a rule for each product. If it is a positive value, then it is the actual number that should be used as the product serial number. Otherwise, the value is bitwise inverted and interpreted using bits. See **rxnstring2sernocode** and **molfiindserno**.

pr dintersurf is a list of product placement rules for intersurface reactions, meaning bimolecular reactions in which the two reactants are on two different surfaces. These are not allowed typically. However, they are allowed if this vector is allocated. Each value corresponds to a product. The rules are: 1 implies that this product should be placed on the surface, or relative to the surface, of the first reactant, and 2 implies that this product should be placed on the surface, or relative to the surface, of the second reactant. If the reaction has no products, then the vector can be allocated for a single element, which is supposed to have value 0 to indicate that intersurface reactions are allowed.

logserno and **logfile** are for logging individual reactions as they occur, which is done to the file named **logfile**. **logserno** is a list of serial numbers for which this reaction should be logged. If it is an empty list, then all instances of this reaction are logged.

rate is the reaction rate constant, measured in whichever unit system that the user is using for other aspects of the configuration file. The general rate units are $\text{molecules} * \text{volume}^{(\text{order}-1)} / \text{time}$. The precise meaning of **rate** depends on the order of the reaction. The actual reaction rate is multiplied by **multiplicity**, which is here in case the same reaction arises multiple times through network generation. **bindrad2**, which only applies to order 2 and higher reactions, is the squared binding radius of the reactants. **prob** is, roughly, the reaction probability per time step. For zeroth order reactions, **prob** is the expectation number of reactions per time step in the entire simulation volume; for first order reactions, **prob** is the probability of a reactant reacting during one time step; and for second order reactions, **prob** is the probability of a reaction occurring between two reactants that have already diffused closer than their binding radius. **chi** is the diffusion-limited fraction for this reaction, meaning that it is the simulated steady-state reaction rate divided by the diffusion-limited rate that would occur with these parameters if the time step were equal to zero. **tau** is the characteristic time for the reaction. It is calculated from the other reaction parameters and, for order 2 reactions, from the initial concentrations of the molecules. The information in **tau** is completely redundant with information that is elsewhere.

rparamt is the type of the reversible parameter and **rparam** is the value of the reversible parameter. **unbindrad** applies to all reactions that have exactly 2 products; it is the unbinding radius of the products. **pr dpos** is a list of product position displacements from the reaction position.

See the description for `RxnSetRevparam`. `disable` is a flag that is 0 for normal operation and 1 if a reaction is disabled, meaning that it isn't run. One use of this is that some reactions should run in lattice space only and not in particle space, so these are disabled in particle space. `cmpt` is the compartment that a reaction occurs in, or `NULL` if it occurs everywhere. Conformational spread reactions, identified with `rparamt` equal to `RPconfspread` have reverse reaction rates that are not accounted for during rate calculations and the products are placed in the exact same places as the reactants.

Unimolecular reaction rates are surprisingly complicated. For a single reaction channel, they are simple. For multiple channels, they use the formula that is given in Andrews and Bray, 2004 for each reaction rate. Then, each probability is divided by one minus the sum of the prior probabilities to account for the fact that what's wanted is the conditional probability that a reaction happens, given that prior reactions did not happen. An alternate and possibly better method is used for surface actions, where the probability of each individual event is not stored, but instead the cumulative probability for the events is stored. Both methods are accurate.

5.4.3 reaction superstructure

```
typedef struct rxnsuperstruct {
    enum StructCond condition; // structure condition
    struct simstruct *sim;     // simulation structure
    int order;                 // order of reactions listed: 0, 1, or 2
    int maxspecies;            // maximum number of species
    int maxlist;               // copy of maximum number of molecule lists
    int *nrxn;                 // number of rxns for each reactant set [i]
    int **table;               // lookup table for reaction numbers [i][j]
    int maxrxn;                // allocated number of reactions
    int totrxn;                // total number of reactions listed
    char **rname;              // names of reactions [r]
    rxnptr *rxn;               // list of reactions [r]
    int *rxnmollist;           // live lists that have reactions [ll]
} *rxnssptr;
```

The reaction superstructure, `rxnsuperstruct`, is a structure that is used for all of the reactions that are accounted for by the simulation, of a given order. Thus, there may be one for zeroth order reactions, another for first order reactions, and a third for second order reactions. Higher order reactions may be supported as well, although they are not currently. `condition` is the current condition of the superstructure and `sim` is a pointer to the simulation structure that owns this superstructure. `order` is the order of the reactions that are listed in this superstructure and `maxspecies` is simply a copy of the `maxspecies` value from the simulation structure. `maxlist` is the number of molecule lists that are assumed for, and thus contributes to the size of, `rxnmollist`.

`nrxn` is the number of reactions that are defined for a certain reactant code; conversions between reactant lists and reactant codes may be performed by the functions `rxnpackident` and `rxnunpackident`. `table` is a lookup table with which one inputs the reactant code (`[i]`) and the reaction number for that code (`[j]`, which is 0 to `nrxn[i]-1`), and is given a reaction number; `table` is always symmetric with respect to reactant identities, which only applies to order 2 and higher reactions. Empty molecules are included in these lists, accessed with `nrxn[0]` and `table[0]`, where the former should always equal 0 and the latter should always be `NULL`. The reactions are listed next. `maxrxn` is the number of reactions of this order that have been allocated, while `totrxn` is the total number of reactions of this order that are currently defined. `rname` and `rxn`, which may be indexed from 0 to `totrxn-1`, are the list of reaction names, and the respective reactions,

5 C Code: structures and functions

respectively. `rxnmolllist`, which has $maxlist^{order}$ elements where `maxlist` is listed above, is a list of flags that indicate which molecule lists, or molecule list combinations, need to be checked to find reactions of this order.

5.4.4 packed species identities

Several of the structure elements use packed values, which can be performed with `rxnpackident` and similar functions. Alternatively, they can be done directly according to the following scheme:

item	examples	order = 1	order = 2
identity	<code>nrxn[i], table[i]</code>	<code>i1</code>	<code>i1*maxspecies+i2</code>
state	<code>permit[ms]</code>	<code>ms1</code>	<code>ms1*MSMAX1+ms2</code>
live list	<code>rxnmolllist[l1]</code>	<code>l11</code>	<code>l11*maxlist+l12</code>

See the Wildcards section of the Code Design section for an explanation of how reactions are set up and expanded.

5.4.5 reaction functions

enumerated types

`enum RevParam rxnstring2rp(char *string)`

Converts string to enumerated `RevParam` type. This reads either single letter inputs or full word inputs. Unrecognized inputs are returned as `RPnone`.

`char *rxn2string(enum RevParam rp, char *string)`

Converts `RevParam` enumerated type variable `rp` to a word in `string` that can be displayed. Unrecognized inputs, as well as `RPnone`, get returned as "none".

`enum SpeciesRepresentation rxnstring2sr(char *string)`

Converts string to enumerated `Species Representation` type. This reads as many letters of the string as are given, seeing if they match or not. Unrecognized inputs are returned as `SRnone`.

`char *rxnsr2string(enum SpeciesRepresentation sr, char *string)`

Converts enumerated `SpeciesRepresentation` variable `sr` to a word in `string` that can be displayed. Unrecognized inputs, as well as `SRnone`, get returned as "none".

low level utilities

`int readrxnname(simptr sim, const char *rname, int *orderptr, rxnptr *rxnpt, listptrv *vlistptr, int rxntype)`

Using a reaction name in `rname`, this looks for it in one of several places, set by `rxntype` until it finds it. If `rxntype` is 1, this looks in the current list of reaction names, working with increasing reaction orders. If it is found, it returns the reaction order in `orderptr`, a pointer to the reaction in `rxnpt`, and the reaction number directly; if not, it returns -1. If `rxntype` is 3, this searches the list of rules to see if there is a reaction rule with the same reaction name. If so, this returns a pointer to the template reaction in that rule in `rxnpt`, creating one if necessary, the order in `orderptr`, and the rule number directly. If `rxntype` is 2, this searches the current list of reaction names, working with increasing order numbers. If this finds one or more names, all of the same order, that all start with `rname` and are followed

5.4 Reactions (functions in *smolrxn.c*)

by an underscore and then a number, then this lists those reaction pointers, cast as `void*`s, in `vlistptr`. In this case, the function returns 0; it also returns the first reaction that it found in `rxnpt`. In all cases, this returns -1 if no reaction is found to match `rname`. This also returns -2 for failure to allocate memory (which is only possible for `rxntype` 2 or 3).

`int rxnpackident(int order, int maxident, int *ident)`

Packs a list of order identities that are listed in `ident` into a single value, which is returned. `maxident` is the maximum number of identities, from either the reaction superstructure or the simulation structure.

`void rxnunpackident(int order, int maxident, int ipack, int *ident)`

Unpacks a packed identity that is input in `ipack` to order individual identities in `ident`. `maxident` is the maximum number of identities, from the reaction superstructure or the simulation structure.

`enum MolecState rxnpackstate(int order, enum MolecState *mstate)`

Packs of list of `order` molecule states that are listed in `mstate` into a single value, which is returned.

`void rxnunpackstate(int order, enum MolecState mspack, enum MolecState *mstate)`

Unpacks a packed molecule state that is input in `mspack` to order individual states in `mstate`.

`int rxnreactantstate(rxnptr rxn, enum MolecState *mstate, int convertb2f)`

Looks through the reaction `permit` element to see if the reaction is permitted for any state or state combination. If not, it returns 0; if so, it returns 1. Also, if the reaction is permitted and if `mstate` is not NULL, the “simplest” permitted state is returned in `mstate`. Preference is given to `MSsoln` and `MSbsoln`, with other states investigated afterwards. If `convertb2f` is set to 1, any returned states of `MSbsoln` are converted to `MSsoln` before the function returns. This function always returns a single state in `mstate` (or `MSnone` if the reaction is not permitted at all), and never `MSall` or `MSsome`.

`int rxnallstates(rxnptr rxn)`

Returns 1 if the listed reaction is permitted for all reactant states and 0 if not.

`int findreverserxn(simpstr sim, int order, int r, int *optr, int *rptr)`

Inputs the reaction defined by order `order` and reaction number `r` and looks to see if there is a reverse reaction. All molecule states for the input reaction that can react with reaction `r` are considered. If there is a direct reverse reaction, meaning the products of the input reaction (including states), are themselves able to react to form the reactants of the input reaction (with states that can produce reaction `r`), then the function returns 1 and the order and reaction number of the reverse reaction are pointed to by `optr` and `rptr`. If there is no direct reverse reaction, but the products of the input reaction are still able to react, the function returns 2 and `optr` and `rptr` point to the first listed continuation reaction. The function returns 0 if the products do not react with each other, if there are no reactants, or if there are no products. -1 is returned for illegal inputs. Either or both of `optr` and `rptr` are allowed to be sent in as NULL values if the respective pieces of output information are not of interest.

`int rxnisprod(simpstr sim, int i, enum MolecState ms, int code)`

Determines if a molecule with identity `i` and state `ms` is the product of any reaction, of any order, returning 1 if so and 0 if not. `ms` can include `MSbsoln`. If `code` is 0, there are no

5 C Code: structures and functions

additional conditions. If `code` is 1, the molecule also has to be displaced from the reaction position (i.e. either `confspread` or the unbinding radius is non-zero) in order to qualify.

`long int rxnstring2sernocode(char *pattern, int prd)`

Converts a string pattern to a serial number code. If the pattern is a positive number, then this is the serial number code. Otherwise, a bitcode is computed from the pattern, the bitcode is bitwise inverted, and that value is returned as the serial number code, which is then a negative number. See `molfindserno` for a description of the bitcode. This function is written in a brute force fashion because it was simpler to write and to read, but could have been written more cleverly. Returns 0 for failure, a negative value if a bitcode is used, or a positive number for an actual serial number.

`char *rxnsernocode2string(long int pserno, char *pattern)`

Converts a serial number code to a pattern string. This assumes that `pattern` is allocated and is large enough for the result (maximum of 7 characters for codes or the number of characters for a long int, which is probably a few more than 7). See the descriptions for `rxnstring2sernocode` and `molfindserno`.

memory management

`rxnptr rxnalloc(int order)`

Allocates and initializes a reaction structure of order `order`. The reaction has `order` reactants, a `permit` element that is allocated, no products, and most parameters are set to -1 to indicate that they have not been set up yet.

`void rxnfree(rxnptr rxn)`

Frees a reaction structure.

`rxnssptr rxnssalloc(rxnssptr rxnss, int order, int maxspecies)`

Allocates and initializes a reaction superstructure of order `order` and for `maxspecies` maximum number of species (the same value that is in the molecule superstructure). The superstructure is left with `nrxn` and `table` allocated but with no reactions. This function may be called more than once, which is useful for increasing `maxspecies`. On the first call, enter `rxnss` as NULL and enter it with the existing value on subsequent calls. `maxspecies` may not be decreased. Returns a pointer to the reaction superstructure on success, or NULL on inability to allocate memory.

`void rxnssfree(rxnssptr rxnss)`

Frees a reaction superstructure including all component reactions.

`int rxnexpandmaxspecies(simptr sim, int maxspecies)`

Expands the `maxspecies` value for all existing reaction superstructures, allocating memory as needed. These values should be kept synchronized with the master `maxspecies` in the molecule superstructure, so this is called whenever the master one changes. Returns 0 for success or, if memory could not be allocated, 1 plus the order of the superstructure where the failure occurred.

data structure output

`void rxnoutput(simptr sim, int order)`

Displays the contents of a reaction superstructure for order `order`, as well as all of the

5.4 Reactions (functions in *smolrxn.c*)

component reactions. It also does some other calculations, such as the probability of geminate reactions for the products and the diffusion and activation limited rate constants.

void writereactions(simpstr sim, FILE *fptr)

Writes all information about all reactions to the file **fptr** using a format that can be read by Smoldyn. This allows a simulation state to be saved.

int checkrxnparams(simpstr sim, int *warnptr)

Checks some parameters of reactions to make sure that they are reasonable. Prints warning messages to the display. Returns the total number of errors and, if **warnptr** is not NULL, the number of warnings in **warnptr**.

parameter calculations

int rxnsetrate(simpstr sim, int order, int r, char *erstr)

Sets the internal reaction rate parameters for reaction **r** of order **order**. These parameters are the squared binding radius, **bindrad2**, and the reaction probability, **prob**. Zero is returned and **erstr** is unchanged if the function is successful. Possible other return codes are: 1 for a negative input reaction rate (implies that this value has not been defined yet, which is not necessarily an error; other parameters are not modified), 2 for order 1 reactions for which different reactant states would have different reaction probabilities, 3 for confspread reactions that have a different number of reactants and products, 4 for non-confspread bimolecular reactions that have non-diffusing reactants, or 5 for a reaction probability that is out of range.

For zeroth order reactions, **rxn->prob** is the expectation number of molecules that should be produced in the entire simulation volume during one time step, which is $rate \cdot dt \cdot volume$.

For first order reactions, **rxn->prob** is the conditional probability of a unimolecular reaction occurring for an individual reactant molecule during one time step, where the condition is that any previously listed possible reactions for this reactant were not chosen. First, this computes the unconditional probability. For this, if there is only one reaction possible, then the reaction probability is $prob = 1 - \exp(-rate \cdot dt)$. However, other reaction channels affect the probability because, over a finite time step, the reactant may get used up before it has a chance to react in this reaction. The solution is in Andrews and Bray, 2004, eq. 14, which is $prob = (rate/sum)[1 - \exp(-sum \cdot dt)]$, where *sum* is the sum of all of the reaction channel rates. Next, this probability is conditioned as follows. Consider a reactant, A, which reacts to product B₁ with probability p_1 , to product B₂ with probability p_2 , to product B₃ with probability p_3 , etc. The algorithm is: the first path is taken with probability p'_1 ; if that is not taken, then the next path is taken with probability p'_2 ; if that is not taken, then the next path is taken with probability p'_3 , etc. Since there is no prior condition, $p_1 = p'_1$. The probability that path 2 is actually taken is $p_2 = p'_2(1 - p_1)$ because it is the probability that event 2 happens and that path 1 was not chosen. The probability that path 3 is actually taken is $p_3 = p'_3(1 - p_1)(1 - p_2)$. Thus, for example, $p'_3 = p_3/(1 - p_1)(1 - p_2)$. Here, **product** is the probability that prior paths were not taken and **prob** is the probability that a certain path is taken (e.g. p_1 or p_2).

For second order reactions, **rxn->bindrad2** is the squared binding radius of the reactants, found from **bindingradius**. In this case, the reverse parameter is accounted for in the reaction rate calculation if there is a direct reverse reaction and if it is appropriate (see the discussion of "Binding and unbinding radii," and the description for **findreverserxn**). The requested

5 C Code: structures and functions

rate is doubled if the two reactants are the same due to the fact that there are half as many possible interactions in this case than if the two reactants are different. This is explained in the User's manual: "Consider a situation with 1000 A molecules and 1000 B molecules. Despite the fact that each A molecule has about 1000 potential collision partners, whether the reactants are A + A or A + B, there are twice as many A-B collisions as A-A collisions. This is because each A-A pair can be counted in either of two ways, but is still only a single possible collision. To achieve the same reaction rate for A + A reactants as for A + B, despite the fact that there are fewer collisions, Smoldyn uses a larger binding radius for the former." Smoldyn also doubles the requested reaction rate if one of the species is membrane-bound and the other is in solution. In this case, it is because the membrane occludes half of the binding volume, with the result that only half of it is accessible to binding. Thus, Smoldyn simulates the reaction with only half of the rate that would result if both species were in solution, which means that the requested rate needs to be doubled to get the correct rate.

`int rxnsetrates(simptr sim, int order, char *erstr)`

Sets internal reaction rate parameters for all reactions of order `order`. The return value of the function is -1 for correct operation. If errors occur, the reaction number where the error was encountered is returned and an error string is written to `erstr`, which should be pre-allocated to size `STRCHAR`. This function simply calls `rxnsetrate` for each reaction.

`int rxnsetproduct(simptr sim, int order, int r, char *erstr)`

Sets the initial separations for the products of reaction `r` of order `order`. This uses the `rparamt` and `rparam` elements of the reaction to do so, along with other required values such as the binding radius and parameters from any reverse reaction. The `unbindrad` and `prdpos` elements are set up here. If `rpart` is either `RPoffset` or `RPfixed`, then it is assumed that the product positions have already been set up; they are not modified again by this routine. Otherwise, they are set here, described next. This returns 0 for success or any of several error codes for errors. For each error, a message is written to `erstr`, which needs to have been pre-allocated to size `STRCHAR`. The error codes aren't listed here because they aren't used by the functions that call this one; in other words, they're irrelevant. For warnings, the function returns 0 but a warning string is written to `erstr`.

If the `rparamt` value is for `RPoffset` or `RPfixed`, then `prdpos` is not changed; otherwise it is. If there are 0 products, then everything is an error except for `RPnone`, `RPirrev`, and `RPconfspread` because nothing else makes sense. If there is 1 product, then the product position is set to the 0 vector for all `rparamt` values that make sense because there is no unbinding radius to worry about. However, most `rparamt` values don't make sense. If there are 2 products, the value depends on the `rparamt` value and on whether the reaction is reversible or not. In all cases, only the x coordinate of the `prdpos` vector is set to a non-zero value (the other coordinates are not touched at all, and nor should they be used). The following table is for reactions with 2 products.

<code>rparamt</code>	<code>rparam</code>	<code>unbindrad</code>	<code>prdpos</code>
<code>RPnone</code>	-	0 for irrev., error for rev.	0
<code>RPirrev</code>	-	0	0
<code>RPconfspread</code>	-	0	0
<code>RPbounce</code>	σ_u	σ_u	sum is σ_u
	-1	-1	sum is σ_b
	-2	-2 (?)	0 (?)

5.4 Reactions (functions in *smolrxn.c*)

RPpgem	ϕ	σ_u	difference is σ_u
RPpgemmax	ϕ_{max}	σ_u	difference is σ_u
RPpgemmaxw	ϕ_{max}	σ_u	difference is σ_u
RPratio	σ_u/σ_b	σ_u	difference is σ_u
RPunbindrad	σ_u	σ_u	difference is σ_u
RPpgem2	ϕ	σ_u	difference is σ_u
RPpgemmax2	ϕ_{max}	σ_u	difference is σ_u
RPratio2	σ_u/σ_b	σ_u	difference is σ_u
RPoffset	-	set from <code>prdpos</code>	unchanged
RPfixed	-	set from <code>prdpos</code>	unchanged

* For **RPbounce**, both product positions are positive and are set for a cumulative vector length of the unbinding radius if the **rparam** value is positive and to a cumulative vector length of the binding radius if **rparam** is negative.

int rxnsetproducts(simpstr sim, int order, char *erstr)

Sets initial separations for products of all reactions of order **order**. This returns -1 for success and the reaction number for failure. Upon failure, this also returns **erstr**, which needs to have been preallocated with size **STRCHAR**, with an error message. See the discussion in the section called “Binding and unbinding radii” for more details.

double rxnccalcrate(simpstr sim, int order, int r, double *pgemptr)

Calculates the macroscopic rate constant using the microscopic parameters that are stored in the reaction data structure. All going well, these results should exactly match those that were requested initially, although this routine is useful as a check, and for situations where the microscopic values were input rather than the mass action rate constants.

For unimolecular reactions, this reverses the computations that are made in **rxnsetrate**; see that description for more details. **ms1** is a state for which this reaction is permitted. Next, this goes through all reactions of the same reactant and state. For each, this computes the unconditional probability of the reaction being chosen, in **prob**. This is the product of the conditional probability of it being chosen (**rxn2->prob**), where the condition is that prior reactions were not chosen, and the probability that prior reactions were not chosen (**product**). This also adds up the unconditional probabilities in **sum**. Addition is appropriate here because a reaction can only happen along a single reaction channel, so these are exclusive probabilities. At the end, **sum** is the total probability that a reaction happens. Then, inversion of the equation, $sum = 1 - \exp(-ratesum \cdot dt)$ gives the sum of all the rate constants. From this, the rate constant for this particular reaction is computed.

For bimolecular reactions that are reversible, the routine calculates rates with accounting for reversibility if the reversible parameter type of the reverse reaction is: **RPpgem**, **RPpgemmax**, **RPpgemmaxw**, **RPratio**, or **RPunbindrad** and not otherwise (this list was changed for version 2.56, to make it agree with the Smoldyn User’s Manual). A value of -1 is returned if input parameters are illegal and a value of 0 is returned if the microscopic values for the indicated reaction are undefined (j0). If the input reaction has a reverse reaction or a continuation reaction, and **pgemptr** is not input as **NULL**, then ***pgemptr** is set to the probability of geminate recombination of the products; if there is no reverse or continuation reaction, its value is set to -1.

void rxnccalctau(simpstr sim, int order)

5 C Code: structures and functions

Calculates characteristic times for all reactions of order `order` and stores them in the `rxn->tau` structure elements. These are ignored for 0th order reactions, are $1/k$ for first order reactions, and are $[A][B]/[k([A]+[B])]$ for second order reactions. The actual calculated rate constant is used, not the requested ones. For second order, the current average concentrations are used, which does not capture effects from spatial localization or concentration changes. For bimolecular reactions, if multiple reactant pairs map to the same reaction, only the latter ones found are recorded. Also, all molecule states are counted, which ignores the `permit` reaction structure element.

structure set up

```
void rxnsetcondition(simptr sim, int order, enum StructCond cond, int upgrade)
```

Sets the reaction superstructure condition, for order `order`, to `cond`, if appropriate. Set `order` to the desired reaction order, or to -1 for all reaction orders. Set `upgrade` to 1 if this is an upgrade, to 0 if this is a downgrade, or to 2 to set the condition independent of its current value. If the condition is downgraded, this also downgrades the simulation structure condition.

```
int RxnSetValue(simptr sim, char *option, rxnptr rxn, double value)
```

Sets certain options of the reaction structure for reaction `rxn` to `value`. If `sim` is not NULL, this then downgrades the reaction condition to `SCLists` to cause recomputation of molecule lists to check for reactions and also reaction simulation parameters. Returns 0 for success, 2 for unknown option, or 4 for an illegal value (e.g. a negative rate). In most cases, the value is set as requested, despite the error message.

If `option` is “rate”, the `rate` element is set; if `option` is “multiplicity”, the `multiplicity` element is set; if `option` is “multiplicity++”, the `multiplicity` element is incremented; if `option` is “rateadd”, the value is added to the current reaction rate (this has been largely or completely superseded by the `multiplicity` element); if `option` is “confspreadrad”, the reaction type is made `confspread` and the squared binding radius is set; if `option` is “bindrad”, the squared binding radius is set; if `option` is “prob”, the probability is set; if `option` is “disable”, the `disable` element is set (used for hybrid simulation in which reactions in one part of space might be disabled in other parts of space).

```
int RxnSetValuePattern(simptr sim, const char *option, const char *pattern, const enum MolecState *rctstate, const enum MolecState *prdstate, double value, int oldnresults, const rxnptr template)
```

Same as `RxnSetValue` except for one or more reactions that are described in `pattern` rather than with a pointer to a reaction. As before, `option` is the parameter to be set and `value` is the value it is to be set to. This also needs the reactant and product states, which are not part of the `pattern`. `oldresults` is the point of the index list where this should start setting values. This steps through all of the reactions listed in `pattern`, starting with `oldresults` and continuing to the end of the list, and calls `RxnSetValue` for each one. This function checks to see if a reaction has had its rate set before; if not, it sets the rate to `value` and if so, it adds `value` to the existing rate. Returns 0 for success, 2 for unknown option, 4 for an illegal value (e.g. a negative rate), or 5 for a reaction that is listed in the pattern but does not actually exist.

```
int RxnSetRevparam(simptr sim, rxnptr rxn, enum RevParam rparamt, double rparam, int prd, double *pos, int dim)
```

5.4 Reactions (functions in *smolrxn.c*)

Sets the reversible parameter type and the appropriate reversible parameters for reaction **rxn**. This function is called by the `product_placement` portion of `simreadstring`. The parameter type, `rxn->paramt`, is set to `rpart`. If `rpart` requires a single value, which is stored in `rxn->rparam`, it is sent in with `rparam`. Otherwise, for fixed and offset parameter types, send in the product number that is being altered with `prd`, the vector with `pos`, and the system dimensionality with `dim`. This returns 0 for success, 1 as a warning that the reversible parameter type has been set before (except for offset and fixed types, where many different products need to be set), 2 for parameters that are out of bounds, 3 for an unrecognized `rparamt`, 4 for `prd` out of bounds, or 5 for a missing `pos` vector.

<code>rparamt</code>	<code>rparam</code>	<code>prd</code>	<code>pos</code>
<code>RPhnone</code>	-	-	-
<code>RPirrev</code>	-	-	-
<code>RPconfspread</code>	-	-	-
<code>RPbounce</code>	σ_u or -1	-	-
<code>RPpgem</code>	ϕ	-	-
<code>RPpgemmax</code>	ϕ_{max}	-	-
<code>RPpgemmaxw</code>	ϕ_{max}	-	-
<code>RPratio</code>	σ_u/σ_b	-	-
<code>RPunbindrad</code>	σ_u	-	-
<code>RPpgem2</code>	ϕ	-	-
<code>RPpgemmax2</code>	ϕ_{max}	-	-
<code>RPratio2</code>	σ_u/σ_b	-	-
<code>RPOffset</code>	-	product number	relative position
<code>RPfixed</code>	-	product number	relative position

If `method` is `RPbounce`, then a negative number for the `parameter` indicates default bounce behavior, which is that molecules are separated by an amount that is equal to their previous overlap.

`void RxnCopyRevParam(simpstr sim, rxnptr rxn, const rxnptr template)`

Copies the reverse reaction parameters from the `template` reaction to the `rxn` reaction. This doesn't do any checking. All inputs are required.

`void RxnSetPermit(simpstr sim, rxnptr rxn, int order, enum MolecState *rctstate, int value)`

Sets the `permit` element of reaction `rxn`, which has order `order`, for the states that are included in `rctstate` to value `value`. `value` should be 0 to set permissions to forbidden or 1 to set permissions to permitted. Each item of `rctstate` may be an individual state or may be `MSall`. Other values are not allowed (and are not caught here). This does not affect other `permit` elements. If `order` is 2 and both reactants are the same, this automatically makes the `permit` matrix symmetric (e.g. if input states are `MSall` and `MSfront`, respectively, permissions will also be set for the pair `MSfront` and `MSall`).

`void RxnSetCmpt(rxnptr rxn, compartptr cmpt)`

Sets the `cmpt` element of the `rxn` reaction to compartment `cmpt`. This does no checking, and assigns regardless of whether `cmpt` is NULL or not.

`void RxnSetSurface(rxnptr rxn, surfaceptr srf)`

5 C Code: structures and functions

Sets the `srf` element of the `rxn` reaction to surface `srf`. This does no checking, and assigns regardless of whether `srf` is `NULL` or not.

`int RxnSetPrdSerno(rxnptr rxn, long int *prdserno)`

Sets the product serial number list for reaction `rxn` to the list given in `prdserno`. If the product serial number list had not been allocated previously, it is allocated here. Returns 0 for success or 1 for inability to allocate memory.

`int RxnSetIntersurfaceRules(rxnptr rxn, int *rules)`

Sets the product intersurface rules for reaction `rxn` to the list given in `rules`. If the product intersurface list had not been allocated previously, it is allocated here. Returns 0 for success or 1 for inability to allocate memory.

`int RxnSetRepresentationRules(rxnptr rxn, int order, const enum`

`SpeciesRepresentation *rctrep, const enum SpeciesRepresentation *prdrep)`

Sets lattice versus particle representation for reaction reactants and products, for use in overlapping space hybrid simulation. Enter the reaction in `rxn`, the reaction order in `order`, and the desired representation for the reactants and products in the vectors `rctrep` and `prdrep`. Enter the first element of `rctrep` as `SRfree` to free the data structures. This function copies the contents of `rctrep` and/or `prdrep` to the arrays of the same name in the reaction structure.

`int RxnSetLog(simptr sim, char *filename, rxnptr rxn, listptrli list, int turnon)`

Sets reaction logging, which occurs upon reaction occurrence. Enter `filename` as the name of the file that reaction data should be logged to. This gets stored in `rxn->logfile`. Enter `rxn` as the reaction that should get logged, or as `NULL` for all reactions in the simulation. Enter `list` as a list of molecule serial numbers that should be logged, or as a list with only the value -1 if all molecules serial numbers should be logged. Finally, enter `turnon` as 1 to turn on logging and as 0 to turn off logging. In the latter case, the `filename` entry is ignored, the `rxn` entry is the reaction for which logging should be turned off, and `list` is the list of serial numbers for which logging should be turned off. Returns 0 for success, 1 for inability to allocate memory, or 2 as a warning that prior logfile was overwritten. If logging is turned off for all serial numbers, this erases the stored logfile name as well, meaning that it would need to be re-entered if logging is wanted again.

`rxnptr RxnAddReaction(simptr sim, char *rname, int order, int *rctident, enum MolecState *rctstate, int nprod, int *prdident, enum MolecState *prdstate, compartptr cmpt, surfaceptr srf)`

Adds a reaction to the simulation, including all necessary memory allocation. `rname` is the name of the reaction, `order` is the order of the reaction, and `nprod` is the number of products. `rctident` and `rctstate` are vectors of size `order` that contain the reactant identities and states, respectively. Likewise, `prdident` and `prdstate` are vectors of size `nprod` that contain the product identities and states. This returns the just added reaction for success and `NULL` for inability to allocate memory. This allocates reaction superstructures and reaction structures, and will enlarge any array, as needed. This function can also be used sequentially for reactants and products: first call it with reactants and 0 for `nprod`; next time, call it with the correct `order`, `NULL` for both reactant inputs, and the full product information.

5.4 Reactions (functions in *smolrxn.c*)

`rxnptr RxnTestRxnExist(simptr sim, int order, const char *rname, const int *rctident, const enum MolecState *rctstate, int nprod, const int *prdidnt, const enum MolecState *prdstate, int exact)`
Tests if a reaction already exists. Set `exact` to 1 if this should just test to see if the reaction name is already in use. If not, this checks to see if some other reaction has the same `order` value, the same root of the reaction name, which is entered in `rname`, the same reactants as `rctident`, the same `order` reactant states as `rctstate`, the same `nprod` products as `prdidnt`, and the same `nprod` product states as `prdstate`. Returns the reaction if the reaction already exists and NULL if not.

`int RxnAddReactionPattern(simptr sim, const char *rname, const char *pattern, int oldnresults, enum MolecState *rctstate, enum MolecState *prdstate, compartptr cmpt, surfaceptr srf, int isrule, rxnptr *rxnpt)`
Adds one or more reactions to the simulation, where the reaction is defined in the `pattern` string. Enter `rname` as the name of the reaction if there is just one reaction, and as the root of the name if there are multiple reactions. The `pattern` string should have been created in the `molstring2pattern` function; it should have the reactants as space-separated words, then a newline character, then the products as space-separated words. Enter `oldnresults` as -1 if the current value of the index variable should be used to determine how many results from index have been considered already and as a number greater than or equal to zero to specify how many results from index have been considered already. Enter `rctstate` as a list of reactant states, where this list should have the same number of elements as the number of reactants in `pattern`, and enter `prdstate` as a list of product states, where this list should have the same number of elements as the number of products in `pattern`. `cmpt` and `srf` are optional; they give the compartment that the reaction will occur in and the surface that it will occur on, if the reaction should be restricted in these ways. Set `isrule` to 1 to indicate that this is a rule, meaning that any products that have not been declared beforehand will be created here, or to 0 to indicate that this is not a rule. If `rxnpt` is entered as non-NULL and if there is just one reaction created here, then it will be returned pointing to the new reaction. If more than one reaction is created here, it will be returned as NULL.

This function checks the current status of the pattern to see how much of it has been updated previously. Then, it loops over all reactions that have not been added to the simulation yet. If it encounters a reaction that already exists, it simply skips over it.

Returns 0 for success, -1 for inability to allocate memory, -2 if no wildcards were entered and one or more of the match species is unknown, -3 if a substitute species is unknown and the substitute species had no wildcards in it, -4 if the match string included more words than allowed by this function (which is 4 currently), -5 if a trial match string was too long to fit in `STRCHAR` characters (even if this wasn't actually a match), -6 if species generation failed, -11 for inability to allocate memory, -12 for missing ' ' operand, -13 for missing & operand, -15 for mismatched braces, or -20 for a destination pattern that is incompatible with the matching pattern (i.e. it has to have either 1 destination or the same number of destination options as pattern options), or -30 for failure to add the reaction.

`rxnptr RxnAddReactionCheck(simptr sim, char *rname, int order, int *rctident, enum MolecState *rctstate, int nprod, int *prdidnt, enum MolecState *prdstate, compartptr cmpt, surfaceptr srf, char *erstr)`
This is a simple wrapper for `RxnAddReaction`. Before it calls `RxnAddReaction` though, it checks as many of the input parameters as possible to make sure that they are reasonable.

5 C Code: structures and functions

If they are not reasonable, it returns NULL and an error message in `erstr`, which should be allocated to size STRCHAR.

`int loadrxn(simptr sim, ParseFilePtr *pfpptr, char *line2, char *erstr)`

Loads a reaction structure from an already opened disk file described with `pfpptr`. If successful, it returns 0 and the reaction is added to `sim`. Otherwise it returns 1 and error information in `pfpptr`. If a reaction structure of the same order has already been set up, this function can use it and add more reactions to it. It can also allocate and set up a new structure, if needed. This need for this function has been largely superceded by functionality in `loadsim`, but this is kept for backward compatibility.

`int rxnupdateparams(simptr sim)`

Sets reaction structure parameters for the simulation time step. Return values are 0 for success, 1 for an error with setting either rates or products (and output to `stderr` with an error message), or 2 if the reaction structure was not sufficiently set up beforehand.

`int rxnupdatelists(simptr sim, int order)`

Sets the `rxnmollist` element of the reaction superstructure of order `order`. If one already exists, it is freed and then reallocated; otherwise it is just allocated. Afterwards, this function goes through all reactants of the superstructure, including their `permit` values, and registers their respective molecule lists in the `rxnmollist` array. Returns 0 for success, 1 for failure to allocate memory, 2 for a requested order that is greater than 2 (which is the highest that this function can handle), or 3 for molecules not being set up sufficiently.

`int rxnsupdate(simptr sim)`

Sets up reactions from data that have already been entered. This sets the reaction rates, sets the reaction product placements, sets the reaction `tau` values, and sets the molecule list flags. Returns 0 for success, 1 for failure to allocate memory, 2 for a Smoldyn bug, 3 for molecules not being set up sufficiently, 4 for an error with setting either rates or products (in this case, an error message is displayed to `stderr`), or 5 if the reaction structure was not sufficiently set up. This may be run at at start-up or afterwards.

reaction parsing function

`int rxnparsereaction(simptr sim, const char *word, char *line2, char *errstr)`

Parses reaction statement in configuration file. This code was in the “reaction” statement section of `simreadstring` function but became too long and complicated, so it got moved to its own function. It’s only called by `simreadstring`.

core simulation functions

`int doreact(simptr sim, rxnptr rxn, moleculeptr mptr1, moleculeptr mptr2, int l11, int m1, int l12, int m2, double *pos, panelptr rxnpl)`

Executes a reaction that has already been determined to have happened. `rxn` is the reaction and `mptr1` and `mptr2` are the reactants, where `mptr2` is ignored for unimolecular reactions, and both are ignored for zeroth order reactions. `l11` is the live list of `mptr1`, `m1` is its index in the master list, `l12` is the live list of `mptr2`, and `m2` is its index in the master list; if these don’t apply (i.e. for 0th or 1st order reactions, set them to -1 and if either `m1` or `m2` is unknown, again set the value to -1. If there are multiple molecules, they need to be in the same order as they are listed in the reaction structure (which is only important for confspread reactions and for a completely consistent panel destination for reactions between two surface-bound

5.4 Reactions (functions in *smolrxn.c*)

molecules). The `pos` and `rxnpl` inputs are only looked at for 0th order reactions; for these, they need to be a random position for the reaction to occur, and the panel if any.

Reactants are killed, but left in the live lists. Any products are created on the dead list, for transfer to the appropriate live list by the `molsort` routine. Molecules that are created are put at the reaction position, which is the average position of the reactants weighted by the inverse of their diffusion constants, plus an offset from the product definition. The cluster of products is typically rotated to a random orientation. If the displacement was set to all 0's (recommended for non-reacting products), the routine is fairly fast, putting all products at the reaction position. If the `rparam` character is `RPfixed`, the orientation is fixed and there is no rotation. Otherwise, a non-zero displacement results in the choosing of random angles and vector rotations. If the system has more than three dimensions, only the first three are randomly oriented, while higher dimensions just add the displacement to the reaction position. The function returns 0 for successful operation and 1 if more molecules are required than were initially allocated. This function lists the correct box in the box element for each product molecule, but does not add the product molecules to the molecule list of the box.

For the bounce product placement, the general rule is that reactants 0 and 1 are copied over to the first two products and further products are placed at the reaction position. During the computation of the first product, `v1` is the vector from `mptr1` to `mptr2` and `dist`, or d , is computed as the length of this vector. If the distance equals zero, then this function pretends that the reactants were separated by σ_b along the x-axis because that makes as much sense as anything else. Next, the scaling factor `x` is computed. If `rparam` is positive, meaning fixed unbinding radius, then `v1` needs to be lengthened by $\sigma_u - d$. Dividing this by σ_u which is the cumulative length of the `prdpes` values, gives the scaling factor as $1 - d/\sigma_u$. Finally, this is divided by d because vector `v1` has length d . If `rparam` is negative, meaning new separation is equal to prior overlap, then $\sigma_b - d$ is the prior overlap, $\sigma_b + (\sigma_b - d)$ is the new separation, and `v1` needs to be lengthened by $\sigma_b + (\sigma_b - d) - d$. Simplifying and then dividing by σ_b , which is the cumulative length of the `prdpes` values, gives the scaling factor as $2 - 2d/\sigma_b$. As before, this is finally divided by length d .

Reaction logging is fairly simple but is described anyhow. Nothing happens if `rxn->logserno` is NULL. Otherwise, `dorxnlog` is set to 0 for not logging and to 1 for logging. Then, it changes to 2 as the logging is in progress, or to -1 if logging failed due to a bad file name.

`int zeroreact(simptr sim)`

Figures out how many molecules to create for each zeroth order reaction and then tells `doreact` to create them. It returns 0 for success or 1 if not enough molecules were allocated initially.

`int unireact(simptr sim)`

Identifies and performs all unimolecular reactions. Reactions that should occur are sent to `doreact` to process them. The function returns 0 for success or 1 if not enough molecules were allocated initially.

`int morebireact(simptr sim, rxnptr rxn, moleculeptr mptr1, moleculeptr mptr2, int ll1, int m1, int ll2, enum EventType et, double *vect)`

Given a probable reaction from `bireact`, this checks for compartment or surface reactions, orders the reactants, checks for reaction permission, moves a reactant in case of periodic boundaries, increments the appropriate event counter, and calls `doreact` to perform the reaction. The return value is 0 for success (which may include no reaction) and 1 for

5 C Code: structures and functions

failure. The `vect` input is only considered here, and must be non-NULL, if the event type is `ETrxn2wrap`; in this case, it is the vector from the current position of `mpt1` to the current position of `mpt2`.

`int bireact(simptr sim, int neigh)`

Identifies likely bimolecular reactions, sending ones that probably occur to `morebireact` for permission testing and reacting. `neigh` tells the routine whether to consider only reactions between neighboring boxes (`neigh=1`) or only reactions within a box (`neigh=0`). The former are relatively slow and so can be ignored for qualitative simulations by choosing a lower simulation accuracy value. In cases where walls are periodic, it is possible to have reactions over the system walls. The function returns 0 for success or 1 if not enough molecules were allocated initially.

5.5 Rules (functions in `smolrule.c`)

Rules were part of the reaction superstructure through version 2.46 but were then moved to their own superstructure in version 2.47. As with other portions of Smoldyn, rules are organized with a superstructure that contains information about all of the rules, and then with individual rule structures that contain information about the individual rules. The rules superstructure is only created if necessary. Each rule is a partially parsed set of instructions that tell Smoldyn what definitions to make whenever the rules are expanded, which may be before the simulation runs or on-the-fly during the simulation.

Reaction rules are typically only useful in combination with wildcard characters. See the Wildcards section of the Code Design chapter for the definitive description of the reaction rules.

5.5.1 enumerated types

The only enumerated type in the rules definitions is the `RuleType`:

```
#define RTMAX 4
enum RuleType {RTreaction, RTdifc, RTdispsize, RTcolor, RTnone};
```

The constant `RTMAX` is the number of enumerated elements. The elements are the different types of rules, where each rule is used to specify a certain type of definition, which may be to create reactions, set diffusion coefficients, etc.

5.5.2 rule structure

```
typedef struct rulestruct {
    struct rulesuperstruct *rules; // pointer to superstructure
    enum RuleType ruletype;        // type of rule
    char *rulename;                // pointer to name of rule
    char *rulepattern;             // pattern for the rule
    int *ruledetails;              // list of rule states and restrictions
    double rulerate;               // rate constant
    rxnptr rulerxn;                // template reaction for reaction rules
} *ruleptr;
```

The `rules` element points to the superstructure that owns this rule. The `ruletype` element tells what type of rule this is. The `rulename` element is the name of the rule which, for reaction rules, is the name of the reaction as entered by the user without any suffix. The rule structure does

5.5 Rules (functions in *smolrule.c*)

not own the `rulename` pointers but instead they are copied over from the rule superstructure. The `rulepattern` is the pattern for the rule, including for a reaction rule. For example, a pattern for a bimolecular reaction has a single-species pattern for the first reactant, which is a species name possibly with wildcard characters, a space, a second single-species pattern, a newline character, and then space-separated product species names. The `ruledetails` element contains an array of integers. For reaction rules, these are the species states cast as integers for the reactants and products in their sequence, then the compartment number if the reaction is restricted to a specific compartment, and then the surface number if the reaction is restricted to a single surface. The `rulerate` element is, for reaction rules, the reaction rate entered by the user. The `rulernxn` element contains a reaction template for reaction rules.

The following table shows the contents of the different portions of the data structure for different types of rules.

type	rate	detailsi	detailsf
RTreaction	rate constant	react. states, prod. states, compart., surf.	NULL
RTdifc	diff. coefficient	0: state	NULL
RTdifm	0	0: state	diff. matrix
RTdrift	0	0: state	drift vector
RTsurfdrift	0	0: state, 1: surface, 2: panel shape	drift vector
RTmollist	0	0: state, 1: list number	NULL
RTdispsize	display size	0: state	NULL
RTcolor	0	0: state	color vector
RTsurfaction	0	0: state, 1: surface, 2: face, 3: action	NULL
RTsurfrate	rate constant	0: state, 1: surface, 2: ms1, 3: ms2, 4: new species	NULL
RTsurfrateint	probability	0: state, 1: surface, 2: ms1, 3: ms2, 4: new species	NULL

5.5.3 rule superstructure

```
typedef struct rulesuperstruct {
    struct simstruct *sim;           // simulation structure
    int maxrule;                    // allocated size of rule list
    int nrule;                      // actual size of rule list
    char **rulename;                // list of rule names
    ruleptr *rule;                  // list of rules
    int ruleonthe-fly;              // for expanding rules on the fly
} *rulespptr;
```

The rule superstructure includes a list of rule names and a list of the actual rules, each with allocated size `maxrule` and actual size `nrule`. The `rulename` list contains strings with the rule names, which, for reactions, are simply the reaction names entered by the user, without any suffix. The `rule` list contains pointers to individual rules.

These lists are maintained by the `RuleAddRule` function. These rules do not need updating during the simulation. Instead, they are simply stored here and referred to when updating is required. The `ruleonthe-fly` element is initialized to 0. It is then set to 1 if on-the-fly rule generation is desired.

5.5.4 rule functions

enumerated types

5 C Code: structures and functions

enum RuleType rulestring2rt(const char *string)

Converts string to enumerated RuleType type. This reads full word inputs. Unrecognized inputs are returned as RTnone.

char *rulert2string(enum RuleType rt, char *string)

Converts enumerated RuleType type to string, which needs to be pre-allocated. Unrecognized rule types are converted to "none". The string is returned to simplify function cascading.

memory management

ruleptr rulealloc()

Allocates a new rule structure. All pointers are initialized to NULL, the ruletype is initialized to RTnone and the rulerate element is initialized to -1. The rulepattern and ruledetails are allocated in RuleAddRule.

void rulefree(ruleptr rule)

Frees a rule structure, including its contents.

rulespstr rulesalloc(rulespstr rules, int maxrule)

Allocates or expands the size of a rule superstructure. Send in the current rule superstructure in as rules if there is one, or NULL if there isn't one. Send in the desired value for the allocated number of rules in maxrule. This creates a new superstructure if there wasn't one initially, allocates and initializes the rule and rule name lists, copies over any prior data, frees any prior data, and returns the rule superstructure. Returns NULL if out of memory.

void rulesfree(rulespstr rules)

Frees a rule superstructure, including all of its contents.

data structure output

void ruleoutput(simpstr sim)

Outputs information about the rules, including the rule superstructure and all of the individual rules.

void writerules(simpstr sim, FILE *fpstr)

Not written yet. This will write the rules to a text file using Smoldyn input format.

int checkruleparams(simpstr sim, int *warnptr)

This checks to make sure that rule parameters are reasonable. There are no checks at all here yet.

structure set up

int RuleAddRule(simpstr sim, enum RuleType type, const char *rname, const char *pattern, const enum MolecState *rctstate, const enum MolecState *prdstate, double rate, const int *detailsi, const double *detailsf)

Adds a rule to the list of rules. Enter the rule type in type. For reaction rules, enter the rule name, which is the root of the reaction name and the entire reaction name if only one reaction results, in rname. Enter the rule pattern in pattern, the species or reactant state(s) in rctstate, any product states in prdstate, and the rate constant in rate. Enter additional integer details in detailsi and additional floating point details in detailsf.

5.5 Rules (functions in *smolrule.c*)

This function stores the rule in the rule superstructure. This function creates and/or expands the rule list as needed. It then creates a new rule by copying in the rule name, copying in the rule pattern, and storing the reactant state, product state, compartment index, and surface index in the rule details, in that order. Returns 0 for success or 1 for inability to allocate memory.

The exact inputs depend on the rule type as follows. Note that these are inputs to this function, and not the storage of details in the data structure. In the data structure, the first elements of the **detailsi** vector include the species states and then the last elements are the same as the ones given here.

type	rate	detailsi	detailsf
RTreaction	rate constant	0: compartment, 1: surface	NULL
RTdifc	diff. coefficient	NULL	NULL
RTdifm	0	NULL	diff. matrix
RTdrift	0	NULL	drift vector
RTsurfdrift	0	0: surface, 1: panel shape	drift vector
RTmollist	0	0: list number	NULL
RTdispsize	display size	NULL	NULL
RTcolor	0	NULL	color vector
RTsurfaction	0	0: surface, 1: face, 2: action	NULL
RTsurfrate	rate constant	0: surface, 1: ms1, 2: ms2, 3: new species	NULL
RTsurfrateint	probability	0: surface, 1: ms1, 2: ms2, 3: new species	NULL

core simulation functions

int RuleExpandRules(simp_{tr} sim, int iterations)

Expands the rules by **iterations** iterations. Also, **iterations** can be set to -1 for expanding all rules until everything is up to date, -2 to set up on-the-fly expansion but not to actually do it, or -3 to do on-the-fly expansion if necessary. This goes through the rules sequentially. If there are no rules, this returns error code -41, which isn't actually an error. For each rule, it gets the rule information, including when it was last updated. It then calls **RxnAddReactionPattern** to update the index variable for the rule pattern and add any new reactions to the simulation. Next, it calls **RxnSetValuePattern** to set the reaction rates for the new reactions. This part of the function still needs work. Finally, it sees if all of the rules are up to date, terminating if so.

Returns 0 for success or the following errors which arise from **RxnAddReactionPattern**: -1 for inability to allocate memory, -2 if no wildcards were entered and one or more of the match species is unknown, -3 if a substitute species is unknown and the substitute species had no wildcards in it, -4 if the match string included more words than allowed by this function (which is 4 currently), -5 if a trial match string was too long to fit in **STRCHAR** characters (even if this wasn't actually a match), -6 if species generation failed, -11 for inability to allocate memory, -12 for missing ' ' operand, -13 for missing & operand, -15 for mismatched braces, or -20 for a destination pattern that is incompatible with the matching pattern (i.e. it has to have either 1 destination or the same number of destination options as pattern options), or -30 for failure to add the reaction. Other return code are -40 for a bug in **RxnSetValuePattern** and, of actual use, -41 for expansion requested but there are no rules to expand.

5.6 Surfaces (functions in smolsurf.c)

Surfaces are organized with a surface superstructure that contains not much more than just a list of surfaces and their names. Each of these surfaces, defined with a surface structure, has various properties that apply to the whole surface, such as its color on the front and back faces, how it is drawn, and how it interacts with diffusing molecules. A surface structure also includes lists of panels that comprise the surface. These panels may be rectangular, triangular, spherical, or other shapes. A single surface can contain many panels of multiple shapes.

5.6.1 Surface geometry

The table below lists the types of panels and key aspects of how they are stored internally. Panel locations and sizes, plus some drawing information, are given with sets of **dim**-dimensional points, stored in the **point** element. There are **npts** points for a panel, listed below, where **npts** depends on both the panel shape and the system dimensionality. Additionally, each panel has a **dim**-dimensional **front** vector, which contains information about the direction that the panel faces. In some cases, such as for triangles, this is the normal vector to the surface and is redundant with the information in the points. In others, it contains additional information. For example, for spheres, only one element of **front** is used, and it is used to tell if the front of the panel is on the inside or outside of the sphere, which cannot be known from just the list of points. In the table below, **p** is used for point, and **f** is used for front.

Table: Properties of panels

1D	2D	3D
rectangles, ps = PSrect		
npts = 1	npts = 4	npts = 8
p[0][0] = location	p[0][0...1] = start	p[0...3][0...2]
	p[1][0...1] = end	= corners
	parallel to an axis	parallel to an axis
	front is on right	front has CCW winding
f[0] = ±1	f[0] = ±1	f[0] = ±1
(+ for facing +0)	(+ for facing +axis)	(+ for facing +axis)
f[1] = 0 (perp. axis)	f[1] = perp. axis (0, 1)	f[1] = perp. axis (0, 1, 2)
f[2] = undefined	f[2] = parallel axis	f[2] = axis parallel
	p[2] = normal for end 0	to edge from point 0 to point 1
	p[3] = normal for end 1	p[4] = normal for 0-1 edge
		p[5] = normal for 1-2 edge
		p[6] = normal for 2-3 edge
		p[7] = normal for 3-0 edge
triangles, ps = PStri		
npts = 1	npts = 4	npts = 6
p[0][0] = location	p[0][0...1] = start	p[0...2][0...2]
	p[1][0...1] = end	= corners
	front is on right	front has CCW winding
f[0] = ±1	f[0...1] = normal vect.	f[0...2] = normal vect.
(+1 for facing +0)	p[2] = normal for end 0	p[3] = normal for 0-1 edge
	p[3] = normal for end 1	p[4] = normal for 1-2 edge
		p[5] = normal for 2-0 edge

spheres, ps = PSsph		
npts = 2	npts = 2	npts = 2
p[0][0] = center	p[0][0...1] = center	p[0][0...2] = center
p[1][0] = radius	p[1][0] = radius	p[1][0] = radius
	p[1][1] = slices	p[1][1] = slices
		p[1][2] = stacks
f[0] = ± 1	f[0] = ± 1	f[0] = ± 1
(+ for front outside)	(+ for front outside)	(+ for front outside)
f[1...2] = undefined	f[1...2] = undefined	f[1...2] = undefined
cylinders, ps = PScyl		
	npts = 5	npts = 5
not allowed	p[0][0...1] = start center	p[0][0...2] = start center
	p[1][0...1] = stop center	p[1][0...2] = stop center
	p[2][0] = radius	p[2][0] = radius
		p[2][1] = slices
		p[2][2] = stacks
	f[0...1] = norm. right vect.	
	f[2] = ± 1	f[2] = ± 1
	(+ for front outside)	(+ for front outside)
		f[0...1] = undefined
	p[3] = normal for end 0	p[3] = normal for end 0
	p[4] = normal for end 1	p[4] = normal for end 1
hemispheres, ps = PShemi		
	npts = 3	npts = 3
not allowed	p[0][0...1] = center	p[0][0...2] = center
	p[1][0] = radius	p[1][0] = radius
	p[1][1] = slices	p[1][1] = slices
		p[1][2] = stacks
	p[2][0...1] = outward vect.	p[2][0...2] = outward vect.
	f[0] = ± 1	f[0] = ± 1
	(+ for front outside)	(+ for front outside)
	f[1...2] = undefined	f[1...2] = undefined
disks, ps = PSDisk		
	npts = 2	npts = 2
not allowed	p[0][0...1] = center	p[0][0...2] = center
	p[1][0] = radius	p[1][0] = radius
		p[1][1] = slices
	f[0...1] = normal vect.	f[0...2] = normal vect.
	f[2] = undefined	

To add a new panel shape, several things need to be done. Add the panel shape name to `PanelShape` and increment the `#define` constant `PSMAX`, which are defined in the `smoldyn.h` header file. Define the panel `point` and `front` values in the table above. Add the new panel shape to the following functions (and maybe others): `psstring2ps`, `ps2psstring`, `panelpoints`, `loadsurface` (panel input section), `surfaceoutput`, `panelside`, `surfacearea`, `panelrandpos`, `lineXpanel`, `fixpt2panel`, `surfacereflect`, `surfacejump`, and `panelinbox`. Also, add the panel shape to `RenderSurfaces` in `smoldyn.c`. Most of these are relatively easy, although some math likely needs to be done for a couple of them. Finally, check and document.

5.6.2 Molecule-surface interactions

Molecule-surface interactions arise when a molecule collides with a surface, or when a surface-bound molecule undergoes a spontaneous state change, such as desorption. Collisions can arise both for solution-phase molecules or surface-bound molecules; in the latter case, a molecule bound to surface A diffuses along that surface and then collides with surface B, which intersects surface A.

Surface interactions can be certain or probabilistic. The former interactions, which the user enters with the **action** configuration file statement, always happen immediately upon interaction. These certain interactions are: reflect, transmit, absorb, jump, and port. Probabilistic interactions occur with certain probabilities either upon interaction or at each time step. The user enters the rates of these interactions with the **rate** statement. This statement is also used for spontaneous transition rates of surface-bound molecules.

Internally, both the **action** and **rate** elements of the surface data structure refer to both collision interactions and spontaneous state changes of surface-bound molecules. The “face” index of these elements is either **PFfront** or **PFback** for collisions, or is **PFnone** for surface-bound state changes.

There are several ways of describing surface interactions. One can use a verb, such as reflect, transmit, or adsorb, or one can list the beginning molecule state, the surface interaction face, and the ending molecule state. Or, one can just list the beginning and ending molecules states, with pseudo-states, plus a third state when necessary (see the **surfsetrate** function and the **rate** configuration file statement). Smoldyn uses all of these methods which means that interconversions become necessary. The following table lists the states used in the action details data structure and their meanings. Conversions are given in later tables.

interaction class	forward states			action
	ms1	face1	ms2	
collision from solution state	soln	front	fsoln	reflect
	”	”	bsoln	transmit
	”	”	bound	adsorb
	”	back	fsoln	transmit
	”	”	bsoln	reflect
	”	”	bound	adsorb
impossible	”	none	any	
collision from bound state	bound	front	fsoln	reflect
	”	”	bsoln	transmit
	”	”	bound	hop
	”	”	bound’	hop
	”	back	fsoln	transmit
	”	”	bsoln	reflect
	”	”	bound	hop
	”	”	bound’	hop
action from bound state	”	none	fsoln	desorb
	”	”	bsoln	desorb
	”	”	bound	no
	”	”	bound’	flip

For the most part, surface-bound molecules cannot be absorbed, jumped, or ported, using the same surface. The exception is that if a surface-bound molecule in its **MSfront** or **MSback** state diffuses onto a new surface panel, and the new panel has jump behavior for its **MSsoln** or **MSbsoln**

states, then the molecule is jumped.

If a surface-bound molecule collides with another surface, then the result is typically dictated by the “action” statement and data structure contents, as described above. However, if the panel that it collides with has been declared to be a neighbor of the panel that the molecule is on, then the molecule does not behave according to the action given above. Instead, it can “hop”, meaning that it has a 50% chance of moving to the new panel and 50% chance of staying on its current panel, or it can “stay”, meaning that it stays on its current panel. These options are stored in the surface **neighhop** data structure element, where 1 indicates hopping and 0 indicates staying. Staying is the default. The hopping option is provided so that molecules can diffuse between neighboring surface panels, even if the panels don’t meet at their edges (e.g. a collection of spheres).

Note that absorption, jumping, and porting will have time-step dependent behaviors; from Andrews, *Phys. Biol.*, 2009, the absorption/ jumping/ porting coefficient is about $0.86s/\Delta t$, where s is the rms step length and Δt is the time step.

5.6.3 Surface data structures

```
#define PSMAX 6                                // maximum number of panel
      shapes
enum PanelFace {PFfront, PFback, PFnone, PFboth};
enum PanelShape {PSrect, PStri, PSsph, PScyl, PShemi, PSdisk, PSall,
      PSnone};
enum SrfAction {Sareflect, SAtrans, SAabsorb, SAjump, SAport, SAMult, SAno
      , SAnone, SAadsorb, SArevdes, SAirrevdes, SAflip};
enum DrawMode {DMno=0, DMvert=1, DMedge=2, DMve=3, DMface=4, DMvf=5, DMef
      =6, DMvef=7, DMnone};
enum SMLflag {SMLno=0, SMLdiffuse=1, SMLreact=2, SMLsrfbound=4};
```

Panel faces can be front or back, and there are also enumerations for both and none. For version 2.19, I changed the enumeration sequence to put **PFnone** before **PFboth**. This sequence is important because the surface **action** and **rate** elements are allocated for the first three enumerated panel faces, but not more.

Panel shapes are enumerated with **PanelShape**, of which there are **PSMAX** shapes, plus enumerations for all and none, which can be useful as function arguments.

Not all surface actions apply to all circumstances. For example, reflect, transmit, absorb, and jump only apply to collisions between diffusing molecules and surfaces (plus, diffusing molecules in front and back states can jump as well, although the jump action is not assigned to these states). “no” applies to surface-bound molecules, meaning that they are static and don’t change over time; in contrast, those that might change are labeled as **SAMult**, meaning that there are multiple possible outcomes at each time step. “none” does not imply “no action”, but means instead “none of the other options”. “port” is an action for exporting molecules to other simulators, such as MOOSE. The actions **SAadsorb**, **SArevdes** (reversible desorption), **SAirrevdes** (irreversible desorption), and **SAflip** (change of surface-bound state), are returned by the **surfaction** function, but are not options that the user can choose. This is because they are, in a sense, sub-actions within the **SAMult** option and they cannot be chosen for exclusive use.

DrawMode lists the drawing modes for polygons or other surfaces. **vert** or **v** are for vertices, **edge** or **e** is for edges, and **face** or **f** are for faces. Where multiple options are listed, Smoldyn is supposed to draw multiple methods simultaneously, although I don’t believe that it supports this at present. Numbers are explicitly specified in **DrawMode** because there are distinct bits for vertex, edge, or face, which allows them to be extracted from the code using bitwise logic operations.

5 C Code: structures and functions

SMLflag (stands for surface molecule list) lists binary flags for the molecule lists for which surface checking is required. For example, if the **SMLdiffuse** flag is set for some molecule list, then surfaces need to be checked for all molecules in that list after diffusion occurred. Similarly, those with **SMLreact** set need to be checked after reactions occur, because they might have been placed across a surface during product placement. Those with **SMLsrfbound** set are surface bound molecules that might be able to desorb or flip orientation.

```
typedef struct surfactionstruct {
    int *srfnewspec;           // surface convert mol. species [ms]
    double *srfrate;           // surface action rate [ms]
    double *srfprob;           // surface action probability [ms]
    double *srfcumprob;        // surface cumulative probability [ms]
    int *srfdatasrc;           // surface data source [ms]
    double *srfrevprob;        // probability of reverse action [ms]
} *surfactionptr;
```

The surface action structure collects together details for surface actions. It is typically only allocated if a surface action for a specific molecular species, molecular state, and surface interaction face is of type **SAMult**, meaning that multiple possible outcomes are possible. This arises for the “rate” Smoldyn statement, and not for the “action” statement. In that case, this structure records the rate at which conversion can take place to each of the possible output states, the probability of each of these transitions for each time step, the cumulative probabilities (used for efficient simulation), the new species that should be created upon transition (usually the same as the current species, but not necessarily), and the source of the interaction rate data. Each of the vectors in the surface action structure is allocated to size **MSMAX1**, meaning that **MSbsoln** is an allowed outcome state. The **srfdatasrc** value is initialized to 0, is set to 1 if the user entered an interaction rate, is set to 2 if the user entered an interaction probability, or is set to 3 if the user entered a certain action (see next paragraph) with the “action” statement. Note that each of these vectors refers to the new state so, for example, **srfnewspec**, lists the new species that should happen upon changing to each of the possible new states.

This data structure is also allocated for simple surface actions, arising for the Smoldyn “action” statement if there is a new species. In this case, the entire data structure is ignored, except only for one element of the **srfnewspec** vector. The element is not the new state, as it is normally, but the originating state.

```
typedef struct panelstruct {
    char *pname;               // panel name (reference, not owned)
    enum PanelShape ps;        // panel shape
    struct surfacestruct *srf;  // surface that owns this panel
    int npts;                   // number of defining points
    double **point;             // defining points, [number][d]
    double **oldpoint;          // prior defining points, [number][d]
    double front[DIMMAX];       // front parameters, which depend on the
                                // shape
    double oldfront[DIMMAX];     // prior front parameters, which depend
                                // on the shape
    struct panelstruct *jumpp[2]; // panel to jump to, if appropriate [face]
    enum PanelFace jumpf[2];     // face to jump to, if appropriate [face]
    int maxneigh;               // maximum number of neighbor panels
    int nneigh;                 // number of neighbor panels
    struct panelstruct **neigh;  // list of neighbor panels [p]
    double *emitterabsorb[2];    // absorption for emitters [face][i]
} *panelptr;
```

5.6 Surfaces (functions in *smolsurf.c*)

pname is a pointer to the panel name, which is contained in the surface structure; this memory is owned by the surface, not by the panel. **ps** is the panel shape. **srf** is a pointer to the surface that owns this panel; it would be called a **surfaceptr**, except that a **surfaceptr** isn't declared until later. **npts** is the number of **dim**-dimensional points that are allocated for this panel. **point** and **front** have meanings that depend on the panel shape and on the dimensionality, described in the preceding table. **jumpp** and **jumpf** are used for periodic boundary conditions and jumping molecules; these are the panel and face that a molecule will be sent to if it collides with the front or back face of this panel. **maxneigh** is the number of neighbor references that are allocated, **nneigh** is the number of neighboring panels and **neigh** is the list of neighboring panels. These neighboring panels may be within the same surface or on a different surface. These are used for diffusion of surface-bound molecules. Neighbors are only allocated as necessary (by the **surfsetneighbors** function). **emitterabsorb[face][i]** is the panel absorption probability for face **face** and species **i** to account for emitters (see the user's manual). It is only allocated if necessary.

The **oldpoint** and **oldfront** elements are for moving surfaces. They are used to show where a surface came from. They are used only when a surface is about to be moved and during the moving process.

```
typedef struct surfacestruct {
    char *sname;                // surface name (reference, not owned)
    struct surfacesuperstruct *srfss; // owning surface superstructure
    int selfindex;              // index of self
    enum SrfAction ***action;    // action for molecules [i][ms][face]
    surfactionptr ***actdetails; // action details [i][ms][face]
    int neighhop;               // whether molecules hop between neighbors
    double fcolor[4];           // RGBA color vector for front
    double bcolor[4];           // RGBA color vector for back
    double edgepts;             // thickness of edge for drawing
    unsigned int edgestipple[2]; // edge stippling [factor, pattern]
    enum DrawMode fdrawmode;    // polygon drawing mode for front
    enum DrawMode bdrawmode;    // polygon drawing mode for back
    double fshiny;              // front shininess
    double bshiny;              // back shininess
    int maxpanel[PSMAX];        // allocated number of panels [ps]
    int npanel[PSMAX];          // actual number of panels [ps]
    char **pname[PSMAX];        // names of panels [ps][p]
    panelptr *panels[PSMAX];    // list of panels [ps][p]
    struct portstruct *port[2]; // port, if any, for each face [face]
    double totarea;             // total surface area
    int totpanel;               // total number of panels
    double *areatable;          // cumulative panel areas [pindex]
    panelptr *paneltable;       // sequential list of panels [pindex]
    int *maxemitter[2];         // maximum number of emitters [face][i]
    int *nemitter[2];           // number of emitters [face][i]
    double **emitteramount[2];  // emitter amounts [face][i][emit]
    double ***emitterpos[2];    // emitter positions [face][i][emit][d]
    int *maxmol;                // allocated size of live lists [ll]
    int *nmol;                  // number of molecules in live lists [ll]
    moleculeptr **mol;          // live molecules on the surface [ll][m]
} *surfaceptr;
```

selfindex is the index of the surface in the surface structure; this is useful in case the surface is sent to a function using a pointer and the function needs to know which surface it is. **action** lists the the actions that happen to molecules for each state. It is allocated to size **[maxspecies][MSMAX][3]**,

5 C Code: structures and functions

which means that it does not allow the `MSbsoln` state and also that it allows `PFfront`, `PFback`, or `PFnone` panel faces. An action may be, for example, `SArereflect` or `SAno`, if these actions always happen and is `SAmult` if there are multiple possible actions that could happen at any particular time step. This vector has `maxident` elements times `MSMAX` elements to account for each molecule species and then each state. Only for those that have type `SAmult`, is the data pointed to by `actdetails` relevant. As described above, these details list the action rates, probabilities, and other values for transitions between states for each molecular species. `neighhop` tells whether surface-bound molecules should be allowed to hop to neighboring panels upon collision or whether they should stay on the originating panel.

`fcolor` and `bcolor` are the colors of the front and back of the surface in the order: red, green, blue, alpha; each has a value between 0 and 1. `edgepts` is the thickness of edges in points for drawing, which applies to all drawing situations except for 3D and when the surface faces are rendered. `fdrawmode` and `bdrawmode` describe how the surface front and back should be drawn. Not all options apply to 1D and 2D simulations. `fshiny` and `bshiny` are shininess values for OpenGL surface rendering. `maxpanel` and `npanel` are the number of panels that are allocated or used, respectively, for each of the panel shapes. `panelname` lists a name for each panel. `panels` are lists of pointers to the panels for the possible shapes. Note that every panel within a surface has the same drawing scheme and the same interaction with molecules. `port` points to the port structure that applies to each face, if any.

Considering all surface panels, the total surface area is `totarea` and there are `totpanel` panels. These panels are listed sequentially in the list `paneltable`, and their cumulative areas are listed in `areatable`. Thus, for example, `areatable[0]` is the area of the first panel and `areatable[totpanel-1]` is equal to `totarea`. These four elements are set up in `surfacesupdate`.

Although it is a rather specialized function, surfaces can be configured to absorb molecules with coefficients that yield concentrations that are the same as those for unbounded systems (see the user's manual). This configuration relies on the definitions of point "emitters". For face `face` and species `i`, `maxemitter[face][i]` emitters are allocated, of which `nemitter[face][i]` are actually used. These emitters have amount `emitteramount[face][i][emit]` (`emit` is the emitter index) and are at locations `emitterpos[face][i][emit][d]`, where `d` is the index for the dimensionality.

For Smoldyn version 2.50, I added molecule lists in surfaces, so surfaces know what molecules are adsorbed to them. These lists are in the `maxmol`, `nmol`, and `mol` lists. Molecules are assigned to these lists in the `assignmolecs` function in `smolboxes.c`.

```
typedef struct surfacesuperstruct {
    enum StructCond condition;    // structure condition
    struct simstruct *sim;        // simulation structure
    int maxspecies;               // maximum number of molecular species
    int maxsrf;                   // maximum number of surfaces
    int nsrf;                     // number of surfaces
    double epsilon;               // max deviation of surface-point from
        surface
    double margin;                // panel margin away from edge
    double neighdist;             // neighbor distance value
    char **snames;                // surface names [s]
    surfaceptr *srflist;          // list of surfaces [s]
    int maxmollist;               // number of molecule lists allocated
    int nmollist;                 // number of molecule lists used
    enum SMLflag *srfmollist;     // flags for molecule lists to check [ll]
} *surfacessptr;
```

This is the superstructure for surfaces. `condition` is the current condition of the superstructure

5.6 Surfaces (functions in *smolsurf.c*)

and **sim** is a pointer to the simulation structure that owns this superstructure. **maxspecies** is a copy of **maxspecies** from the molecule superstructure, and is the allocated size of the surface action, rate, and probability elements. **maxsrf** and **nsrf** are the number of surfaces that are allocated and defined, respectively. **epsilon** is a distance value that is used when fixing molecules to panels; if a molecule is already within **epsilon** of a panel and on the correct side, no additional moving is done. **margin** is the distance inside the edge of a panel to which molecules are moved if they need to be moved onto panels. **neighdist** is used for the diffusion of molecules on a surface; if the point where a molecule diffuses off of one panel is within **neighdist**^{1/2} of the closest point on another panel, then the molecule can move to the neighboring panel. **snames** is a list of names for the surfaces. **srflist** is the list of pointers to surfaces. **srfmollist** is a list of flags for which molecule lists need to be checked for surface interactions; **maxmollist** and **nmollist** are local copies of **sim->mols->maxlist** and **sim->mols->nlist**, and are used to read the **srfmollist** element. The **SMLflag** enumerated values are or-ed together in these elements.

It was surprisingly difficult to get surfaces to work well enough that diffusing molecules did not leak through reflective panels. Because of that, the code is written unusually carefully, and in ways that are not necessarily obvious, so be careful when modifying it. For example, round-off error differences between two different but mathematically identical ways of calculating a molecule distance from a surface can easily place the molecule on the wrong side of a surface panel.

If a molecule is exactly at a panel, it is considered to be at the back side of the panel. Initially, I defined direct collisions as collisions in which the straight line between two points crosses a surface, whereas an indirect collision is one in which the straight line does not cross a surface but it was determined with a random number that the Brownian motion trajectory did contact the surface. Indirect collisions proved to slow down the program significantly, greatly complicate the code development, and provided minimal accuracy improvements, so I got rid of them. Now, only direct collisions are detected and dealt with.

5.6.4 Function interdependence

Following is a partial listing of what functions call what other functions. This is incomplete but could be completed relatively easily if necessary.

```
surfreadstring
  name
    surfaddsurface
      surfenablesurfaces if needed
        surfacessalloc
          surfacealloc
            emittersalloc
  action
    surfsetaction
      surfsetcondition to SCparams
  rate or rate_internal
    surfsetrate
      surfaceactionalloc if needed
      surfsetcondition to SCparams
  color
    surfsetcolor
```

5 C Code: structures and functions

- thickness
 - surfsetedgepts
- stipple
 - surfsetstipple
- polygon
 - surfsetdrawmode
- shininess
 - surfsetshiny
- max_panels (deprecated function)
 - surfsetmaxpanel
 - panelsalloc
- panel
 - surfaddpanel
 - panelsalloc if needed
 - emittersalloc
 - surfsetcondition to SClists
 - boxsetcondition to SCparams
- jump
 - surfsetjumppanel
- neighbors
 - surfsetneighbors, allocates space as needed
- unbounded_emitter
 - surfaddemitter
 - emittersalloc if needed
 - surfsetcondition to SCparams

surfsetepsilon

- surfenablesurfaces if needed, see above

surfsetmargin

- surfenablesurfaces if needed, see above

surfsetneighdist

- surfenablesurfaces if needed, see above

5.6.5 Surface functions

enumerated types

enum PanelFace surfstring2face(char *string)

Converts panel face **string** to an enumerated panel face type. Input strings are “front” for the front, “back” for the back, or “all” or “both” for both sides. Also, initial portions of these strings, such as “f” or “fro” are sufficient. Other inputs result in **PNone**.

char *surface2string(enum PanelFace face, char *string)

Converts enumerated panel **face** to a string, in **string**, which must be pre-allocated. Output strings are “front”, “back”, “both”, or “none”. **string** is returned to allow for function nesting.

enum SrfAction surfstring2act(char *string)
 Converts action *string* to an enumerated action type. Input strings are the same as the *SrfAction* strings. Initial portions of strings are sufficient. Unknown strings result in *SAnone*. This cannot return the actions *SAadsorb*, *SAirrevdes*, *SArevdes*, or *SAflip*, because the user cannot enter them. They are "sub-actions" of *SAmult*.

char *surfact2string(enum SrfAction act, char *string)
 Converts enumerated surface action *act* to a string, in *string*, which must be pre-allocated. Output strings are "reflect", "transmit", etc. *string* is returned to allow for function nesting.

enum PanelShape surfstring2ps(char *string)
 Converts panel shape *string* to an enumerated panel shape type. Input strings are the same as *PanelShape* strings.

char *surfps2string(enum PanelShape ps, char *string)
 Converts enumerated panel shape to a string. Output strings are abbreviated shape names, such as "rect" to designate a rectangle. Also, *PSall* and *PSnone* result in the strings "all" and "none". *string* is returned to allow for function nesting.

enum DrawMode surfstring2dm(char *string)
 Converts drawing mode *string* to an enumerated drawing mode type. Input strings are *DrawMode* names. Unrecognized input results in *DMnone*.

char *surfdm2string(enum DrawMode dm, char *string)
 Converts enumerated drawing mode to a string. Output strings are abbreviated drawing mode names. *string* is returned to allow for function nesting.

low level utilities

int readsurfacename(simptr sim, char *str, enum PanelShape *psptr, int *pptr)
 Reads the first word of string *str* to parse the surface name and an optional panel name, which are entered in the format *surface:panel*. Returns the surface index directly and, if the pointers are not NULL, returns the panel shape in *psptr* and the panel index in *pptr*. Returns the surface index, or -1 if *str* is missing, -2 if no surfaces have been defined in the current simulation, -3 if the name string cannot be read, -4 if the surface name is unknown, or -5 if the surface is "all". In *psptr* and *pptr* are returned, respectively: *PSnone* and -1 if the panel is not given, *PSnone* and -2 if the surface is "all" and the panel is something else (this is an error), *PSnone* and -3 if the panel name is unknown, *PSall* and -5 if the panel is "all", and otherwise the panel shape and panel number. These outputs were changed 4/24/12, while developing Smoldyn 2.27.

panelptr readpanelname(simptr sim, surfaceptr srf, const char *str)
 Reads the name of a panel in *str*. If this name already includes a surface, using the format *surface : panel*, then *srf* is ignored. Otherwise, *srf* is assumed to be the surface that owns the panel. Returns a pointer to the panel, or NULL if the panel wasn't found.

int srfoincidentpanels(panelptr pnl1, panelptr pnl2)
 Tests two panels to see if they are essentially identical to each other, meaning defined at the exact same position. Returns 0 if not, 1 if they are completely identical, 2 if their defining points are sequentially rotated (e.g. *pnl1* has triangle points *a*, *b*, and *c* and *pnl2* has triangle points *b*, *c*, *a*), and 3 if the two panels have front and back flipped from each other. Also returns -1 for an error (e.g. the panels have different shapes).

5 C Code: structures and functions

int panelpoints(enum PanelShape ps, int dim)

Returns the number of point elements that need to be allocated for a panel of shape **ps** and for system dimensionality **dim**. These numbers are the same as those listed in the table above. 0 is returned for inputs that don't make sense (e.g. **PSall**) or for shapes that are not permitted in the requested dimension.

int surfpanelparams(enum PanelShape ps, int dim)

Returns the number of numerical parameters that the user needs to enter to define a panel of shape **ps** and in a **dim** dimensional system. 0 is returned for inputs that don't make sense or for shapes that are not permitted in the requested dimension.

void panelmiddle(panelptr pnl, double *middle, int dim, int onpanel)

Returns the middle of panel **pnl** in the **dim**-dimensional vector **middle**; **dim** is the system dimensionality. For spheres, hemispheres, and cylinders, the middle point is the actual center location if **onpanel** is 0, which is enclosed by the panel but not on it; for these, set **onpanel** to 1 for **middle** to be returned as a point on the panel, although it will no longer be in the middle. **onpanel** is ignored for rectangles, triangles, and disks. If **onpanel** is 1: for spheres, the returned point is directly to the positive *x* direction from the sphere center; for cylinders, the returned point is as close as possible to the center point; and for hemispheres, the returned point is the center of the on-panel locations.

double panelarea(panelptr pnl, int dim)

Returns the area of panel **pnl**; **dim** is the system dimensionality, as always.

double surfacearea(surfaceptr srf, int dim, int *totpanelptr)

Returns the total area of surface **srf**; **dim** is the system dimensionality. If **totpanelptr** is not NULL, it is returned with the total number of panels in the surface. This function calculates the area, rather than just returning the value that is in the **totarea** surface element.

double surfacearea2(simpstr sim, int surface, enum PanelShape ps, char *pname, int *totpanelptr)

Returns the area of one or more panels. For the area of a single panel, the inputs are the surface number, the panel shape, and the panel name. The area is returned (number of points for 1-D, line length for 2-D, and area for 3-D). If **totpanelptr** is sent in as not NULL, it will point to the integer 1 on return. For multiple panels, set any or all of the inputs to "all" using: a negative number for surface, **PSall** for **ps**, and/or "all" for **pname**; **totpanelptr** will point to the number of panels included in the sum. For example, if surface is a positive number, **ps** is **PSall**, and panel is "all", then the total area of all panels of the specified surface is found. Or, if surface is negative, **ps** is **PSall** and **pname** is "endcap" then the area is found for all panels named "endcap", regardless of their surface or shape. If no panels match the input description, 0 is returned and **totpanelptr** is left pointing to a 0.

void panelrandpos(panelptr pnl, double *pos, int dim)

Returns a random position, in **pos**, on the surface of panel **pnl**, in a **dim** dimensional system. The result might be on either side of the panel.

panelptr surfrandpos(surfaceptr srf, double *pos, int dim)

Returns a random position, in **pos**, on the surface **srf**, in a **dim** dimensional system. The result might be on either side of the surface. The return value is a pointer to the panel that the point is in, or NULL if the surface has no panels.

5.6 Surfaces (functions in *smolsurf.c*)

`int issurfprod(simptr sim, int i, enum MolecState ms)`

Determines if molecule identity `i` and state `ms` is the product of a surface action, accounting for all surfaces. Returns 1 if so and 0 if not. `ms` can be `MSbsoln`. This should work after surfaces have been loaded and either before or after they have been set up. This does not return 1 if molecule `i` and state `ms` only participates in surface interactions, but is not produced by them. For example, if a molecule type simply reflects off of a surface with no species or state change, then this doesn't count. On the other hand, if a molecule type is produced when a molecule of a different state adsorbs to a surface, then this does count.

`int srfsamestate(enum MolecState ms1, enum PanelFace face1, enum MolecState ms2, enum MolecState *ms3ptr)`

Determines if a molecule in state `ms2` is in the same state as it was in `ms1`, and returns 1 if so and 0 if not. Also, if `ms3ptr` is not NULL, this returns the state that is the same as the `ms1` and `face1` information in the value pointed to by `ms3ptr`. This returns values from the following table.

interaction class	forward states			action	return values	
	ms1	face1	ms2		function	*ms3ptr
collision from solution state	soln	front	fsoln	reflect	1	fsoln
	"	"	bsoln	transmit	0	fsoln
	"	"	bound	adsorb	0	fsoln
	"	back	fsoln	transmit	0	bsoln
	"	"	bsoln	reflect	1	bsoln
	"	"	bound	adsorb	0	bsoln
impossible	"	none	any		0	none
collision from bound state	bound	front	fsoln	reflect	1	fsoln
	"	"	bsoln	transmit	0	fsoln
	"	"	bound'	hop	0	fsoln
	"	back	fsoln	transmit	0	bsoln
	"	"	bsoln	reflect	1	bsoln
	"	"	bound'	hop	0	bsoln
action from bound state	"	none	fsoln	desorb	0	bound
	"	"	bsoln	desorb	0	bound
	"	"	bound	no	1	bound
	"	"	bound'	flip	0	bound

`void srfreverseaction(enum MolecState ms1, enum PanelFace face1, enum MolecState ms2, enum MolecState *ms3ptr, enum PanelFace *face2ptr, enum MolecState *ms4ptr, enum SrfAction *actionptr)`

This function simply takes in a surface interaction and returns what the reverse interaction would be. It does not consider the specifics of individual surfaces, any interaction rates, or any other details. Instead, it simply inverts the table of surface interactions that is presented above. More specifically, given that some molecule starts with state `ms1`, interacts with `face1` of a surface, and then ends in state `ms2`, this finds the reverse surface action. In this reverse action, the molecule starts in state `ms3`, interacts with `face2`, and ends in state `ms4`. These latter parameters are pointed to by `ms3ptr`, `face2ptr`, and `ms4ptr`, respectively. Additionally, the type of action is returned in `actionptr`.

In concept, the reverse action is that `ms3` should equal `ms2` and `ms4` should equal `ms1`, although

5 C Code: structures and functions

it's rarely this simple. The reason is that the starting states cannot include **MSbsoln**, whereas the end states can include it, and also the actions have to allow for surface-bound molecules to interact with other surfaces that they cross. The following table shows the input and output values.

interaction class	forward states			action	reverse states		
	ms1	face1	ms2		ms3	face2	ms4
collision from solution state	soln	front	fsoln	reflect	soln	front	fsoln
	"	"	bsoln	transmit	soln	back	fsoln
	"	"	bound	adsorb	bound	none	fsoln
	"	back	fsoln	transmit	soln	front	bsoln
	"	"	bsoln	reflect	soln	back	bsoln
	"	"	bound	adsorb	bound	none	bsoln
impossible	"	none	any	none	none	none	none
collision from bound state	bound	front	fsoln	reflect	bound	front	fsoln
	"	"	bsoln	transmit	bound	back	fsoln
	"	"	bound'	no (hop)	<i>bound'</i>	<i>both</i>	<i>bound</i>
	"	back	fsoln	transmit	bound	front	bsoln
	"	"	bsoln	reflect	bound	back	bsoln
	"	"	bound'	no (hop)	<i>bound'</i>	<i>both</i>	<i>bound</i>
action from bound state	"	none	fsoln	rev. desorb	soln	front	bound
	"	"	bsoln	rev. desorb	soln	back	bound
	"	"	bound'	flip	bound'	none	bound

The italicized rows for the “reverse states” columns indicate that a bound-state molecule collided with a new surface and then hopped to this new surface. In this case, it's easy to know the beginning and ending states, but it is impossible to know the face of the reverse action. Also, the kinetics of the reverse process cannot be ascertained without knowing which surface the molecule was initially bound to. For this reason, a **face2** return value of **PFboth** indicates that the molecule hopped from one surface to another, which means that the interaction face of the reverse action cannot be determined without further information. This might be interpreted as an error condition. Finally, if **ms1** equals **MSsoln** and **face1** equals **PFnone**, this is an error; in this case, the function returns “none” for all variables.

```
void srfrystate2index(enum MolecState ms, enum MolecState ms1, enum MolecState
ms2, enum MolecState *ms3ptr, enum PanelFace *faceptr, enum MolecState
*ms4ptr)
```

Converts between the format that the **rate** configuration file statement and the **surfsetrate** function input, called tristate format, and the index format that the action details data structure uses. This function converts according to the following table. Rows that are listed in italics may be forbidden as input combinations, but still make logical sense and so they are converted here.

interaction class	tristate format			action	index format		
	ms	ms1	ms2		ms3	face	ms4
collision from	soln/none	soln	soln	<i>reflect</i>	soln	front	fsoln
	"	"	bsoln	transmit	soln	front	bsoln
	"	"	bound	adsorb	soln	front	bound

5.6 Surfaces (functions in *smolsurf.c*)

solution state	"	bsoln	soln	transmit	soln	back	fsoln
	"	"	bsoln	<i>reflect</i>	soln	back	bsoln
	"	"	bound	adsorb	soln	back	bound
action from bound state	"	bound	soln	desorb	bound	none	fsoln
	"	"	bsoln	desorb	bound	none	bsoln
	"	"	bound	<i>no change</i>	bound	none	bound
	"	"	bound'	flip	bound	none	bound'
collision from bound state	bound	soln	soln	<i>reflect</i>	bound	front	fsoln
	"	"	bsoln	transmit	bound	front	bsoln
	"	"	bound	hop	bound	front	bound
	"	"	bound'	hop	bound	front	bound'
	"	bsoln	soln	transmit	bound	back	fsoln
	"	"	bsoln	<i>reflect</i>	bound	back	bsoln
	"	"	bound	hop	bound	back	bound
	"	"	bound'	hop	bound	back	bound'
action from bound state	"	bound	soln	desorb	bound	none	fsoln
	"	"	bsoln	desorb	bound	none	bsoln
	"	"	bound	<i>no change</i>	bound	none	bound
	"	"	bound'	flip	bound	none	bound'
impossible	"	bound'	any	<i>nonsense</i>	none	none	none

```
void srfindex2tristate(enum MolecState ms3, enum PanelFace face, enum MolecState
ms4, enum MolecState *msptr, enum MolecState *mslptr, enum MolecState
*ms2ptr)
```

Inverse function as `srfrtristate2index`. This function uses the same table as shown for `srfrtristate2index`, but in reverse. There are two ways that actions from bound states can be described using the tristate format, of which the former uses the default value of `ms` where it equals `MSsoln` and the latter uses the better value, which equals `ms1`. This function inverts to the better tristate values.

```
int srffcompareaction(enum SrfAction act1, surfactionptr details1, enum SrfAction
act2, surfactionptr details2)
```

Compares the activity levels of `act1` and `act2`, returning -1 if `act1` is "more active", 1 if `act2` is more active, and 0 if they are the same. If they have exactly the same activity levels, then this returns 0. If both `act1` and `act2` equal `SAmult`, then this turns to the details of these actions, entered in `details1` and `details2`, to determine which is more active. Here, the first difference that is encountered determines which is more active. For the actions, the "activity" of an action is ordered as: `SAtans` < `SAmult` < `SArelect` < `SAjump` < `SAabsorb` < `SAport`.

memory management

```
surfaceactionptr surfaceactionalloc(int species)
```

Allocates a surface action structure for storing action details, and returns a pointer to it, or NULL if memory could not be allocated. Initializes most values to 0 or equivalent. Initializes `srffnewspec` to `species`.

```
void surfaceactionfree(surfaceactionptr actdetails)
```

Frees a surface action structure and its data.

5 C Code: structures and functions

int panelsalloc(surfaceptr srf, int dim, int maxpanel, int maxspecies, enum PanelShape ps)

Allocates **maxpanels** of shape **ps** for the surface **srf**; **srf** cannot be NULL but must be a surface. The **srf** element of the panels are set to **srf**. In **srf**, the correct **maxpanel** entry is set to **maxpanel**, the **npanel** entry is unchanged, and the proper list of panels are allocated and cleared. Also, panel names are created, each of which is set to the panel number, as a default. All points are set to all zeros. The function returns 1 for success and 0 for failure to allocate memory; if it fails, it does not do a good job of freeing working memory. The default jump destination is to the opposite side of the same panel.

This function may be called multiple times. It should be called when more panels of shape **ps** are needed. It *should not* be called if the only change is a larger value of **maxspecies**; in this case, a simple call to **surfacealloc** will take care of all required updates. This function calls **emittersalloc** to take care of the **emitterabsorb** panel element.

void panelfree(panelptr pnl)

Frees a single panel and all of its substructures (but not **srf**, because that's a reference and is not owned by the panel). This is called by **surfacefree** and so should not need to be called externally.

int emittersalloc(surfaceptr srf, enum PanelFace face, int oldmaxspecies, int maxspecies)

Allocates basic space for emitters (used for concentrations that match those for unbounded diffusion). This allocates the **srf->maxemitter** and **nemitter** arrays, as well as the first levels of the **emitteramount**, and **emitterpos** arrays. It also allocates all of the panel **emitterabsorb** arrays. Returns 0 for success or 1 for inability to allocate memory; in the latter case, this will create memory leaks.

There are three reasons to call this function: to set up the basics of emitters (send in **oldmaxspecies** as 0), to address emitter issues for new panels that had not been declared when emitters were set up originally (send in **oldmaxspecies** equal to **maxspecies**), or to allow for a larger **maxspecies** value (send in **oldmaxspecies** as whatever the prior **maxspecies** value was).

int surfexpandmollist(surfaceptr srf, int newmax, int ll) Allocates the surface molecule lists and expands them as needed. Returns 0 for success or 1 for failure to allocate memory. Enter **newmax** as the new maximum size for list number **ll**. Enter **ll** as -1 to set the number of molecule lists to **newmax** rather than to expand a single list.

surfaceptr surfacealloc(surfaceptr srf, int oldmaxspecies, int maxspecies, int dim)

Allocates a surface structure, and sets all elements to initial values. **maxspecies** is the maximum number of molecular species, which is used for allocating **action** and **actdetails**, as well as some emitter things. Colors are set to all 0's (black), but with alpha values of 1 (opaque); polygon modes are set to face if **dim** is 3, and to edge otherwise; **edgepoints** is set to 1; **action** is set to for transmitting (**SAttrans**) for the solution elements and to "no" (**SAno**) for the surface-bound elements. Panels and emitters are allocated here, although they are updated if necessary. This is called by **surfaccessalloc** and so should not need to be called externally.

This function can be called multiple times. The two times when it might need to be called are when the surface does not exist and needs to be allocated, in which case send in **srf** as NULL

5.6 Surfaces (functions in *smolsurf.c*)

and `oldmaxspecies` as 0, or when the surface does exist and `maxspecies` is being updated. In this case, send in the existing structure in `srf` and send in the prior number of maximum species in `oldmaxspecies`.

If this function is unable to allocate adequate memory, it returns `NULL`. If this is the case, it does not change any pre-existing `srf` data structure. In this case (which generally should not arise), it does not fully free the memory allocated here, leading to memory leaks.

`void surfacefree(surfaceptr srf, int maxspecies)`

Frees a surface, including all substructures and panels in it. This is called by `surfacessfree` and so should not need to be called externally. `maxspecies` is the number of molecule identities that the system was allocated with.

`surfacessptr surfacessalloc(surfacessptr srfss, int maxsurface, int maxspecies, int dim)`

Allocates a surface superstructure for `maxsurface` surfaces, as well as all of the surfaces. Each surface name is allocated to an empty string of `STRCHAR` (256) characters. Each surface is allocated for `maxspecies` species (a value of 0 is allowed). This function may be called more than once. On the first call, send in `srfss` as `NULL`, and the surface superstructure pointer will be returned; if it fails to allocate memory, it will return `NULL`. On subsequent calls, send in the existing surface superstructure pointer in `srfss` and the superstructure will be expanded as needed for the new larger `maxsurface` and/or `maxspecies` values (they may not be shrunk). In this case, the function will return the same pointer that was sent in, or `NULL` if it could not allocate memory; in the latter case, the original superstructure is unchanged. If this function is unable to allocate adequate memory (which generally should not arise), it does not fully free the memory allocated here, leading to memory leaks.

`void surfacessfree(surfacessptr srfss)`

Frees a surface superstructure pointed to by `srfss`, and all contents in it, including all of the surfaces and all of their panels.

data structure output

`void surfaceoutput(simptr sim)`

Prints out information about all surfaces, including the surface superstructure, each surface, and panels in the surface.

`void writesurfaces(simptr sim, FILE *fptr)`

Writes all information about all surfaces to the file `fptr` using a format that can be read by Smoldyn. This allows a simulation state to be saved.

`int checksurfaceparams(simptr sim, int *warnptr)`

Checks some surface parameters. Many more should be checked as well, although those haven't been written yet.

structure set up

`int surfenablesurfaces(simptr sim, int maxsurf)`

Allocates and sets up the surface superstructure for a maximum of `maxsurf` surfaces. This function may be called multiple times. If the input parameters are different at one call than from the previous call, then the surface superstructure is updated to the new parameters. Enter `maxsurf` as -1 to indicate default operation, meaning that nothing is done if the surface

5 C Code: structures and functions

superstructure was allocated previously, or that 5 surfaces are allocated otherwise. Returns 0 for success (or nothing was done because it was done previously), 1 for failure to allocate memory, or 2 if `sim` is undefined.

int surfexpandmaxspecies(simptr sim, int maxspecies)

Expands the number of molecular species that the surfaces can work with, allocating memory as needed. Surfaces and molecules should be kept synchronized, so this should be called whenever the master maxspecies changes. Returns 0 for success or 1 for failure to allocate memory.

surfaceptr surfaddsurface(simptr sim, char *surfname)

Adds a surface called `surfname` to the simulation. This enables surfaces and/or allocates more surfaces, as necessary. A pointer to the surface is returned, or NULL is returned for failure to allocate memory. If `surfname` was already declared as a surface, then this function simply returns a pointer to the existing surface.

void surfsetcondition(surfaceptr surfss, enum StructCond cond, int upgrade)

Sets the surface superstructure condition to `cond`, if appropriate. Set `upgrade` to 1 if this is an upgrade, to 0 if this is a downgrade, or to 2 to set the condition independent of its current value. If the condition is downgraded, this also downgrades the simulation structure condition.

int surfsetepsilon(simptr sim, double epsilon)

Sets the `epsilon` value in the surface superstructure. Returns 0 for success, 2 if the surface superstructure did not exist and could not be created, or 3 for an illegal requested value (≤ 0).

int surfsetmargin(simptr sim, double margin)

Sets the `margin` value in the surface superstructure. Returns 0 for success, 2 if the surface superstructure did not exist and could not be created, or 3 for an illegal requested value (≤ 0).

int surfsetneighdist(simptr sim, double neighdist)

Sets the neighbor distance value in the surface superstructure. Returns 0 for success, 2 if the surface superstructure did not exist and could not be created, or 3 for an illegal requested value (≤ 0).

int surfsetneighhop(surfaceptr srf, int neighhop)

Sets the `neighhop` element of the surface structure to the value entered. Returns 0 unless `srf` is NULL, in which case it returns 1.

int surfsetcolor(surfaceptr srf, enum PanelFace face, double *rgba)

Sets the color vector for face `face` of surface `srf` to `rgba`. Any face value is allowed, including both and none. Returns 0 for success, 1 if no surface was entered, or 2 if one or more of the color or alpha values is out of range (between 0 and 1 inclusive).

int surfsetedgepts(surfaceptr srf, double value)

Sets the drawing thickness for surface `srf` to `value`. Returns 0 for success, 1 if no surface was entered, or 2 if `value` is negative.

int surfsetstipple(surfaceptr srf, int factor, int pattern)

Sets the stippling pattern for drawing surface `srf`. `factor`, which needs to be at least 1, is the repeat distance for the entire stippling pattern and `pattern`, which needs to be between

5.6 Surfaces (functions in *smolsurf.c*)

0 and 0xFFFF inclusive, is the pattern to be used. Enter either or both as negative values to not set that parameter. Returns 0 for success, 1 if the surface is undefined, or 2 if inputs are out of range.

int surfsetdrawmode(surfaceptr srf, enum PanelFace face, enum DrawMode dm)
Sets the surface drawing mode of face **face** of surface **srf** to **dm**. Any face value is allowed. The only drawing mode that is not allowed is **DMnone**. Returns 0 for success, 1 if the surface is undefined, or 2 if the drawing mode is out of range.

int surfsetshiny(surfaceptr srf, enum PanelFace face, double shiny)
Sets the shininess of face **face** of surface **srf** to **shiny**. Any face value is allowed. The shininess needs to be between 0 and 128, inclusive. Returns 0 for success, 1 if the surface is undefined, or 2 if the shininess is out of range.

int surfsetaction(surfaceptr srf, int ident, const int *index, enum MolecState ms, enum PanelFace face, enum SrfAction act, int newident)
Sets the **action** element of surface **srf**. **ident** is the molecule species, which should be a valid species index, -5 to indicate all species, or 0 to indicate a group of species, perhaps with a wildcard. If value 0 is used, then **index** should be sent in as the **index** element that gets returned by **molstring2index1**. **ms** is the species state, which should be a regular species state (not **MSbsoln**) or should be **MSall** to indicate all states. **face** is the interaction face of the surface, if any; it is allowed to adopt any value, including **PFfront**, **PFback**, **PFboth** (to indicate both front and back), or **PFnone** to indicate actions for surface-bound species that do not collide with other surfaces. **act** is the desired action. **newident** is the new identity of the molecule; set it to -1 for no new identity. Returns 0 for success, 1 if **ident** is out of range, 2 if **ms** is out of range, or 3 if the requested action is not permitted for the indicated interaction face.

int surfsetrate(surfaceptr srf, int ident, const int *index, enum MolecState ms, enum MolecState ms1, enum MolecState ms2, int newident, double value, int which)
Sets the **srfrate** or **srfprob** element of the action details of surface **srf**, along with the **srfnewspec** and **srfdatasrc** elements. See the data table for the **srftristate2index** function for the possible input combinations of **ms**, **ms1**, and **ms2** and what they mean (this function uses the tristate input format). Enter **ident** as a positive number to indicate a specific species, -5 to indicate all species, or -6 to indicate wildcard selected species (channel 0). Typically, **newident** will be the same as **ident**, although it can be different for a species change at the surface. **value** is the desired rate or probability value. Enter **which** as 1 to set the **srfrate** element and as 2 to set the **srfprob** element. Returns 0 for success, 1 if **ident** is out of range (equal to 0), 2 if **ms** is out of range, 3 if **ms1** is out of range, 4 if **ms2** is out of range, 5 if **newident** is out of range, 6 if **value** is out of range, or -1 if memory could not be allocated for the surface action details data structure.

int surfsetmaxpanel(surfaceptr srf, int dim, enum PanelShape ps, int maxpanel)
Sets the maximum number of panels of shape **ps** for surface **srf** to **maxpanel**. The system dimensionality is **dim**. This function may be called multiple times. It allocates memory as needed.

int surfaddpanel(surfaceptr srf, int dim, enum PanelShape ps, char *string, double *params, char *name)

5 C Code: structures and functions

Adds or modifies a panel of shape **ps** to surface **srf**, in a **dim** dimensional system. **string** lists any text parameters for the panel, which in practice is only a single word that gives the orientation of a rectangle panel (e.g. “+0” or “-y”). **params** lists the numerical parameters for the panel location, size, and drawing characteristics. The number of necessary parameters can be found from the **surfpanelparams** function and the specific parameters, which depend on the panel shape and the system dimensionality, are described in the User Manual. **name** is an optional parameter; if it is included and is not an empty string, the panel is named **name**. If this panel name was already used by a panel of the same shape, then this function overwrites that panel’s data with the new data. If the name was already used by a panel with a different shape, then this creates an error, and if the name was not used before, then a new panel is created. To use default panel naming, send in **name** as either NULL or as an empty string.

This function returns 0 for success, -1 for inability to allocate memory, 1 for missing surface, 2 for **ps** out of range, 3 for unable to parse **string**, 4 for drawing slices or stack are zero or negative, 5 for cylinder ends are at same location, 6 for hemisphere outward pointing vector has zero length, 7 for a zero or negative radius, 8 for a normal vector with zero length, or 9 if the panel name was used before for a panel with a different shape.

This function calculates all of the geometrical data that are stored with the panel, such as the normal direction and the edge normals.

```
int void surftransformpanel(panelptr pnl, int dim, double *translate, double
    *origin, double *expand)
```

Performs translation and expansion of a surface panel, while maintaining the panel shape as much as possible. Enter the translation vector in **translate**, the origin coordinates about which the expansion should be performed in **origin**, and the expansion values for the different coordinates in **expand**. If expansion should be isotropic, enter the same values for all components of the **expand** vector. Each component should be 1 for no expansion, less than 1 for retraction, and greater than 1 for expansion. Negative values indicate inversion. This function updates all panel **point** and **front** values. This lowers the surface condition, box condition, and compartment condition to indicate that some recomputation is needed. Note that Smoldyn cannot distort spheres, hemisphere, cylinders, or disks, so anisotropic expansion is only done approximately.

Most of this function is quite straightforward. Each panel point value that represents an actual coordinate is transformed through translation and expansion. Normals are recomputed as needed. Note that all computations are done in-place, which is usually easy, but some care is required to make sure that one is using the correct values. Several functions transform the radius of a panel. The math for this is as follows. Consider a 2D system, a radius vector (often given as **radiusv** in the code), **r**, and suppose that the expansion vector is (e_x, e_y) . Also assume this radius vector starts at the expansion origin. In this case, the transformed vector endpoint will be at $(e_x r_x, e_y r_y)$. This means that the new radius will be $(e_x^2 r_x^2 + e_y^2 r_y^2)^{1/2}$. If the old radius vector has length r and is parallel to the unit vector **n**, then the new radius is $r(e_x^2 n_x^2 + e_y^2 n_y^2)^{1/2}$. For 3D objects, anisotropic expansion would normally distort shapes away from being perfect cylinders, hemispheres, spheres, or disks. However, Smoldyn does not support such distorted shapes, so this function just picks specific radius vectors and uses those.

```
void surftranslatepanel(panelptr pnl, int dim, double *translate)
```

Translates panel **pnl** by amount given in **translate**. No superstructure conditions are

5.6 Surfaces (functions in *smolsurf.c*)

changed in this function. This function is similar to `surftransformpanel` but it much simpler.

```
void surfupdateoldpos(surfaceptr srf, int dim)
```

Copies the contents of each panel `point` element to the panel `oldpoint` element, and also each `front` element to the panel `oldfront` element. Use this when a surface is about to be moved and the `oldpoint` and `oldfront` elements should represent the prior location.

```
void surftranslatesurf(surfaceptr srf, int dim, double *translate)
```

Translates surface `srf` by amount given in `translate`. The box superstructure condition is downgraded to `SCparams`.

```
int surfsetemitterabsorption(simptr sim)
```

Sets emitter absorption probabilities for panels based on emitter information in the surface structures. Returns 0 for success or 1 if one or more of the distances between emitters and a surface panel was zero (which leads to divide-by-zero errors).

```
int surfsetjumppanel(surfaceptr srf, panelptr pnl1, enum PanelFace face1, int  
bidirect, panelptr pnl2, enum PanelFace face2)
```

Sets up jumping between face `face1` of panel `pnl1` and face `face2` of panel `pnl2`, for surface `srf`. Jumping is set up to be unidirectional, from `pnl1` to `pnl2` if `bidirect` equals 0 and is set up to go in both directions if `bidirect` equals 1. This only sets the panel jumping indices and does not assign jumping actions to any molecules. Returns 0 for success, 1 for no surface, 2 for no `pnl1`, 3 for `face1` out of range, 4 for `bidirect` out of range, 5 for an error with `pnl2` including it having a different shape from `pnl1` or it equaling `pnl1`, and 6 for `face2` out of range.

```
double srfcalcrate(simptr sim, surfaceptr srf, int i, enum MolecState ms1, enum  
PanelFace face, enum MolecState ms2)
```

Calculates the actual rate for the interaction of a molecule of type `i` and state `ms1` interacting with face `face` of surface `srf`, and ending up in state `ms2`. This uses the `srf->actdetails[i][ms][face]->srfrprob[ms2]` data to calculate the actual conversion rate, and accounts for reversible interactions as much as possible. All cases are calculated assuming steady-state behavior, and all are found using the `SurfaceParam.c` function `surfacefacerate`. Returned rates will be between 0 and `MAX_DBL`. Error codes are returned with negative numbers: -1 indicates that the input situation is impossible (i.e. `ms1=MSsoln` and `face=PFnone`), or that input data are unavailable (i.e. the surface action isn't `SAmult` or the action details aren't recorded for this action); -2 indicates that the interaction probabilities haven't been computed yet in the action details structure, probably because the appropriate set up routine hasn't been called yet; and -3 indicates that reflection coefficients were requested, which cannot be computed here. See `srfcalcprob`.

```
double srfcalcprob(simptr sim, surfaceptr srf, int i, enum MolecState ms1, enum  
PanelFace face, enum MolecState ms2)
```

Calculates the surface interaction probability for the interaction of a molecule of type `i` and state `ms` interacting with face `face` of surface `srf`, and ending up in state `ms2`. This uses the `srf->actdetails[i][ms][face]->srfrate[ms2]` data to calculate the conversion probability, and accounts for reversible interactions as much as possible. All cases are calculated assuming steady-state behavior, and all are found using the `SurfaceParam.c` function `surfaceprob`. Returned probabilities will be between 0 and 1, inclusive, or an error code. Error codes are returned with negative numbers: -1 indicates that input data

5 C Code: structures and functions

are unavailable (i.e. the surface action isn't `SAmult` or the action details aren't recorded for this action); -2 indicates that the rate is listed as being negative, which is impossible; and -3 indicates that reflection probabilities were requested, which cannot be computed here. See `srfcalsrate`.

`int surfsetneighbors(panelptr pnl, panelptr *neighlist, int nneigh, int add)`

Adds or removes neighbors to or from a panel's list of neighbors. `pnl` is the panel whose neighbor list should be modified, `neighlist` is a list of neighboring panels to be added or removed, `nneigh` is the number of neighbors that are listed in `neighlist`, and `add` is 1 if those listed in `neighlist` should be added, or 0 if they should be removed. For addition, neighbors are not added again if they are already in the list. If all neighbors should be removed, send `neighlist` in as `NULL`. This allocates space as needed. It returns 0 for success or 1 if not enough space could be allocated. For optimal memory allocation, it's slightly better if many neighbors are added at once in a single function call, rather than one neighbor per function call.

`int surfaddemitter(surfaceptr srf, enum PanelFace face, int i, double amount, double *pos, int dim)`

Adds an emitter to a surface so that it can be used for simulating unbounded diffusion. This takes care of any necessary memory allocating. `srf` is the surface that the emitter is being added to, `face` is the surface face, `i` is the species number, `amount` is the emitter amount, flux, or weighting, `pos` is the `dim`-dimensional position of the emitter, and `dim` is the system dimensionality. Returns 0 for success or 1 if this is unable to allocate memory. This does not calculate the panel absorption probabilities, but does allocate space for them, if needed.

`surfaceptr surfreadstring(simptr sim, surfaceptr srf, char *word, char *line2, char *erstr)`

Reads and processes one line of text from the configuration file, or some other source, for the surface pointed to by `srf`. If the surface is not known or has not been defined yet, then set `srf` to `NULL`. The first word of the line should be sent in as `word` and the rest sent in as `line2`. If this function is successful, it returns the surface pointer and it does not change the contents of `erstr`; if not, it returns -1 and it writes an error message to `erstr`.

`int loadsurface(simptr sim, ParseFilePtr *pfpptr, char *line2, char *erstr)`

`loadsurface` loads a surface from an already opened disk file pointed to with `fp`. `lctrptr` is a pointer to the line counter, which is updated each time a line is read. If successful, it returns 0 and the surface is added to the surface superstructure in `sim`, which should have been already allocated. Otherwise it returns the updated line counter along with an error message. If a surface with the same name (entered by the user) already exists, this function can add more panels to it. It can also allocate and set up a new surface. If this runs successfully, the complete surface structure is set up, with the exception of box issues. If the routine fails, any new surface structure is freed.

`int surfupdateparams(simptr sim)`

Sets the surface interaction probabilities within the surface action structure, calling `srfcalsprob` to do the actual calculations. This function loops through all possible interactions and sets those that are needed. Returns 0 for success or 2 if the molecules aren't adequately set up.

This function loops over all surfaces and then within each surface, it loops over all possible interactions several times. In the first pass, it computes (with `srfcalsprob`) and assigns

the interaction probabilities. In the next pass, it adds up all the probabilities for exiting each particular state; if those probabilities add up to more than one, then all are scaled down to make the sum equal to one, and if the sum is less than one, then the remaining probability is the probability of staying in the originating state. The next pass through adds up probabilities to compute cumulative probabilities, which are the ones that are actually used in a simulation. The final pass through records the reverse probabilities by simply looking them up in the newly updated action details structure, which accounting for changes to new species.

At the very end of this function, `surfsetemitterabsorption` is called to set up effective emitter stuff.

```
int surfupdatelists(simptr sim)
```

Sets up surface molecule lists, area lookup tables, and action probabilities. If calculated probabilities exceed 1 or add up to more than 1, they are adjusted as needed, although this may affect simulation results. No warnings are returned about these possible problems, so they should be checked elsewhere. Returns 0 for success, 1 for inability to allocate memory, or 2 for molecules not being sufficiently set up beforehand. This function may be called at setup, or later on during the simulation.

```
int surfupdate(simptr sim)
```

Sets up or updates surface data structures.

core simulation functions

```
enum PanelFace panelside(double* pt, panelptr pnl, int dim, double *distptr, int strict, int useoldpos)
```

Returns the side of the panel `pnl` that point `pt` is on. If `strict` is 0, then `PFback` is returned if the point is exactly at the panel, while if `strict` is 1 then `PFnone` is returned if the point is exactly at the panel. In general, this should not be set for strict use when getting panel faces and should be for strict use for assistance in setting panel faces. If `distptr` is sent in as a non-NULL pointer, its contents will be set to the distance that `pt` is away from the infinite panel, with a positive number for the front side and negative or zero for the back. The values returned by this function define the side that `pt` is on, so should be called for other functions that care. The distance value that is returned by this function is also used to determine the face; thus, if two calls with different `pt` values return the exact same distance value, then the same face will be returned. Set `useoldpos` to 1 if the old point and front values should be used for the panel position rather than the current point and front values.

```
void panelnormal(panelptr pnl, double *pos, enum PanelFace face, int dim, double *norm)
```

Returns, in `norm`, the normal vector for the panel `pnl`, that points outwards from the `face` side. If this is a curved panel, such as a sphere or a cylinder, then `pos` is the position on the surface for which the local normal should be computed. If `face` is not equal to `PFfront` or `PFback`, then it is assumed that the front side is desired.

```
int lineXpanel(double *pt1, double *pt2, panelptr pnl, int dim, double *crsspt, enum PanelFace *face1ptr, enum PanelFace *face2ptr, double *crossptr, double *cross2ptr, int *veryclose, int useoldpoint)
```

This determines if the line from `pt1` to `pt2` crosses the panel `pnl`, using a `dim` dimensional system. These input variables are not changed by this function; other variables are for output

5 C Code: structures and functions

only such that prior values are not looked at by the function. With those, **crsspt** must be a **dim** dimensional vector and the others may be **NULL** or can be pointers to single values that will be overwritten. The panel includes all of its edges. 1 is returned if the line crosses the panel at least once and 0 if it does not. If it does not cross, all return values except **veryclose** are undefined. If it crosses, **crsspt** will be the coordinates of the crossing, **face1** will be the side of the panel that is first impacted, **face2** will be the side that is towards **pt2**, and the contents of **crossptr** will be the crossing position on the line, which is between 0 and 1 inclusive. The type of crossing can be determined by looking at the returned **face** values. (1) If **face1!=face2**, then the line crosses the panel exactly once and the contents of **cross2** are undefined. (2) if **face1==face2**, then the line crosses the panel exactly twice (implying a curved panel) for which the first crossing will be recorded in **crsspt** and **crossptr** and the latter in **cross2ptr**. Suppose the line crosses a hemisphere in such a way that it enters the corresponding sphere where the hemisphere is open and departs where it is a panel. In this case, the **face1** value will equal the inside face (unlike the value returned by **panelside**). In contrast, if it enters on the closed side and exits on the open side, both **face1** and **face2** will equal the outside. The same rules apply for cylinders.

veryclose, if it is entered as non-**NULL**, is returned equal to 1 if **pt1** is within **VERYCLOSE** distance units of the panel, as 2 if **pt2** is this close to the panel, as 3 if both points are this close to the panel, and as 0 if neither point is this close to the panel. These codes can be used to determine if round-off error is likely to be a problem.

Set **useoldpos** if the **oldpoint** and **oldfront** elements should be used for the panel position rather than the current location given in **point** and **front**.

While **crsspt** will be returned with coordinates that are very close to the panel location, it may not be precisely at the panel, and there is no certainty about which side of the panel it will be on; if it matters, fix it with **fixpt2panel**. Similar rules apply for the contents of **crossptr** and **cross2ptr**.

Each portion of this routine does the same things, and usually in the same order. Crossing of the infinite panel is checked, the crossing value is calculated, the crossing point is found, and finally it is determined if intersection actually happened for the finite panel. For hemispheres and cylinders, if intersection does not happen for the first of two possible crossing points, it is then checked for the second point.

```
int lineexitpanel(double *pt1, double *pt2, panelptr pnl, int dim, double
    *pnledgept, int *exitside)
```

This finds where the line segment that goes from **pt1** to **pt2** exits the panel **pnl**. Although it probably isn't essential, this function assumes that the line segment is nearly co-planar with its local position on the panel (or parallel to the panel for 2D). It returns the answer as coordinates in **pnledgept**. This works in 1, 2, or 3-D. This returns the side or axis of the panel that the line exits through in **exitside**. See below. Returns 0 for success or 1 if **pt1** is identical to **pt2**, which is an error.

For the most part, this function calls other functions for each of the possibilities. An exception is for 2D cylinders, which are really just 2 parallel lines. In this case, this function computes the dot product between the vector from a cylinder end to **pt1** and the cylinder normal vector, using the sign of the dot product to determine which line to consider. Then, it calls **Geo_LineExitLine2** using this line. This function was new in Smoldyn 2.37.

1D	2D	3D
rectangles, ps = PSrect		
1	1 for p[0] 2 for p[1]	1 for p[0] - p[1] 2 for p[1] - p[2] 3 for p[2] - p[3] 4 for p[3] - p[0]
triangles, ps = PStri		
1	1 for p[0] 2 for p[1]	1 for p[0] - p[1] 2 for p[1] - p[2] 3 for p[2] - p[3]
spheres, ps = PSsph		
1	1	1
cylinders, ps = PScyl		
not allowed	1 for p[0] 2 for p[1]	1 for p[0] 2 for p[1]
hemispheres, ps = PShemi		
not allowed	1 for cw end 2 for ccw end	1
disks, ps = PSDisk		
not allowed	1 for end cw of normal 2 for end ccw of normal	1

```
void paneledgenormal(panelptr pnl, double *pnledgept, int dim, int edgenum,
double *normal)
```

Returns the outward-pointing unit normal vector to the edge of panel **pnl**, where this normal is plane-parallel to the local panel surface and perpendicular to the panel edge. In 2D, this normal is parallel to the panel at its edge. This normal points out of the panel. Send in **pnledgept** as a point that is at the edge of the panel (not always used) and **edgenum** as the number of the edge (not always used). See the table for **lineexitpanel** for the correct values to use for this input. The result is returned in the vector **normal**. If **edgenum** is entered as zero, then this implies that the point is not at a panel edge but in the center of the panel somewhere. In this case, this returns a vector that is parallel to the surface at **pnledgept**. It is not specified here in which direction this vector points, except simply that it is parallel to the surface at the point **pnledgept**.

Most portions of this function are fairly straightforward. This function only works for 2D and 3D, because it doesn't make sense in 1D. This function was new in version 2.37.

The sphere portion is written under the assumption that **edgenum** is equal to zero, because there are no edges to this surface.

```
int ptinpanel(double *pt, panelptr pnl, int dim)
```

Determines if the point **pt** is inside the finite panel **pnl**, returning 1 if so and 0 if not. Here, inside only means that the point is within the volume that is swept out perpendicular to the plane of the panel, and says nothing about the position of the point relative to the plane of the panel.

This function is nearly identical to the portion of **lineXpanel** that checks whether the position **crsspt** is within the panel or not.

5 C Code: structures and functions

`enum SrfAction surfaction(surfaceptr srf, enum PanelFace face, int ident, enum MolecState ms, int *i2ptr, enum MolecState *ms2ptr)`

Returns the surface action that should happen to a molecule of type `ident` and state `ms` that interacts with face `face` of surface `srf`. `ms` needs to be a real molecule state, meaning that it is not allowed to be `MSbsoln`. If the state of `ident` should be changed, then the new state is returned in `ms2ptr`, if `ms2ptr` is not NULL (`ms2ptr` may be returned pointing to `MSbsoln`, as well as to `MSsoln`). If the species of `ident` should be changed, then the new species number is returned in `i2ptr`, if `i2ptr` is not NULL; if it should not be changed, then `i2ptr` is returned pointing to the same value as `ident`. This function does not return `SAmult`; instead, it specifies what should happen in detail, including `SAadsorb` for adsorption, `Sarevdes` for reversible desorption, `SAirrevdes` for irreversible desorption, and `SAflip` for on-surface state change.

`int rxnXsurface(simptr sim, moleculeptr mptr1, moleculeptr mptr2)`

Returns 1 if a potential bimolecular reaction between `mptr1` and `mptr2` is across a non-transparent surface, and so cannot actually happen. Returns 0 if a reaction is allowed. Using the diffusion coefficients of the two molecules, this calculates the reaction location and then determines which molecules need to diffuse across which surfaces to get to that location. If the molecules can diffuse across the necessary surfaces, then the reaction is allowed, and not otherwise. This routine does not allow reactions to occur across jump surfaces. Also, it does not look for jump paths that go from `mptr1` to `mptr2`. Surface-bound molecules that are in their “up” or “down” state are assumed to be accessible from both sides of the surface, whereas those that are in the “front” or “back” states are accessible from only one side.

`void getpanelnormal(double *pt, panelptr pnl, int dim, double *norm)`

Finds the local panel normal for panel `pnl` at position `pt`, returning the vector in `norm`. The result points towards the front side of the panel. This only considers the infinite plane of the panel, while ignoring its boundaries (similarly, hemispheres are considered to be identical to spheres and cylinders are considered to be infinitely long).

`void fixpt2panel(double *pt, panelptr pnl, int dim, enum PanelFace face, double epsilon)`

Fixes the point `pt` onto the face `face` of panel `pnl`. Send in `face` equal to `PNone` if `pt` should be moved as close as possible to `pnl`. If it should also be on the front or back face of the panel, as determined by `panelside`, then send in `face` equal to `Pffront` or `Pfback`, respectively. Before moving, if `pt` is already on the proper face and its distance is less than or equal to `epsilon`, it is not moved; setting `epsilon` to 0 ensures moving. This function first moves `pt` to the panel in a direction normal to the local panel surface and then nudges `pt` as required to get it to the proper side. This only considers the infinite plane of the panel, while ignoring its boundaries (similarly, hemispheres are considered to be identical to spheres and cylinders are considered to be infinitely long).

`int fixpt2panelnocross(simptr sim, double *pt, panelptr pnl, int dim, enum PanelFace face, double epsilon)`

This is very similar to `fixpt2panel`, in that it moves point `pt` to the correct side of panel `pnl` based on the value of `face`. The difference is that this also ensures that no other panels are crossed in the process. If other panels are crossed, then this keeps on adjusting the point until no other panels are crossed. The function usually returns 0 but will return 1 if it tried 20 times to find a suitable fixed position and was unable to do so; in this latter case, the

5.6 Surfaces (functions in *smolsurf.c*)

point is not moved at all. Any crossings of the panel `pnl`, and all of its neighboring panels are ignored.

`void movept2panel(double *pt, panelptr pnl, int dim, double margin)`

This moves the point `pt` to the nearest location that it is over the panel `pnl`, and also inside the edge by distance `margin`. This means that `pt` is not moved into the plane of the panel, which is done by `fixpt2panel`, but is moved parallel to the plane of the panel.

`int closestpanelpt(panelptr pnl, int dim, double *testpt, double *pnlpt, double margin)`

Finds the closest point that is on panel `pnl` to the test point `testpt` and returns it in `pnlpt`. `testpt` and `pnlpt` should not point to the same memory address. This returns an integer which is 0 if the closest panel point is not on the panel edge, and which is the edge number if the closest panel point is on an edge. The edge numbers are listed in the table following the `lineexitpanel` description. If the test point is exactly at the panel edge, then the edge number is returned.

If the closest panel point is within `margin` of the edge or a corner, even if the test point is technically within the panel, then the coordinates at the edge or corner are returned.

`double closestsurfacept(surfaceptr srf, int dim, double *testpt, double *pnlpt, panelptr *pnlptr, boxptr bptr)`

Finds the closest point that is on the surface `srf` to the test point `testpt` and returns it in `pnlpt` if `pnlpt` is not NULL.. This also returns the panel that that point is on, in `pnlptr`, if `panelptr` is not NULL. It returns the distance between `testpt` and `pnlpt`, which should always be positive. If this returns a negative number, that means that the surface has no panels. If `bptr` is entered as NULL, then this searches over all panels on the listed surface, but if `bptr` is entered pointing to a box, then this only searches the panels that are in that box. `pnlpt` is allowed to be the same pointer as `testpt`.

`void movemol2closepanel(simptr sim, moleculeptr mptr)`

Checks to see if surface-bound molecule `mptr` is within the area of the finite panel `mptr->pnl` (i.e. over or under the panel, ignoring the position relative to the plane of the panel). If it isn't, this sees if `mptr` is over a neighboring panel and if so, this puts `mptr->pos` on the neighboring panel at the correct location. At the end, and regardless of whether the molecule changed panels or not, the molecule position is fixed to the correct face of its panel using `fixpt2panel`. If the molecule needs to be moved parallel to the plane of a panel, whether back to its original panel or onto a new panel, then it is inset from the edge by distance `margin`. This function should only be called if the system dimensionality is 2 or 3.

This function is called by the `diffuse` function for putting surface-bound molecules in the correct places after they are diffused. It is also called by the `dosurfinteract` function after a surface-bound molecule reflects off of another surface. It was re-written in July-September of 2015, for Smoldyn version 2.37, because Boris Slepchenko discovered that its prior design produced inaccurate diffusion. In the new version, the molecule is projected onto its infinite panel according to the local panel normal. If it is then within the finite panel, the function is done. Otherwise, the function determines the coordinates where the molecule's trajectory exits the panel, it picks a random neighboring panel that is close to these coordinates to be the new panel, and the molecule's trajectory is rotated over to the new panel. If there is no new panel, then the molecule bounces off of the edge of the current panel. The function repeats this process until the full molecule trajectory is used up.

5 C Code: structures and functions

```
void surfacereflect(moleculeptr mptr, panelptr pnl, double *crsspt, int dim, enum
    PanelFace face)
```

This bounces the molecule `mptr` off of the face `face` side of panel `pnl`. Elastic collisions are performed, which should work properly for any shape panel and any dimensionality. For flat panels, elastic collisions also apply to Brownian motion. It is assumed that the molecule travels from some point (not given to this function, and irrelevant) to `mptr->pos`, via a collision with the panel at location `crsspt`, where `crsspt` is either exactly at the panel or is slightly on impact side of the panel. The molecule `pos` element is set to the new, reflected, position, which will always be on the face side of the panel.

```
int surfacejump(moleculeptr mptr, panelptr pnl, double *crsspt, enum PanelFace
    face, int dim)
```

This performs a jump for molecule `mptr` that hit panel `pnl` on face `face`. The contact location is input in `crsspt`, which needs to be very close to the panel but does not have to be on the proper side. This looks up the jump destination and translates both the `crsspt` value and the molecule position in `mptr->pos` to represent this jump. On return, `crsspt` is on the destination face of the destination panel. The molecule `pos` element will always be returned on the destination face side of the destination panel. For the most part, this function only allows jumps between panels with the same shape, the same dimensions, and the same orientation. Exceptions are that sphere, hemisphere, and cylinder radii are allowed to differ between origin and destination panels. This function works for any molecule state. Returns 0 if no jump happened (i.e. `pnl->jumpp[face]` was NULL or `pnl->jumpf[face]` wasn't either `PFfront` or `PFback`) and returns 1 if a jump happened.

Internally, for each surface, a few things are calculated. `delta` is the jump offset (jumped position minus current position) and `dir` is the relative orientation of the panels (1 if parallel, -1 if antiparallel). The jump offset is then added to `crsspt` and to `mptr->pos`, and subtracted from `mptr->posoffset`; the final thing means that `pos+posoffset` is always the diffused to position, and does not include any jumps. At the end, the `crsspt` and molecule position are finalized.

```
int dosurfinteract(simptr sim, moleculeptr mptr, int ll, int m, panelptr pnl,
    enum PanelFace face, double *crsspt)
```

Performs interaction between molecule and surface for a collision that is known to have happened or for possible interaction from a surface-bound state. This converts, kills, reflects, adsorbs, desorbs, etc. molecules, as appropriate. This function is called in three situations: (1) by `checksurfaces` when a molecule is found to have diffused across a surface, in which case `pnl` is the panel that was diffused across and `face` is the fact that was diffused into, (2) by `checksurfacebound` when a molecule is surface-bound, and thus might be able to desorb or flip, in which case `pnl` is the panel to which the molecule is bound and `face` is `PFnone`, and (3) by `checksurfaces1mol` which is used by the lattice hybrid software to check the surfaces for a single molecule; in this case, `ll` and `m` are set to -1. Quite possibly, the code would be better if these uses were separated into independent functions, but that's not the case at the moment. This function is sometimes called twice for the same molecule during a single time step, if that molecule is both surface-bound and it diffuses across a different panel, but that should not lead to incorrect rates due to the different input parameters for the two calls.

At the beginning of the function, a few lines take care of special cases. The `isneigh` test, which only applies to situation (1) above, determines if the molecule is surface-bound and it diffused across a neighboring panel. If this test returns true, then the molecule has equal odds

5.6 Surfaces (functions in *smolsurf.c*)

of staying on the same panel or crossing to the new panel. This is used so that molecules can diffuse from, say, one sphere to a neighboring and intersecting sphere, despite the fact that neither sphere has an open edge. This special case doesn't really address a proper surface interaction, but instead enables diffusion from panel to panel on a single surface. The next test is also for situation (1) and is for collision with unbounded emitter type surfaces. If neither special case holds, the action is gotten from **surfaction**. This function again copes with the two separate situations by testing whether **face** is **PFnone**.

On return, **crsspt** will be on the same side of the surface as the molecule. Returns 1 if the molecule does not need additional trajectory tracking (e.g. it's absorbed) and 0 if it might need additional tracking (e.g. it's reflected). This function does not consider opposite-face actions. For example, if the front of a surface is transparent and the back is absorbing, an impact on the front will result in the molecule being transmitted to the far side, and not being absorbed.

int checksurfaces1mol(simp_{ptr} sim, molecule_{ptr} mptr, double crossminimum)

Essentially identical to **checksurfaces** function, but just for a single molecule. Also, and very importantly, this assumes that a molecule's trajectory starts at **mptr->via** and not at **mptr->posx**. The reason is that this function was designed for checking surfaces after a molecule hit a port, so it's for the surfaces that are after the "via" position. If the molecule's list specifies that the molecule should be in a port buffer, even if it isn't there already, then the live list is not updated further to reflect new surface interactions. However, the molecule will still change states, get killed, etc. as appropriate. The **crossminimum** value is used to indicate that a crossing should be ignored unless its value is greater than the **crossminimum** value. As usual, the value is the distance along the molecule's straight-line trajectory where the crossing occurs, where 0 is the starting point and 1 is the ending point. This check adds a distance of **VERYCLOSE** to the **crossminimum** value to prevent unintentional cross determinations due to round-off error.

Returns 0.

int checksurfaces(simp_{ptr} sim, int ll, int reborn)

Takes care of interactions between molecules and surfaces that arise from diffusion. Molecules in live list **ll** are considered; if **reborn** is 1, only the reborn molecules of list **ll** are considered. This transmits, reflects, or absorbs molecules, as needed, based on the panel positions and information in the molecule **posx** and **pos** elements. Absorbed molecules are killed but left in the live list with an identity of zero, for later sorting. Reflected molecules are bounced and their **posx** values represent the location of their last bouncing point. This function does not rely on molecules being properly assigned to boxes, and nor does it assign molecules to boxes afterwards. However, it does rely on the panels being properly assigned to boxes. If multiple surfaces are coincident, only the last one is effective. Returns error code of 0.

This code cycles through many molecules. For each molecule, it first copies **posx** to **via** (**posx** never changes in this function, while **via** will point to the last surface crossing location). Then it goes through the boxes that lie along the trajectory between **via** and **pos**. At each one, it goes through the panels in that box and asks if the trajectory crosses that panel. At the end of this scan, **crossmin** is the relative position along the trajectory, from 0 to 1, for the closest panel crossing, **pnlmin** is that panel, **facemin** is the panel face that was hit, **crssptmin** is the physical location of that crossing, and **crossmin2** is the relative position along the trajectory for the second closest panel crossing. Then, typically, this calls **dosurfinteract** with the molecule, **pnlmin**, **facemin**, and **crssptmin**, which takes care of the interaction and **via** is

5 C Code: structures and functions

set to the crossing point. This then repeats with the new trajectory, from `via` to the updated value of `pos` until there are no more crossings.

This function includes two “hacks.” First, if a molecule has over 100 surface interactions during the same diffusion step, this function decides that something has gone wrong, and it simply puts the molecule back to where it started and moves on to deal with the next molecule. This option addresses the possibility for a molecule to become trapped in an endless loop. In practice, I’ve seen it trigger when a molecule starts very close to the inside of a concave surface and diffuses a long distance parallel to the nearby surface; it then bounces along the surface over and over again as it essentially rolls along, with potentially over 100 bounces. In this case, the warning is incorrect because the molecule was bouncing correctly and not stuck in a loop, and would have ended eventually.

The other “hack” is that this function looks for both the first and second surface crossings along the molecule’s current trajectory. If they differ by less than 10^{-12} (but the difference is greater than zero), this function doesn’t bother dealing with either surface, but puts the molecule back to its last known good position. The idea is that if the difference equals 0, then the last declared surface panel has priority. However, if the difference is negligibly larger than zero, then round-off errors are likely to dominate for calculations, which can cause the molecule to accidentally cross one of the two surfaces. The only time that this hack is likely to become a problem is if the user defines two essentially coincident surfaces, in which case they will become effectively reflective.

```
int checksurfacebound(simptr sim, int ll)
```

Takes care of actions for surface-bound molecules, such as desorption, on-surface orientation flipping, etc. Returns error code of 0.

5.7 Boxes (functions in `smolboxes.c`)

The simulation volume is exactly divided into an array of identical virtual boxes. These allow the simulation to run efficiently because only potential reactions between molecules that are known to be physically close need to be checked, and the same for molecule-surface interactions. In principle, the boxes are fairly simple. In practice though, they complicate the overall code quite significantly. While the boxes are sized to exactly fill the simulation volume, the edge boxes are considered to extend beyond the volume to plus or minus infinity. In this way, all of space is within some box, and points outside of the simulation volume are assigned to the nearest box. Each box has its own box structure.

```
typedef struct boxstruct {  
    int *indx;           // dim dimensional index of the box [d]  
    int nneigh;          // number of neighbors in list  
    int midneigh;        // logical middle of neighbor list  
    struct boxstruct **neigh; // all box neighbors, using sim. accuracy  
    int *wpneigh;        // wrapping code of neighbors in list  
    int nwall;           // number of walls in box  
    wallptr *wlist;      // list of walls that cross the box  
    int maxpanel;        // allocated number of panels in box  
    int npanel;          // number of surface panels in box  
    panelptr *panel;     // list of panels in box  
    int *maxmol;         // allocated size of live lists [ll]  
    int *nmol;           // number of molecules in live lists [ll]
```

```

moleculeptr **mol;           // lists of live molecules in the box [ll][m
    ]
} *boxptr;

```

boxstruct (declared in *smollib.h*) is a structure for each of the virtual boxes that partition space. Each box has a list of its neighbors, in **neigh**, as well as a little information about them. This list extends from 0 to **nneigh-1**. From 0 to **midneigh-1** are those neighbors that logically precede the box, meaning that they are above or to the left, whereas those from **midneigh** to **nneigh-1** logically follow the box. If there are no periodic boundary conditions, the logical order is the same as the address order; however, this is not necessarily true with the inclusion of wrap-around effects. In **wpneigh** is a code for each neighbor that describes in what way it is a neighbor: 0 means that it's a normal neighbor with no edge wrap-around; otherwise pairs of bits are associated with each dimension (low order bits for low dimension), with the bits equal to 00 for no wrapping in that dimension, 01 for wrapping towards the low side, and 10 for wrapping towards the high side. This might be clearer in the *Zn.c* documentation. The neighbors that are listed depend on the requested simulation accuracy:

accuracy	neighbors	wrap-around
<3	none	no
3 to <6	nearest	no
6 to <9	nearest	yes
<9	all	yes

Boxes also have lists of molecules, allocated to size **maxmol[11]** and filled from 0 to **nmol[11]-1**) that correspond to the master molecule lists, and walls (**wlist**, allocated and filled with **nwall** pointers) within them. While the lists are owned by the box, the members of the lists are simply references, rather than implications of ownership. The same, of course, is true of the neighbor list, although the box owns the **wpneigh** list. If wall or neighbor lists are empty, the list is left as **NULL**, whereas the molecule list always has a few spaces in it. Boxes are collected in a box superstructure.

```

typedef struct boxsuperstruct {
    enum StructCond condition; // structure condition
    struct simstruct *sim;     // simulation structure
    int nlist;                 // copy of number of molecule lists
    double mpbox;              // requested number of molecules per box
    double boxsize;            // requested box width
    double boxvol;             // actual box volumes
    int nbox;                  // total number of boxes
    int *side;                 // number of boxes on each side of space
    double *min;               // position vector for low corner of space
    double *size;              // length of each side of a box
    boxptr *blist;             // actual array of boxes
} *boxssptr;

```

boxsuperstruct (declared in *smollib.h*) expresses the arrangement of virtual boxes in space, and owns the list of those boxes and the boxes. **condition** is the current condition of the superstructure and **sim** is a pointer to the simulation structure that owns this superstructure. **nlist** is a copy of the number of molecule lists that are used in the molecule superstructure. This is used here, and the **mol** element of the individual boxes are allocated to this, rather than **maxlist**, because boxes can potentially use up lots of memory, and this saves allocating memory unnecessarily.

Either **mpbox** or **boxsize** are used but not both. Boxes are arranged in a rectangular prism grid and exactly cover all space inside the walls. The structure of the boxes in space is the same as that

5 C Code: structures and functions

of a `dim` rank tensor, allowing tensor indexing routines to be used to convert between box addresses and indices. The box index along the `d`'th dimension of a point with position `x[d]` is

```
indx[d]=(int)((x[d]-min[d])/size[d]);
```

where integer conversion takes care of the truncation. Because of this, a box includes the points that are exactly on the low edge, but not those that are exactly on the high edge. Converting from box index to address is easy with the tensor routine in `Zn.c`, or can also be calculated quickly with the following code fragment, which outputs the box number as `b`,

```
for(b=0, d=0;d<dim;d++) b=side[d]*b+indx[d];
```

Converting the box number to the indices can also be done, but the `Zn.c` routine is easiest for this.

low level utilities

```
void box2pos(simptr sim, boxptr bptr, double *poslo, double *poshi)
```

Given a pointer to a box in `bptr`, this returns the coordinate of the low and/or high corners of the box in `poslo` and `poshi`, respectively. They need to be pre-allocated to the system dimensionality. If either point is unwanted, enter `NULL`. This requires that the `min` and `size` portions of the box superstructure have been already set up.

```
boxptr pos2box(simptr sim, double *pos)
```

`pos2box` returns a pointer to the box that includes the position given in `pos`, which is a `dim` size vector. If the position is outside the simulation volume, a pointer to the nearest box is returned. This routine assumes that the entire box superstructure is set up.

```
void boxrandpos(simptr sim, double *pos, boxptr bptr)
```

Returns a uniformly distributed random point, in `pos`, that is within the box `bptr`.

```
int panelinbox(simptr sim, panelptr pnl, boxptr bptr)
```

Determines if any or all of the panel `pnl` is in the box `bptr` and returns 1 if so and 0 if not. For most panel shapes, this is sufficiently complicated that this function just calls other functions in the library file `Geometry.c`.

```
int boxaddmol(moleculeptr mptr, int ll)
```

Adds molecule `mptr`, which belongs in live list `ll`, to the box that is pointed to by `mptr->box`. Returns 0 for success and 1 if memory could not be allocated during box expansion.

```
void boxremovemol(moleculeptr mptr, int ll)
```

Removes molecule `mptr` from the live list `ll` of the box that is pointed to by `mptr->box`. Before returning, `mptr->box` is set to `NULL`. If the molecule is not in the box that's listed, then this doesn't try removing it; this result is fine if the molecule is actually in no box at all (which can happen) but is probably a bug if the molecule is in the wrong box.

```
boxptr boxscansphere(simptr sim, const double *pos, double radius, boxptr bptr, int *wrap)
```

Allows the calling function to scan over all of the boxes that are at least partially within `radius` distance of the point `pos`. This accounts for system edges and periodic boundaries. Call this at the first time in a scan with `bptr` equal to `NULL`, which instructs the function to set up internal static variables, and in subsequent scans with `bptr` not equal to `NULL`. This returns boxes repeatedly for the scan, and then returns `NULL` when the scan is complete. At

present, it is not possible to skip ahead in the scan using the **bptr** input, although this could be implemented if it would be useful.

If the system has periodic boundaries, this returns boxes that are found upon wrapping, possibly returning the same box multiple times if it is within **radius** of **pos** in multiple ways. The amount of wrapping on each axis is returned in **wrap**, which needs to be entered as a non-NULL vector.

This function was rewritten for version 2.61 because it didn't correctly account for system volumes that didn't start at (0, 0, 0) and/or that didn't use periodic boundary conditions.

memory management

boxptr boxalloc(int dim, int nlist)

boxalloc allocates and minimally initializes a new **boxstruct**. Lists allocated are **indx**, which is size **dim**, and **maxmol**, **nmol**, and **mol**, each of which are size **nlist**. **nlist** may be entered as 0 to avoid allocating the latter lists. No molecule spaces are allocated.

int expandbox(boxptr bptr, int n, int ll)

Expands molecule list **ll** within box **bptr** by **n** spaces. If **n** is negative, the box is shrunk and any molecule pointers that no longer fit are simply left out. This function may be used if the initial list size (**bptr->maxmol[ll]**) was zero and can also be used to set the list size to zero. The book keeping elements of the box are updated. The function returns 0 if it was successful and 1 if there was not enough memory for the request.

int expandboxpanels(boxptr bptr, int n)

Expands the list of panels in box **bptr** by **n** spaces. If $n \leq 0$, this function ignores it, and does not shrink the box. This updates the **maxpanel** element. Returns 0 for success and 1 for insufficient memory.

void boxfree(boxptr bptr, int nlist)

Frees the box and all of its lists, although not the structures pointed to by the lists. **nlist** is the number of live lists.

boxptr *boxesalloc(int nbox, int dim, int nlist)

boxesalloc allocates and initializes an array of **nbox** boxes, including the boxes. **dim** is the system dimensionality and **nlist** is the number of live lists. There is no additional initialization beyond what is done in **boxalloc**.

void boxesfree(boxptr *blist, int nbox, int nlist)

Frees an array of boxes, including the boxes and the array. **nlist** is the number of live lists.

boxssptr boxssalloc(int dim)

Allocates and initializes a superstructure of boxes, including arrays for the **side**, **min**, and **size** members, although the boxes are not added to the structure, meaning that **blist** is set to NULL and **nbox** to 0.

void boxssfrees(boxssptr boxs)

Frees a box superstructure, including the boxes.

data structure output

5 C Code: structures and functions

void boxoutput(boxssptr boxs, int blo, int bhi, int dim)

This displays the details of virtual boxes in the box superstructure **boxs** that are numbered from **blo** to **bhi-1**. To continue to the end of the list, set **bhi** to -1. This requires the system dimensionality in **dim**.

void boxssoutput(simpstr sim)

Displays statistics about the box superstructure, including total number of boxes, number on each side, dimensions, and the minimum position. It also prints out the requested and actual numbers of molecules per box.

int checkboxparams(simpstr sim, int *warnptr)

Checks and displays warning about various box parameters such as molecules per box, box sizes, and number of panels in each box.

structure set up

void boxsetcondition(boxssptr boxs, enum StructCond cond, int upgrade)

Sets the box superstructure condition to **cond**, if appropriate. Set **upgrade** to 1 if this is an upgrade, to 0 if this is a downgrade, or to 2 to set the condition independent of its current value. If the condition is downgraded, this also downgrades the simulation structure condition.

int boxsetsize(simpstr sim, const char *info, double val)

Sets the requested box size. **info** is a string that is “molperbox” for the **mpbox** element, or is “boxsize” for the **boxsize** element, and **val** is the requested value. If the box superstructure has not been allocated yet, this allocates it. Returns 0 for success, 1 for failure to allocate memory, 2 for an illegal value, or 3 for the system dimensionality has not been set up yet.

int boxesupdateparams(simpstr sim)

Creates molecule lists for each box and sets both the box and molecule references to point to each other.

int boxesupdatelists(simpstr sim)

Sets up a superstructure of boxes, and puts some things in them boxes, including wall references. It sets up the box superstructure, then adds indices to each box, then adds the box neighbor list along with neighbor parameters, then adds wall references to each box. The function returns 0 for successful operation, 1 if it was unable to allocate sufficient memory, 2 if required things weren't set up yet. This function can be very computationally intensive.

int boxesupdate(simpstr sim)

Sets up or updates box data structures.

core simulation functions

boxptr line2nextbox(simpstr sim, double *pt1, double *pt2, boxptr bptr)

Given a line segment which is defined by the starting point **pt1** and the ending point **pt2**, and which is known to intersect the virtual box pointed to by **bptr**, this returns a pointer to the next box along the line segment. If the current box is also the final one, NULL is returned. Virtual boxes on the edge of the system extend to infinity beyond the system walls, so this function accurately tracks lines that are outside of the system volume.

5.8 Compartments (functions in *smolcompart.c*)

This function would be simple, except that it has to cope with a few fairly rare exceptions. In particular, if one of the points is exactly in line with a corner between boxes (**flag** equals 1 or 2), it has to deal with it correctly, which is what the **if(flag)** portion copes with. As mentioned above, boxes include their low sides, but not their high sides. Thus if the destination is towards a corner point that is not in the current box, then the line should first visit a direction that is increasing, followed by a direction that is decreasing; if multiple directions increase, the trajectory moves diagonally, and if all directions are decreasing (**flag** equals 2) then the trajectory also moves diagonally. This code is messy, but I think it works.

Near the end of the function is a line that checks if **crsmin==1.01**. This is true only if there is no next box, despite the fact that **pos2box**, in the first line, said that there was one. It can arise from round-off error.

```
int reassignmolecs(simptr sim, int diffusing, int reborn)
```

Reassigns molecules to boxes. If **diffusing** is 1, only molecules in lists that include diffusing molecules (**sim->mols->diffuselist**) are reassigned; otherwise all lists are reassigned. If **reborn** is 1, only molecules that are reborn, meaning with **indices** greater than or equal to **topl[11]**, are reassigned; otherwise, entire lists are reassigned. If **reborn** is 1, this assumes that all molecules that are in the system are also in a box, meaning that the **box** element of the molecule structure lists a box and that the **mol** list of that box lists the molecule. Molecules are arranged in boxes according to the location of the **pos** element of the molecules. Molecules outside the set of boxes are assigned to the nearest box. If more molecules belong in a box than actually fit, the number of spaces is doubled using **expandbox**. The function returns 0 unless memory could not be allocated by **expandbox**, in which case it fails and returns 1.

This function was modified on 1/15/16 so that if **reborn** is 1, then molecules are only moved if they are in the wrong places and if so, then they are removed from the old box and placed in the new box, but if **reborn** is 0, then all boxes are cleared out and molecules are assigned from scratch. This is a much faster routine than it was before, especially if boxes have a lot of molecules in them.

5.8 Compartments (functions in *smolcompart.c*)

Compartments are regions of volume that are bounded by surfaces. They do not include their bounding surfaces. They have no function in the performance of the simulation, but are useful for input, output, and communication with MOOSE. Compartments do not maintain a record of what they contain, but instead they define a set of rules that make it possible to test whether objects are inside or outside of the compartment. Compartments may be disjoint and they may overlap each other.

The inside of a compartment is defined to be all points from which one can draw a straight line to one of the “inside-defining points” without crossing any bounding surface. For example, to create a spherical compartment, one would define a spherical surface as the boundary and some point inside the sphere (the center, or any other internal point) to be the inside-defining point. In addition, compartments can be combined with previously defined compartments with logic arguments. Thus, for example, a cell cytoplasm compartment can be defined with the logic equation: equal to the cell compartment and not the nucleus compartment.

5.8.1 Data structures

5 C Code: structures and functions

```
enum CmpLogic {CLEqual, CLequalnot, CLand, CLor, CLxor, CLandnot, CLornot, CLnone};

typedef struct compartstruct {
    struct compartsuperstruct *cmptss; // compartment superstructure
    char *cname;                       // compart. name (reference, not owned)
    int nsrf;                          // number of bounding surfaces
    surfaceptr *surflist;              // list of bounding surfaces [s]
    int npts;                          // number of inside-defining points
    double **points;                   // list of inside-defining points [k][d]
    int ncmptl;                        // number of logic compartments
    struct compartstruct **cmptl;     // list of logic compartments [cl]
    enum CmpLogic *clsym;              // compartment logic symbol [cl]
    double volume;                     // volume of compartment
    int maxbox;                        // maximum number of boxes in compartment
    int nbox;                          // number of boxes inside compartment
    boxptr *boxlist;                   // list of boxes inside compartment [b]
    double *boxfrac;                   // fraction of box volume that's inside [b]
    double *cumboxvol;                 // cumulative cmpt. volume of boxes [b]
} *compartptr;
```

The volume of a compartment is initialized to 0, and is also reset to 0 whenever its definition changes. This indicates that it needs to be updated.

```
typedef struct compartsuperstruct {
    enum StructCond condition; // structure condition
    struct simstruct *sim;     // simulation structure
    int maxcmpt;               // maximum number of compartments
    int ncmpt;                 // actual number of compartments
    char **cnames;             // compartment names
    compartptr *cmptlist;      // list of compartments
} *compartssptr;
```

This structure contains information about all of the compartments. `condition` is the current condition of the superstructure and `sim` is a pointer to the simulation structure that owns this superstructure.

enumerated types

```
enum CmpLogic compartstring2cl(char *string)
```

Converts compartment logic symbol string to an enumerated compartment logic type. Input strings can be: “equal”, “equalnot”, “and”, “or”, “xor”, “andnot”, or “ornot”. Anything else results in `CLnone`.

```
char *compartcl2string(enum CmpLogic cls, char *string)
```

Converts enumerated compartment logic type to a string, in `string`, which must be pre-allocated. Output strings are “equal”, “equalnot”, “and”, “or”, “xor”, “andnot”, “ornot” or “none”. `string` is returned to allow for function nesting.

low level utilities

```
int posincompart(simptr sim, double *pos, compartptr cmpt, int useoldpos)
```

Tests if position `pos` is in compartment `cmpt`, returning 1 if so and 0 if not. This includes composed compartment logic tests. It does not use the compartment box list. This function

5.8 Compartments (functions in *smolcompart.c*)

is quite efficient for surfaces with few panels, but inefficient if surfaces have lots of panels. `useoldpos` tells the function to use the old surface panel position variables `oldpoint` and `oldfront` rather than the current ones.

`int compartrandpos(simptr sim, double *pos, compartptr cmpt)`

Returns a random position, in `pos`, within compartment `cmpt`. Returns 0 and a valid position, unless a point cannot be found, in which case this returns 1.

memory management

`compartptr compartalloc(void)`

Allocates memory for a compartment. All arrays are set to NULL, and not allocated. Returns the compartment or NULL if unable to allocate memory.

`void compartfree(compartptr cmpt)`

Frees a compartment, including all of its arrays.

`compartssptr compartssalloc(compartssptr cmptss, int maxcmpt)`

Allocates a compartment superstructure as well as `maxcmpt` compartments. Space is allocated and initialized for compartment names. Returns the compartment superstructure or NULL if unable to allocate memory. This function may be called multiple times in order to space for additional compartments. See `surfacessalloc`.

`void compartssfree(compartssptr cmptss)`

Frees a compartment superstructure, including all compartments and everything within them.

data structure output

`void compartoutput(simptr sim)`

Displays all important information about all compartments to stdout.

`void writecomparts(simptr sim, FILE *fptr)`

Prints information about all compartments to file `fptr` using a format that allows the compartments to be read as a configuration file.

`void checkcompartparams(simptr sim)`

This checks a few compartment parameters.

structure set up

`void compartsetcondition(compartssptr cmptss, enum StructCond cond, int upgrade)`

Sets the compartment superstructure condition to `cond`, if appropriate. Set `upgrade` to 1 if this is an upgrade, to 0 if this is a downgrade, or to 2 to set the condition independent of its current value. If the condition is downgraded, this also downgrades the simulation structure condition.

`int compartenablecomparts(simptr sim, int maxcmpt)`

Enables compartments in the simulation by allocating the compartment superstructure to hold `maxcmpt` compartments and setting the necessary condition state. This function may be called multiple times. Returns 0 for success or 1 if memory could not be allocated. Function is analogous to `surfenablesurfaces`.

5 C Code: structures and functions

compartptr compartaddcompart(simptr sim, char *cmptname)

Adds a compartment named **cmptname** to **sim**. This allocates all necessary memory, including the superstructure if needed. Returns a pointer to the compartment for success or NULL for failure. Function is analogous to **surfaddsurface**.

int compartaddsurf(compartptr cmpt, surfaceptr srf)

Adds surface **srf** to the compartment **cmpt**. This increments **nsrf** and appends the surface to **srflist**. Returns 0 for success, 1 if memory could not be allocated, and 2 if the surface was already in the list (in which case it is not added again).

int compartaddpoint(compartptr cmpt, int dim, double *point)

Adds point **point** to the compartment **cmpt**, in a **dim** dimensional system. This increments **npts** and appends the point to **points**. Returns 0 for success and 1 if memory could not be allocated.

int compartaddcmpt1(compartptr cmpt, compartptr cmpt1, enum CmpLogic sym)

Add logically composed compartment **cmpt1**, which is composed with symbol **sym**, to the compartment **cmpt**. This increments **ncmpt1** and appends the new logic compartment to **cmpt1**. Returns 0 for success, 1 if memory could not be allocated, or 2 if **cmpt** and **cmpt1** are the same, which is not allowed.

int compartupdatebox(simptr sim, compartptr cmpt, boxptr bptr, double volfrac)

Updates the listing of box **bptr** in compartment **cmpt**, according to the rule that boxes should be listed if any portion of them is within the compartment and should not be listed if no portion is within the compartment. This also updates the **cumboxvol** and volume structure elements as needed. If the fraction of the box within the compartment is known, including 0, enter it in **volfrac**. If it is unknown and should be calculated, enter -1 in **volfrac**. If the fraction is unknown and should be unchanged if the box was already in the compartment and calculated if the box wasn't in the compartment, then enter -2 in **volfrac**. This returns 0 for no change, 1 for box successfully added, 2 for box successfully removed, 3 for box was already listed but volume was updated, or -1 for failure to allocate memory. If the volume of the box within the compartment needs to be calculated, this calculates it with a hard-coded value of 100 random trial points. Memory is allocated as needed.

The following table lists the return values, which is useful for understanding them and for reading through the function. The former value for each pair is for the actual volume fraction, in **volfrac2**, equal to 0 and the latter is for the actual volume fraction >0.

		value of volfrac		
		-2	-1	0 to 1
bptr was in cmpt	yes	0 / 0	2 / 0, 3	2 / 0, 3
(bc <= cmpt -> nbox)	no	0 / 1	0 / 1	0 / 1

compartptr compartreadstring(simptr sim, compartptr cmpt, char *word, char *line2, char *erstr)

Reads and processes one line of text from the configuration file, or some other source, for the compartment **cmpt**. If the compartment is not known, then set **cmpt** to NULL. The first word of the line should be sent in as **word** and the rest sent in as **line2**. If this function is successful, it returns the compartment and it does not change the contents of **erstr**; if not, it returns NULL and writes an error message to **erstr**.

```
int loadcompart(simptr sim, ParseFilePtr *pfp, line2, char *erstr)
```

Loads a compartment, or information for an already existing compartment, from an already opened configuration file. This is used to fill in basic compartment details. However, it does not address any of the box information. Returns 0 for success and 1 for an error; error messages are returned in *erstr*.

```
int compartmentsupdateparams(simptr sim)
```

Sets up the boxes and volumes portions of all compartments. Returns 0 for success, 1 for inability to allocate sufficient memory, or 2 for boxes not set up before compartments. This function may be run during initial setup, or at any time afterwards. It is computationally intensive.

```
int compartmentsupdatelists(simptr sim)
```

Does nothing. This function is here for future expansion, and to keep similarity between different Smoldyn modules.

```
int compartmentsupdate(simptr sim)
```

Sets up or updates all portions of compartment data structures.

core simulation functions

```
void comparttranslate(simptr sim, compartptr cmpt, int code, double *translate)
```

Translates compartment *cmpt* by the displacement given in *translate*. The different bits of *code* tell which attributes of the compartment should be translated. If *code*&1, then the bounding surfaces of the compartment (that are listed in the compartment definition, not including those that are implied through the logic statements) are translated. If *code*&2, then the molecules that are bound to the bounding surfaces of the compartment (which are listed in the compartment definition) are translated. If *code*&4, then the molecules that are inside the compartment are translated. If *code*&8, then molecules that get bumped into by the moving surfaces get translated.

This has a number of flaws with the external molecules. The algorithm is that any molecule that would get bumped into gets translated by the same amount as the compartment, except if its surface action is “transmit” for both surface faces. After being translated, the function checks for its collisions with any surfaces, dealing with them as required. One problem is that this ignores most surface actions, except only for a slight check for transmit vs. anything else. Another problem is that molecules will end up on the wrong side of the moving surface if they get squeezed between the moving surface and a static surface. Here, what happens is the molecule is moved because of the moving surface, then it bounces off of the static surface back towards where it came from, and then the moving surface moves over it. In some ways, this is unavoidable because there is no good solution for what to do with squeezed molecules. However, it seems that it could be improved.

5.9 Ports (functions in *smolport.c*)

Ports are data structures for importing and exporting molecules between a Smoldyn simulation and another simulation. In particular, they are designed for the incorporation of Smoldyn into MOOSE, but they could also be used to connect multiple Smoldyn simulations or for other connections.

A port is essentially a surface and a buffer. Smoldyn molecules that hit the porting surface are put into the buffer for export. Alternatively, molecules may be added to the Smoldyn simulation

5 C Code: structures and functions

at the porting surface by other programs. To perform porting within Smoldyn, a command called `transport` will move molecules from one port to another. This can be used to test porting, or, hopefully, to transport molecules between multiple Smoldyn simulations.

As much as possible, the code for ports is very analogous to the code for compartments.

5.9.1 Data structures

```
typedef struct portstruct {
    struct portsuperstruct *portss; // port superstructure
    char *portname;                 // port name (reference, not owned)
    surfaceptr *srf;                // porting surface (ref.)
    enum PanelFace face;            // active face of porting surface
    int llport;                     // live list number for buffer
} *portptr;

typedef struct portsuperstruct {
    enum StructCond condition; // structure condition
    struct simstruct *sim;     // simulation structure
    int maxport;               // maximum number of ports
    int nport;                 // actual number of ports
    char **portnames;          // port names
    portptr *portlist;         // list of ports
} *portsspstr;
```

This structure contains information about all ports. `condition` is the current condition of the superstructure and `sim` is a pointer to the simulation structure that owns this superstructure.

memory management

`portptr portalloc(void)`

Allocates memory for a port. Pointers are set to `NULL` and `llport` is set to `-1`. Returns the port or `NULL` if unable to allocate memory.

`void portfree(portptr port)`

Frees a port.

`portsspstr portssalloc(portsspstr portss, int maxport)`

Allocates a port superstructure, if needed, as well as `maxport` ports. Space is allocated and initialized for port names and port lists. Returns the port superstructure or `NULL` if unable to allocate memory. This function can be called repeatedly to expand the number of ports.

`void portssfree(portsspstr portss)`

Frees a port superstructure, including all ports.

data structure output

`void portoutput(simpstr sim)`

Displays all important information about all ports to `stdout`.

`void writeports(simpstr sim, FILE *fpstr)`

Prints information about all ports to file `fpstr` using a format that allows the ports to read as a configuration file.

```
int checkportparams(simptr sim, int *warnptr)
```

This checks a few port parameters.

structure set up

```
void portsetcondition(portssptr portss, enum StructCond cond, int upgrade)
```

Sets the port superstructure condition to `cond`, if appropriate. Set `upgrade` to 1 if this is an upgrade, to 0 if this is a downgrade, or to 2 to set the condition independent of its current value. If the condition is downgraded, this also downgrades the simulation structure condition.

```
int portenableports(simptr sim, int maxport)
```

Enables ports in the simulation by allocating the port superstructure to hold `maxport` ports and setting the necessary condition state. This function may be called multiple times. Returns 0 for success or 1 if memory could not be allocated. Function is analogous to `surfenablesurfaces`.

```
portptr portaddport(simptr sim, char *portname, surfaceptr srf, enum PanelFace face)
```

Adds, or updates, a port to the port superstructure. If `portname` is not the name of an existing port, a new port is defined, the `nport` element of the superstructure is incremented, and this name is copied over for the new port. Alternatively, if `portname` has already been defined, it is used to reference an existing port. Either way, the port surface is set to `srf` if `srf` is not NULL, the port face is set to `face` if `face` is not `PFnone`, and the address of the port is returned.

```
portptr portreadstring(simptr sim, portptr port, char *word, char *line2, char *erstr)
```

Reads and processes one line of text from the configuration file, or some other source, for the port `port`. If the port is not known, set `port` to NULL. The first word of the line should be sent in as `word` and the rest sent in as `line2`. If this function is successful, it returns the port and it does not change the contents of `erstr`; if not, it returns NULL and it writes an error message to `erstr`.

```
int loadport(simptr sim, ParseFilePtr *pfpptr, char* line2, char *erstr)
```

Loads a port, or information for an already existing port, from an already opened configuration file. This is used to fill in basic port details. However, it does not assign a molecule buffer to the port. Returns 0 for success and 1 for an error; error messages are returned in `erstr`.

```
int portsupdateparams(simptr sim)
```

Does nothing. This function is here for future expansion, and to maintain similarity between different Smoldyn modules.

```
int portsupdatelists(simptr sim)
```

Sets up the molecule buffers for all ports. Returns 0 for success, 1 for inability to allocate sufficient memory, or 2 for molecules not set up sufficiently.

```
int portsupdate(simptr sim)
```

Sets up or updates all port data structure components.

core simulation functions

```
int portgetmols(simptr sim, portptr port, int ident, enum MolecState ms, int remove)
```

Returns the number of molecules of type `ident` (use -1 for all species) and state `ms` (`MSall` is allowed) that are in the export buffer of port `port`. If `remove` is 1, this kills those molecules, so that they will be returned to the dead list at the next sorting; otherwise they are left in the port. The intention is that molecules that are gotten from the export list with this function are then added to MOOSE or another simulator.

```
int portputmols(simptr sim, portptr port, int nmol, int ident, int *species,  
double **positions, double **positionsx)
```

Adds `nmol` molecules of state `MSsoln` to the simulation system at the porting surface of port `port`. If `species` is non-NULL, then it needs to be an `nmol` length list of species numbers, which are used for the molecule species and `ident` is ignored; alternatively, if `species` is NULL, then all molecules will have species `ident`. Likewise, if `positions` is non-NULL, then it needs to be a list of molecule positions that will be used for the new molecule positions. If `positions` is NULL, then molecules are placed randomly on the porting surface. If `positionsx` is NULL, then molecule old positions are fixed to the panel, and otherwise molecule old positions are set to these values. This returns 0 for success, 1 for insufficient available molecules, 2 for no porting surface defined, 3 for no porting surface face defined, and 4 for no panels on porting surface.

```
int porttransport(simptr sim1, portptr port1, simptr sim2, portptr port2)
```

Transports molecules from `port1` of simulation structure `sim1` to `port2` of simulation structure `sim2`. `sim1` and `sim2` may be the same and `port1` and `port2` may be the same. This is designed for testing ports or for coupled Smoldyn simulations that communicate with ports.

5.10 Lattices (functions in smollattice.c)

A lattice is a region of space in which molecules do not have precise spatial locations, but are compartmentalized to lattice subvolumes. At present, lattices use axis-aligned rectangular subvolumes, but they could at some point be unstructured meshes (as in URDME). In the lattice region of space, molecules are currently represented as discrete individuals, in each subvolume. At some point though, this could be extended to continuous concentrations, to enable PDE simulation methods. Lattice regions of space interface to particle-based regions of space through Smoldyn's ports.

Martin Robinson added lattice functionality to Smoldyn during 2013. This documentation is my best understanding of his code, which I don't promise is correct. The lattice code is in several subdirectories of `source`. The `NextSubVolume` directory contains Martin's code for the next subvolume simulation method. The `vtk` directory contains a couple of VTK wrapper code files, which I assume are used to interface Martin's code to the VTK visualization software. I suspect that Martin wrote those wrapper files. Finally, the lattice code is integrated into the rest of Smoldyn using the `smollattice.c` file, which compiles with the rest of Smoldyn. The code in `smollattice.c` is always compiled, regardless of compiling options, but if the `LATTICE` configuration variable is undefined, then all calls to the `nsv` code are disabled.

5.10.1 Data structures

```

enum LatticeType {LATTICEnone, LATTICEnsv, LATTICEpde};

typedef struct latticestruct {
    struct latticesuperstruct *latticess; // lattice superstructure
    char *latticename; // lattice name (reference, not owned)
    enum LatticeType type; // type of lattice
    double min[DIMMAX]; // lower spatial boundaries
    double max[DIMMAX]; // upper spatial boundaries
    double dx[DIMMAX]; // lattice lengthscale (subvolume width)
    char btype[DIMMAX]; // boundary type (r)eflective or (p)eriodic
    portptr port; // interface port (ref.)
    int **convert; // convert to particle at port, 0 or 1 [lat.species][face]
    // ??
    int maxreactions; // maximum number of reactions
    int nreactions; // number of reactions
    rxnptr *reactionlist; // list of reactions
    int *reactionmove; // 0 or 1 for moving reactions
    maxsurfaces??
    nsurfaces??
    surfacelist??
    int maxspecies; // maximum number of species
    int nspecies; // number of species
    int *species_index; // species indices
    int *maxmols; // allocated size of molecule list [lat.species]
    int *nmols; // number of individual molecules [lat.species]
    double*** mol_positions; // molecule positions [lat.species][nmols][dim]
    NextSubvolumeMethod* nsv; // nsv class
    NextSubvolumeMethod* pde; // pde class
} *latticeptr;

typedef struct latticesuperstruct {
    enum StructCond condition; // structure condition
    struct simstruct *sim; // simulation structure
    int maxlattice; // maximum number of lattices
    int nlattice; // actual number of lattices
    char **latticenames; // lattice names
    latticeptr *latticelist; // list of lattices
} *latticessptr;

```

These structures contain information about all lattices. Starting with the lattice superstructure, `condition` is the current condition of the superstructure and `sim` is a pointer to the simulation structure that owns this superstructure. These, and other superstructure elements are completely standard. Space is allocated for `maxlattice` lattices, of which `nlattice` are actually defined and used. Their names are in `latticenames` and they are pointed to by pointers in `latticelist`.

In the lattice structure, of type `latticestruct`, the `latticess` pointer points to the owning superstructure and `latticename` points to the lattice name. Each lattice is either type `LATTICEnsv` if it is for discrete particles or `LATTICEpde` if it is for continuous concentrations. There is also the enumeration `LATTICEnone`, which is helpful for specifying neither of the other options but is not valid for actual lattices. `min` and `max` vectors represent the lower and upper spatial boundaries of the lattice space; this space is divided into subvolumes that each have width given with the `dx` vector. `btype` is a vector with a character for each coordinate that equals 'r' for reflective boundaries, 'p'

5 C Code: structures and functions

for periodic boundaries, and ‘u’ for undefined.

The `port` element points to the port that divides the lattice part of space from the particle-based part of space. The `maxreactions` and `nreactions` elements tells how many reactions are allocated and are being used here. These reactions, listed in `reactionlist` are simply pointers to Smoldyn’s normal reactions. This means that there is no need for additional reaction data structures here. `reactionmove` is either 0 for 1 for each reaction, where 1 means that it should be moved and 0 means that it should not be moved; a moved reaction is only implemented on the lattice side of space and is set to a rate of 0 on the particle side of space. `nspecies` is the number of species that are declared for the lattice region of space. It may be different from `nspecies` values in the particle side of space, defined in the rest of Smoldyn. `species_index` is a list of species indices from the particle side of space; it is a look-up table that connects lattice species numbers to particle side species numbers. `maxmols` is the allocated size of the molecule list and `nmols` is the number of molecules for each species, indexed with the lattice-side indices. `mol_positions` are the positions of the molecules, indexed as the lattice-side species number, molecule number, and dimensional coordinate. I don’t really understand this because I didn’t think these molecules had precise positions. The `nsv` and `pde` pointers point to classes in the core NSV and PDE code (the PDE code isn’t included yet).

5.10.2 Functions

memory management

`latticeptr latticealloc(int dim)`

Allocates memory for a lattice assuming a `dim` dimensional system. Pointers are set to NULL and values are set to 0 or other defaults. Returns the lattice or NULL if unable to allocate memory.

`void latticefree(latticeptr lattice)`

Frees a lattice and all of the memory in it.

`int latticeexpandreactions(latticeptr lattice, int maxrxns)`

Expands the number of reactions in a lattice to `maxrxns`, allocating new memory and freeing old memory as needed. This addresses the memory in `reactionlist` and `reactionmove`. Returns 0 for success or 1 for failure.

`int latticeexpandspecies(latticeptr lattice, int maxspecies)`

Expands the number of species in a lattice to `maxspecies`, allocating new memory and freeing old memory as needed. This addresses the memory in `species_index` and `nmols`. It also addresses the memory in `mol_positions`. Returns 0 for success or 1 for failure.

`int latticeexpandmols(latticeptr lattice, int species, int maxmols, int dim)`

Expands the `mol_positions` list for the lattice species number `species` so that it will have `maxmols` spaces in it, each with `dim` values.

`latticessptr latticessalloc(latticessptr latticess, int maxlattice)`

Allocates a lattice superstructure, if needed, as well as `maxlattice` lattices. Space is allocated and initialized for lattice names and lattice lists. Returns the lattice superstructure or NULL if unable to allocate memory. This function can be called repeatedly to expand the number of lattices.

```
void latticesfree(latticesptr lattices)
    Free a lattice superstructure and all of its contents.
```

data structure output

```
void latticeoutput(simpstr sim)
    Outputs the contents of the lattice superstructure and all of the member lattices to the display.
    This shows essentially everything except for the molecule positions.
```

```
void writelattices(simpstr sim, FILE *fptr)
    Writes the contents of the lattice superstructure and all of the member lattices to file fptr
    in a format that enables it to be read in again.
```

```
int checklatticeparams(simpstr sim, int *warnptr)
    Checks the lattice parameters, sending output to the simLog function along with the
    appropriate error level.
```

structure set up

```
void latticesetcondition(latticesptr lattices, enum StructCond cond, int upgrade)
```

Sets the lattice superstructure condition to *cond*, if appropriate. Set *upgrade* to 1 if this is an upgrade, to 0 if this is a downgrade, or to 2 to set the condition independent of its current value. If the condition is downgraded, this also downgrades the simulation structure condition.

```
int latticeenablelattices(simpstr sim)
    Enables simulation with lattices. This function may be called more than once but doesn't do
    anything on repeat calls. This allocates the superstructure and 1 lattice. It also initializes the
    NSV code.
```

```
latticeptr latticeaddlattice(simpstr sim, latticeptr *latptr, const char
    *latticename, const double *min, const double *max, const double *dx, const
    char *btype, enum LatticeType type)
    Adds a new lattice or modifies an existing lattice. This enables lattices, creates the
    superstructure, and allocates memory as needed. Send in latptr with a pointer to the
    lattice if it has already been created and now needs to be modified, or NULL if a new lattice
    is wanted. In the latter case, this pointer is returned with the address of the new lattice.
    All of the lattice parameters that can be set here are optional, meaning that NULL can be
    entered to choose to not set the parameter. If entered, the lattice name is set to latticename,
    the minimum boundary coordinates to min, the maximum boundary coordinates to max, the
    lattice length parameters to dx, the boundary types to btype, and the lattice type to type.
    In the last case, enter LATTICENone to not set this parameter. Returns 0 for success or 1 for
    inability to allocate memory.
```

```
int latticeaddspecies(latticeptr lattice, int ident, int *index)
    Adds one or more Smoldyn species to an existing lattice, allocating new memory as needed.
    If index is NULL, this sets the species_index element of the lattice structure to the ident
    value that is entered (the index of the Smoldyn species) and this also sets the number of
    molecules of this species on the lattice to 0. Returns 0 for success, 1 for inability to allocate
    memory, or 2 if this species was already in the lattice. If index is not NULL, then this adds
```

5 C Code: structures and functions

each of the listed species to the lattice, returning 0 for success or 1 for inability to allocate memory; in this case any duplicate species are just ignored.

`int latticeaddrxn(latticeptr lattice, rxnptr reaction, int move)`

Adds a reaction to an existing lattice, allocating new memory as needed. This puts the reaction at the end of the current reaction list and sets its “move” status to the `move` value that is entered. Returns 0 for success, 1 for inability to allocate memory, or 2 if the reaction was already in the lattice.

`int latticeaddmols(latticeptr lattice, int nmol, int i, double *poslo, double *poshi, int dim)`

Adds `nmol` molecules of Smoldyn species `i` to lattice `lattice`. These molecules are randomly positioned in the axis-aligned bounding box defined by the `poslo` and `poshi` vectors. The space has `dim` dimensions. This allocates memory as needed. This also adds the species index `i` to the lattice if it’s not already there. Returns 0 for success or 1 for failure to allocate memory.

`void latticeaddport(latticeptr lattice, portptr port)`

Sets the port for a lattice. A lattice only works with a single port, so there’s no need to allocate memory, check for prior status, etc. Instead, this simply sets the port variable.

`int latticeaddconvert(latticeptr lattice, int ident, int *index, enum PanelFace face, int convert)`

Sets the lattice `convert` element for the single Smoldyn species `ident`, or for the group of Smoldyn species in `index`, and the face `face` to the value in `convert`. This is for lattice-molecules that collide with the given face side of the port, to indicate if they should convert to particle molecules or stay as lattice molecules. Returns 0 for success or the identity of the Smoldyn species if the Smoldyn species has not been imported to the lattice.

`latticeptr latticereadstring(simpstr sim, ParseFilePtr`

`pfp, latticeptr lattice, const char *word, char *line2)` This reads and processes user input, much like `portreadstring` and similarly named functions. As always, `sim` is the simulation pointer, `pfp` is the parse file pointer, which contains the file information, `lattice` is the pointer to the current lattice being worked on or NULL if there is no current lattice, `word` is the first word of the input string, and `line2` is the remainder of the input string. This returns a pointer to the current lattice being worked on if successful and NULL if there is an error; in the latter case, the error is first sent to `simParseError`.

`int loadlattice(simpstr sim, ParseFilePtr *pfp, char* line2)` Loads a lattice using input from the user. This function is called by `loadsims` when file parsing reaches “start_lattice”. This is essentially identical in structure to `loadport`.

`int latticesupdateparams(simpstr sim)`

Updates the parameters of all lattices, meaning that it sends various parameters from the lattice structures to the NSV code. This sends the port details, including the porting rectangle panel locations, dimensions, and orientations. This also sends the list of reactions to the NSV code. Returns 0 for success or 1 for inability to allocate memory.

`int latticesupdatelists(simpstr sim)`

Updates the more fundamental parameters of all lattices, meaning that it sends them from the lattice structures to the NSV code. First, this deletes any existing nsv data. Then, this

5.10 Lattices (functions in *smollattice.c*)

sends the lattice dimensions, the species, and the numbers of molecules of each species. This deletes the local molecule data afterwards.

int latticesupdate(simptr sim)

Updates the lattice superstructure and all member lattices as needed, while updating the condition as appropriate. This function is essentially identical to ones with similar names in other code modules.

core simulation functions

int latticeruntimestep(simptr sim)

Runs the lattice NSV code for one simulation time step for each lattice. Returns 0.

NSV functions, in nsvc.cpp

void nsv_init()

Initializes the Kairos simulation engine.

NextSubvolumeMethod* nsv_new(double* min, double* max, double* dx, int n)

Allocates and initializes a new nsv type data structure, and then returns it. **min** and **max** are the low and high corners of space, **dx** is the lattice spacing on each coordinate, and **n** is the lattice dimensionality.

void nsv_delete(NextSubvolumeMethod* nsv)

Deletes an nsv type data structure.

void nsv_print(NextSubvolumeMethod* nsv, char bufferptr)**

Prints out information about the nsv data structure to the buffer, which is simply a pointer to an unallocated string. The buffer needs to be freed afterwards.

void nsv_add_interface(NextSubvolumeMethod* nsv, int id, double dt, double *start, double *end, double *norm, int dim)

Adds information about a porting surface panel to the lattice **nsv**. This panel has to be a rectangle shape. It is used for conversion of a particular species at this surface. **id** is the Smoldyn species index (not the lattice species index) of the species that should be ported at the interface. **dt** is the Smoldyn time step (same for the lattice). **start** is a **dim**-dimensional list of the rectangle low-value coordinates and **end** is a list of the rectangle high-value coordinates. **norm** is the **dim**-dimensional surface normal vector that points towards the particle side of the interface and away from the lattice-side of the interface. Finally, **dim** is the simulation dimensionality.

void nsv_add_species(NextSubvolumeMethod* nsv, int id, double D, char *btype, int dim)

Adds a new species to the lattice **nsv**. **id** is the Smoldyn species number and **D** is the solution-phase diffusion coefficient of this species. **btype** is a vector with the boundary type on each axis, where the options are 'r' for reflective and 'p' for periodic. **dim** is the lattice dimensionality.

extern void nsv_add_surface(NextSubvolumeMethod* nsv, surfacestruct* surface)

Adds a surface to the lattice **nsv**. **surface** is a pointer to a Smoldyn surface that should be added to the lattice.

5 C Code: structures and functions

`void nsv_add_reaction(NextSubvolumeMethod* nsv, rxnstruct *reaction)`

Adds a reaction to the lattice `nsv`. `reaction` is a pointer to a Smoldyn reaction that should be added to the lattice.

`void nsv_integrate(NextSubvolumeMethod* nsv, double dt, portstruct *port, latticestruct *lattice)`

Runs the lattice simulation over `dt` amount of time. `nsv` is the lattice, `port` is the port that is between the lattice and the Smoldyn simulation, and `lattice` is the Smoldyn lattice data structure.

`void nsv_add_mol(NextSubvolumeMethod* nsv, int id, double* pos, int dim)`

Adds one molecule to lattice `nsv`. The molecule has Smoldyn species `id` and `dim`-dimensional location vector `pos`. `dim` is the lattice dimensionality.

`int nsv_get_species_copy_numbers(NextSubvolumeMethod* nsv, int id, const int **copy_numbers, const double** positions)`

Gets the copy number of a species in the `nsv` data structure. Enter `id` as the Smoldyn species number (not the lattice species number), `copy_numbers` as a pointer to an unallocated integer array and `positions` as an unallocated double array. Both of these need to be freed afterwards.

5.11 Filaments (functions in `smolfilament.c`)

Filaments are 1-dimensional structures that are designed to represent cytoskeletal filaments, such as actin or microtubules, or DNA. They are intended to be much longer than they are wide.

The functions here are designed for filaments in either 2D or 3D space. However, the vectors and matrices are allocated and set up for 3D space, regardless of the actual space, with the z -coordinate set up so that the filament always stays within the x, y plane; for example, a direction cosine matrix for 2D space is still 9 elements, but the elements involving z coordinates are all set to either 0 or 1 as appropriate. Despite this, most code treats 2D and 3D separately.

5.11.1 Data structures declared in `smoldyn.h`

Filaments are defined in a hierarchical fashion. At the top level, there is one filament superstructure which contains a list of filament types that are used in the simulation. Below that are multiple filament type structures. For example, actin in the cytoplasm might be one type; another might be double-stranded DNA, and yet another is single-stranded DNA. Each filament type contains multiple individual filaments, which is the next level down, each of which looks and behaves identically. Each filament is composed of a list of segments. Segments can vary in lengths and thicknesses. Each filament also contains a data structure of “working data” that is used for dynamics integration.

Enumerations

```
enum FilamentDynamics {FDnone, FDeuler, FDRK2, FDRK4, FDeulermat,
    FDimplicit};
```

The filament dynamics enumeration describes what algorithms to use for its dynamics. They are summarized below.

Dynamics	Description
None	filament is rigid and does not move
Euler	Euler ODE integration
RK2	2nd order Runge-Kutta integration
RK4	4th order Runge-Kutta integration
EulerMat	Euler integration with matrix method
ImplicitOld	Euler implicit integration with original method
Implicit	Euler implicit integration

Segments

```
typedef struct segmentstruct {
    struct filamentstruct* fil;           // owning filament
    int index;                           // self index along filament
    double *xyzfront;                    // Coords. for segment front
    double *xyzback;                     // Coords. for segment back
    double len;                          // segment length
    double thk;                          // thickness of segment
    double ypr[3];                       // relative ypr angles
    double qrel[4];                      // relative rotation quaternion
    double qabs[4];                      // absolute rotation quaternion
} * segmentptr;
```

Segments start by listing the filament that owns them. Then they have their own position within the filament in `index`; a pointer to the same segment would be found with `segment->fil->segments[segment->index]` and the result would be identical to `segment`.

`xyzfront` is the x, y, z coordinate of the segment starting point and `xyzback` is the x, y, z coordinate of the segment end point. These vectors are not owned by the segment but point to vectors in the `nodes` element of the owning filament. The back of one segment is the identical vector (same memory address) as the front of the next segment. `len` is the segment length and `thk` is the segment thickness.

`ypr` is the relative yaw, pitch, and roll angle for the segment relative to the orientation of the prior segment. It is largely ignored for the first segment, whose orientation is instead given with `seg0dcm` in the filament data structure. `qrel` and `qabs` are rotation quaternions that express the segment's orientation, relative to the prior segment and in the system coordinate system respectively. Identical information is contained in `ypr`, `qrel`, and `qabs`, but the redundancy is here for faster computation.

For 2D simulations, only the first 2 values are used in the coordinates vectors (`xyzfront` and `xyzback`) and the third coordinate should always equal 0; only the first value (yaw) is used in the `ypr` angles and the other two should always equal 0; and only the first and last elements are used in the quaternions.

Segment lengths can vary but standard lengths are constant over an entire filament. This means that segments all have roughly equal lengths. I initially thought that it could be useful to have large standard length variations to focus computation on more important regions, but this is sufficiently hard to program that it's not worth it.

Filament working data

```
typedef struct filworkstruct {
    struct filamentstruct *fil;           // owning filament
```

5 C Code: structures and functions

```

double ***wnodes;           // working nodes (3,nseg+1,3)
double **wroll;             // working rolls (3,nseg)
double **flatnodes;         // flattened nodes (2,(2or4)*nseg)
double **flatforces;        // flattened forces (2,(2or4)*nseg)
double **wseg0up;           // seg. 0 up vector (3,3)
double **forces;            // forces on nodes (nseg+1)
double *torques;            // list of segment torques (nseg)
sparsematrix forcemat;      // force matrix for implicit int.
double thermtime;           // time for thermal forces
double **thermforce;        // thermal forces on nodes
} *filworkptr

```

Filament dynamics integration requires a substantial amount of working data, which are stored here. None of these data define the filament state at the end of a time step, but this is only scratch space for use by the dynamics functions. `fil` points to the filament that owns this data structure. `wnodes` is 3 lists of nodes; number 0 is not owned by this data structure, but is simply a pointer to `fil->nodes`, whereas numbers 1 and 2 are owned by this data structure. `wroll` is 3 lists of segment roll values; the first points to `fil->roll` and the other two are owned by this data structure. `flatnodes` is 2 lists of flattened node and roll data; for 2D, the list is $[x_{0x}, x_{0y}, x_{1x}, x_{1y}, \dots, x_{nx}, x_{ny}]$ and for 3D the list is $[x_{0x}, x_{0y}, x_{0z}, \tau_0, x_{1x}, x_{1y}, x_{1z}, \tau_1, \dots, x_{nx}, x_{ny}, x_{nz}]$. `flatforces` is identical to `flatnodes` but is for forces instead of node positions. It has the same layout and same vector sizes. `wseg0up` is 3 lists; number 0 is not owned here but points to `fil->seg0up`, and the other two are working copies to represent the segment 0 up vector. `forces` is used to store the cumulative forces on each node. Similarly, the `torques` element is used to store the roll torque acting on each segment (along its axis). `forcemat` is a flattened matrix of forces that is used for matrix-method Euler integration and implicit integration. `thermtime` is the simulation time for the most recent calculation of thermal forces (they need to be computed only once per time step to avoid giving the integrators conflicting information) and `thermforce` is a list of thermal forces on the nodes.

Filaments

```

typedef struct filamentstruct {
    struct filamenttypestruct* filtype; // owning filament type
    char* filename;                     // filament name (ref, not owned)
    int maxseg;                          // number of segments allocated
    int nseg;                            // number of segments
    segmentptr* segments;               // array of segments (nseg)
    double **nodes;                     // list of nodes (nseg+1)
    struct filamentworkstruct *filwork; // working data for dynamics
    double **nodesx;                    // old node locations (nseg+1)
    double *roll;                       // segment roll values (nseg)
    double *nodemobility;               // extra node mobility (nseg+1)
    double seg0up[3];                   // segment 0 up vector
    struct filamentstruct* frontend;    // what front attaches to
    struct filamentstruct* backend;     // what back attaches to
    int maxbranch;                      // max branches off this filament
    int nbranch;                        // num branches off this filament
    int* branchspots;                   // segments where branches are
    struct filamentstruct** branches;   // list of branching filaments
    int maxsequence;                    // allocated sequence characters
    int nsequence;                      // number of sequence characters
    char* sequence;                     // sequence code
}

```

```
} * filamentptr;
```

Each filament data structure represents a single unbranched chain of segments. Multiple filaments of the same type (e.g. two microtubules) are stored in multiple data structures. In the filament data structure, **filtype** points to the filament type that owns this filament and **filname** points to the name of this filament; names are assigned by default if not assigned by the user.

Next, **maxseg** and **nseg** give the allocated and actual numbers of segments. The segments, in **segments**, are an array, where the index of the front is equal to 0. The segment endpoints are stored in **nodes**, which is allocated to be **maxseg+1** long by 3 elements wide for the *x*, *y*, and *z* values (regardless of system dimension). The **filwork** element points to a data structure for dynamics integration data.

The **nodesx** vector contains node positions at the end of the previous time step, enabling determination of surface crossing or other events during the time step. The **roll** vector contains the roll values for individual segments. It is only here for convenience, and I'm not sure it's a good design, but it's here for now. It is redundant with the segment **ypr[2]** values, which are the definitive ones, so any function that writes to **roll** must also write to **ypr[2]** and any function that reads from **roll** must first make sure that it equals **ypr[2]**. The reason for redundant data is that it can be easier to work with either a list of roll values or lists of segment rotation values, depending on the situation. The **nodemobility** vector is a list of unitless mobility factors for each node, which are multiplied by the filament type mobility value to give the node's total mobility; this is useful for fixing particular nodes in space. The **seg0up** vector represents the unit "up" vector for segment 0 (which is perpendicular to the segment itself).

Filaments are allowed to branch. If the current filament is a branch off another filament, then **frontend** points to the filament that the front end attaches to and/or **backend** points to the filament that the back end attaches to; these are **NULL** if they are free ends. Additionally, this filament knows about the filaments that branch off this one, which are listed in **maxbranch**, **nbranch**, and **branches**. The branched filaments branch off at the locations listed in **branchspots**.

Along the length of each filament are equally spaced sequence codes, which represent things like DNA sequence or modification state of proteins. They are stored in an array with parameters **maxsequence** and **nsequence**, and then the list **sequence**. This array is analogous to the segments array, but may be larger or smaller because there's no necessary correlation between sequence code and segments.

Filament types

```
typedef struct filamenttypestruct {
    struct filamentsuperstruct* filss; // filament superstructure
    char* ftype;                       // fil type name (ref, not owned)
    enum FilamentDynamics dynamics;    // dynamics for the filament
    double bundlevalue;                // microfilaments in bundle
    double color[4];                   // filament color
    double edgethick;                  // thickness of edge for drawing
    unsigned int edgestipple[2];       // edge stippling [factor, pattern]
    enum DrawMode drawmode;            // polygon drawing mode
    double shiny;                      // shininess
    double drawforcescale;             // scale factor for force arrows
    double drawforcecolor[4];          // color for force arrows
    double stdlen;                     // minimum energy segment length
    double stdypr[3];                  // minimum energy bend angle
}
```

5 C Code: structures and functions

```
double klen;           // force constant for length
double kypr[3];        // force constant for angle
double kT;             // thermodynamic temp, [0,inf)
double treadrate;      // treadmilling rate constant
double mobility;       // mobility
double filradius;      // segment radius
int maxface;           // filament faces allocated
int nface;             // number of filament faces
char** facename;       // list of face names
double facetwist;      // twisting rate of faces
int maxfil;            // maximum number of filaments
int nfil;              // actual number of filaments
int autonamenum;       // next number for autoname
filamentptr* fillist;  // list of filaments
char** filenames;     // names of filaments
} * filamenttypeptr;
```

Filament types describe various types of filaments, such as actin, microtubule, etc. All filaments of the same type have the same graphical display, force constants, dynamic simulation methods, etc. It's reasonable to have multiple types that represent a single biological type. For example, there could be actin in cytoplasm, actin in nucleoplasm, and an actin bundle.

The `filss` element points to the filament superstructure that owns this filament type, and the name of the type is in `fname`. The type of simulation dynamics is in `dynamics`, where the options are listed above. The `bundlevalue` parameter gives the number of microfilaments that are bundled together in this particular type.

In the graphical display parameters, `color` is the filament color, with red, green, blue, and alpha values. `edgepts` is the edge thickness for drawing, `edgestipple` is the stippling code for the edge stippling, if any. `drawmode` is the polymer drawing mode, which can be face, edge, vertex, or a combination of these. `shiny` is the shininess for graphical display. These graphical display statements are very similar to those used for surfaces. The `drawforcescale` parameter is for drawing force arrows to show what the forces are at each node; it is set to 0 for no arrows and otherwise the force value is multiplied by this scaling number to determine the arrow length. The `drawforcecolor` parameter is for the colors of these arrows.

Filament mechanics information is in the next elements. `stdlen` is the standard segment length, meaning the segment length at minimum energy, where segments are neither stretched nor compressed. `stdypr` is the minimum energy relative ypr angle. `klen` is the stretching force constant. Its value is < 0 for an infinite force constant, meaning that the segment lengths are fixed. `kypr` is the bending force constant. For any element, its value is 0 for zero force constant, meaning a freely jointed chain, and < 0 for an infinite force constant, meaning a fixed bending angle. `kT` is the thermodynamic energy. This parameter isn't ideal here because it allows a different temperature for different components in the simulation, but it's here for now. `treadrate` is the treadmilling rate constant, with 0 for no treadmilling, a positive number for the back growing and a negative number for the front growing. `mobility` is the mobility of the nodes for computing dynamics. `filradius` is the radius of segments.

Next, `maxfil` and `nfil` give the number of allocated and actual filaments of this type, and they are listed in `fillist`. The filaments have names in `filnames`. Names can be assigned automatically, in which case they are named as the filament type name appended with an integer; this integer is found from `autonamenum`.

Filament superstructure

```
typedef struct filamentsuperstruct {
    enum StructCond condition; // structure condition
    struct simstruct* sim;     // simulation structure
    int maxtype;               // maximum number of filament types
    int ntype;                 // actual number of filament types
    char** ftnames;            // filament type names
    filamenttypeptr* filtypes; // list of filament types
} * filamentssptr;
```

The superstructure contains a list of filament types. It starts with its condition element in `condition` and a pointer up the heirarchy to the simulation structure in `sim`. The list of filament types is allocated with `maxtype` entries, has `ntype` entries, and has list `ftnames` for the type names and `filtypes` is the actual list of types.

5.11.2 Filament math - 3D conformations

*This is the first of several sections on filament math, but these aren't being maintained because I'm now doing math in a different document, called *FilamentNotes.tex*.*

There are several ways to describe 3D rotations, including with various conventions of Euler angles, with direction cosine matrices, and with quaternions. Smoldyn uses the yaw-pitch-roll Euler angle convention for relative angles between segments, which it also uses for force and energy computations. I developed most of the filament theory using direction cosine matrices, which are relatively straightforward and widely used in the physics literature, and are described in this subsection. However, I then replaced the direction cosine matrices with quaternions, described below, because they are likely to be faster computationally. At this point, I am continuing to do theory with dcms and code with quaternions.

This subsection uses direction cosine matrices to relate filament relative angles (**A**, was `dcm` in the code), absolute angles (**B**, was `adcm` in the code), and positions (**x**, `xyz`). See also Goldstein Chapter 4 (page 128) and Appendix B (page 608), Andrews, 2004 (rotational averaging paper), and Andrews, 2014 (filament paper). Equations are given here for 2D and 3D systems.

Filament bending is described by the Tait-Bryan convention for Euler angles:

name	symbol	perpendicular axis	positive direction	positive unit vector
2D: yaw	ϕ	z	counter-clockwise	$\hat{z}$
3D: yaw	ϕ	z	turn left	$\hat{z}$
pitch	θ	y	turn down	$\hat{y}$
roll	ψ	x	tilt right	$\hat{x}$

Note that pitch is defined so that positive pitch is rotation downward. This is the opposite of many conventions, but is used here so that the rotation vector is along the positive $\hat{y}$ vector while using a right-handed coordinate system. Figure 5.2 shows these rotations using a standard right-handed coordinate system. The pitch rotation angle is negative, with the result that the rotation vector is along the rotated $-\hat{y}$.




Figure 5.1: Two filament segments in 3D, showing positive yaw, negative pitch, and positive roll.

Consider a direction cosine matrix (dcm), Φ . It can be expanded as:

$$\text{2D: } \Phi = \begin{bmatrix} \Phi_{Xx} & \Phi_{Xy} \\ \Phi_{Yx} & \Phi_{Yy} \end{bmatrix} \quad (5.1)$$

$$\text{3D: } \Phi = \begin{bmatrix} \Phi_{Xx} & \Phi_{Xy} & \Phi_{Xz} \\ \Phi_{Yx} & \Phi_{Yy} & \Phi_{Yz} \\ \Phi_{Zx} & \Phi_{Zy} & \Phi_{Zz} \end{bmatrix} \quad (5.2)$$

There are two interpretations for this dcm. It can convert between two frames of reference for a static object, which is called passive rotation, or it can rotate an object in a fixed reference frame, which is called active rotation. In passive rotation, shown in Figure 5.2, a vector doesn't move, but is simply being expressed in different coordinate systems. Here, X , Y , and Z are the unit vectors of the “lab frame”, meaning the absolute coordinates of the simulation, and x , y , and z are the unit vectors of the “molecule frame”, such as the local coordinates of a filament segment. The dcm expresses the dot products of these unit vectors. As a result, each column of the matrix gives the lab frame coordinates of each unit vector of the molecule frame, and each row of the matrix gives the molecule frame coordinates of each unit vector of the lab frame. It is unitary, meaning that all its eigenvalues equal 1, and its transpose is its inverse, so $\Phi^T \Phi = \mathbf{1}$, where $\mathbf{1}$ is the identity matrix.




Figure 5.2: Diagram for 2D coordinate rotation. Black is lab frame, blue is molecule frame, and red is a vector that just happens to be aligned with the lab frame.

Here's the dcm in more detail, now shown for Φ as a product of yaw, pitch, and roll

transformations. See Goldstein pp. 135 and 609.

$$\begin{aligned}
\text{2D: } \Phi &= \mathbf{Z}(\phi) = \begin{bmatrix} c\phi & s\phi \\ -s\phi & c\phi \end{bmatrix} \\
\text{3D: } \Phi &= \mathbf{X}(\psi)\mathbf{Y}(\theta)\mathbf{Z}(\phi) \\
&= \begin{bmatrix} 1 & 0 & 0 \\ 0 & c\psi & s\psi \\ 0 & -s\psi & c\psi \end{bmatrix} \begin{bmatrix} c\theta & 0 & -s\theta \\ 0 & 1 & 0 \\ s\theta & 0 & c\theta \end{bmatrix} \begin{bmatrix} c\phi & s\phi & 0 \\ -s\phi & c\phi & 0 \\ 0 & 0 & 1 \end{bmatrix} \\
&= \begin{bmatrix} c\theta c\phi & c\theta s\phi & -s\theta \\ s\psi s\theta c\phi - c\psi s\phi & s\psi s\theta s\phi + c\psi c\phi & c\theta s\psi \\ c\psi s\theta c\phi + s\psi s\phi & c\psi s\theta s\phi - s\psi c\phi & c\theta c\psi \end{bmatrix} \tag{5.3}
\end{aligned}$$

Right-multiplying the direction cosine matrix with a vector that's expressed in the lab frame returns the coordinates of the same vector but now using the molecule frame (i.e. this is passive rotation for a molecule frame that is rotated from the lab frame by the given angles). For example, suppose $\mathbf{r}_L$ is a vector expressed in the lab frame. Then,

$$\mathbf{r} = \Phi \mathbf{r}_L \tag{5.4}$$

gives the same vector but expressed in the molecule frame. As a particularly simple vector (shown with the red arrow in Figure 5.2), consider 2D and suppose $\mathbf{r}_L = [1, 0]^T$. In the molecule frame, which is rotated counter-clockwise by ϕ from the lab frame, then $\mathbf{r} = [c\phi, -s\phi]^T$.

In the code and most of my math, the absolute direction cosine matrix, meaning for conversion between the simulation reference frame and the local segment frame, is called $\mathbf{B}$. Right multiplying the dcm with a lab-frame vector, here $\mathbf{p}_L$, gives the molecule frame, leading to $\mathbf{p} = \mathbf{B}\mathbf{p}_L$. This expression can be transposed and/or multiplied on both sides by $\mathbf{B}^T$ to yield (non-transposed vectors are column vectors and transposed ones are row vectors)

$$\mathbf{p} = \mathbf{B}\mathbf{p}_L \quad \mathbf{p}_L = \mathbf{B}^T \mathbf{p} \tag{5.5}$$

$$\mathbf{p}^T = \mathbf{p}_L^T \mathbf{B}^T \quad \mathbf{p}_L^T = \mathbf{p}^T \mathbf{B} \tag{5.6}$$

The Sphere.c library has functions for these multiplications. The transpose of the dcm listed above performs an active rotation of a vector in a constant coordinate system.

The front end of a filament is the front end of segment 0, which is at $\mathbf{x}_0$. In its own coordinate system, this segment lies along direction $\hat{\mathbf{x}} = [1, 0, 0]^T$. In the lab coordinate system, this same vector is $\mathbf{B}_0^T \hat{\mathbf{x}}$. From this, the back of segment 0, which is also the front of segment 1, is at

$$\mathbf{x}_1 = \mathbf{x}_0 + l_0 \mathbf{B}_0^T \hat{\mathbf{x}}$$

Alternatively, focusing on only the displacement,

$$\Delta \mathbf{x}_0 = \mathbf{x}_1 - \mathbf{x}_0 = l_0 \mathbf{B}_0^T \hat{\mathbf{x}}$$

The relative rotation matrix $\mathbf{A}$ expresses the rotation of each segment relative to the previous one. Also, the reference orientation is assumed to be parallel to the lab x -axis. Thus,

$$\begin{aligned}
\mathbf{B}_0 &= \mathbf{A}_0 \\
\mathbf{B}_i &= \mathbf{A}_i \mathbf{B}_{i-1} = \mathbf{A}_i \mathbf{A}_{i-1} \cdots \mathbf{A}_0 \\
\mathbf{B}_i^T &= \mathbf{B}_{i-1}^T \mathbf{A}_i^T = \mathbf{A}_0^T \mathbf{A}_1^T \cdots \mathbf{A}_i^T \tag{5.7}
\end{aligned}$$

5 C Code: structures and functions

To make sense of this, consider the orientation of the i th segment. It is $\hat{\mathbf{x}}$ in its own reference frame and $\mathbf{A}_i^T \hat{\mathbf{x}}$ in the reference frame of the prior segment; continuing this process, it is $\mathbf{A}_0^T \mathbf{A}_1^T \cdots \mathbf{A}_i^T \hat{\mathbf{x}}$, which is $\mathbf{B}_i^T \hat{\mathbf{x}}$, in the lab frame. The locations of the segment front ends for subsequent segments are

$$\begin{aligned}\mathbf{x}_{i+1} &= \mathbf{x}_i + l_i \mathbf{B}_i^T \hat{\mathbf{x}} \\ \Delta \mathbf{x}_i &= \mathbf{x}_{i+1} - \mathbf{x}_i = l_i \mathbf{B}_i^T \hat{\mathbf{x}}\end{aligned}\tag{5.8}$$

To compute a filament's node positions from angle data, we start with segment 0. This segment's orientation can be set initially with ypr angles, but those are not used when computing dynamics because the yaw and roll angles become undefined when the segment is vertical, and are extremely sensitive to the segment's orientation when it is nearly vertical. Instead, the segment 0 orientation is stored in the `seg0dcm` direction cosine matrix; this is both $\mathbf{A}_0$ and $\mathbf{B}_0$ in the analytical math. For subsequent segments, we know the ypr angles, so we compute the $\mathbf{A}_i$ matrices from eq. 5.3. From those, we find the $\mathbf{B}_i$ matrices from eq. 5.7. Then, we find the node positions from eq. 5.8. These computations are the same for 2D and 3D.

Going backward from node positions to ypr angles is more difficult. Starting with 2D systems, the only things that are needed are absolute and relative yaw angles, since pitch and roll are always zero. For segment, i , which goes from $\mathbf{x}_i$ and $\mathbf{x}_{i+1}$ and has spatial difference vector $\Delta \mathbf{x}_i = \mathbf{x}_{i+1} - \mathbf{x}_i$, the segment's absolute angle is

$$\phi = \text{atan2}(\Delta x_{i,y}, \Delta x_{i,x})$$

This uses the `atan2` function, which is the same as the arctangent, but keeps track of the quadrant; this uses $\Delta x_{i,x}$ and $\Delta x_{i,y}$ for the x and y axes. The result is between $-\pi$ and π , (inclusive at both ends). Note that this fails if both arguments of the function are equal to zero. The code doesn't store the absolute angle, but this is converted to an absolute direction cosine matrix that is stored. The difference between this angle and the angle for the prior segment is the relative angle, but that is no longer necessarily within $[-\pi, \pi]$, so that angle needs to be clamped to this range. Testing shows that this works well, so long as the mobility value isn't so high as to create unstable integration.

The 3D computation does not use absolute angles. Also, the node positions are not sufficient input information but one also needs to know the roll value for each segment, except for segment 0. These roll values give the roll amount relative to the prior segment's orientation. For segment 0, I tried using a roll value relative to the absolute $\hat{\mathbf{x}}$ vector, but that creates discontinuities when the segment is nearly vertical. Instead, the solution I'm choosing is to record an upward vector, which is in the system reference frame, for segment 0. Ideally, this vector would be perpendicular to the segment and would have unit length, in which case it is the $\hat{\mathbf{z}}$ vector in the segment 0 reference frame. Deviations from this ideal are not a problem because they can be easily fixed.

For other segments, consider segment i , which has displacement $\Delta \mathbf{x}_i$ and roll ψ_i . This displacement is in the lab coordinate system and the roll is in the coordinate system of the prior segment. From 5.8,

$$\Delta \mathbf{x}_i = l_i \mathbf{B}_i^T \hat{\mathbf{x}} = l_i \mathbf{B}_{i-1}^T \mathbf{A}_i^T \hat{\mathbf{x}}$$

Rearrange to

$$\mathbf{A}_i^T \hat{\mathbf{x}} = \frac{1}{l_i} \mathbf{B}_{i-1} \Delta \mathbf{x}_i\tag{5.9}$$

We know the right hand side, and also the roll ψ_i , and want to compute the left hand side. The matrix product on the left expands to local ypr angles, which are in the frame of the $i - 1$ segment.

5.11 Filaments (functions in *smolfilament.c*)

Define $\Delta \mathbf{x}'_i$ as the normalized and rotated version of $\Delta \mathbf{x}_i$ into the $i - 1$ segment frame, with $\Delta \mathbf{x}'_i = \mathbf{B}_{i-1} \Delta \mathbf{x}_i / l_i$. These yield

$$\mathbf{A}^T(\phi_i, \theta_i, \psi_i) \hat{\mathbf{x}} = \begin{bmatrix} c\theta_i c\phi_i \\ c\theta_i s\phi_i \\ -s\theta_i \end{bmatrix} = \frac{1}{l_i} \mathbf{B}_{i-1} \Delta \mathbf{x}_i = \Delta \mathbf{x}'_i = \begin{bmatrix} \Delta x'_{i,x} \\ \Delta x'_{i,y} \\ \Delta x'_{i,z} \end{bmatrix}$$

Solving the three scalar equations for the two unknowns, θ_i and ϕ_i yield

$$\begin{aligned} \phi_i &= \text{atan2}(\Delta x'_{i,y}, \Delta x'_{i,x}) \\ \theta_i &= -\text{asin} \frac{\Delta x'_{i,z}}{l_i} \end{aligned} \quad (5.10)$$

We still have ψ_i from beforehand, still in the coordinate system of the $i - 1$ segment. Combine these three values, ϕ_i , θ_i , and ψ_i , with eq. 5.3 to get $\mathbf{A}_i$. This can be done by first computing the ypr angles and substituting them into eq. 5.3. Alternatively, one can substitute in the analytical values from eq. 5.10 and then simplify. Using the latter approach (and abbreviating $\Delta x'_{i,x}$, $\Delta x'_{i,y}$, and $\Delta x'_{i,z}$ as x' , y' , and z'),

$$\mathbf{A} = \begin{bmatrix} x' & y' & z' \\ \frac{-y' c\psi - x' z' s\psi}{\sqrt{x'^2 + y'^2}} & \frac{x' c\psi - y' z' s\psi}{\sqrt{x'^2 + y'^2}} & \sqrt{x'^2 + y'^2} s\psi \\ \frac{y' s\psi - x' z' c\psi}{\sqrt{x'^2 + y'^2}} & \frac{-x' s\psi - y' z' c\psi}{\sqrt{x'^2 + y'^2}} & \sqrt{x'^2 + y'^2} c\psi \end{bmatrix}$$

While we used ypr angles to find this solution, they do not appear here (except for ψ) because this is the unique solution to the stated problem of finding $\mathbf{A}_i$ from $\Delta \mathbf{x}'_i$ and ψ_i , independent of approach. The top row is equal to x' , y' , and z' because this comes directly from the equation $\mathbf{A}^T \hat{\mathbf{x}} = \Delta \mathbf{x}'_i$. This top row not hypersensitive to particular values of x' , y' , or z' . However, the first two columns of the other two rows are highly sensitive to small changes in x' or y' if the segment is nearly vertical (in the reference frame of the prior segment), so that $x'^2 + y'^2 \approx 0$. This is an unavoidable consequence of how the roll value is defined.

Next,

$$\mathbf{B}_i = \mathbf{A}_i \mathbf{B}_{i-1}$$

This finishes the calculations for segment i , and then walk forward to the next segment.

5.11.3 Filament math - quaternions

Quaternions work in relatively similar ways (see the online article “Rotation Quaternions, and How to Use Them” by Rose, 2015). A quaternion, $\mathbf{q}$ has 4 values, (q_0, q_1, q_2, q_3) . If it represents a 3D point in space, then $\mathbf{q} = (0, x, y, z)$ and if it represents a rotation, then all four values can be non-zero and they have total magnitude of 1. The quaternion $\mathbf{q} = (1, 0, 0, 0)$ is the rotation identity quaternion, which produces no rotation. The inverse or complex conjugate quaternion, $\mathbf{q}^* = (q_0, -q_1, -q_2, -q_3)$, performs the opposite rotation as the original one.

Quaternions correspond to direction cosine matrices. For quaternion $\mathbf{q}$, the corresponding direction cosine matrix is as follows (from Rose). Note that the unitary nature of a rotation quaternion ($q_0^2 + q_1^2 + q_2^2 + q_3^2 = 1$) results in different ways of expressing the diagonal elements.

$$\mathbf{Q} = \begin{bmatrix} 1 - 2q_2^2 - 2q_3^2 & 2q_1q_2 - 2q_0q_3 & 2q_1q_3 + 2q_0q_2 \\ 2q_1q_2 + 2q_0q_3 & 1 - 2q_1^2 - 2q_3^2 & 2q_2q_3 - 2q_0q_1 \\ 2q_1q_3 - 2q_0q_2 & 2q_2q_3 + 2q_0q_1 & 1 - 2q_1^2 - 2q_2^2 \end{bmatrix} \quad (5.11)$$

5 C Code: structures and functions

Consider rotation by ypr angles (ϕ, θ, ψ) . This converts to quaternion $\mathbf{q}$ with the following equation (this equation is from Rose but with reversed signs because theirs is for active rotation, whereas this work assumes passive rotation). The equation on the right performs the opposite conversion (again, equation from Rose, but with sign changes)

$$\mathbf{q} = \begin{bmatrix} \frac{\psi}{2} \frac{\theta}{2} \frac{\phi}{2} + \frac{\psi}{2} \frac{\theta}{2} \frac{\phi}{2} \\ -\frac{\psi}{2} \frac{\theta}{2} \frac{\phi}{2} + \frac{\psi}{2} \frac{\theta}{2} \frac{\phi}{2} \\ -\frac{\psi}{2} \frac{\theta}{2} \frac{\phi}{2} - \frac{\psi}{2} \frac{\theta}{2} \frac{\phi}{2} \\ -\frac{\psi}{2} \frac{\theta}{2} \frac{\phi}{2} + \frac{\psi}{2} \frac{\theta}{2} \frac{\phi}{2} \end{bmatrix} \quad \mathbf{ypr} = \begin{bmatrix} \text{atan2}(-2q_0q_3 + 2q_1q_2, q_0^2 + q_1^2 - q_2^2 - q_3^2) \\ \text{asin}(-2q_0q_2 - 2q_1q_3) \\ \text{atan2}(-2q_0q_1 + 2q_2q_3, q_0^2 - q_1^2 - q_2^2 + q_3^2) \end{bmatrix} \quad (5.12)$$

To check these equations, I computed from ypr angles (ϕ, θ, ψ) to quaternion and then to dcm, and arrived at eq. 5.3. I also converted from ypr angles to quaternion and back again.

For passive rotation of position quaternion $\mathbf{p}$ between the lab frame and molecule frame using rotation quaternion $\mathbf{q}$,

$$\mathbf{p} = \mathbf{q}^* \mathbf{p}_L \mathbf{q} \quad \mathbf{p}_L = \mathbf{q} \mathbf{p} \mathbf{q}^*$$

The former equation is for “forward” rotation and is equivalent to the matrix equation $\mathbf{p} = \mathbf{B} \mathbf{p}_L$, whereas the latter equation is for “inverse” rotation and is equivalent to the matrix equation $\mathbf{p}_L = \mathbf{B}^T \mathbf{p}$. I checked this by converting ypr angles (ϕ, θ, ψ) to a dcm with eq. 5.3 and using that to rotate vector $[x, y, z]^T$, and also converting ypr angles to a quaternion with eq. 5.12 and rotating the same vector; the results were the same.

For a filament, suppose $\hat{\mathbf{x}}$ is the quaternion version of the unit vector on the x -axis, and $\mathbf{b}_i$ is the absolute rotation quaternion for segment i . Then, the segment displacement is

$$\Delta \mathbf{x}_i = \mathbf{x}_{i+1} - \mathbf{x}_i = l_i \mathbf{b}_i \hat{\mathbf{x}} \mathbf{b}_i^*$$

From this, it can be seen that quaternions for multiple rotations combine as

$$\mathbf{b}_i = \mathbf{b}_{i-1} \mathbf{a}_i$$

$$\mathbf{b}_i^* = \mathbf{a}_i^* \mathbf{b}_{i-1}^*$$

These are opposite the equations for matrices, where $\mathbf{B}_i = \mathbf{A}_i \mathbf{B}_{i-1}$.

Converting from node positions and roll values to quaternions is similar to the case with direction cosine matrices. The direct approach, which will be shown to fail, starts with the equation

$$\Delta \mathbf{x}_i = l_i \mathbf{b}_i \hat{\mathbf{x}} \mathbf{b}_i^* = l_i \mathbf{b}_{i-1} \mathbf{a}_i \hat{\mathbf{x}} \mathbf{a}_i^* \mathbf{b}_{i-1}^*$$

Defining $\Delta \mathbf{x}'_i$ as the vector for the i th segment in the reference frame of the $i-1$ segment, and also normalized, this equation rearranges and expands to

$$\Delta \mathbf{x}'_i = \begin{bmatrix} \Delta x'_{i,x} \\ \Delta x'_{i,y} \\ \Delta x'_{i,z} \end{bmatrix} = \frac{1}{l_i} \mathbf{b}_{i-1}^* \Delta \mathbf{x}_i \mathbf{b}_{i-1} = \mathbf{a}_i \hat{\mathbf{x}} \mathbf{a}_i^* = \begin{bmatrix} 1 - 2a_2^2 - 2a_3^2 \\ 2a_1a_2 + 2a_0a_3 \\ 2a_1a_3 - 2a_0a_2 \end{bmatrix}$$

Note that the vector on the right is the same as the first column of eq. 5.11, much as was seen before with dcms. In principle, one could solve this equation for the four a values, and also include the roll value somehow. However, the solution seems difficult (and is beyond Mathematica), and I’m also not sure how to include the roll value.

Another approach, which does work, is to use the prior solution for the ypr values, from eq. 5.10 and the previously known roll value. Then substitute these solutions into eq. 5.12 to compute the quaternion. Doing so is straightforward numerically. It’s also easy analytically, although the equations are very long and don’t clean up. Thus, one needs to compute the ypr angles ϕ , θ , and ψ first, and then use those to compute the quaternion. This works well, for the most part, but is problematic when x and y are small, meaning that the filament segment is nearly vertical.

5.11.4 Filament math - segment 0 orientation

The segment 0 orientation cannot be described with ypr angles because its orientation values become highly sensitive on the precise orientation when the segment is nearly vertical. Furthermore, the segment orientation also cannot be described as a $\Delta \mathbf{x}_0$ displacement in combination with a roll value, because the roll angle is still highly sensitive when the segment is nearly vertical.

The solution is as follows. First, compute all forces. Move nodes in response. Apply rolls in response. However, for segment 0, the roll is not just an angle value. Instead, this segment has a perpendicular vector that defines the local up direction. This vector is rotated for segment 0 rolling.

5.11.5 Filament math - filament stretching forces

For filament stretching forces, consider a segment, i . Its segment vector is $\Delta \mathbf{x}_i = \mathbf{x}_{i+1} - \mathbf{x}_i$, as usual. Also, its length is $l_i = |\mathbf{x}_i|$. For standard length l° , the extension force is

$$F = -k_{stretch}(l_i - l^\circ)$$

This is the extension force of the segment, meaning that a positive force implies that the segment wants to be longer and a negative force implies that the segment wants to be shorter. The product of this with the segment's unit vector yields the force on the $i + 1$ node, and then the negative of that yields the force on the i node,

$$\begin{aligned} \mathbf{F}_{i+1} &= F \frac{\Delta \mathbf{x}_i}{l_i} \\ \mathbf{F}_i &= -F \frac{\Delta \mathbf{x}_i}{l_i} \end{aligned}$$

Note that $\mathbf{F}_i$ is a *dim*-dimensional vector. Also, i is the segment number (from 0 to `nseg-1`), while the subscript on $\mathbf{F}$ is the node number. Slight rearrangement gives the scaled stretch force for segment i as

$$s_i = -k_{stretch} \frac{l_i - l^\circ}{l_i}$$

and then the forces on the adjacent nodes from stretching in segment i as

$$\begin{aligned} \mathbf{F}_i &= -s_i \Delta \mathbf{x}_i \\ \mathbf{F}_{i+1} &= s_i \Delta \mathbf{x}_i \end{aligned} \tag{5.13}$$

This is implemented in the code `filAddStretchForces`.

A slightly different version is implemented in `filAddStretchForceMat`. Here, each matrix row represents a change of a state variable over time and each matrix column represents the state variables that this change depends on. In other words,

$$\dot{\mathbf{x}}(t) = \mu \mathbf{F} \cdot \mathbf{x}(t)$$

where $\mathbf{F}$ is the matrix, which I call the force matrix. Now, $\dot{\mathbf{x}}$ is a linear function of $\mathbf{x}$. This linear relationship can also be expressed in the force vectors by expanding eq. 5.13 and rearranging slightly to give

$$\begin{aligned} \mathbf{F}_i &= s_i \mathbf{x}_i - s_i \mathbf{x}_{i+1} \\ \mathbf{F}_{i+1} &= -s_i \mathbf{x}_i + s_i \mathbf{x}_{i+1} \end{aligned} \tag{5.14}$$

5 C Code: structures and functions

These vectors and the force matrix are flattened in the code. For example, the 2D version of $\mathbf{x}(t)$ is $[x_{0,x}, x_{0,y}, x_{1,x}, x_{1,y}, \dots]$, where the first index is the node number and the second is the Cartesian coordinate. With flattening, the 2D version of eq. 5.14, which again is the force created by stretching of segment i , becomes

$$\begin{bmatrix} F_{i,x} \\ F_{i,y} \\ F_{i+1,x} \\ F_{i+1,y} \end{bmatrix} = \begin{bmatrix} s_i & 0 & -s_i & 0 \\ 0 & s_i & 0 & -s_i \\ -s_i & 0 & s_i & 0 \\ 0 & -s_i & 0 & s_i \end{bmatrix} \cdot \begin{bmatrix} x_{i,x} \\ x_{i,y} \\ x_{i+1,x} \\ x_{i+1,y} \end{bmatrix}$$

The 3D version is identical, but has the z coordinate as well. Again, these are the forces from stretching or compression of segment i . If there is stretching or compression for other segments, then those forces are added to the force matrix. This equation is implemented in `filAddStretchForceMat`. The following matrix equation shows the stretching forces on the 4 nodes of a 2D 3-segment filament.

$$\mathbf{F} = \begin{bmatrix} s_0 & 0 & -s_0 & 0 & 0 & 0 & 0 & 0 \\ 0 & s_0 & 0 & -s_0 & 0 & 0 & 0 & 0 \\ -s_0 & 0 & s_0 + s_1 & 0 & -s_1 & 0 & 0 & 0 \\ 0 & -s_0 & 0 & s_0 + s_1 & 0 & -s_1 & 0 & 0 \\ 0 & 0 & -s_1 & 0 & s_1 + s_2 & 0 & -s_2 & 0 \\ 0 & 0 & 0 & -s_1 & 0 & s_1 + s_2 & 0 & -s_2 \\ 0 & 0 & 0 & 0 & -s_2 & 0 & s_2 & 0 \\ 0 & 0 & 0 & 0 & 0 & -s_2 & 0 & s_2 \end{bmatrix} \cdot \begin{bmatrix} x_{0,x} \\ x_{0,y} \\ x_{1,x} \\ x_{1,y} \\ x_{2,x} \\ x_{2,y} \\ x_{3,x} \\ x_{3,y} \end{bmatrix}$$

Again, generalizing to 3D is trivial.

5.11.6 Filament math - filament bending forces

For filament forces from bending angles, consider a pair of segments, labeled as $i - 1$ and i . These segment vectors, pointing from front to back as normal, are $\Delta \mathbf{x}_{i-1}$ and $\Delta \mathbf{x}_i$, where

$$\begin{aligned} \Delta \mathbf{x}_{i-1} &= \mathbf{x}_i - \mathbf{x}_{i-1} \\ \Delta \mathbf{x}_i &= \mathbf{x}_{i+1} - \mathbf{x}_i \end{aligned} \quad (5.15)$$

The segment nodes are $i - 1$ for the front-most node, i for the middle node, and $i + 1$ for the back-most node. See Figure 5.3. The absolute dcm for the first segment is $\mathbf{B}_{i-1}$ and the relative dcm for the second segment is $\mathbf{A}_i$.




Figure 5.3: Indices for a pair of segments. Also, forces and torques.

Supposing the first segment were fixed and one applied force $\mathbf{F}_{i+1}$ to the back end node, this would create torque at the center node of

$$\boldsymbol{\tau}_i = \Delta \mathbf{x}_i \times \mathbf{F}_{i+1}. \quad (5.16)$$

5.11 Filaments (functions in *smolfilament.c*)

In 2D, this torque is along the z -axis. In 3D, this torque is necessarily perpendicular to $\Delta \mathbf{x}_i$ (and $\mathbf{F}_{i+1}$) but may not be perpendicular to $\Delta \mathbf{x}_{i-1}$. Alternatively, the second segment could be fixed and one applies force $\mathbf{F}_{i-1}$ to the first node, which would create torque at the center node of

$$\boldsymbol{\tau}'_i = -\Delta \mathbf{x}_{i-1} \times \mathbf{F}_{i-1} \quad (5.17)$$

The negative sign is included because the $\Delta \mathbf{x}_{i-1}$ vector points toward node i instead of away from it.

In practice, the center node exerts a three-dimensional torque from effective yaw, pitch, and roll springs. These torques, which are computed below, then create forces at the adjacent nodes. Inverting the torque equations (see Wikipedia “Cross product”) gives the forces on the nodes that arise from the torques. This inversion is not mathematically unique but can also have components that are parallel to the segments, but those are omitted here because we know from the physical situation that the induced forces must be perpendicular to the segment vectors. The induced forces are

$$\begin{aligned} \mathbf{F}_{i+1} &= \frac{\boldsymbol{\tau}_i \times \Delta \mathbf{x}_i}{l_i^2} \\ \mathbf{F}_{i-1} &= -\frac{\boldsymbol{\tau}'_i \times \Delta \mathbf{x}_{i-1}}{l_{i-1}^2} \end{aligned} \quad (5.18)$$

Because the $\boldsymbol{\tau}_i$ and $\boldsymbol{\tau}'_i$ torques are created from the same yaw, pitch, and roll springs, and they push on the two segments in opposite directions, they are equal to each other but have opposite signs, so

$$\boldsymbol{\tau}'_i = -\boldsymbol{\tau}_i$$

Consider, for example, the situation shown in Figure 5.3 and suppose this is in 2D. Consistent with eqs. 5.16 and 5.17, if there were a force pushing down on the right segment, then the torque $\boldsymbol{\tau}_i$ would point into the page and if there were a force pushing down on the left segment, then the torque $\boldsymbol{\tau}'_i$ would point out of the page, as shown in the diagram. Likewise, these torques would create the forces listed in eqs. 5.18.

For 3D, in addition to creating forces at the end nodes, the torque created at node i by the springs also induces torques into the two segments around their axes. These segment torques are given as $\tau_{R,i-1}$ for the first segment and $\tau_{R,i}$ for the second segment (no bold face because these are scalars; also the R subscript indicates that these are roll torques, as opposed to the total torques denoted by $\boldsymbol{\tau}_i$). These induced roll torques are the projections of the spring torques onto the segment axes,

$$\begin{aligned} \tau_{R,i} &= \frac{\boldsymbol{\tau}_i \cdot \Delta \mathbf{x}_i}{l_i} \\ \tau_{R,i-1} &= \frac{\boldsymbol{\tau}'_i \cdot \Delta \mathbf{x}_{i-1}}{l_{i-1}} \end{aligned} \quad (5.19)$$

This shows that spring torques that are perpendicular to the segments create forces at the end nodes and parallel spring torques create torques along the segments.

Considering the yaw spring first, which is only thing that matters for 2D, it creates torques

$$\tau_{\phi,i} = -k_{\phi} \Delta \phi_i \hat{z} = -\boldsymbol{\tau}'_{\phi,i} \cdot \hat{z}. \quad (5.20)$$

Here and below, we work in the coordinate system of the former segment (the $i-1$ segment), so $\hat{x}$, $\hat{y}$ and $\hat{z}$ are unit vectors for the former segment’s coordinate system. In this coordinate system,

$$\Delta \mathbf{x}_{i-1} = l_{i-1} \hat{x}$$

5 C Code: structures and functions

$$\Delta \mathbf{x}_i = l_i \mathbf{A}_i^T \hat{x}$$

Also, $\Delta \phi_i = (\phi_i - \phi^\circ)$ is the yaw spring displacement. Although not used right away, we also define $\Delta \theta_i = (\theta_i - \theta^\circ)$ and $\Delta \psi_i = (\psi_i - \psi^\circ)$. This yaw spring torque creates forces

$$\begin{aligned} \mathbf{F}_{i+1} &= \frac{-k_\phi \Delta \phi_i \hat{z} \times \Delta \mathbf{x}_i}{l_i^2} = \frac{-k_\phi \Delta \phi_i}{l_i} \hat{z} \times (\mathbf{A}_i^T \hat{x}) \\ \mathbf{F}_{i-1} &= \frac{-k_\phi \Delta \phi_i \hat{z} \times \Delta \mathbf{x}_{i-1}}{l_{i-1}^2} = \frac{-k_\phi \Delta \phi_i}{l_{i-1}} \hat{z} \times \hat{x} = \frac{-k_\phi \Delta \phi_i}{l_{i-1}} \hat{y} \end{aligned} \quad (5.21)$$

These forces are helpful to think about and are useful for the 2D case. Note that they are in the reference frame of the $i-1$ segment. However, they aren't directly useful for the 3D case. For that, we'll add up all the torques first.

Repeating what was written above in eq. 5.20, the yaw spring creates torques (in the segment $i-1$ reference frame)

$$\boldsymbol{\tau}_{\phi,i} = -k_\phi \Delta \phi_i \hat{z} = -\boldsymbol{\tau}'_{\phi,i}.$$

Next, the pitch spring creates torques

$$\boldsymbol{\tau}_{\theta,i} = -k_\theta \Delta \theta_i \mathbf{Z}^T(\phi_i) \hat{y} = -\boldsymbol{\tau}'_{\theta,i}, \quad (5.22)$$

where $\mathbf{Z}^T(\phi_i)$ represents the yaw active rotation matrix. The $\mathbf{Z}^T(\phi_i)$ matrix is here because yaw is performed first, and then pitch, so pitch deviations create a torque that's along the rotated y -axis and it needs to be rotated to the segment $i-1$ reference frame. The roll spring creates torques

$$\boldsymbol{\tau}_{\psi,i} = -k_\psi \Delta \psi_i \mathbf{Y}^T(\theta_i) \mathbf{Z}^T(\phi_i) \hat{x} = -\boldsymbol{\tau}'_{\psi,i}, \quad (5.23)$$

where $\mathbf{Y}^T(\theta_i)$ represents the pitch rotation matrix. It's here because yaw is performed first, then pitch, and then roll, so roll deviations create a torque that's along the rotated x -axis and, again, it needs to be converted back to the segment $i-1$ reference frame. See Figure 5.1. Adding up these torques leads to (math is in Sphere2.nb, from the Sphere.c library)

$$\begin{aligned} \boldsymbol{\tau}_i &= -k_\phi \Delta \phi_i \hat{z} - k_\theta \Delta \theta_i \mathbf{Z}^T(\phi_i) \hat{y} - k_\psi \Delta \psi_i \mathbf{Y}^T(\theta_i) \mathbf{Z}^T(\phi_i) \hat{x} \\ &= -k_\phi \Delta \phi_i \hat{z} - k_\theta \Delta \theta_i (-s\phi_i \hat{x} + c\phi_i \hat{y}) - k_\psi \Delta \psi_i (c\theta_i c\phi_i \hat{x} + s\phi_i \hat{y} - c\phi_i s\theta_i \hat{z}) \\ &= \begin{bmatrix} -k_\psi \Delta \psi_i c\theta_i c\phi_i + k_\theta \Delta \theta_i s\phi_i \\ -k_\theta \Delta \theta_i c\phi_i - k_\psi \Delta \psi_i s\phi_i \\ -k_\phi \Delta \phi_i + k_\psi \Delta \psi_i c\phi_i s\theta_i \end{bmatrix} \end{aligned} \quad (5.24)$$

In the vector form, the terms are arranged so that the dominant one for small bending angles is listed first.

This torque induces forces and segment roll torques. It's simplest to convert to the lab coordinate system for this. The total torque in the lab coordinate system due to the bending springs at node i is

$$\boldsymbol{\tau}_{L,i} = \mathbf{B}_{i-1}^T \boldsymbol{\tau}_i \quad (5.25)$$

Continuing in the lab frame (I'm dropping the L subscript, but it's implied), the forces on the

adjacent nodes due to bending at node i are

$$\begin{aligned}
\mathbf{F}_{i+1} &= \frac{\boldsymbol{\tau}_i \times \Delta \mathbf{x}_i}{l_i^2} \\
\mathbf{F}_{i-1} &= \frac{\boldsymbol{\tau}_i \times \Delta \mathbf{x}_{i-1}}{l_{i-1}^2} \\
\mathbf{F}_i &= -(\mathbf{F}_{i+1} + \mathbf{F}_{i-1}) \\
\tau_{R,i} &= \frac{\boldsymbol{\tau}_i \cdot \Delta \mathbf{x}_i}{l_i} \\
\tau_{R,i-1} &= -\frac{\boldsymbol{\tau}_i \cdot \Delta \mathbf{x}_{i-1}}{l_{i-1}}
\end{aligned} \tag{5.26}$$

These equations are implemented in `filAddBendForces`.

Equations 5.26 can also be expressed in a different manner. Here is the *total* force on node i from bending at nodes $i-1$, i , and $i+1$, respectively

$$\begin{aligned}
\mathbf{F}_{tot.,i} &= \frac{\boldsymbol{\tau}_{i-1} \times \Delta \mathbf{x}_{i-1}}{l_{i-1}^2} - \left(\frac{\boldsymbol{\tau}_i \times \Delta \mathbf{x}_i}{l_i^2} + \frac{\boldsymbol{\tau}_i \times \Delta \mathbf{x}_{i-1}}{l_{i-1}^2} \right) + \frac{\boldsymbol{\tau}_{i+1} \times \Delta \mathbf{x}_i}{l_i^2} \\
&= \frac{\boldsymbol{\tau}_{i-1} - \boldsymbol{\tau}_i}{l_{i-1}^2} \times \Delta \mathbf{x}_{i-1} + \frac{\boldsymbol{\tau}_{i+1} - \boldsymbol{\tau}_i}{l_i^2} \times \Delta \mathbf{x}_i \\
&= \mathbf{b}_i \times \Delta \mathbf{x}_i - \mathbf{b}_{i-1} \times \Delta \mathbf{x}_{i-1} \\
\text{where } \mathbf{b}_i &\equiv \frac{\boldsymbol{\tau}_{i+1} - \boldsymbol{\tau}_i}{l_i^2}
\end{aligned} \tag{5.27}$$

The interpretation of $\mathbf{b}_i$ (“b” for bending) is that it’s the torque applied to segment i by the springs at both of its ends. Continuing the manipulation leads to

$$\begin{aligned}
\mathbf{F}_{tot.,i} &= \mathbf{b}_i \times (\mathbf{x}_{i+1} - \mathbf{x}_i) - \mathbf{b}_{i-1} \times (\mathbf{x}_i - \mathbf{x}_{i-1}) \\
&= \mathbf{b}_{i-1} \times \mathbf{x}_{i-1} - (\mathbf{b}_i + \mathbf{b}_{i-1}) \times \mathbf{x}_i + \mathbf{b}_i \times \mathbf{x}_{i+1}
\end{aligned} \tag{5.28}$$

For 2D and 3D, the cross product $\mathbf{b} \times \mathbf{x}$ expands to

$$\begin{aligned}
\mathbf{b} \times \mathbf{x} &= \begin{bmatrix} -b_z x_y \\ b_z x_x \end{bmatrix} = \begin{bmatrix} 0 & -b_z \\ b_z & 0 \end{bmatrix} \begin{bmatrix} x_x \\ x_y \end{bmatrix} \\
\mathbf{b} \times \mathbf{x} &= \begin{bmatrix} b_y x_z - b_z x_y \\ b_z x_x - b_x x_z \\ b_x x_y - b_y x_x \end{bmatrix} = \begin{bmatrix} 0 & -b_z & b_y \\ b_z & 0 & -b_x \\ -b_y & b_x & 0 \end{bmatrix} \begin{bmatrix} x_x \\ x_y \\ x_z \end{bmatrix}
\end{aligned}$$

Using these expansions, the total bending force on node i can be expanded from eq. 5.28 to express $\mathbf{F}_{tot.,i}$ with flattened vectors and matrices in 2D and 3D as

$$\mathbf{F}_{tot.,i} = \left[\begin{array}{cc|cc|cc} 0 & -b_{i-1} & 0 & b_{i-1} + b_i & 0 & -b_i \\ b_{i-1} & 0 & -b_{i-1} - b_i & 0 & b_i & 0 \end{array} \right] \cdot \begin{bmatrix} x_{i-1,x} \\ x_{i-1,y} \\ \hline x_{i,x} \\ x_{i,y} \\ \hline x_{i+1,x} \\ x_{i+1,y} \end{bmatrix}$$

5 C Code: structures and functions

$$= \left[\begin{array}{ccc|ccc} 0 & -b_{i-1,z} & b_{i-1,y} & 0 & b_{i-1,z} + b_{i,z} & -b_{i-1,y} - b_{i,y} \\ b_{i-1,z} & 0 & -b_{i-1,x} & -b_{i-1,z} - b_{i,z} & 0 & b_{i-1,x} + b_{i,x} \\ -b_{i-1,y} & b_{i-1,x} & 0 & b_{i-1,y} + b_{i,y} & -b_{i-1,x} - b_{i,x} & 0 \end{array} \right] \cdot \left[\begin{array}{ccc} 0 & -b_{i,z} & b_{i,y} \\ b_{i,z} & 0 & -b_{i,x} \\ -b_{i,y} & b_{i,x} & 0 \end{array} \right] \cdot \left[\begin{array}{c} x_{i-1,x} \\ x_{i-1,y} \\ x_{i-1,z} \\ x_{i,x} \\ x_{i,y} \\ x_{i,z} \\ x_{i+1,x} \\ x_{i+1,y} \\ x_{i+1,z} \end{array} \right] \quad (5.29)$$

These equations are implemented in `filAddBendForceMat`.

In addition to the forces, the *total* roll torque on segment i can also be computed from eqs. 5.26, giving

$$\begin{aligned} \tau_{tot.,i} &= \frac{\tau_i \cdot \Delta \mathbf{x}_i}{l_i} - \frac{\tau_{i+1} \cdot \Delta \mathbf{x}_i}{l_i} \\ &= \frac{\tau_i - \tau_{i+1}}{l_i} \cdot \Delta \mathbf{x}_i \end{aligned}$$

As before, this torque is zero for 2D because the torques are parallel to the z -axis and the $\Delta \mathbf{x}_i$ displacements are in the x, y plane. Define the $\mathbf{b}'_i$ vector in close analogy to the $\mathbf{b}_i$ vector defined above,

$$\mathbf{b}'_i = \frac{\tau_{i+1} - \tau_i}{l_i} = l_i \mathbf{b}_i$$

Then,

$$\tau_{tot.,i} = \mathbf{b}'_i \cdot (\mathbf{x}_i - \mathbf{x}_{i+1})$$

Expressing this in matrix form gives the following for 3D,

$$\tau_{tot.,i} = \left[\begin{array}{ccc|ccc} b'_{i,x} & b'_{i,y} & b'_{i,z} & -b'_{i,x} & -b'_{i,y} & -b'_{i,z} \end{array} \right] \cdot \left[\begin{array}{c} x_{i,x} \\ x_{i,y} \\ x_{i,z} \\ x_{i+1,x} \\ x_{i+1,y} \\ x_{i+1,z} \end{array} \right]$$

In practice, I'm interleaving the forces and torques because this maintains the band-diagonal structure of the matrix. There are $4n + 3$ rows to the matrix for a filament with n segments. This comes from $3(n + 1)$ rows for the forces on the $n + 1$ nodes and n rows for the segment roll torques. There are also $4n + 3$ columns for the matrix, with $3(n + 1)$ columns for the node positions and n columns for the segment roll values. None of the equations given here are linearly dependent on the segment roll values, so all of the corresponding matrix columns are equal to zero. That is, there is dependence on these values, through the τ_i bend torques (eq. 5.24), but that dependence is not expressed in the force matrix. Nevertheless, this isn't a problem, because the matrix is added to the identity matrix for use, which then makes the columns not empty. At the end, this means that $\psi_{new} = \psi_{old}$ for each of the roll values, so there's no update here, but then there is an update when the nodes positions are converted to bending angles.

5.11 Filaments (functions in smolfilament.c)

Here is the full stretch plus bend matrix for a 3 segment filament:

$$\begin{bmatrix} s_0 & b_{0z} & -b_{0y} & 0 & -s_0 & -b_{0z} & b_{0y} & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ -b_{0z} & s_0 & b_{0x} & 0 & b_{0z} & -s_0 & -b_{0x} & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ b_{0y} & -b_{0x} & s_0 & 0 & -b_{0y} & b_{0x} & -s_0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ b_{0x} & b_{0y} & b_{0z} & 0 & -b_{0x} & -b_{0y} & -b_{0z} & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ \hline -s_0 & -b_{0z} & b_{0y} & 0 & s_0 + s_1 & b_{0z} + b_{1z} & -b_{0y} - b_{1y} & 0 & -s_1 & -b_{1z} & b_{1y} & 0 & 0 & 0 & 0 & 0 & 0 \\ b_{0z} & -s_0 & -b_{0x} & 0 & -b_{0z} - b_{1z} & s_0 + s_1 & b_{0x} + b_{1x} & 0 & b_{1z} & -s_1 & -b_{1x} & 0 & 0 & 0 & 0 & 0 & 0 \\ -b_{0y} & b_{0x} & -s_0 & 0 & b_{0y} + b_{1y} & -b_{0x} - b_{1x} & s_0 + s_1 & 0 & -b_{1y} & b_{1x} & -s_1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & b_{1x} & b_{1y} & b_{1z} & 0 & -b_{1x} & -b_{1y} & -b_{1z} & 0 & 0 & 0 & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & 0 & -s_1 & -b_{1z} & b_{1y} & 0 & s_1 + s_2 & b_{1z} + b_{2z} & -b_{1y} - b_{2y} & 0 & -s_2 & -b_{2z} & b_{2y} & 0 & 0 \\ 0 & 0 & 0 & 0 & b_{1z} & -s_1 & -b_{1x} & 0 & -b_{1z} - b_{2z} & s_1 + s_2 & b_{1x} + b_{2x} & 0 & b_{2z} & -s_2 & -b_{2x} & 0 & 0 \\ 0 & 0 & 0 & 0 & -b_{1y} & b_{1x} & -s_1 & 0 & b_{1y} + b_{2y} & -b_{1x} - b_{2x} & s_1 + s_2 & 0 & -b_{2y} & b_{2x} & -s_2 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & b_{2x} & b_{2y} & b_{2z} & 0 & -b_{2x} & -b_{2y} & -b_{2z} & 0 & 0 \\ \hline 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & -s_2 & -b_{2z} & b_{2y} & 0 & s_2 & b_{2z} & -b_{2y} & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & b_{2z} & -s_2 & -b_{2x} & 0 & -b_{2z} & s_2 & b_{2x} & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & -b_{2y} & b_{2x} & -s_2 & 0 & b_{2y} & -b_{2x} & s_2 & 0 & 0 \end{bmatrix}$$

5.11.7 Filament math - filament thermal forces

To start with, consider the thermal distribution for filament stretches and bends. For a 2D filament with only two segments, it has bending angle ϕ , intrinsic bending angle ϕ° , and bending force constant k_ϕ . Its mean bending energy relates to the thermal energy from the equipartition theorem as

$$\frac{k_\phi}{2} \langle (\phi - \phi^\circ)^2 \rangle = \frac{k_B T}{2}$$

Rearrange to get the rms ϕ angle as

$$\phi_{rms} = \sqrt{\frac{k_B T}{k_\phi}}$$

The same approach applies to stretching and other bending angles for 3D. In addition, define R as the filament thickness, where 1 is the standard thickness for which the force constants are measured. I'm going to assume that the force constants get multiplied by the thickness. Based on this, the stretching and bending thermal distributions are

$$l_{rms} = \sqrt{\frac{k_B T}{R k_{len}}} \quad (5.30)$$

$$\phi_{rms} = \sqrt{\frac{k_B T}{R k_\phi}} \quad \theta_{rms} = \sqrt{\frac{k_B T}{R k_\theta}} \quad \chi_{rms} = \sqrt{\frac{k_B T}{R k_\chi}} \quad (5.31)$$

The next goal is to compute the thermal force. In 2D, the torque on bend i with angle ϕ_i is

$$\tau_i = -k_\phi (\phi_i - \phi^\circ)$$

Squaring and averaging this squared torque over many possible bends leads to

$$\langle \tau_i^2 \rangle = k_\phi^2 \langle (\phi_i - \phi^\circ)^2 \rangle$$

Simplifying and substituting from eq. 5.31 gives

$$\tau_{i,rms} = k_\phi \phi_{i,rms} = \sqrt{k_\phi k_B T}$$

This torque exerts perpendicular forces on the preceding and following nodes with $F_{i-1} = \tau/l_{i-1}$ and $F_{i+1} = \tau/l_i$, where l_i is the i 'th segment length. There are also forces applied to this node from the preceding and following torques. From eq. 5.27, there are four total terms for the bending force on node i . The standard deviation for the sum of four Gaussian distributed random variables

5 C Code: structures and functions

is the square root of the sum of the squared standard deviations. Using that, the rms forces on node i are

$$F_{i,rms} = 2 \frac{\sqrt{k_\phi k_B T}}{l^\circ} \quad (5.32)$$

For 3D, equivalent forces should apply for the pitch degree of freedom, and a similar calculation should yield the rms roll torques (I think the only difference is that there's no l in the denominator). In addition, the stretch degrees of freedom add yet more forces. I am not convinced that these will return the correct answer. It only makes physical sense for the thermal forces to have the same distributions on the 3 axes, and to be the same at every node, but it's not clear to me that this will yield the same Boltzmann distribution statistics given in eqs. 5.30 and 5.31.

Here is a second approach for computing thermal forces. A bead that diffuses due to Brownian motion moves due to forces exerted on it according to the overdamped Langevin equation (see 5.34), so its motion parallel to the x -axis over time step Δt is

$$\Delta x = \mu F_x \Delta t$$

where μ is the mobility and F_x is the x component of the applied force. Repeating this motion over and over again leads to diffusion, so we compute the diffusion coefficient. Over time step Δt , the bead's rms displacement in 1D relates to the diffusion coefficient as

$$\Delta x_{rms} = \sqrt{2D\Delta t}.$$

Substituting and rearranging leads to

$$D = \frac{\Delta x_{rms}^2}{2\Delta t} = \frac{\mu^2 F_{x,rms}^2 \Delta t}{2}$$

Also, the Einstein relation relates the diffusion coefficient to the temperature, T , and mobility with

$$D = \mu k_B T$$

Combining these equations and solving for the force leads to

$$\begin{aligned} \frac{\mu^2 F_{x,rms}^2 \Delta t}{2} &= \mu k_B T \\ F_{x,rms} &= \sqrt{\frac{2k_B T}{\mu \Delta t}} \end{aligned} \quad (5.33)$$

This is the rms thermal force on each axis for a diffusing bead. It's not entirely physical, due to the presence of the Δt term, but should yield the correct diffusion coefficient. Filament nodes can be modeled as beads, so these are also the forces on the nodes. The only way for the prior force result and this one to both be correct is if the mobility is related to the force constant and time step.

5.11.8 Filament math - filament dynamics

These and other forces push on nodes, causing the nodes to move. Also, these and other torques create rotations in the segments. The displacements are the product of the forces, mobilities, and the time step Δt , which is essentially the overdamped Langevin equation. The displacement motion of a node over one timestep is

$$\Delta \mathbf{x} = \mu \mathbf{F} \Delta t \quad (5.34)$$

where μ is the mobility coefficient.

Additionally, for 3D, the rotation of segment i over one timestep is

$$\Delta\psi = \mu\tau_{R,i}\Delta t \quad (5.35)$$

This uses the same mobility value, μ as for displacement, although this is probably not accurate because rotational friction is likely to be less than displacement friction. However, I'll assume it's the same for now. Note that $\Delta\psi$ is independent of reference frame, but is just the rotation of this one segment around its axis during this time step.

These are described for the Euler method in which the displacement is directly proportional to the time step. The Runge-Kutta method is preferable but requires more force calculations. I had thought that my machine learning method might be better yet, but testing showed that it's worse than Euler for complex problems (see my document *BeadDynamicsNotes.pdf*).

Here is the Euler method math

$$\mathbf{x}(t + \Delta t) = \mathbf{x}(t) + \mu\Delta t\mathbf{F}(t, \mathbf{x})$$

Suppose the force fuction is expressed as a linear function of $\mathbf{x}(t)$,

$$\mathbf{F}(t, \mathbf{x}) = \mathbf{F}_0(t) + \mathbf{F}_1(t)\mathbf{x}(t)$$

Here, $\mathbf{F}_0(t)$ is a vector and $\mathbf{F}_1(t)$ is a matrix. With this, the Euler method can be rewritten as

$$\mathbf{x}(t + \Delta t) = \mathbf{x}(t) + \mu\Delta t(\mathbf{F}_0(t) + \mathbf{F}_1(t)\mathbf{x}(t))$$

The semi-implicit method is the same, but the $\mathbf{x}$ value on the right hand side is evaluated at $t + \Delta t$ rather than at t . The fully implicit method would also evaluate the forces at the later time as well, but that's harder to accomplish. Here's the semi-implicit method,

$$\mathbf{x}(t + \Delta t) = \mathbf{x}(t) + \mu\Delta t(\mathbf{F}_0(t) + \mathbf{F}_1(t)\mathbf{x}(t + \Delta t))$$

Some rearranging leads to

$$\begin{aligned} [\mathbf{1} - \mu\Delta t\mathbf{F}_1(t)]\mathbf{x}(t + \Delta t) &= \mathbf{x}(t) + \mu\Delta t\mathbf{F}_0(t) \\ \mathbf{x}(t + \Delta t) &= [\mathbf{1} - \mu\Delta t\mathbf{F}_1(t)]^{-1}(\mathbf{x}(t) + \mu\Delta t\mathbf{F}_0(t)) \end{aligned}$$

I've been trying to compute $\mathbf{F}_1(t)$ above, but I'm having a hard time for the bending component. My new proposal is that I do it numerically instead.

5.11.9 Filament simulation approach

The steps in filament simulation are:

(1) Compute the forces on the nodes and torques on the segments using the data stored in `fil->nodes` and `segment->ypr`. Performed by `filComputeForces`. The forces (`fil->forces`) are stored in the lab reference frame and the segment torques (`fil->torques`) are stored in the local segment reference frame. These forces and segment torques arise from several independent sources. (a) Segment stretching and compression creates forces on both the front and back nodes for each segment and are given by eq. 5.13. These don't affect segment torques. These are computed in `filAddStretchForces`. (b) There are ypr springs at each segment junction that create forces at that particular node, as well as at both neighboring nodes. They are computed by first finding the 3D local torque with eq. 5.24 (in `filBendTorque`), which is then converted to the lab frame with

5 C Code: structures and functions

5.25 (still in `filBendTorque`). Combining this torque with the segment displacements yields the node forces and segment torques given in eq. 5.26. These are computed in `filAddBendForces`. (c) There are also forces from external influences, such as surfaces, which are not included yet. Another force is at filament ends from branches (d) Thermal forces can be computed from either eq. 5.33 or 5.32, but I don't know which is better; the latter one is implemented in `filAddThermalForces` but is not included in the total force calculation yet. Although not done yet, it might be helpful to characterize forces by their stiffness, computed as the derivative of the force with respect to the displacement.

(2) Perform node displacement and segment rotation, which is performed by `filEulerDynamics`. The nodes move in direct proportion to the applied forces, given in eq. 5.34. Also, the segment roll values move in direct proportion to the applied segment torques, given in eq. 5.35. These latter values are stored in the segment ypr vectors, although recognizing that the yaw and pitch values are no longer current.

(3) Convert node positions, which are in the lab frame, to local ypr angles and also compute both dcm matrices. This is performed in `filNodes2Angles`. To do so, walk forward along a filament, computing each segment displacement in the prior segment's reference frame with the left side of 5.9. Then, use these displacements to compute local θ_i and ϕ_i values from eqs. 5.10 and without changing each ψ_i , thus yielding the $\mathbf{A}_i^T$ matrices on the right side of eq. 5.9. Multiply $\mathbf{A}_i \mathbf{B}_{i-1}$ to get $\mathbf{B}_i$, from eq. 5.7. Then, walk forward to the next segment.

5.11.10 Filament math - numbers

According to my 2014 paper [1], a filament in the Angle-Biased Chain (ABC) model with mean segment lengths of b and pitch and yaw force constants of k_θ and k_ϕ has a persistence length of

$$P = \frac{2bk_\theta k_\phi}{k_B T (k_\theta + k_\phi)}$$

Suppose this filament is rotationally isotropic, so $k_\theta = k_\phi$. This simplifies the persistence length to

$$P = b \frac{k_\phi}{k_B T}$$

Rearranging yields

$$k_\phi = k_B T \frac{P}{b}$$

Note that the bending force constants k_ϕ and k_θ have units of energy.

For values, $k_B T \approx 4.28 \cdot 10^{-21}$ J = 4.28 zJ for $T = 37^\circ\text{C}$. Suppose I am modeling actin, which has a persistence length of 17 μm , and I am using 1 μm segments. Its bending force constant would be 72.76 zJ.

5.11.11 Filament status

At present (5/22/25), it's possible to define basic filaments and to run their dynamics. Dynamics methods include Euler, RK2, RK4, an old version of implicit Euler that doesn't work, and a new version of implicit Euler that does work. The data structures allow for filament branching but this isn't implemented in the code yet. Thermal forces are ad hoc and not quantitatively correct. Implicit dynamics are functional but could be made more efficient (many more bending forces are evaluated than should be necessary when building the numerical matrix and also stretching forces and thermal forces should not be treated numerically since they can be done analytically). Also,

dynamics need validation. A minor issue that needs fixing is that the implicit integration methods don't account for local mobility factors. The next major steps are to add filament branching and filament interactions. Also, it might be good to allow for non-uniform standard segment lengths.

For documentation, there is actually too much detail on the theory here in this documentation. I have moved all theory to my FilamentNotes.tex document, so all work should be there and only the essential results and equations should be here.

5.11.12 Functions.

As much as possible, functions are declared locally rather than in the *smoldynfuncs.h* header file. This simplifies the code reading because it clarifies which functions might be called externally versus those that are only called internally. In practice, many functions are listed both places, but they shouldn't be.

Enumerations (No changes to data structures.)

```
char *filFD2string(enum FilamentDynamics fd, char *string)
```

Local. Converts filament dynamics type to string. Returns "none" for unrecognized dynamics type. Writes result to *string* and also returns it directly.

```
enum FilamentDynamics filstring2FD(const char *string)
```

Local. Converts filament dynamics string to enumerated dynamics type. Returns *FDnone* for unrecognized input.

Low level utilities (No changes to data structures; no math.)

```
int filReadFilName(const simptr sim, const char *str, filamenttypeptr
    *filtypeptr, filamentptr *filptr, char *filename)
```

Reads a filament name in *str* and returns the filament type in *filtypeptr*, the filament in *filptr*, and the filament name in *filename*. This accepts a name in several possible formats. (1) It allows the format "filtype:filename", where a colon separates the type name from the filament name; if the filament type is recognized but not the filament name, then this sends back the filament type and returns -5 (see below for return codes). (2) If the filament type is known, then it can be sent in using *filtypeptr*, and only that type will be searched for the filament name. (3) If the filament type is not known, then this searches all filament types and, if found, returns the type in *filtypeptr* and the filament in *filptr*.

Note that *filtypeptr* and *filptr* cannot be NULL and, the value of *filtypeptr* must equal NULL if the filament type is unknown or to a filament type if the type is known (*filptr* does not need to be initialized). *filename* is allowed to equal NULL if the filament name is not of interest. Returns 0 for success, -1 for no *str*, -2 for no filaments have been defined yet, -3 for *str* is unreadable, -4 for unrecognized filament type name, or -5 for unrecognized filament name.

```
int filGetFilIndex(const simptr sim, const char *name, int *ftptra)
```

Returns the index of the filament named *name* or -1 if not found; also, returns the filament type in *ftptra*. This searches through all filament types for the filament with this name. If multiple filaments with the same name are found, this returns -2. *ftptra* should be sent in pointing to a memory address, but the contents are unimportant. Example: *filGetFilIndex(sim, "fil1", &ft)*.

Memory management (Allocate, free, and initialize data structures.)

`segmentptr filSegmentAlloc()`

Allocates memory for a single segment and initializes this segment. Returns a pointer to this segment, or NULL if memory could not be allocated. The segment length is initialized to 0 to indicate that it hasn't been set up yet, and the thickness is initialized to 1, which is good default value. `xyzfront` and `xyzback` are set to NULL. Ypr angles are initialized to 0 and direction cosine matrices to unit matrices.

`void filSegmentFree(segmentptr segment)`

Frees memory for a single segment.

`filamentworkptr filWorkAlloc(filamentptr fil, int maxseg)`

Allocates a data structure for filament dynamics data. Enter the filament in `fil` and the maximum number of segments that are currently allowed in `maxseg`. All working data are allocated to this filament size. The function sets up all necessary internal pointers. This function can be called repeatedly to enlarge the working data (but old data are not kept if so).

`void filWorkFree(filamentworkptr filwork, int maxseg)` Frees working data and the data structure.

`filamentptr filAlloc(filamentptr fil, int maxseg, int maxbranch, int maxsequence)`

Allocates or expands a filament, returning a pointer to the result. This does not shrink any elements of filaments. Enter `maxseg` to expand the number of segments, `maxbranch` to expand the number of branches off of this filament, and `maxsequence` to expand the sequence list. Any of these can be set to 0 to ignore this component. This returns NULL if memory could not be allocated but does not free things, which can cause a minor memory leak, but is better than risking a segfault and easier than a sophisticated analysis of what can and cannot be freed.

`void filFree(filamentptr fil)`

Frees a filament and all of the data structures in it.

`filamenttypeptr filamentTypeAlloc(filamenttypeptr filtype, int maxfil, int maxface)`

Allocates or expands memory for a filament type structure and initializes it. Returns a pointer to this filament type, or NULL if memory could not be allocated. Enter `filtype` as an existing one if it is to be expanded, `maxfil` as the number of filament spaces to be allocated, and `maxface` as the number of face names to be allocated. This only expands the filament or face lists if the requested sizes are larger than the current sizes. When putting new filaments into the list, they are allocated here and initialized, but have no segments.

`void filamentTypeFree(filamenttypeptr filtype)`

Frees memory for a single filament type.

`filamentssptr filssAlloc(filamentssptr filss, int maxtype)`

Allocates or expands a filament superstructure for a maximum size of `maxtype` filament types.

`void filssFree(filamentssptr filss)`

Frees a filament superstructure and all of the structures within it.

Data structure output (No changes to data structures.)

```
void filOutput(const filamentptr fil)
```

Outputs all of the key information about a filament to the display, including essential information about each segment.

```
void filtypeOutput(const filamenttypeptr filtype)
```

Outputs all of the key information about a filament type to the display, along with information about each filament.

```
void filssOutput(const simptr sim)
```

Outputs all of the key information about a filament superstructure, and all of the filaments in it, to the display.

```
void filWrite(const simptr sim, FILE *fptr)
```

Writes filament information to a file in Smoldyn format for loading in again later on. NOT WRITTEN YET.

```
int filCheckParams(const simptr sim, int *warnptr)
```

Checks filament parameters to make sure they are all reasonable. I've started writing this but much more is needed.

Filaments (Data structures are changed.)

```
int filSetNodeMobility(filamentptr fil, int node, double value)
```

Sets the unitless node mobility for node `node` to `value`. Enter `node` as 0 for the front end of the filament, as 1 for the next node, etc. Also, `node` can be entered as `-1` for the back of the filament, `-2` as the preceding node, etc. Returns 0 for success or 1 for `node` out of range.

Filament type (Data structures are changed.)

```
int filtypeSetParam(filamenttypeptr filtype, const char *param, int index, double value)
```

Sets various filament type parameters. The parameter name is `param` and it is set to the value `value`. If this parameter has multiple indices, enter the index to be changed in `index`, or enter `index` as `-1` to set all of the indices at once to the same value. Returns 0 for success or 2 for out of bounds value.

<code>param</code>	Description	<code>index</code>	Range
<code>stdlen</code>	standard length	n/a	$[0, \infty)$
<code>stdypr</code>	standard ypr angles	0,1,2	$[-\pi, \pi]$
<code>klen</code>	stretching force constant	n/a	$-1, [0, \infty)$
<code>kypr</code>	bending force constants	0,1,2	$-1, [0, \infty)$
<code>kT</code>	thermodynamic temperature	n/a	$[0, \infty)$
<code>treadrate</code>	treadmill rate constant	n/a	$(-\infty, \infty)$
<code>mobility</code>	filament node mobility	n/a	$(0, \infty)$
<code>bundle</code>	bundle value	n/a	$(0, \infty)$
<code>radius</code>	filament radius	n/a	$(0, \infty)$
<code>facetwist</code>	face twist rate	n/a	$(-\infty, \infty)$

```
int filtypeSetColor(filamenttypeptr filtype, const double *rgba)
```

Set the 4-value color vector of the filament type to the values entered in `rgba`. Returns 0

5 C Code: structures and functions

unless one of the entered values is outside of the range $[0, 1]$, in which case this does not change the current color and returns an error code of 2.

int `filtypeSetEdgePts(filamenttypeptr filtype, double value)`

Sets the drawing thickness of the filament type to the value entered in `value`. Returns 0 for success or 2 if the entered value is less than 0, which is an error.

int `filtypeSetStipple(filamenttypeptr filtype, int factor, int pattern)`

Sets the stippling factor and pattern for the filament type to the entered values. Enter a negative number to not set the given parameter. The factor should be a positive value and the pattern should be a value between 0 and 0xFFFF. Returns 0 for success or 2 if values are out of bounds.

int `filtypeSetDrawmode(filamenttypeptr filtype, enum DrawMode dm)`

Sets the filament type drawing mode to the entered value. Returns 0 for success or 2 if the input value is unrecognized (i.e. equal to `DMnone`).

void `filtypeSetDrawForceArrows(filamenttypeptr filtype, double scale, const double *rgba)`

Sets the `drawforcescale` and/or `drawforcecolor` elements. Enter `scale` as 0 or a positive value to set this parameter and negative to ignore it; a value of 0 implies that no arrows are drawn. The `rgba` vector is ignored if it is NULL and is otherwise copied into `drawforcecolor`. Each value should be between 0 and 1, inclusive.

int `filtypeSetShiny(filamenttypeptr filtype, double shiny)`

Sets the shininess for the filament type to the entered value, which should be within $[0, 128]$. Returns 0 for success or 2 if the value is out of bounds.

int `filtypeSetDynamics(filamenttypeptr filtype, enum FilamentDynamics fd)`

Sets the dynamics for the filament type to `fd`.

int `filtypeAddFace(filamenttypeptr filtype, const char* facename)`

Adds a face name to the filament, enlarging the list of face names if required. Returns 0 for success and -1 for failure to allocate memory.

filamenttypeptr `filAddFilamentType(simptr sim, const char *ftname)`

Adds a new filament type to the simulation, which is named `ftname`, returning a pointer to the new filament type. Both `sim` and `ftname` need to be defined. If there is already a filament type with this name, a pointer to it is returned. Returns a pointer to the new filament type or NULL for failure, meaning that memory could not be allocated.

Filament superstructure (Superstructure is changed and/or lower levels are updated.)

void `filSetCondition(filamentssptr filss, enum StructCond cond, int upgrade)`

Local. This function sets the condition of the filament superstructure. Set `upgrade` to 1 if this is an upgrade, to 0 if this is a downgrade, or to 2 to set the condition independent of its current value. If the resulting condition is lower than the `sim` condition, it downgrades the `sim` condition.

int `filEnableFilaments(simptr sim)`

Allocates a filament superstructure, adds it to the simulation structure, and sets the filament condition to `SCLists`. This does not create any filament types or other items that are lower in the hierarchy.

int filUpdateParams(simptr sim)

Does nothing currently. This function will compute any simulation parameters necessary from user parameters.

int filUpdateLists(simptr sim)

Does nothing currently. This function will compute any simulation list changes necessary from user parameters.

int filUpdate(simptr sim)

Updates filaments as needed and upgrades the condition. Returns 0 for success or other error codes passed through from `filupdatelists` and `filupdateparams`.

User input (Data structures are changd.)

filamenttypeptr filtypeReadString(simptr sim, ParseFilePtr pfp, filamenttypeptr filtype, const char *word, char *line2)

Reads one line of user input within a filament type block.

filamentptr filReadString(simptr sim, ParseFilePtr pfp, filamentptr fil, const char *word, char *line2)

Reads one line of user input within a filament block.

int filLoadType(simptr sim, ParseFilePtr *pfpPtr, char *line2)

Reads multiple lines of user input from within a filament type block, calling `filtypeReadString` with each line.

int filLoadFil(simptr sim, ParseFilePtr *pfpPtr, char *line2)

Reads multiple lines of user input from within a filament block, calling `filReadString` with each line.

Sequences (Data structures may be changed.)

int filSetSequence(filamentptr fil, char *str)

Sets the filament sequence to `str`, overwriting any previous sequence. This allocates memory and updates pointers as needed.

Computations on segments (No changes to data structures.)

double filRandomLength(const filamenttypeptr filtype, double thickness, double sigmamult)

Local. Returns a random segment length using the mechanics parameters given in filament type `filtype`. `thickness` is the thickness of the new segment and `sigmamult` is multiplied by the normal standard deviation of the length. If $k_{len} < 0$, this returns the standard length; otherwise, the returned length is Gaussian distributed, with mean equal to `fil->lstd` and standard deviation, which is constrained to be always positive. From eq. 5.30 it is equal to

$$\sigma_{mult} \sqrt{\frac{kT}{R * k_{len}}}$$

double *filRandomAngle(const filamenttypeptr filtype, int n, double thickness, double sigmamult, double *angle)

5 C Code: structures and functions

Local. Returns a random bending angle, which is a relative ypr angle, in **angle** (which needs to be allocated with 3 elements) and directly. This is for filament type **filtype** and segment number **n**; this **n** number is only checked to see if it's zero or non-zero; in the latter case, angles are chosen from angle bending forces. Angles are for a segment with thickness **thickness** and the computed standard deviation is multiplied by **sigmamult**. For **n>0**, if k_{ypr} is less than 0 for a particular coordinate, this returns the standard bending angle, if it is equal to 0, this returns a uniformly distributed bending angle, and if it is greater than 0, this returns a mean of ypr_{std} and standard deviation of (see eq. 5.31

$$\sigma_{mult} \sqrt{\frac{kT}{R * k_{ypr}}}.$$

For 2D space, this sets the last two elements of **angle** to 0. The Gaussian distributions, for 2D and 3D angles, are truncated so that the angle stays within appropriate bounds. For yaw angles, in both 2D and 3D, this is from $-\pi$ to π ; for pitch angles, this is from $-\pi/2$ to $\pi/2$, and for roll angles, this is from $-\pi$ to π .

For segment 0, angles are chosen uniformly over all possibilities. For 2D, the returned angle, which is in the first element only (yaw) is relative to the system $\hat{x}$ vector.

Computations on filaments (No changes to data structures.)

double *filGetPositionOnFil(const filamentptr fil, double frac, double *pos)

Returns the coordinates of the point that is **frac** fraction of the pathlength along the filament, putting the results in **pos** and also returning them directly. This computes the total filament pathlength, then finds the segment that is **frac** distance along the filament, and then interpolates within the segment to find the coordinates. This function finds the position as a fraction of the actual pathlength, and not of the relaxed pathlength.

double filStretchEnergy(const filamentptr fil, int seg1, int seg2)

Computes the stretching energy for filament **fil**. Enter **seg1** as the first segment to compute from, or as -1 to indicate that it should start at the front of the filament, and **seg2** as one more than the last segment to compute to, or as -1 to indicate that calculation should end at the end of the filament. The equation is

$$E_{stretch} = \sum_s \frac{R_s k_{len} (l_s - l^o)^2}{2}$$

double filBendEnergy(const filamentptr fil, int seg1, int seg2)

Computes the bending energy for filament **fil**. Enter **seg1** as the first segment to compute from, or as -1 to indicate that it should start at the front of the filament, and **seg2** as the last segment to compute to, or as -1 to indicate that calculation should end at the end of the filament. Here, **seg2** is inclusive of the last segment; the bends that are included here start at the back of **seg1** and continue up to the front of **seg2**; there are one fewer bends in a filament than the number of segments. The equation is

$$E_{bend} = \sum_{s=1}^n \frac{R_{s-1} + R_s}{2} \frac{k_y (a_y - a_{y,std})^2 + k_p (a_p - a_{p,std})^2 + k_r (a_r - a_{r,std})^2}{2}$$

5.11 Filaments (functions in *smolfilament.c*)

The first term in the sum computes the average thickness of the segment in front of and behind the bend. The other term is the squared bending angle on each coordinate (for 3D; only the first term is included for 2D).

void filBendTorque(const filamentptr fil, int node, double *torque)

Computes the bending torque vector for node number `node`, which is the front of segment `node`, in filament `fil`. This uses the data stored in `segment->ypr`. The `torque` vector needs to be pre-allocated. This uses eq. 5.24 and then rotates the torque with the absolute quaternion of the prior segment to express the torque in absolute coordinates, which is how it's returned (in matrix notation, $\tau_L = \mathbf{B}^T \tau$). This accepts `node` values of 0 and `nseg`, returning torques of zero in both cases.

Filament updating (Filament is unchanged but data structures are updated.)

void filNodes2Angles(filamentptr fil, int nodemin, int nodemax)

Computes angle information for a filament, which are this function's output, from its node positions and roll values (`segment->ypr[2]`), which are the input. In more detail, the input information is in the segment front and back elements (`xyzfront` and `xyzback`, which are the same memory elements as in `fil->nodes`) and the segment roll values (`segment->ypr[2]`). It computes segment lengths (`len`), ypr values (`ypr[0]` and `ypr[1]`), local quaternions (`qrel`), and absolute quaternions (`qabs`).

Enter `nodemin` as -1 (or ≤ 0) to start at the front and `nodemax` as -1 (or $\geq n_{seg}$) to end at the back. Otherwise, enter these values for the first and last nodes, inclusive, that should be updated by this function. For example, if node i is moved, then the absolute quaternions change for segments $i - 1$ and i , and the angles and relative quaternions change for nodes $i - 1$, i , and $i + 1$, so run this with `nodemin` equal to $i - 1$ and `nodemax` equal to $i + 1$.

void filCopyRollToYpr2(filamentptr fil)

Copies from the `fil->roll` vector to the segment `ypr[2]` elements.

void filCopyYpr2ToRoll(filamentptr fil)

Copies from the segment `ypr[2]` elements to the `fil->roll` vector.

void filCopyWorkingNodes(filamentptr fil, int from, int to)

This is used for the RK2 and RK4 integration methods, and maybe elsewhere. It copies values from the `wnodes[from]` element of the filament working structure to the `wnodes[to]` element in order to save the values. For 3D filaments, it also copies the first segment "up" vector and roll values. If `from` is equal to 0 (and in 3D), then this also copies the `segment->ypr[2]` value, which is the segment roll value, into the `fil->roll` vector.

void filFlattenNodes(filamentptr fil, int from, int to)

This is used for matrix Euler integration and implicit integration. It copies the node position data in the `wnodes[from]` element of the filament working structure to the `flatnodes[to]` element of the working structure. For 3D, this also copies roll data from the `wroll[from]` element of the working structure to `flatnodes[to]`. See Section 5.11.6. If `from` is equal to 0 (and in 3D), then this also copies the `segment->ypr[2]` value, which is the segment roll value, into the `fil->roll` vector.

void filUnflattenNodes(filamentptr fil, int from, int to)

Used for matrix Euler integration and implicit integration. This copies the node position and

5 C Code: structures and functions

roll (for 3D) data in the `flatnodes[from]` element of the filament working structure to the `wnodes[to]` and `wroll[to]` elements of the working structure. See Section 5.11.6. For 3D, this also rotates the `seg0up[from]` vector using the segment 0 roll vector to create the new `seg0up[to]` vector.

`void filFlattenForces(filamentptr fil,int to,int nodemin,int nodemax)`

Used for numerical implicit integration. This copies the force data in the `forces` element of the filament working structure to the `flatforces[to]` element of the working structure. For 3D, this also copies torque data from the `torques` element of the working structure to `flatforces[to]`. This copies force data for the nodes from `nodemin` to `nodemax`, inclusive. Enter `nodemin` as -1 (or ≤ 0) to start at the front and `nodemax` as -1 (or $\geq n_{seg}$) to go to the back.

Filament manipulation (Filament is changed.)

`void filArrayShift(filamentptr fil, int shift)`

Shifts the segment and nodes lists in the filament either to higher indices or to lower indices (indices in the array, not the segment `index` element). Enter `shift` as a positive number to increase all the indices, as a negative number to decrease all the indices, and as 0 to not move at all. If `shift` is made too large or too small, such that the filament would move off the end of the allocated memory, then it is automatically fixed so that the filament just gets shifted to the end of the allocated memory instead. Thus, an easy way to make sure that the filament is at the end of the allocated memory is to enter `shift` as $\pm \text{fil} \rightarrow \text{maxseg}$.

`int filAddSegment(filamentptr fil, const double *x, double length, const double *angle, double thickness, char endchar)`

Adds a segment to filament `fil`. If this is the first segment, then `x` becomes the starting location of the filament; otherwise, `x` is ignored. `length` is the length of the segment, `angle` is the relative angle of the segment facing along the filament from front to back (for the first segment, this is the absolute angle), `thickness` is the segment thickness, and `endchar` should be set to 'b' to add to the back of the filament or 'f' to add to the front of the filament. For segments added to the front, `angle` is the new angle from the new first monomer to the next monomer, facing towards the back. If the first segment is added to the back, then `angle` is the angle of the new segment off of the coordinate system. If the first segment is added to the front, then `angle` is the angle from the new segment to the coordinate system. The `angle` vector must have 3 elements in it, even for 2D simulations; in the latter case, the latter two elements should be set to 0. Returns 0 for success and 1 for out of memory.

For the first segment: $\mathbf{x}_0 = \mathbf{x}$, $\mathbf{x}_1 = \mathbf{x}_0 + l_0 \mathbf{B}_0^T \cdot \hat{\mathbf{x}}$

For segments added to the back, with index i : $\mathbf{ypr}_i = \mathbf{angle}$, $\mathbf{A}_i = DCM(\mathbf{ypr}_i)$, $\mathbf{B}_i = \mathbf{A}_i \cdot \mathbf{B}_{i-1}$, $\mathbf{x}_{i+1} = \mathbf{x}_i + l_i \mathbf{B}_i^T \cdot \hat{\mathbf{x}}$

For segments added to the front, with index i (typically equal to 0): $\mathbf{B}_i = \mathbf{A}_{i+1}^T \cdot \mathbf{B}_{i+1}$, $\mathbf{A}_i = \mathbf{B}_i$, $\mathbf{ypr}_i = YPR(\mathbf{B}_i)$, $\mathbf{x}_i = \mathbf{x}_{i+1} - l_i \mathbf{B}_i^T \cdot \hat{\mathbf{x}}$

`int filAddRandomSegments(filamentptr fil,int number,const char *xstr,const char *ystr,const char *zstr,const char *phistr, const char* thtstr,const char* psistr,double thickness)`

Add `number` of random segments to the back filament `fil`, while ignoring any surfaces or other potential collisions. Enter `xstr`, `ystr`, and `zstr` as strings for the starting position coordinates of a new filament, and `phistr`, `thtstr`, and `psistr` as strings for the starting

5.11 Filaments (functions in *smolfilament.c*)

angle of a new filament using Tait-Bryan (yaw-pitch-roll) angles. Here, math equations are allowed, using Smoldyn variables, or use the ‘u’ character to indicate uniform starting position within the system volume or angle in spherical coordinates. **thickness** is the thickness of the segments being added. For 2D simulations, the z , ϕ , and χ values are automatically set to zero, regardless of what is entered in **zstr**, **phistr**, and **chistr**.

Returns 0 for success, 2 for syntax error in position strings, or 3 for syntax error in angle strings.

```
int filAddOneRandomSegment(simpstr sim, filamentptr fil, const double *x, double
    thickness, char endchar, int constraints)
```

Adds one segment to the **endchar** end of **fil** using a random length and bending angle. The **x** input is ignored unless the filament has 0 current segments. The new segment has thickness **thickness**. The **constraints** input is supposed to be a collection of flags for which constraints the addition needs to account for. For now, the only **constraints** options are 0 for no constraints at all, and 1 for the new segment should not cross any surface panel. If the addition is constrained, this function tries up to **FILMAXTRIES** times to add a segment that does not disobey the constraints. Returns 0 for success, 1 for out of memory, and 2 for failure to place a segment that obeys the constraints.

```
void filTreadmill(simpstr sim, filamentptr fil, char endchar)
```

This performs one step of treadmilling for filament **fil** at the **endchar** end. It adds one segment at this end and removes one segment from the opposite end.

```
int filRemoveSegment(filamentptr fil, char endchar)
```

Removes one segment or one bead from either the front or back end of a filament. Specify the end in **endchar** with ‘f’ for front and ‘b’ for back. Returns 0 for success or -1 for a filament that already had no segments.

Removing a segment from the back requires only shortening the filament by one segment. Removing one from the front requires copying the absolute dcm or quaternion into the relative dcm or quaternion for the new front segment, and then translating this into the ypr angles.

```
void filTranslate(filamentptr fil, const double *vect, char func)
```

Translates an entire filament in space without rescaling or rotation. Enter **vect** with a 3D vector and **func** with ‘=’ for translate to the given position, with ‘-’ for to subtract the value of **vect** from the current position, and with ‘+’ to add the value of **vect** to the current position. For 2D simulations, **vect**[2] needs to equal 0.

```
void filTranslateNode(filamentptr fil, int node, const double *vect, char func)
```

Translates node **node**. Enter **func** as ‘=’ to set the value to **vect**, as ‘-’ to subtract **vect**, or as ‘+’ to add **vect**. This calls **filNodes2Angles** to update the angles and quaternions as needed.

```
int filLengthenSegment(filamentptr fil, int seg, double length, char endchar,
    char func)
```

This function modifies the length of a single segment of index **seg** in filament **fil**. The **endchar** input can be **b** or **f** and gives which end of the segment moves, along with the entire rest of that end of the filament. The **func** input can be **=**, **+**, or **-** and gives whether the segment length should be made equal to **length**, increased by **length**, or decreased by **length**. Returns 0 for success or an error code of 1 if the new segment length would be 0 or negative.

5 C Code: structures and functions

This function computes the change of length, and then multiplies that by the local segment orientation to create the shift vector `xdelta`. This is added to the following or preceding nodes, depending on `endchar`.

```
int filChangeThickness(filamentptr fil, int seg, double thick, char func)
```

This function modifies the thickness of a single segment of index `seg` in filament `fil`. The `func` input can be `=`, `+`, or `-` and gives whether the segment thickness should be made equal to `thick`, increased by `thick`, or decreased by `thick`. Returns 0 for success or an error code of 1 if the new segment thickness would be negative.

```
void filRotateVertex(filamentptr fil, int seg, const double *angle, char endchar, char func)
```

This function rotates part of the filament about the vertex that is at the front of segment `seg` by ypr angle `angle`. Enter `endchar` as `b` to have the back of the filament move and as `f` to have the front move. Enter `func` as `=` to set the current relative rotation angles (`segment->dcm`) to the input ones, as `+` for increased rotation by the given angles and as `-` for decreased rotation by these angles.

This function makes the requested modification to the segment's relative quaternion. Then, it propagates the changes either backward or forward as needed. During propagation, new segment positions are computed from neighbor positions, rather than being rotated independently, in order to avoid round-off errors.

```
int filCopyFilament(filamentptr filto, const filamentptr filfrom, const filamenttypeptr filtype)
```

This function copies nearly all of the values from `filfrom` to `filto`, overwriting any prior data in `filto` in the process. It does not copy the name. Also, it assigns the filament type to `filtype` if this value is not NULL, and to the `filfrom` filament type otherwise. Returns 0 for success, 1 for out of memory, or 2 for illegal function inputs.

```
filamentptr filAddFilament(filamenttypeptr filtype, const char *filename)
```

This adds a filament to the filament type `filtype`, expanding the list if needed. If there is already a filament with this name, a pointer to it is returned. If `filename` is NULL, this generates an automatic name for the new filament; otherwise, it names the new filament to `filename`. Returns a pointer to the new filament or NULL for failure, arising from failure to allocate memory.

Filament interactions (Data structures are changed.)

```
int filSegmentXSurface(const simptr sim, const segmentptr segment, panelptr *pnlptr)
```

Tests if segment `segment` crosses any surface in the simulation. Returns 0 if not and 1 if so. If so, then this returns a pointer to the panel that is crossed in `pnlptr`. Many more details about the crossing could be returned, such as where along the segment the crossing is, if there are multiple crossings with curved panels, etc., but these are not being returned at present. This function scans through the possible surface panels and checks each with `lineXpanel`. THIS IS VERY INEFFICIENT AND SHOULD USE BOXES.

```
int filSegmentXFilament(const simptr sim, const segmentptr segment, filamentptr *filptr)
```

Tests if segment `segment` crosses any filament in the simulation. Returns 0 if not and 1 if so.

If so, then this returns a pointer to the filament that is crossed in `filptr`. More details about the crossing could be returned, such as which segment within the filament, but this is not supported at present. This function scans through all filaments in the simulation and tests for crossing for each one, meaning that the segment radius overlaps the radius of the other filament (this uses the `thk` element as the radius). This ignores possible overlaps with the self-segment, along with its nearest neighbor on either side. THIS IS VERY INEFFICIENT AND SHOULD USE BOXES.

Force computation (Only writing to `forces`, `torques`, `forcemat` elements.)

`void filAddStretchForces(filamentptr fil,int nodemin,int nodemax)`

Computes filament stretching (and compression) forces using the node positions in `fil->nodes`. This adds forces to any prior values in the `fil->forces` vector, which is in the system reference frame. These forces do not affect segment torsion. They are computed from eq. 5.13. Set `nodemin` to -1 (or ≤ 0) to start at the filament front and set `nodemax` to -1 (or $\geq n_{seg}$) to end at the filament back. Otherwise, forces are computed for segment numbers `nodemin` to `nodemax-1`, inclusive. For example, if node i is moved, then segments $i-1$ and i change lengths, so evaluate the new forces with `nodemin` as $i-1$ and `nodemax` as $i+1$.

`void filAddStretchForceMat(filamentptr fil)`

Computes filament stretching (and compression) forces using the node positions in `fil->nodes`. This adds forces to any prior values in the `filwork->forcemat` matrix. They are computed from eq. 5.14.

`void filAddThermalForces(filamentptr fil,int nodemin,int nodemax)`

Computes thermal forces acting on filament nodes and adds them to any prior values in the `fil->forces` vector. This stores the forces in the `filwork->thermforce` vector so that they only get computed once per time step. This function uses forces from eq. 5.32, but it's clearly incorrect; in part, this uses the yaw force constant only. Also, there should be thermal torsion as well, which should be added to the `fil->torques` vector, but is not computed at present. As usual, enter `nodemin` as the starting node or -1 (or ≤ 0) for the filament front and enter `nodemax` as the ending node, inclusive, or as -1 (or $\geq n_{seg}$) for the filament back. NEEDS WORK.

`void filAddBendForces(filamentptr fil,int nodemin,int nodemax)`

Computes filament bending forces due to ypr springs using the ypr values in `segment->ypr`. This adds forces to any prior values in the `fil->forces` vector, which is in the system reference frame. For 3D filaments, bending forces can create torsion for individual segments, which are added to any prior values in `fil->torques`. These results are computed from eqs. 5.26.

As usual, enter `nodemin` as the starting node or -1 (or ≤ 0) for the filament front and enter `nodemax` as the ending node or as -1 (or $\geq n_{seg}$) for the filament back. *However*, this only computes bending forces arising from bends that are between `nodemin` and `nodemax`, exclusive of the endpoints, which are then applied to nodes from `nodemin` to `nodemax`, inclusive of the endpoints. This may not be as useful as it could be, so this function needs work.

`void filAddBendForceMat(filamentptr fil)` Computes filament bending forces due to ypr springs using the ypr values in `segment->ypr`, adding the forces to any prior forces in the `filwork->forcemat` matrix. These forces are computed from eq. 5.29. This matrix is

5 C Code: structures and functions

technically correct but of minimal use because the equations are not linearized correctly and so it is not lead to stable implicit integration.

void filComputeForces(filamentptr fil,int nodemin,int nodemax)

Computes all forces and segment torques acting on filament **fil**. This clears the **filwork->forces** and **filwork->torques** vectors and then calls **filAddStretchForces**, **filAddBendForces**, and potentially other functions to compute the complete forces acting on the filament. The current version is adequate for basic operation but will need expansion for surface interactions, molecule interactions, thermal interactions, etc.

Enter **nodemin** as the starting node or -1 (or ≤ 0) for the filament front and enter **nodemax** as the ending node or as -1 (or $\geq n_{seg}$ for the filament back. This computes all forces within this range, inclusive of the endpoints.

void filComputeForceMatrix(filamentptr fil)

Computes all forces and segment torques acting on filament **fil**, putting the results in the **filwork->forcemat** matrix.

Dynamics functions

void filStepDynamics(filamentptr fil,double dtmu,int in,int out)

Computes one step for explicit dynamics simulations, which is essentially an Euler step but is also used for RK2 and RK4 functions. **fil** is the filament and **dtmu** is the time step times the mobility for this particular step; this is used as is for Euler and divided by various values for RK2 and RK4. This works with the **filwork** data structure and reads from the **in** node positions, and related data, and writes to the **out** node positions, and related data. This adds $\mu\Delta t\mathbf{F}$ to the input node positions. For 3D, this also uses the torque of the first segment to update its orientation, and uses the torques of other segments to update their orientations. See Section 5.11.9.

void filEulerDynamics(simptr sim,filamenttypeptr filtype)

Integrates filament dynamics over one time step using the Euler method. This loops over all filaments of the filament type. For each one, it computes the forces using **filComputeForces** which get stored in **fil->forces**. It then updates the node positions with the **filStepDynamics** function. At the end, it converts the node positions back to angles.

Euler

force = stretch + bend forces

nodes += μdt force

ypr from **nodes**

void filRK2Dynamics(simptr sim,filamenttypeptr filtype)

Integrates filament dynamics over one time step using the RK2 method. This loops over all filaments of the filament type.

RK2

(1) **force** = stretch + bend forces

nodes1 = **nodes**

nodes += $\mu dt/2$ force

5.12 BioNetGen (functions in smolbng.c)

```
ypr from nodes
(2) force = stretch + bend forces
nodes = nodes1 +  $\mu dt$  force
ypr from nodes
```

```
void filRK4Dynamics(simptr sim,filamenttypeptr filtype)
```

Integrates filament dynamics over one time step using the RK4 method. This loops over all filaments of the filament type.

```

RK4
-----
(1) force = stretch + bend forces
nodes1 = nodes
nodes +=  $\mu dt/2$  force
nodes2 = nodes1 +  $\mu dt/6$  force
ypr from nodes
(2) force = stretch + bend forces
nodes = nodes1 +  $\mu dt/2$  force
nodes2 +=  $\mu dt/3$  force
ypr from nodes
(3) force = stretch + bend forces
nodes = nodes1 +  $\mu dt$  force
nodes2 +=  $\mu dt/3$  force
ypr from nodes
(4) force = stretch + bend forces
nodes = nodes2 +  $\mu dt/6$  force
ypr from nodes
```

```
void filEulerMatDynamics(simptr sim,filamenttypeptr filtype)
```

Integrates filament dynamics over one time step using the Euler method, but with matrix approach rather than a direct approach. This loops over all filaments of the filament type. This is supposed to return identical results to the `filEulerDynamics` function, but is used as a matrix test system for implicit integration.

```
void filImplicitDynamics(simptr sim,filamenttypeptr filtype)
```

Integrates filament dynamics over one time step using the implicit Euler method. This loops over all filaments of the filament type. This function works but doesn't yield stable results.

```
int filDynamics(simptr sim)
```

Runs all dynamics for filaments. At present, this supports treadmilling and conformational dynamics using the Euler, RK2, RK4, and EulerMat integration methods.

5.12 BioNetGen (functions in smolbng.c)

BioNetGen is a separate software tool that expands biochemical reaction network rules to form complete biochemical reaction networks. Weiren Cui wrote a short perl program, called b2s, which reads BioNetGen output (.net file) and converts that into Smoldyn input. Inspired by it, I wrote a more elaborate interpreter into the main Smoldyn source code (in C) so that Smoldyn can read these files directly. The plan is that Smoldyn will be distributed with the basic BioNetGen perl

5 C Code: structures and functions

program. Then, when Smoldyn encounters rules within an input file, it will automatically call BioNetGen, BioNetGen will expand the rules into a .net file, and Smoldyn will read the .net file into Smoldyn structures.

Data structures declared in smoldyn.h

```
typedef struct bngstruct {
    struct bngsuperstruct *bngss; // bng superstructure
    char *bngname;                // bng name
    int bngindex;                 // index of this bng structure

    int maxparams;                // maximum number of numeric parameters
    int nparams;                 // actual number of numeric parameters
    char **paramnames;            // names of parameters [index]
    char **paramstrings;          // strings for parameter values [index]
    double *paramvalues;          // actual parameter values [index]

    int maxmonomer;               // maximum number of monomers
    int nmonomer;                // actual number of monomers
    char **monomernames;          // names of monomers [index]
    int *monomercount;            // monomer count work space [index]
    double *monomerdifc;          // diffusion coefficient of monomer [index]
    double *monomerdisplaysize;   // display size of monomer [index]
    double **monomercolor;        // color of monomer [index][RGB]
    enum MolecState *monomerstate; // default monomer state [index]
    int bngmaxsurface;            // local copy of nsurface
    enum SrfAction ***monomeraction; // monomer surface actions [index][srf][face]
    surfactionptr ***monomeractdetails; // monomer action details [index][srf][face]

    int maxbspecies;              // maximum number of bng species
    int nbspecies;                // actual number of bng species
    char **bsplongnames;          // complete bng species names [index]
    char **bspshortnames;         // shortened bng species names [index]
    enum MolecState *bspstate;     // default species state [index]
    char **bspcountstr;           // strings for initial bng species counts [index]
    double *bspcount;             // actual initial bng species counts [index]
    int *spindex;                 // smoldyn index of this species [index]

    int maxbrxns;                 // maximum number of bng reactions
    int nbrxns;                   // actual number of bng reactions
    char **brxnreactstr;          // strings for reactants [index]
    char **brxnprodstr;           // strings for products [index]
    char **brxnratestr;           // strings for reaction rates [index]
    int **brxnreact;              // reactant bng species indices [index][rct]
    int **brxnprod;               // product bng species indices [index][prd]
    int *brxnorder;               // order of bng reaction [index]
    int *brxnnpod;                // number of products of bng reaction [index]
    rxnptr *brxn;                 // pointer to this reaction [index]
} *bngptr;

typedef struct bngsuperstruct {
    enum StructCond condition; // structure condition
```

5.12 BioNetGen (functions in smolbng.c)

```

struct simstruct *sim;          // simulation structure
int maxbng;                    // maximum number of bng networks
int nbng;                      // actual number of bng networks
char **bngnames;               // names of bng networks
bngptr *bnglist;               // list of bng networks
} *bngssptr;

```

Superstructure description. As usual, **condition** tells about whether this data structure is up-to-date or not and **sim** points to the simulation structure. **BNG2path** is a string with the path information to the BNG2 program, along with the name of the program (usually BNG2.pl). This superstructure allows for multiple bng networks, the idea being that different macromolecular complexes might be described with separate sets of rules. Each bng network has a name listed in **bngnames** and is listed in **bnglist**. This structure is allocated for **maxbng** possible bng networks, of which **nbng** are currently defined.

Bng structure definition. The bng structure includes a pointer to the bng superstructure that owns it in **bngss**, a pointer to its name in **bngname**, and a copy of its index number in **bngindex**. The **unirate** and **birate** elements are multipliers for unimolecular and bimolecular reaction rates, respectively, which are useful due to the possibility of different units in the BNG file.

To understand the monomer elements, each complex is composed of a collection of monomers, also called seed species (BioNetGen), mols (Molecularizer), or subunits. It is helpful for this parsing software to have a list of these monomers, so they are stored here. The monomers that are listed here are inferred from the species names. There are **nmonomer** monomers, of a total of **maxmonomer** spaces for them. Each monomer has a name. The **monomercount** element is not for storing data, but is purely workspace for the **bngparsespecies** function. Monomer diffusion coefficients, in **monomerdifc** are retrieved from the **sim->mols->difc** data location. The **monomerdisplaysize** and **monomercolor** elements are also retrieved from the **sim->mols** superstructure. The **monomerstate** element stores the default state of the monomer, which is used to compute default states for species and then the states for reactions. The rule is that higher value states take priority over lower value ones (so if a species includes some monomers with default state of **MSsoln** and one of state **MSup**, then the ‘up’ state will win and that becomes the default state for the species). These monomer lists are sorted so that monomers are listed in alphabetical order. The **bngmaxsurface**, **monomeraction**, and **monomeractdetails** elements are for molecule-surface interactions. The **bngmaxsurface** value is the allocated size of the following two elements, which is made equal to **srfss->nsrf** when things are updated, but might be smaller in the meantime. The next two elements store information about monomer interactions with the surface. The general rule is that more action takes priority over less action.

Most other structure elements correspond to those of the BioNetGen .net file. The species list has allocated size **maxbspecies** and actual size **nbspecies** (the ‘b’ part of these terms is for BioNetGen, to differentiate these lists from the regular Smoldyn lists). Each species is one that is listed in the .net file. The listed name is put in the **bsplongnames** element and the count from the file is put in the **bspcountstr** element. The name is parsed to monomers and then simplified to a short species name, in **bspshortnames**. The default state is also from the monomers, and is stored in **bspstate**. The species count is parsed into the **bspcount** element. The index of this species within the **sim->mols** data structure is stored in **spindex**.

The reaction list has allocated size **maxbrxns** and actual size **nbrxns**. The reaction data read in from the .net file is separated into reactants in **brxnreactstr**, products in **brxnprodstr**, and the reaction rate in **brxnratestr**. Those items are parsed into bionetgen species numbers (as opposed to Smoldyn species numbers) in the **brxnreact** and **brxnprod** vectors. The reaction has order **brxnorder**, and number of products **brxnnpod**. After the reaction is added to the main Smoldyn

5 C Code: structures and functions

program, the pointer to the reaction data structure for this reaction is stored in **brxn**.

Function declarations. As much as possible, the smolbng functions are declared locally rather than in the smoldynfuncs.h header file. This simplifies the code reading because it clarifies which functions might be called externally versus those that are only called internally. Below, all functions are labeled as either “Local” or “Global” to indicate this status.

Memory management functions

void bngallocsurface(bngptr bng, int maxsurface)

This allocates memory for the surface action elements for the monomers. Enter **bng** with the bng to be updated and **maxsurface** for the desired number of surface spaces (which should be equal to **srfs->nsrf**). This function allocates all of the surface spaces for any monomers that don’t have any yet. It also expands the number of surface spaces for monomers that were allocated when there were fewer surfaces. This doesn’t return anything except for the updated **bng** structure. If the **bng->bngmaxsurface** element wasn’t updated, then this function wasn’t able to allocate memory.

bngptr bngalloc(bngptr bng, int maxparams, int maxbspecies, int maxbrxns)

Local. Allocates or expands a bng structure, returning a pointer to it. Enter **bng** as NULL to allocate a new structure, or as an existing pointer to expand lists within the data structure. Spaces for parameters, species, and reactions are allocated if the entered **max** value is larger than the current one. Returns a pointer to the new or previously existing structure on success, or NULL on failure.

void bngfree(bngptr bng)

Local. Frees a bng structure, including all of its contents.

bngssptr bngssalloc(bngssptr bngss, int maxbng)

Local. Allocates or expands a bngss structure, returning a pointer to it. Enter **bngss** as NULL to allocate a new structure, or as an existing pointer to expand lists within the data structure. Spaces for bng names and structures are allocated if the entered **maxbng** value is larger than the current one. Returns a pointer to the new or previously existing structure on success, or NULL on failure.

void bngssfree(bngssptr bngss)

Local. Frees a bng superstructure, including all of its contents.

Data structure output

void bngoutput(simptr sim)

Global. Outputs the contents of the bng superstructure and all bng structures to the display.

int checkbngparams(simptr sim, int *warnptr)

Global. Checks that the bng parameters are reasonable. This does very little at present.

Structure set up - bng

void bngsetcondition(bngssptr bngss, enum StructCond cond, int upgrade)

Local. This function sets the condition of the bng superstructure. Set upgrade to 1 if this is an upgrade, to 0 if this is a downgrade, or to 2 to set the condition independent of its current value.

```
int bngenablebng(simptr sim, int maxbng)
```

Local. This function enables simulation with bng capability. It allocates and initializes the bng superstructure if it doesn't already exist. Send in **maxbng** with -1 for automatic behavior, which is 1 bng structure, or with some other value to specify the number of bng structures. Returns 0 on success or 1 for failure to allocate memory.

```
bngptr bngaddbng(simptr sim, const char *bngname)
```

Local. Adds a bng structure to the bng superstructure, returning a pointer to it. The new bng structure is named **bngname**. This enables bng function if it hasn't been enabled already. If a bng structure already exists with the same name, then a pointer to that bng is returned.

```
int bngsetparam(bngptr bng, char *parameter, double amount)
```

Local. Sets the bng structure parameter called **parameter** to the value **amount**. The only options for **parameter** are "unimolecular_rate" and "bimolecular_rate". Returns 0 for success, 1 for illegal **parameter** string, and 2 for illegal **amount** value.

```
int bngsetBNG2path(bngptr bng, char *path)
```

Local. Sets the bng superstructure BNG2path element to the string entered in **path**. Returns 0.

Structure set up - parameters

```
int bngparseparameter(bngptr bng, int index)
```

Local. Parses the parameter value string for the parameter with index **index**, and puts the result into the parameter value element. Returns 0 for success or 1 for failure. If there is a failure, then its description can be found using **strmatherror**.

```
int bngaddparameter(bngptr bng, const char *name, const char *string)
```

Local. Adds a parameter, or modifies an existing parameter, in a bng structure. Send in the parameter name in **name**. If a parameter of this name doesn't already exist, one is created. The value of **string**, which is allowed to be NULL is copied into the **bng->paramstrings** element for parsing to a value. This function sends the new parameter to **bngparseparameter** for parsing. Returns the index of the parameter for success, -1 for failure to allocate memory, or -2 for a math parsing error; in the last case, the error message can be found using **strmatherror**.

Structure set up - monomers

```
int bngaddmonomer(bngptr bng, const char *name, enum MolecState ms)
```

Local. Adds a monomer, or modifies an existing monomer, in a bng structure. Send in the monomer name in **name**. If the default state of this monomer is known, enter it in **ms**; if not, then enter **ms** as **MSsoln**. If a monomer of this name doesn't already exist, one is created. It is added to the list while maintaining alphabetical order. The monomer counts are not changed. Returns the monomer index for success, -1 for failure to allocate memory, or -2 for an invalid monomer name. This sets the monomer diffusion coefficient and color to values in the main Smoldyn code if it is available. In doing so, it looks first for a Smoldyn species name that is the same as the monomer name. If it doesn't find one, it looks for a Smoldyn species name that has two dots and where the portion before the first name is the same as the monomer name (e.g. "X.1.0").

```
int bngsetmonomerdifc(bngptr bng, char *name, double difc)
```

Local. Sets the diffusion coefficient the monomer called **name** to **difc**. This also adds the

5 C Code: structures and functions

monomer to the list if it wasn't there already, and initializes the diffusion coefficient. Returns 0 for success, -1 if out of memory, or -2 for invalid monomer name.

int bngsetmonomerdisplaysize(bngptr bng, char *name, double displaysize)

Local. Sets the display size of the monomer called **name** to **displaysize**. This also adds the monomer to the list if it wasn't there already, and initializes the diffusion coefficient. Returns 0 for success, -1 if out of memory, or -2 for invalid monomer name.

int bngsetmonomercolor(bngptr bng, char *name, double *color)

Local. Sets the color of the monomer called **name** to **color**, which is a 3-element vector. This also adds the monomer to the list if it wasn't there already, and initializes the diffusion coefficient. Returns 0 for success, -1 if out of memory, or -2 for invalid monomer name.

int bngsetmonomerstate(bngptr bng, char *name, enum MolecState ms)

Local. Sets the state of the monomer called **name** to **ms**. This also adds the monomer to the list if it wasn't there already, and initializes the diffusion coefficient. Returns 0 for success, -1 if out of memory, or -2 for invalid monomer name.

Structure set up - species

int bngmakeshortname(bngptr bng, int index, int totalmn, int hasmods)

Local. Generates a short name for a bspecies, saving it in **bng->bspshortnames**. Send in **index** as the index of this species, **totalmn** as the total number of monomers in this species, and **hasmods** as 1 if the species has at least one modification site and 0 if not. The **monomercount** array needs to be prepared as well. If **totalmn** is 1 and the species has no modification sites, then this species is just a simple unmodifiable monomer; in this case, the short name is the name of the monomer with no suffix. Otherwise, a short name is created which concatenates each monomer name and the number of copies of that monomer, for each monomer, and then adds an isomer code to the end. Always returns 0. This function is only called by **bngparsespecies**. If the short name length as defined here would be longer than the maximum string length, **STRCHAR**, then the length is truncated at a value that is slightly less than this and this species is distinguished from others that have the same name with the isomer number. Thus, errors cannot arise from excessively long strings.

enum MolecState bngmakedefaultstate(bngptr bng, int index, int totalmn)

Local. Generates and returns the default state for a bspecies, which has index **index** and total monomers **totalmn**. The **monomercount** array needs to be prepared as well. If the species is also in Smoldyn, then that value is retrieved. If the species has only 1 monomer, the state is the state of that monomer. Otherwise, the state is the greatest state of the monomer states (although with a little re-ordering, so that the priority is **MSsoln**, **MSbsoln**, **MSfront**, **MSback**, **MSup**, **MSdown**. This function is only called by **bngparsespecies**.

double bngmakedifc(bngptr bng, int index, int totalmn)

Local. Generates and returns the diffusion coefficient for a bspecies, which has index **index** and total monomers **totalmn**. The **monomercount** array needs to be prepared as well. If the species is also in Smoldyn, then that value is retrieved. If the species has only 1 monomer, then the diffusion coefficient is that of the monomer. Otherwise, it is a weighted average of the monomer diffusion coefficients using the equation $D_{species} = (\sum_i D_i^{-3})^{-1/3}$ where $D_{species}$ is the diffusion coefficient of the species and D_i is the diffusion coefficient of the i th monomer within the species. This function is only called by **bngparsespecies**.

```
double bngmakedisplaysize(bngptr bng, int index, int totalmn)
```

Local. Generates and returns the display size for a bspecies, which has index `index` and total monomers `totalmn`. The `monomercount` array needs to be prepared as well. If the species is also in Smoldyn, then that value is retrieved. If the species has only 1 monomer, then the display size is the display size of that monomer. Otherwise, it is a weighted average of the monomer display sizes using the equation $S_{species} = (\sum_i S_i^3)^{1/3}$ where $S_{species}$ is the display size of the species and S_i is the display size of the i th monomer within the species. This function is only called by `bngparsespecies`.

```
int bngmakecolor(bngptr bng, int index, int totalmn, double *color)
```

Local. Generates and returns the color for a bspecies, which has index `index` and total monomers `totalmn`. The `monomercount` array needs to be prepared as well. If the species is also in Smoldyn, then that value is retrieved. If the species has only 1 monomer, then the color is the color of that monomer. Otherwise, it is a weighted average of the monomer colors. This averaging is weighted by the display sizes of each monomer. This function is only called by `bngparsespecies`.

```
void bngmakesurfaction(bngptr bng, int index, int totalmn, enum SrfAction  
**srfaction, surfactionptr **actdetails)
```

Local. Generates and returns the surface actions for a bspecies, which has index `index` and total monomers `totalmn`. The `monomercount` array needs to be prepared as well. If the species is also in Smoldyn, then those values are retrieved. If the species has only 1 monomer, then the surface actions are those of that monomer. Otherwise, it is chosen based on priority, so that greater actions take priority over smaller actions. More precisely, the action ordering is `SAttrans` | `SAmult` | `SAreflect` | `SAjump` | `SAabsorb` | `SAport`. If the only choice is between `SAmult` values, then this goes to the rate values in the action details to decide which has the greater action. This function has barely been tested.

```
int bngparsespecies(bngptr bng, int index)
```

Local. Parses species with index `index` using its longname. This does lots of things as it goes along. It determines if there are any new monomers, and adds those to the monomer list if so. This also generates a short name for the species, its default state, diffusion coefficient, display size, and color. This species is added to the Smoldyn simulation, including all of its attributes. This then parses the species countstring and adds the correct number of molecules to the Smoldyn simulation. Returns 0 for success, -1 for inability to allocate memory, -2 for a longname that cannot be parsed, -3 for an illegal name (e.g. an asterisk in the name), -4 for a count string that cannot be parsed (in this case, the error can be found from `strmathererror`, or -5 if more molecules are requested than the maximum allowed in Smoldyn, set with `maxdlimit`.

```
int bngaddspecies(bngptr bng, int bindex, const char *longname, const char *countstr)
```

Local. Adds a species, or modifies an existing species, in a bng structure. The `bindex` value is used to place the species in the list, meaning that species are not added to the list sequentially, but according to the `bindex` values. This copies `longname` into the `bng->bsplongnames` element and `countstr` into the `bng->bspcountstr` element. Either or both are allowed to be NULL. This calls `bngparsespecies` to parse the species, set up all species parameters, and add it to the Smoldyn simulation. Returns 0 for success or the same error codes as `bngparsespecies` for failure.

Structure set up - reactions

`int bngparsereaction(bngptr bng, int index)`

Local. Parses the reaction with index `index`. This reads the `bng` reactant and product names, deals with reactant and product states, and uses this information to create a new Smoldyn reaction. This also reads the reaction rate and uses it to set the Smoldyn reaction rate. Returns 0 for success, 1 for failure to add the reaction to Smoldyn, or 2 for inability to parse the rate string.

`int bngaddreaction(bngptr bng, int bindex, const char *reactants, const char *products, const`

Local. Adds a reaction, or modifies an existing reaction, in a `bng` structure. The `bindex` value is used to place the reaction in the list, meaning that reactions are not added to the list sequentially, but according to the `bindex` values. This copies `reactants` into the `bng->brxnreactstr` element, `products` into the `bng->brxnprodstr` element, and `rate` into the `bng->brxnratestr` element. Any or all are allowed to be NULL. This calls `bngparsereaction` to perform all parsing. Returns 0 for success, 1 for failure to allocate memory, or 2 for inability to parse the rate string.

Structure set up - reactions

`int bngaddgroup(bngptr bng, int gindex, const char *gname, const char *specieslist)`

Adds a group, created in the `.bnl` file with “begin observables” and in the `.net` file as “begin groups” to the Smoldyn simulation, as a species group. Send in the group index as `gindex`. This value is ignored. Send in the group name as `gname` and the species list as `specieslist`. The species list should be a comma-separated list of species indices, using the BioNetGen indices. Returns 0 for success or 1 for errors (the only error that should ever arise is an out of memory error). This function converts the `bng` species indices to Smoldyn species indices and sends the results to `moladdspeciesgroup` for group creation.

Structure set up - high level functions

`int bngrunBNGL2(bngptr bng, char *filename, char *outname)`

Local. This runs the BNG2.pl program on the BNGL file called `filename`. The output file name is returned in `outname`. Returns 0 for success; 1 for inability to find BNG2.pl software at the stored path location; 2 for missing input file called `filename`; or 3 for no output file generated, due to BNG2.pl terminating because of an error in the input file.

`bngptr bngreadstring(simptr sim, ParseFilePtr pfp, bngptr bng, const char *word, char *line2)`

Local. Reads a line of input. This input can start with the word ‘name’ to give the name of the `bng` structure. Otherwise, it needs to be a line from a BioNetGen `.net` file, or a few other words. Returns a pointer to the `bng` structure for success or NULL for failure.

`int loadbng(simptr sim, ParseFilePtr *pfp, char* line2)`

Reads BioNetGen `.net` file, or set of statements. Returns 0 for success or 1 for failure. This function does not require an end statement to stop reading the file, but an end of file serves as well.

`int bngupdateparams(simptr sim)`

Local. Doesn’t do anything at the moment.

5.13 Complexes (not written yet)

`int bngupdatelists(simptr sim)`

Local. This does nothing, simply returning 0 to indicate success.

`int bngupdate(simptr sim)`

Updates a bng structure, bringing it up to the ok condition. This calls the `bngupdatelists` and `bngupdateparams` functions to carry out the updating tasks.

Core simulation functions

No core simulation functions.

5.13 Complexes (not written yet)

It's becoming increasingly apparent that Smoldyn needs to support macromolecular complexes. This section presents documentation for code that hasn't been written yet in the hopes that this will provide a design for the code once there is time to write it.

Data structures

Each individual complex is listed with a `complexstruct`, and this `complexstruct` lists each of its monomers individually in a `monomerstruct`.

```
typedef struct monomerstruct {
    struct complexstruct *cmplx;      // owning complex superstructure
    moleculeptr *mptr;               // molecule that is this monomer
    double *dispsph;                  // displacement of monomer in r, q, f, x
    double *dispcart;                 // displacement of this monomer in x, y, z
    struct monomerstruct **site;      // monomer at each binding site [bs]
    int *shape;                       // shape of each binding site [bs]
} *monomerptr;

typedef struct complexstruct {
    struct complexsuperstruct *cmplxss; // owning complex superstructure
    int maxmonomer;                     // maximum number of monomers
    int nmonomer;                       // actual number of monomers
    monomerptr *monomers;               // list of component monomers
    double *pos;                        // center of mass position [d]
    double *posx;                       // old COM position [d]
    double *rotation;                   // complex rotation in q, f, x
    double mass;                        // complex mass
    double difc;                        // complex diffusion coefficient
    double difstep;                     // complex rms step length
} *complexptr;

typedef struct complexsuperstruct{
    struct simstruct *sim;              // owning simulation structure
    int maxcomplex;                     // maximum number of complexes
    int ncomplex;                       // actual number of complexes
    complexptr *complexes;              // list of individual complexes
// rules about connectivity
    int *nsites;                        // number of sites on species [i]
    int **nshapes;                       // number of shapes [i][bs]
    double ***shapemass;                 // mass of shape [i][bs][sh]
```

5 C Code: structures and functions

```
char ****shapename;          // shape names [i][bs][sh]
double ****dispsph;          // site displacement in r, q, f, x, [i][bs][
    sh]
// rules for complex names (i.e. explicit-species and species-classes)

} *complexssptr;

Add to moleculesuperstruct: double *mass; // mass of species [i]
```

5.14 Graphics (functions in smolgraphics.c)

Overall, Smoldyn's graphics use is fairly straightforward, although it is nevertheless a little complicated due to the design of the OpenGL glut library. For stand-alone Smoldyn, the program's entry and exit point are in the `main` function, which is in `smoldyn.c`. From here, the program forks to either `smolsimulate` (in `smolsim.c`) if graphics are not used or to `smolsimulategl` (in `smolgraphics.c`) if graphics are used. In the latter case, the OpenGL glut framework is used, in which control is passed from the main program to OpenGL and is not returned again until the user chooses quit from a menu. This leads to a slightly strange program structure, although I have attempted to contain all of the strangeness to a few functions at the end of `smolgraphics.c`.

If graphics are used, then `smolsimulategl` sets up a few graphics options that will apply for the entire simulation, calls `gl2Initialize` with the system boundaries for more overall set up, sets the background color, registers `RenderScene` as the display callback function, registers `TimerFunction` as the simulation callback function, posts a need for a redisplay with `glutPostRedisplay`, and then goes into the black box called `glutMainLoop`. The only graphics done in `TimerFunction` is that it posts a need for graphical update on occasion with `glutPostRedisplay`. Meanwhile, `RenderScene` is just a wrapper for `RenderSim`, which actually draws the entire graphical output. The summary is: initialization is done in `smolsimulategl` and drawing is done by `RenderSim`.

Data structure

```
#define MAXLIGHTS 8;
enum LightParam {LPambient, LPdiffuse, LPspecular, LPposition, LPon, LPoff,
    , LPauto, LPnone};

typedef struct graphicssuperstruct {
    int graphics;          // graphics: 0=none, 1=opengl, 2=good opengl
    int currentit;         // current number of simulation time steps
    int graphicit;         // number of time steps per graphics update
    unsigned int graphicdelay; // minimum delay (in ms) for graphics
        updates
    int tiffit;            // number of time steps per tiff save
    double framepts;       // thickness of frame for graphics
    double gridpts;        // thickness of virtual box grid for
        graphics
    double framecolor[4];  // frame color [c]
    double gridcolor[4];   // grid color [c]
    double backcolor[4];   // background color [c]
    double textcolor[4];   // text color [c]
    int maxtextitems;      // allocated size of item list
    int ntextitems;        // actual size of item list
    char **textitems;      // items to display with text [item]
    double ambient[4];     // global ambient light [c]
```

5.14 Graphics (functions in *smolgraphics.c*)

```
int lightstate[MAXLIGHTS];    // whether light is on or off [lt]
double ambientlight[MAXLIGHTS][4]; // ambient light color [lt][c]
double diffflight[MAXLIGHTS][4]; // diffuse light color [lt][c]
double speclight[MAXLIGHTS][4]; // spectral light color [lt][c]
double lightpos[MAXLIGHTS][3]; // light positions [lt][d]
} *graphicsssptr;
```

The `graphicdelay` value is used during simulation running to slow the simulation down. It is ignored while the simulation is in pause state or is complete, and then a delay of 20 milliseconds is used instead.

enumerated types

`enum LightParam graphicsstring2lp(char *string)`
Converts a string to an enumerated light parameter.

`char *graphicslp2string(enum LightParam lp, char *string)`
Converts an enumerated light parameter to a string. The string is returned.

low level utilities

`int graphicsreadcolor(char **stringptr, double *rgba)`
Reads the text of the string that `stringptr` points to, to find color information. The data are returned in the vector `rgba`, if `rgba` is not sent in as `NULL`. The input data are in a pointer to a string, rather than just a string, so that the string can be advanced to the end of the color information. Upon return, if this function is successful, the contents of `stringptr` points to the first word of the string that follows the color information, or to `NULL` if the end of the string was reached while parsing color information. If `rgba` is supplied, it needs to be allocated to hold at least 4 numbers, which are for the red, green, blue, and alpha color channels, respectively. The string needs to list the color either with three space-separated numbers, each between 0 and 1 inclusive, or with a single word. Word options are: maroon, red, orange, yellow, olive, green, purple, magenta, lime, teal, cyan, blue, navy, black, gray, silver, and white. Other words are not recognized. Following the color information, the string can optionally list the alpha value, as a number between 0 and 1. If alpha is not listed, a default value of 1 is assigned. The function returns 0 for no error, 1 if `string` is missing or is empty, 2 if too few numbers are listed, 3 if one or more of the listed color numbers is out of range (the listed numbers are returned in `rgba`), 4 if a word was given but it isn't recognized, 5 if an alpha value was given but can't be parsed, or 6 if the listed alpha value is out of range.

memory management

`graphicsssptr graphssalloc(void)`
Allocates and initializes the graphics superstructure. No OpenGL stuff is initialized here.

`void graphssfree(graphicsssptr graphss)`
Frees a graphics superstructure.

data structure output

`void graphssoutput(simpstr sim)`
Displays all graphics parameters from the graphics superstructure to stdout. Also displays some information from the `opengl2` library variables, including the TIFF name and TIFF numbering.

5 C Code: structures and functions

void writegraphss(simpstr sim, FILE *fptr)

Writes graphics information to `fptr` as part of a Smoldyn-readable input file.

int checkgraphicsparams(simpstr sim, int *warnptr)

Checks graphics parameters for actual or possible errors. Returns the number of errors directly and returns the number of warnings in `warnptr`, if `warnptr` isn't NULL. At present, this doesn't check anything, but just returns two zeros.

structure setup

void graphicssetcondition(graphicsssptr graphss, enum StructCond cond, int upgrade)

Sets the graphics superstructure condition to `cond`, if appropriate. Set `upgrade` to 1 if this is an upgrade, to 0 if this is a downgrade, or to 2 to set the condition independent of its current value. If the condition is downgraded, this also downgrades the simulation structure condition.

int graphicsenablegraphics(simpstr sim, char *type)

Enables graphics by allocating a graphics superstructure and adding it to the simulation structure. Enter `type` as "none" for no graphics, "opengl" for minimal OpenGL graphics (molecules are square dots), "opengl_good" for reasonably good OpenGL graphics (molecules are solid colored spheres), "opengl_better" for better OpenGL graphics (use of lighting and shininess), or NULL for default enabling which is no change if graphics already exist and basic OpenGL if they don't already exist. Returns 0 for success, 1 for inability to allocate memory, 2 for missing `sim` input, or 3 for an invalid `type` string.

int graphicssetiter(simpstr sim, int iter)

Sets the `graphicit` element of the graphics superstructure, which tells how many simulation iterations should be allowed to pass between graphics renderings. This enables graphics, if needed. Returns 0 for success, 1 for out of memory enabling graphics, 2 for no `sim`, or 3 for an illegal `iter` value (it needs to be at least 1).

int graphicssetdelay(simpstr sim, int delay)

Sets the `graphicdelay` element of the graphics superstructure, which gives the minimum number of milliseconds that should be allowed to elapse between graphics renderings, to keep simulations from running too fast to see. This enables graphics, if needed. Returns 0 for success, 1 for out of memory enabling graphics, 2 for no `sim`, or 3 for an illegal `delay` value (it needs to be at least 0). As described above, this delay value is only used during simulation running; a value of 20 milliseconds is used during pause state or after the simulation is complete.

int graphicssetframethickness(simpstr sim, double thickness)

Sets the `framepts` element of the graphics superstructure, which gives the drawing thickness of the frame that surrounds the simulation volume. This enables graphics, if needed. Returns 0 for success, 1 for out of memory enabling graphics, 2 for no `sim`, or 3 for an illegal `thickness` value (it needs to be at least 0).

int graphicssetframecolor(simpstr sim, double *color)

Sets the `framecolor` elements of the graphics superstructure, which gives the color of the frame that surrounds the simulation volume. Color is a four-element vector (red, green, blue, alpha). This enables graphics, if needed. Returns 0 for success, 1 for out of memory enabling

5.14 Graphics (functions in *smolgraphics.c*)

graphics, 2 for no **sim**, or 3 for one or more illegal **color** values (they need to be between 0 and 1, inclusive).

int graphicssetgridthickness(simptr sim, double thickness)

Sets the **gridpts** element of the graphics superstructure, which gives the drawing thickness of the partitions that separate the virtual boxes. This enables graphics, if needed. Returns 0 for success, 1 for out of memory enabling graphics, 2 for no **sim**, or 3 for an illegal **thickness** value (it needs to be at least 0).

int graphicssetgridcolor(simptr sim, double *color)

Sets the **gridcolor** elements of the graphics superstructure, which gives the color of the partitions that separate the virtual boxes. Color is a four-element vector (red, green, blue, alpha). This enables graphics, if needed. Returns 0 for success, 1 for out of memory enabling graphics, 2 for no **sim**, or 3 for one or more illegal **color** values (they need to be between 0 and 1, inclusive).

int graphicssetbackcolor(simptr sim, double *color)

Sets the **backcolor** elements of the graphics superstructure, which gives the color of the background. Color is a four-element vector (red, green, blue, alpha). This enables graphics, if needed. Returns 0 for success, 1 for out of memory enabling graphics, 2 for no **sim**, or 3 for one or more illegal **color** values (they need to be between 0 and 1, inclusive).

int graphicssettextcolor(simptr sim, double *color)

Sets the **textcolor** elements of the graphics superstructure, which gives the color of text drawn to the graphics window. Color is a four-element vector (red, green, blue, alpha). This enables graphics, if needed. Returns 0 for success, 1 for out of memory enabling graphics, 2 for no **sim**, or 3 for one or more illegal **color** values (they need to be between 0 and 1, inclusive).

int graphicssettextitem(simptr sim, char *itemname)

Adds an item to the list of things that Smoldyn will display to the graphics window. Enter the name of the item as a string in **itemname**. This automatically allocates space for text items as needed. Returns 0 for success, 1 if memory could not be allocated, 2 if the item is not a supported string, or 3 if the item was already listed. Currently supported names are: "time" and species(state) names.

int graphicssetlight(simptr sim, graphicssptr graphss, int lt, enum LightParam ltparam, double *value)

Sets parameters for the lighting portions of the graphics superstructure. If **graphss** is entered as non-NULL, then it is worked with and **sim** is ignored; otherwise, this requires the **sim** input and enables graphics if needed. **lt** is the light number, which should be between 0 and 8, or set **lt** to -1 for the global ambient light source. **ltparam** is the lighting parameter that should be set. It is only allowed to be **LPambient** if **lt** is -1. Otherwise, **ltparam** can be **LPambient**, **LPdiffuse**, **LPspecular** and then value should list the 4 color values. Or, **ltparam** can be **LPposition** and value should list the three light position values. Or, **ltparam** can be **LPon** or **LPoff** to turn the light on or off, in which case **value** is ignored. No checking is done to see that input parameters are legitimate. Returns 0 for success or 1 if memory could not be allocated for the graphics superstructure.

structure update functions

5 C Code: structures and functions

int graphicsupdateinit(simptr sim)

Performs basic graphics initialization. This calls `gl2glutInit` to initialize things and then `gl2Initialize` to create the graphics window and set the viewing coordinates. This function should probably be called only once.

int graphicsupdatelists(simptr sim)

Enables lighting models for the graphics display. This only needs to be called when the user requests the “`opengl.better`” graphics option.

int graphicsupdateparams(simptr sim)

Updates the lighting model parameters for the graphics display. This should be called every time the user changes the lighting model. It also sets the background color.

int graphicsupdate(simptr sim)

Updates the graphics superstructure and upgrades the graphics condition element. This calls `graphicsupdateinit`, `graphicsupdatelists`, and/or `graphicsupdateparams`, depending on the amount of updating required.

core simulation functions

void RenderSurfaces(simptr sim)

Draws all surfaces in the simulation using OpenGL graphics. The 3-D portion of this function needs some work, both to fix 3-D disk drawing, to improve 3-D drawing overall, and for overall cleanup.

void RenderMolecs(simptr sim)

Draws all molecules using OpenGL graphics. Because the molecules are not sorted by type, this function is fairly inefficient; this could be a significant computational burden for movie-making, but shouldn't be for most research purposes.

void RenderText(simptr sim)

Draws any requested text to the OpenGL graphics window.

void RenderSim(simptr sim, int swapbuffers)

Draws the entire graphical output using OpenGL graphics. This calls other functions for most of the work, although it draws the frame and the grid itself. Set the `swapbuffers` input to 0 for this function to not swap graphics buffers when it's done and 1 to swap the buffers. Using 1 is the standard practice but buffer swapping is asynchronous with the rest of the code and only happens on monitor vertical refreshes; thus, I've found that it needs to be set to 0 for recording the output at a specific simulation time to TIFF files.

For some reason, OpenGL calls this function more frequently than it needs to. For this reason, this function doesn't bother running if it's unnecessary, meaning that the simulation time hasn't changed and it was called less than 5 ms earlier.

Top level OpenGL functions Both `RenderScene` and `TimerFunction` are declared locally, rather than in `smoldynfuncs.h`. This makes them invisible outside of this source file. They are callback functions for OpenGL. In addition, the `Sim` variable is declared as a global variable, with the scope of this file. It is here because OpenGL does not allow `void*` pointers to be passed through to all callback functions, so making it a global variable enables the callback functions to access the simulation data structure.

5.15 Simulation structure (functions in smolsim.c)

`smolPostRedisplay(void)`

This is just a wrapper for `glutPostRedisplay`, designed for use elsewhere in the Smoldyn code.

`void RenderScene(void)`

`RenderScene` is the call-back function for OpenGL that displays the graphics. This does nothing but call `RenderSim`.

`void TimerFunction(int state)`

`TimerFunction` is the call-back function for OpenGL that runs the simulation. `state` is positive if the simulation should quit due to a simulation error or normal ending, `state` is negative if the simulation has been over, and `state` is 0 if the simulation is proceeding normally. This also looks at the state defined in the `opengl2` library; if it is 0, the simulation is continuing, if it is 1, the simulation is in pause mode, and if it is 2, the user told the simulation to quit. `oldstate` is the old version of the `gl2State` value. This function runs one simulation time step, posts graphics redisplay flags, and saves TIFF files as appropriate.

oldstate	state	gl2State	meaning	next state
1	-	0	leave pause state	(0 - 0)
0	0	0	continue simulating	(0 =simstep 0)
-	>0	-	stop the simulation	(- -1 -)
-	0	2	" " "	(- -1 2)
0	0/-1	1	enter pause state	(1 - 1)
-	0/-1	-	in pause state or sim is over	(- - -)

`void smolsimulategl(simptr sim)`

`smolsimulategl` initiates the simulation using OpenGL graphics. It does all OpenGL initializations, registers OpenGL call-back functions, sets the global variables to their proper values, and then hands control over to OpenGL. This function returns as the program quits.

5.15 Simulation structure (functions in smolsim.c)

At the highest level of the structures is the simulation structure. This is a large framework that contains information about the simulation that is to be run as well as pointers to each of the component structures and superstructures. It also contains some scratch space for functions to use as they wish.

5.15.1 Data structures

```
#define ETMAX 10
enum SmolStruct {SSmolec, SSwall, SSRxn, SSsurf, SSbox, SScmpt, SSport,
    SScmd, SSmzr, SSsim, SScheck, SSall, SSnone};
enum EventType {ETwall, ETsurf, ETdesorb, ETrxn0, ETrxn1, ETrxn2intra,
    ETrxn2inter, ETrxn2wrap, ETimport, ETexport};

typedef void (*logfnptr)(struct simstruct *, int, const char*, ...);
typedef int (*diffusefnptr)(struct simstruct *);
typedef int (*surfaceboundfnptr)(struct simstruct *, int);
typedef int (*surfacecollisionsfnptr)(struct simstruct *, int, int);
```

5 C Code: structures and functions

```
typedef int (*assignmols2boxesfnptr)(struct simstruct *, int, int);
typedef int (*zeroreactfnptr)(struct simstruct *);
typedef int (*unimolreactfnptr)(struct simstruct *);
typedef int (*bimolreactfnptr)(struct simstruct *, int);
typedef int (*checkwallsfnptr)(struct simstruct *, int, int, boxptr);

typedef struct simstruct {
    enum StructCond condition; // structure condition
    logfnptr logfn;           // function for logging output
    FILE *logfile;            // file to send output
    char *filepath;           // configuration file path
    char *filename;           // configuration file name
    char *flags;               // command-line options from user
    time_t clockstt;           // clock starting time of simulation
    double elapsedtime;        // elapsed time of simulation
    long int randseed;         // random number generator seed
    int eventcount[ETMAX];     // counter for simulation events
    int dim;                   // dimensionality of space.
    double accur;              // accuracy, on scale from 0 to 10
    double time;               // current time in simulation
    double tmin;               // simulation start time
    double tmax;               // simulation end time
    double dt;                 // simulation time step
    int quitatend;             // 1 if quit simulation at end
    rxnssptr rxnss[MAXORDER]; // reaction superstructures
    molssptr mols;             // molecule superstructure
    wallptr *wlist;            // list of walls
    surfacessptr srfs;         // surface superstructure
    boxssptr boxs;             // box superstructure
    compartssptr cmtss;        // compartment superstructure
    portssptr portss;          // port superstructure
    mzzrssptr mzzrss;          // network generation rule superstructure
    void* cmds;                // command superstructure
    graphicssptr graphss;      // graphics superstructure
    threadssptr threads;       // pthreads superstructure
    diffusefnptr diffusefn;    // function for molecule
    // diffusion
    surfaceboundfnptr surfaceboundfn; // function for surface-
    // bound molecules
    surfacecollisionsfnptr surfacecollisionsfn; // function for surface
    // collisions
    assignmols2boxesfnptr assignmols2boxesfn; // function that assigns
    // molecs to boxes
    zeroreactfnptr zeroreactfn; // function for zero order
    // reactions
    unimolreactfnptr unimolreactfn; // function for first order
    // reactions
    bimolreactfnptr bimolreactfn; // function for second order
    // reactions
    checkwallsfnptr checkwallsfn; // function for molecule
    // collisions with walls
} *simptr;
```

ETMAX is the maximum number of event types, which are enumerated with EventType. These

5.15 Simulation structure (functions in *smolsim.c*)

are used primarily for reporting the number of times that various things happened to the user, although they are also used occasionally elsewhere in the code so that certain routines are only done if they are necessary. **SmolStruct** enumerates the different types of superstructures that a simulation can have. It does not appear to be used anywhere in the code, so I'm not sure why I created it.

The list of function pointers defined with typedef statements are used below in the **simstruct**. They make it possible for a simulation to use different collections of core algorithms. The first one, **logfnptr**, is for an external logging function, for use by Libsmoldyn. Others allow the use of either single-threaded or multi-threaded algorithm versions. Also, which has not been done but could be, they could be tuned for better efficiency in, say, a 3D system, rather than being generalists for all system dimensionalities.

simstruct contains and owns all information that defines the simulation conditions, the current state of the simulation, and all other simulation parameters.

The **condition** element is maintained at the lowest level of all superstructure conditions.

The **logfn** is typically set to NULL but can be set to point to a function if text output should be dealt with there instead of in **simLog**. This is a callback function, where this function is called by **simLog** with any possible output, which allows for complex internal error handling. It is not used at present. **logfile** is used by **simLog** for sending all output. The default is stdout. Note that there are also global variables for these purposes called **LogFile** and **LoggingCallback**; the system is that the **sim** variables are used by priority but if they are not available, then the global variables are used.

The **filepath** and **filename** give the configuration file name and **flags** lists the command-line flags that the user supplied.

clockstt is used for the clock value when the simulation starts, and **elapsedtime** is used for storing the simulation run time while the simulation is paused, both of which are for timing simulations.

randseed is the starting random number seed. **eventcount** is a list of counts for each of the enumerated event types.

dim is the system dimensionality and **accur** is the overall simulation accuracy level. Because this has not proven useful, it should be removed at some point, and a version of it should be moved to the box superstructure.

Finally, **time**, **tmin**, **tmax**, and **dt** are the current time, starting time, stopping time, and time step of the simulation, respectively. **tbreak** is the simulation break time, which lets functions that use Libsmoldyn run the simulation for a fixed amount of time and then stop for other operations.

quitatend is a simple flag that is 1 if the simulation should simply quit when done and 0 if not. This is only relevant for simulations with graphics, because the others automatically quit at the end anyhow.

The superstructures are listed next. All of them are optional, where a NULL value simply means that the simulation does not include that feature and an existing superstructure means that the simulation has that feature. The command superstructure is pointed to with a **void*** rather than a **cmdssptr** because the latter is declared in a separate header file and I didn't want to require a dependency between the **smoldyn.h** header file and the **SimCommand.h** header file.

Finally, the simulation structure lists the function pointers for the core simulation algorithms.

5.15.2 Functions

Error handling

5 C Code: structures and functions

```
void simSetLogging(simptr sim, const char *logfile, void (*logFunction)(simptr,
int, const char*, ...))
```

This sets the logging file and logging callback functions. If `sim` is entered, this sets the versions for the simulation structure and if not, this sets the global variables. Enter both `logfile` and `logFunction` as `NULL` to close files, clear both values, and set them to `NULL`. If either or both of `logfile` and `logFunction` are non-`NULL`, this sets them, without disturbing the other one. Files are closed as needed.

```
void simSetThrowing(int corethreshold)
```

Sets the global `ThrowThreshold` threshold for throwing an error. Throwing would normally occur if the error importance values is greater than or equal to this threshold. However, throwing is disabled for now because it conflicts with static linking for `libsmoldyn`, making this function unnecessary.

```
void simLog(simptr sim, int importance, const char* format, ...)
```

All text output should be sent to this function. As a default, it simply prints the output to `stdout`. However, this also sends it elsewhere if the user asked for a different destination. This can also send output to a `Libsmoldyn` host program. Enter `sim` with the simulation structure if possible; if it's not possible, then this function displays the message to `stderr`. Enter `importance` with a value between 0 and 10, as shown below. The `format` and `...` portions are the same format string and arguments that are used for `printf`.

This function starts by printing the message to a string using the variable arguments functions `va_start`, `vsprintf`, and `va_end`. Then, it replaces any unit placeholders with the current working units, if there are any. After that, it sends output to the logging outputs if appropriate. Then it checks the flags and message importance, and sends it to the correct output if appropriate.

This function can output units. It expects unit placeholders to start with a pipe character and then list dimensions. Parentheses around the complete units, which are started after the pipe character, are optional. Using parentheses means that any units are shown within parentheses but lack of units means that parentheses are not shown. The string word that contains the units is allowed to end with a single character of various punctuation (`" , . : ; + - * / < > = ! )"`), which is kept in the output; however, the word needs to have a whitespace character after that punctuation character. Examples: `"%g|L "`, `"%g|L2/T, "`, `"text |(L/T): "`, `"|/T, "`. If the file has working units, then they are substituted into the placeholders, and any parentheses or terminal punctuation are kept. If the file doesn't have working units, then the placeholders and any parentheses are removed, but the terminal punctuation is still kept.

importance	meaning	example	flags and display
0	debugging output		v enables
1	verbose output	box details	v enables
2	normal diagnostics	reaction parameters	q suppresses
3	abbreviated diagnostics	basic species	q suppresses
4	important notices	simulation has ended	q suppresses
5	warnings	time step too big	w suppresses
7	minor errors	command errors	w suppresses
		structure condition not updated	
8	recoverable errors	user input syntax error	always shown

5.15 Simulation structure (functions in *smolsim.c*)

9	error in normal operation	child structures missing parents	always shown
10	unrecoverable errors	bug or memory allocation failure	always shown

`void simParseError(simptr sim, ParseFilePtr pfp)`

Handles parsing errors when reading a file. Send in the parse file pointer in `pfp` with the current status of the parsing. This sends `pfp` off to `Parse_ReadFailure` and then displays the error with appropriate messages using `simLog`.

Enumerated types

`enum SmolStruct simstring2ss(char *string)`

Returns the enumerated simulation structure type that corresponds to the string input. Returns `SSnone` if input is “none” or if it is not recognized.

`char *simss2string(enum SmolStruct ss, char *string)`

Returns the string that corresponds to the enumerated simulation structure input in `string`, which needs to be pre-allocated. The address of the string is returned to allow for function nesting.

`char *simsc2string(enum StructCond sc, char *string)`

Returns the string that corresponds to the enumerated structure condition input in `string`, which needs to be pre-allocated. The address of the string is returned to allow for function nesting.

low level utilities

`double simversionnumber(void)`

Returns the version number of Smoldyn. This reads the `VERSION` string from `smoldyn_config.h` into a double and returns that value. A value of 0 indicates that the reading didn’t work.

`void Simsetrandseed(simptr sim, long int randseed)`

Sets the random number generator seed to `seed` if `seed` is at least 0, and sets it to the current time value if `seed` is less than 0.

memory management

`simptr simalloc(char *root)`

Allocates a simulation structure. Essentially everything, including superstructures, is initialized to 0 or `NULL`. Exceptions are that the `filepath`, `filename`, and `flags` strings are allocated, the random number seed is initialized with a random value, and the command superstructure is allocated. `root` is a required input because it is sent to the command superstructure allocation; it is allowed to be `NULL`.

`void simfree(simptr sim)`

`simfree` frees a simulation structure, including every part of everything in it.

`void simfuncfree(void)`

Frees memory that is allocated by functions within the Smoldyn program and that is only kept track of by the functions themselves, using static variables. This should be called just before program termination.

5 C Code: structures and functions

int simexpandvariables(simptr sim, int spaces)

Expands the number of variables in a simulation structure by **spaces** spaces. This allocates memory, copies over existing variables, and clears new ones as needed.

data structure output

void simoutput(simptr sim)

simoutput prints out the overall simulation parameters, including simulation time information, graphics information, the number of dimensions, what the molecule types are, the output files, and the accuracy.

void simsystemoutput(simptr sim)

Displays information about all components of the simulation to stdout by calling output functions for each superstructure.

void writesim(simptr sim, FILE *fptr)

Writes all information about the simulation structure, but not its substructures, to the file **fptr** using a format that can be read by Smoldyn. This allows a simulation state to be saved.

int checksimparams(simptr sim)

checksimparams checks that the simulation parameters, including parameters of substructures, have reasonable values. If values seem to be too small or too large, a warning is displayed to the standard output, although this does not affect continuation of the program. Returns the number of errors.

structure set up

Initialization procedures are meant to be called once at the beginning of the program to allocate and set up the necessary structures. These routines call memory allocation procedures as needed. **simupdate** is the only one of these routines that should ever need to be called externally, since it calls the other functions as needed.

int simsetpthreads(simptr sim, int number)

Sets the number of pthreads that the simulation should run in to **number**. Send in **number** as 0 for unthreaded mode, which is the default, or as a larger value for threaded operation. A value of 1 will causes multi-threaded operation, but with 1 thread. If threading is requested and the function is able to fulfill it, it returns the number of threads. If threading is not requested, the function returns 0. If threading is requested and the function fails because it's not enabled, then it returns -1, and if it fails because of failure to allocate memory, it returns -2.

void simsetcondition(simptr sim, enum StructCond cond, int upgrade)

Sets the simulation structure condition to **cond**, if appropriate. Set **upgrade** to 1 if this is an upgrade, to 0 if this is a downgrade, or to 2 to set the condition independent of its current value.

int simsetvariable(simptr sim, char *name, double value)

Sets the value of the variable named **name** to **value**. This creates the variable if it doesn't already exist. This also allocates memory for new variables as needed. Returns 0 for success and 1 for out of memory.

5.15 Simulation structure (functions in *smolsim.c*)

int simsetdim(simp_{ptr} sim, int dim)

Sets the simulation dimensionality. Returns 0 for success, 2 if it had already been set (it's only allowed to be set once), or 3 if the requested dimensionality is not between 1 and 3.

int simsettime(simp_{ptr} sim, double time, int code)

Sets the appropriate simulation time parameter to *time*. Enter code as 0 to set the current time, 1 to set the starting time, 2 to set the stopping time, 3 to set the time step, or 4 to set the break time. Returns 0 for success, 1 if an illegal code was entered, or 2 if a negative or zero time step was entered. This function also keeps track of the times that have been set using a static variable called *timedefined*. To see what times have been set, enter code as -1, and this will return a number which is the sum of: 1 for the current time, 2 for the starting time, 4 for the stopping, 8 for the time step, and 16 for the break time. For example, and the return value with 14 to check for the start, stop, and step times.

int simreadstring(simp_{ptr} sim, ParseFilePtr pfp, const char *word, char *line2)

Reads and processes one line of text from the configuration file, or some other source. The first word of the line should be sent in as *word* (terminated by a '\0') and the rest sent in as *line2*. This function may change *line2*. Also send in the "parse file pointer" in *pfp*; this input is optional. Returns 0 for success. On failure, this prints an error message to the global variable *ErrorMessage*, calls *simParseError* to shut down the parsing process (and free *pfp*), and returns 1.

Most numeric input is parsed with the *strmathsscanf* function, which is my version of the standard *sscanf* function. In this version, a %m start to a format specifier means that it should accept math functions and variable names and also might end with units. The function then processes the input to a numeric quantity.

int loadsim(simp_{ptr} sim, const char *fileroot, const char *filename, const char *flags)

loadsim loads all simulation parameters from a configuration file, using a format described above. *fileroot* is sent in as the root of the filename, including all colons, slashes, or backslashes; if the configuration file is in the same directory as Smoldyn, *fileroot* should be an empty string. *filename* is sent in as just the file name and any extension. *sim* is the simulation structure. This routine calls *loadsurface* to load any surfaces. The following things are set up after this routine is completed: all molecule elements except box; all molecule superstructure elements; all wall elements; box superstructure element *mpbox*, but no other elements; no boxes are allocated or set up; all reaction structure elements except *rate2* and the product template position vectors (*pos* in each product); the command superstructure, including all of its elements; and all simulation structure elements except for sub-elements that have already been listed. All new molecules are left in the dead list for sorting later. If the configuration file loads successfully, the routine returns 0. If the file could not be found, it returns 10 and an error message. If an error was caught during file loading, the return value is 10 plus the line number of the file with an error, along with an error message. If there is an error, all structures are freed automatically.

int simupdate(simp_{ptr} sim, char *erstr)

Updates all parts of the simulation structure. This is called on start up by *simUpdateAndDisplay*, and may be called at anytime afterwards. It returns 0 for success or 1 for failure. In the latter case, a string that describes the error should be returned in *erstr*.

5 C Code: structures and functions

```
int simInitAndLoad(const char *fileroot, const char *filename, simptr *smptr, const char *flags)
```

`simInitAndLoad` sets up and loads values for all the structures as well as global variables. This routine calls the other initialization routines, so they do not have to be called from elsewhere. It also displays the status to stdout and calls output routines for each structure, allowing verification of the initialization. Send in `fileroot` and `filename` with strings for the path and name of the input file and send in `smptr` (pointer to a simulation structure) pointing to a NULL. `flags` is a string of command-line flags. `logfile` is an optional entry for the logging filename. This returns 0 for correct operation and 1 for an error. If it succeeds, `smptr` is returned pointing to a fully set up simulation structure. Otherwise, `smptr` is set to NULL and an error message is displayed on stderr. If `smptr` does not point to NULL, then this simply returns without doing anything.

```
int simUpdateAndDisplay(simptr sim)
```

This calls `simupdate` to update all aspects of the simulation structure and then outputs the diagnostics output along with any warnings or errors. Returns 0 for success or 1 for failure.

core simulation functions

```
int simdocommands(simptr sim)
```

Performs all commands that should happen at the current time. This includes commands that should happen before or after the simulation. This function leaves data structures in good shape. Returns 0 to indicate that the simulation should continue, 6 for error with `molsort`, 7 for terminate instruction from `docommand`, or 8 for failed simulation update. These are the same error codes that `simulatetimestep` uses.

```
int simulatetimestep(simptr sim)
```

`simulatetimestep` runs the simulation over one time step. If an error is encountered at any step, or a command tells the simulation to stop, or the simulation time becomes greater than or equal to the requested maximum time, the function returns an error code to indicate that the simulation should stop; otherwise it returns 0 to indicate that the simulation should continue. See the table below. Note that the sequence in which the simulation components are set up is designed carefully and will likely lead to bugs if it is changed.

Error code	Meaning
0	Simulation should continue
1	Simulation completed normally
2	Error with <code>assignmolecs</code> (insufficient memory when expanding boxes)
3	Error with <code>zeroreact</code> (insufficient memory when expanding molecules)
4	Error with <code>unireact</code> (insufficient memory when expanding molecules)
5	Error with <code>bireact</code> (insufficient memory when expanding molecules)
6	Error with <code>molsort</code> (insufficient memory when expanding boxes)
7	Terminate instruction from <code>docommand</code> (e.g. stop command)
8	Failed simulation update
9	Error with <code>diffuse</code>
10	Simulation stopped because the time equals or exceeds the break time
11	Error in filament dynamics
12	Error in lattice simulation
13	Error in reaction network expansion

```
void endsimulate(simptr sim, int er)
```

`endsimulate` takes care of things that should happen when the simulation is complete. This includes executing any commands that are supposed to happen after the simulation, displaying numbers of simulation events that occurred, and calculating the execution time. `er` is a code to tell why the simulation is ending, which has the same values as those returned by `simulatetimestep`. If graphics are used, this routine just returns to where it was called from (which is `TimerFunction`); otherwise, it frees the simulation structure and then returns (to `smolsimulate` and then `main`).

```
int smolsimulate(simptr sim)
```

`smolsimulate` runs the simulation without graphics. It does essentially nothing other than running `simulatetimestep` until the simulation terminates or stops due to reaching the break time.

5.16 Commands (functions in *smolcmd.c*)

5.16.1 Writing commands

Command strings are not parsed, checked, or even looked at during simulation initialization. Instead, they are run by the command interpreter during the simulation. Command routines are given complete freedom to look at and/or modify any part of a simulation structure or substructure. This, of course, also gives commands the ability to crash the computer program, so they need to be written carefully to prevent this. Every command is sent a pointer to the simulation structure in `sim`, as well as a string of command parameters in `line2`.

To write a command, do the following steps, which can be done in any order:

- Write a description of the new command that will go into the reference section of the user's manual.
- In *smolcmd.c*, add a new declaration to the top of the file for the command, which looks like:

```
enum CMDcode cmdname(simptr sim, cmdptr cmd, char *line2);
```
- The first function of *smolcmd.c* is `docommand`. In it, add an `else if` line for the new command. It looks like:

```
else if(!strcmp(word, "name")) return cmdname(sim, cmd, line2);
```
- Write the function for the new command, modeling it on the command functions currently in *smolcmd.c*. See below.
- Proofread the function and test the command.
- Write documentation about the command for this section of this manual.
- Mention the command in the modifications portion of this manual.

Each command is written with a similar structure. As an example, here is `cmdecho`:

5 C Code: structures and functions

```
enum CMDcode cmdecho(simptr sim, cmdptr cmd, char *line2) {
    int er;
    FILE *fptr;
    char *termqt, str[STRCHAR];

    if(line2 && !strcmp(line2, "cmdtype")) return CMDobserve;
    er=scmdgetfptr(sim->cmds, line2, 1, &fptr, NULL);
    SCMDCHECK(er!=-1, "file_name_not_recognized");
    line2=strnword(line2, 2);
    SCMDCHECK(line2=strchr(line2, '"'), "no_starting_quote_on_string");
    strncpy(str, line2+1, STRCHAR-1);
    str[STRCHAR-1]='\0';
    SCMDCHECK(termqt=strchr(str, '"'), "no_terminal_quote_on_string");
    *termqt='\0';
    strbslash2escseq(str);
    scmdfprintf(cmd->cmds, fptr, "%s", str);
    scmdflush(fptr);
    return CMDok; }
```

Every command has essentially the same function call, where the first parameter is the simulation structure, the second parameter is the command structure of the current command, and the third parameter is the remainder of the input line from this command entry in the user's configuration file.

The first line of code checks to see whether `line2` is entered as "cmdtype" (which is never a valid entry from the user) and if so, returns the type of the current command. These types are `CMDcontrol` for commands that control the running of the simulation, `CMDobserve` for commands for observing the system but that do not change it, and `CMDmanipulate` for commands that modify the simulation state but do not have any output. Conditional commands, all of which have an "if" in the name, return the command type of the command that they call if the condition is met (see `cmdifflag`).

The body of the command is after this first line of code. In the body, the command parses `line2` while checking for valid input, does the requested action, flushes any output file (which is opened and closed elsewhere), and then returns `CMDok` to indicate that it terminated successfully. Along the way, any errors are trapped with the `SCMDCHECK` macro. To isolate commands from the user interface, they get file pointers by calling `scmdgetfptr`, they write to the file using `scmdfprintf`, and they flush the file with `scmdflush`.

Commands that output data to file should use the `scmdfprintf` function. This function handles output precision automatically. Also, you should separate data values using the `%`, formatting symbol, which the `scmdfprintf` function converts to either a space or a comma, depending on whether the user wants space-separated vectors or comma-separated vectors.

Commands that read numbers from user input, whether integers or floating point values should not do so with `sscanf` but should use `strmathsscanf` instead. This is a simple replacement for `sscanf` but it evaluates any formulas that the user provides for numerical input. To specify that formula evaluation should be enabled for a specific numerical input, replace the `%i` format symbol with `%mi` and replace `%lg` with `%mlg`. The function call also requires the simulation variable list. These are available as global variables, as `Varnames`, `Varvalues`, and `Nvar`.

Commands that read molecule species should do so using the `string2index1` function. This reads a string for a species name and an optional state, and then returns the species index and the state. Also, if the user did not ask for a single species but for a group of species, then this returns the full list of species in this group. It also returns error codes. The variety of outputs

5.16 Commands (functions in *smolcmd.c*)

would normally be somewhat annoying, which is what **molscan** was written to handle, described next.

A lot of commands scan over all molecules, whether to count the molecules, do other statistics on them, or other things. This looping is already somewhat complicated, and much more so if commands allow the user to enter species names with wildcards or species group names. Furthermore, commands are supposed to support these inputs if at all possible. To simplify the code in the commands, the molecule iteration loop has been written in **molscan**, and that function then calls back to the command to actually do whatever needs to be done with the identified molecules. This leads to a more complicated command structure, of which an example is in **cmdifincmpt**, which calls another command if there are some number of molecules in a specified compartment. Here is part of its listing:

```
enum CMDcode cmdifincmpt(simptr sim, cmdptr cmd, char *line2) {
    ... variable declarations ...
    moleculeptr mptr;
    static compartptr cmpt=NULL;
    static int inscan=0, count=0;

    if(inscan) goto scanportion;
    if(line2 && !strcmp(line2, "cmdtype")) return conditionalcmdtype(sim,
        cmd, 4);

    cmptss=sim->cmptss;
    ... parsing of line2 ...
    cmpt=cmptss->cmptlist[c];

    count=0;
    inscan=1;
    molscan(sim, i, index, ms, cmd, cmdifincmpt);
    inscan=0;
    if((ch=='<' && count<min) || (ch=='=' && count==min) || (ch=='>' &&
        count>min))
        return docommand(sim, cmd, line2);
    return CMDok;

scanportion:
    mptr=(moleculeptr) line2;
    if(posincompart(sim, mptr->pos, cmpt)) count++;
    return CMDok; }
```

When the command is called initially, for running it, **inscan** equals 0 and **line2** is the line to be parsed, so control passes through the first two **if** statements. The command then parses **line2** and sets up the basic variables. Importantly, any variable that will be used with the individual molecules needs to be declared as a static variable. To scan through the molecules, the command sets **inscan** to 1 and runs the scan using **molscan**, while telling **molscan** to call back to itself with each individual molecule. It would be more straightforward if **molscan** called a different function, but that would require much more code and would also make it more difficult to pass information around, so that's why the control returns to the same function. When **molscan** calls back to this function, **inscan** is set to 1, so the flow jumps down to the **scanportion** label, where the individual molecule is processed. The molecule is sent to the command using the **line2** parameter, cast as a **char***. After the molecule is processed, the command returns **CMDok** to indicate that the scan should continue. Finally, all molecules are done, **molscan** returns control back to the main portion

5 C Code: structures and functions

of the command, and the command resets `inscan` back to 0 to indicate that the scan is done. It then finishes what it needs to do and returns `CMDok` to indicate that the task is complete.

5.16.2 Externally accessible function

Not all functions are listed here because many of them don't require any more description than what is already given in the Smoldyn User Manual.

`CMDcode docommand(void *simvd, cmdptr cmd, char *line)`

`docommand` is given the simulation structure in `simvd`, the command to be executed in `cmd`, and a line of text which includes the entire command string. It parses the line of text only into the first word, which specifies which command is to be run, and into the rest of the line, which contains the command parameters. The rest of the line is then sent to the appropriate command routine as `line2`. The return value of the command that was called is passed back to the main program from `docommand`. These routines return `CMDok` for normal operation, `CMDwarn` for an error that does not require simulation termination, `CMDabort` for an error that requires immediate simulation termination, `CMDstop` for a normal simulation termination, and `CMDpause` for simulation pausing.

5.16.3 Individual command functions

simulation control

`enum CMDcode cmdstop(simptr sim, cmdptr cmd, char *line2)`

Returns a value of 2, meaning that the simulation should stop. Any contents of `line2` are ignored.

`enum CMDcode cmdpause(simptr sim, cmdptr cmd, char *line2)`

Causes the simulation to pause until the user tells it to continue. Continuation is effected by either pressing the space bar, if OpenGL is used for graphics, or by pressing enter if output is text only. The return value is 0 for non-graphics and 3 for graphics. Any contents of `line2` are ignored.

`enum CMDcode cmdbeep(simptr sim, cmdptr cmd, char *line2)`

Causes the computer to beep (sent to `stderr`). Any contents of `line2` are ignored.

`enum CMDcode cmdkeypress(simptr sim, cmdptr cmd, char *line2)`

Sets a key press event as though the key had actually been pressed.

`enum CMDcode cmdsetflag(simptr sim, cmdptr cmd, char *line2)`

Sets the command superstructure flag value.

`enum CMDcode cmdsetrandseed(simptr sim, cmdptr cmd, char *line2)`

Sets the random number seed. Negative values indicate that the current time should be used.

`enum CMDcode cmdsetgraphics(simptr sim, cmdptr cmd, char *line2)`

Sets the type of graphics output.

`enum CMDcode cmdsetgraphic_iter(simptr sim, cmdptr cmd, char *line2)`

Sets the number of iterations between each graphics update.

file manipulation

5.16 Commands (functions in *smolcmd.c*)

enum CMDcode cmdoverwrite(simptr sim, cmdptr cmd, char *line2)
Overwrites a prior output file. See the user manual.

enum CMDcode cmdincrementfile(simptr sim, cmdptr cmd, char *line2)
Closes a file, increments the name and opens that one for output. See the user manual.

enum CMDcode cmdsetrandseed(simptr sim, cmdptr cmd, char *line2)
Sets the random number seed.

conditional

enum CMDcode cmdifflag(simptr sim, cmdptr cmd, char *line2)
Runs the command in *line2* depending on value of the command superstructure flag value.

enum CMDcode cmdifprob(simptr sim, cmdptr cmd, char *line2)
Runs the command in *line2* depending on a random number.

enum CMDcode cmdifno(simptr sim, cmdptr cmd, char *line2)
Reads the first word of *line2* for a molecule name and then checks the appropriate simulation live list to see if any molecules of that type exist. If so, it does nothing, but returns 0. If not, it sends the remainder of *line2* to *docommand* to be run as a new command, and then returns 0. It returns 1 if the molecule name was missing or not recognized.

enum CMDcode cmdifless(simptr sim, cmdptr cmd, char *line2)
Identical to *cmdifno*, except that it runs the command in *line2* if there are less than a listed number of a kind of molecules in the appropriate live list.

enum CMDcode cmdifmore(simptr sim, cmdptr cmd, char *line2)
Identical to *cmdifno* except that it runs the command in *line2* if there are more than a listed number of a kind of molecules in the appropriate live list.

enum CMDcode cmdifincmpt(simptr sim, cmdptr cmd, char *line2)
Runs a command depending on the number of molecules in a compartment.

enum CMDcode cmdifchange(simptr sim, cmdptr cmd, char *line2)
Runs a command if the number of some species of molecules changes, where the user can specify various changing criteria. This function uses *i1* to indicate whether the function has been called before and *i2* to store the number of molecules that were counted in the prior invocation.

enum CMDcode cmdif(simptr sim, cmdptr cmd, char *line2)
Runs a command if one value is greater than, less than, or equal to another value.

observation commands

enum CMDcode cmdwarnescapee(simptr sim, cmdptr cmd, char *line2)
Checks for molecules that escaped from the system and displays information about them.

enum CMDcode cmdwarnescapeecmpt(simptr sim, cmdptr cmd, char *line2)
Checks for molecules that escaped from a compartment and displays information about them.

enum CMDcode cmdecho(simptr sim, cmdptr cmd, char *line2)
Echoes a string of text to the filename that is given.

5 C Code: structures and functions

`enum CMDcode cmdevaluate(simptr sim, cmdptr cmd, char *line2)`
Evaluates a math expression (usually a Smoldyn function) and prints the result to the filename that is given.

`enum CMDcode cmdmolcountheadr(simptr sim, cmdptr cmd, char *line2)`
Prints a header line for the `molcount` collection of commands.

`enum CMDcode cmdmolcount(simptr sim, cmdptr cmd, char *line2)`
Reads the output file name from `line2`. Then, to this file, it saves one line of text listing the current simulation time, followed by the number of each type of molecule in the system. This routine does not affect any simulation parameters. It accounts for both particle space and lattice space.

`enum CMDcode cmdmolcountinbox(simptr sim, cmdptr cmd, char *line2)`
Counts and outputs the number of molecules in a specific box.

`enum CMDcode cmdmolcountincmpt(simptr sim, cmdptr cmd, char *line2)`
Counts and outputs the number of molecules in a specific compartment.

`enum CMDcode cmdmolcountincmpts(simptr sim, cmdptr cmd, char *line2)`
Counts and outputs the number of molecules in multiple compartments.

`enum CMDcode cmdmolcountincmpt2(simptr sim, cmdptr cmd, char *line2)`
Counts and outputs the number of molecules in a specific compartment.

`enum CMDcode cmdmolcountonsurf(simptr sim, cmdptr cmd, char *line2)`
Counts and outputs the number of molecules on a specific surface.

`enum CMDcode cmdmolcountspace(simptr sim, cmdptr cmd, char *line2)`
Counts and outputs the number of molecules along a line profile, creating a histogram.

`enum CMDcode cmdmolcountspace2d(simptr sim, cmdptr cmd, char *line2)`
Counts and outputs the number of molecules in a 2D grid, creating a 2D histogram.

`enum CMDcode cmdmolcountspaceradial(simptr sim, cmdptr cmd, char *line2)`
Counts and outputs the number of molecules as a function of radius, creating a histogram.

`enum CMDcode cmdmolcountpolarangle(simptr sim, cmdptr cmd, char *line2)`
Counts and outputs the number of molecules as a function of polar angle, creating a histogram.

`enum CMDcode cmdradialdistribution(simptr sim, cmdptr cmd, char *line2)`
Computes and outputs a radial distribution function. This function is a nice example of scanning over pairs of molecules, where wildcards are allowed. It inputs species 1 into `i1`, `ms1`, and `index1` and species 2 into `i2`, `ms2`, and `index2`. It creates a histogram of separations in `ct=cmd->v1`, which has size `cmd->i1=nbin`. Next, it finds the low and high state numbers (`mslo` and `mshi`) and low and high list numbers (`l1lo` and `l1hi`) for species number 2. Then, it scans over species 1 using the `molscancmd` function, for which the code for each scan is in the `scanportion` section. In that portion, `mptr` is the current molecule of interest, from species 1. The function scans over the boxes that surround `mptr` and, within each box, it scans over all molecules in lists from `l1lo` to `l1hi`, testing to see if they are of type species 2 with the `molismatch` function. If it matches, the separation is computed and included in the histogram. Finally, back in the main function, data analysis is done and results are output.

5.16 Commands (functions in *smolcmd.c*)

`enum CMDcode cmdmolcountspecies(simptr sim, cmdptr cmd, char *line2)`
Counts and outputs the number of molecules of a given species.

`enum CMDcode cmdmolcountspecieslist(simptr sim, cmdptr cmd, char *line2)`
Counts and outputs the number of molecules of multiple species.

`enum CMDcode cmdlistmols(simptr sim, cmdptr cmd, char *line2)`
Reads the output file name from `line2`. To this file, it saves a list of every individual molecule in both live lists of the simulation, along with their positions. This routine does not affect any simulation parameters.

`enum CMDcode cmdlistmols2(simptr sim, cmdptr cmd, char *line2)`
Reads the output file name from `line2`. To this file, it saves the number of times this command was invoked using the `invoke` element of commands, a list of every individual molecule in both live lists of the simulation, along with their positions. This routine does not affect any simulation parameters. Routine originally written by Karen Lipkow and then rewritten by me.

`enum CMDcode cmdlistmols3(simptr sim, cmdptr cmd, char *line2)`
Reads a molecule name and the output file name from `line2`. To this file, it saves the number of times the command was invoked, the identity of the molecule specified, and the positions of every molecule of the specified type. This routine does not affect any simulation parameters.

`enum CMDcode cmdmolpos(simptr sim, cmdptr cmd, char *line2)`
Reads a molecule name and then the output file name from `line2`. To this file, it saves one line of text with the positions of each molecule of the listed identity. This routine does not affect any simulation parameters.

`enum CMDcode cmdtrackmol(simptr sim, cmdptr cmd, char *line2)`
Reads a molecule serial number and an output file name from `line2`. This searches all molecule lists for a molecule with that serial number, meaning an exact match or one part of the two-part molecule serial number is equal to the input value. If it is found, this prints out the time, molecule species name, molecule state, molecule serial number, and whether this molecule is inside or outside of each compartment sequentially. This function is primarily designed for debugging, but could be useful for single molecule tracking too. This routine does not affect any simulation parameters.

`enum CMDcode cmdmolmoments(simptr sim, cmdptr cmd, char *line2)`
Reads a molecule name and then the output file name from `line2`. To this file, it saves in one line of text: the time and the zeroth, first, and second moments of the distribution of positions for all molecules of the type listed. The zeroth moment is just the number of molecules (of the proper identity); the first moment is a dim dimensional vector for the mean position; and the second moment is a dimxdim matrix of variances. This routine does not affect any simulation parameters.

`enum CMDcode cmdsavesim(simptr sim, cmdptr cmd, char *line2)`
Reads the output file name from `line2` and then saves the complete state of the system to this file, as a configuration file. This output can be run later on to continue the simulation from the point where it was saved.

`void cmdmeansqrdispfree(cmdptr cmd)`
A memory freeing routine for memory that is allocated by `cmdmeansqrdisp`.

5 C Code: structures and functions

`enum CMDcode cmdmeansqrdisp(simptr sim, cmdptr cmd, char *line2)`

This calculates the mean square displacements of all molecules of the requested type, based on the difference between their current positions and their positions when the command was first invoked. This uses several of the command memory storage options. `i1` is the number of molecules being tracked and is the size of other arrays; `i2` is 0 if memory and initial values have not been set up, 1 if they have, or 2 if the function failed; `v1` is the list of molecule serial numbers; and `v2` is the list of initial coordinates for each molecule.

`enum CMDcode cmdmeansqrdisp2(simptr sim, cmdptr cmd, char *line2)`

This calculates the mean square displacements of molecules of the requested type, based on the difference between their current or latest positions and their positions when they were first tracked. This uses several of the command memory storage options. `i1` is the maximum number of molecules that can be tracked and is the size of other arrays; `i2` is 0 if memory and initial values have not been set up, 1 if they have, or 2 if the function failed; `i3` is the actual number of molecules being tracked; `v1` is the list of molecule serial numbers; `v2[0]` is a code equal to 0 if the molecule is not being tracked, 1 if it's not being tracked and exists, 2 if it's being tracked, or 3 if it's being tracked and exists; `v2[1, ..., dim]` is the list of initial coordinates for each molecule; and `v2[dim+1, ..., 2*dim]` is the list of current coordinates for each molecule.

`enum CMDcode cmdmeansqrdisp3(simptr sim, cmdptr cmd, char *line2)`

This calculates the effective diffusion coefficient of molecules of the requested type, based on the difference between their current or latest positions and their positions when they were first tracked, as well as their diffusion times. Effective diffusion coefficients are computed by weighting in direct proportion to the molecule lifetime. To do this, effective diffusion coefficients are the sum of squared displacements divided by (2 times the dimensionality times the sum of the molecule lifetimes). This function was modified from `cmdmeansqrdisp2`. This uses several of the command memory storage options. `i1` is the maximum number of molecules that can be tracked and is the size of other arrays; `i2` is 0 if memory and initial values have not been set up, 1 if they have, or 2 if the function failed; `i3` is the actual number of molecules being tracked; `v1` is the list of molecule serial numbers; `v2[0]` is a code equal to 0 if the molecule is not being tracked, 1 if it's not being tracked and exists, 2 if it's being tracked, or 3 if it's being tracked and exists; `v2[1, ..., dim]` is the list of initial coordinates for each molecule; `v2[dim+1, ..., 2*dim]` is the list of current coordinates for each molecule, `v2[2*dim+1]` is the simulation time when tracking began on this molecule; and `f1` is the effective diffusion coefficient from the prior invocation, or -1 on initialization.

`enum CMDcode cmdresidencetime(simptr sim, cmdptr cmd, char *line2)`

This calculates the residence time of molecules of the requested type, based on the difference between the current time and the molecule creation times. This function was modified from `cmdmeansqrdisp3`. This uses several of the command memory storage options. `i1` is the maximum number of molecules that can be tracked and is the size of other arrays; `i2` is 0 if memory and initial values have not been set up, 1 if they have, or 2 if the function failed; `i3` is the actual number of molecules being tracked; `v1` is the list of molecule serial numbers; `v2[0]` is a code equal to 0 if the molecule is not being tracked, 1 if it's not being tracked and exists, 2 if it's being tracked, or 3 if it's being tracked and exists; and `v2[1]` is the simulation time when tracking began on this molecule.

`enum CMDcode cmddiagnostics(simptr sim, cmdptr cmd, char *line2)`

5.16 Commands (functions in *smolcmd.c*)

Displays diagnostics about different data structures to the standard output.

enum CMDcode cmdexecutiontime(simptr sim, cmdptr cmd, char *line2)

Prints simulation time and execution time to file.

enum CMDcode cmdprintLattice(simptr sim, cmdptr cmd, char *line2)

Prints information about the lattices to the a file. This would be better as part of the cmddiagnostics command.

enum CMDcode cmdprintFilament(simptr sim, cmdptr cmd, char *line2)

Prints information about a specific filament to a file.

enum CMDcode cmdprintFilaments(simptr sim, cmdptr cmd, char *line2)

Prints information about all filaments to a file. New runtime command `printFilaments <file>` writes every filament of every type (auto-named daughters included) as a flat node polyline plus its parent name. Independent diagnostic/analysis tool; useful for anyone inspecting filaments, since `filWrite()`/`savesim()` do not serialize filaments (stub).

enum CMDcode cmdwriteVTK(simptr sim, cmdptr cmd, char *line2)

Outputs VTK data for the current simulation state.

enum CMDcode cmdprintdata(simptr sim, cmdptr cmd, char *line2)

Prints out a data array to a file, another data array, or stdout.

system manipulation

enum CMDcode cmdset(simptr sim, cmdptr cmd, char *line2)

Reads `line2` to extract the first word and the rest of the line. These are passed to `simreadstring` where they are interpreted as a configuration file line. Any errors are passed on from `simreadstring`.

enum CMDcode cmdpointsource(simptr sim, cmdptr cmd, char *line2)

Reads `line2` for a molecule name, followed by the number of molecules that should be created, followed by the `dim` dimensional position for them. If all reads well, it creates the new molecules in the system at the appropriate position.

enum CMDcode cmdvolumesource(simptr sim, cmdptr cmd, char *line2)

Reads `line2` for a molecule name, the number of molecules that should be created, and a region. If all reads well, it creates the new molecules in the system in the appropriate volume.

enum CMDcode cmdgaussiansource(simptr sim, cmdptr cmd, char *line2)

Reads `line2` for a molecule name, the number of molecules that should be created, and a region. If all reads well, it creates the new molecules in the system with a Gaussian distribution.

enum CMDcode cmdmovesurfacemol(simptr sim, cmdptr cmd, char *line2)

Probabilistically moves molecules from one surface to another, with an optional state change.

enum CMDcode cmdkillmolinsphere(simptr sim, cmdptr cmd, char *line2)

Reads `line2` for a molecule name and a surface name and then kills all molecules of the given type, that are in spheres of the listed surface. The molecule name and/or the surface name can be "all".

5 C Code: structures and functions

- `enum CMDcode cmdkillmolincmpt(simptr sim, cmdptr cmd, char *line2)`
Reads `line2` for a molecule name and a compartment name and then kills all molecules of the given type, that are in the listed compartment. The molecule name can be “all”.
- `enum CMDcode cmdkillmoloutsidesystem(simptr sim, cmdptr cmd, char *line2)`
Reads `line2` for a molecule name and then kills all molecules of the given type that are outside of the system boundaries. The molecule name and/or the surface name can be “all”.
- `enum CMDcode cmdfixmolcount(simptr sim, cmdptr cmd, char *line2)`
Fixes the copy number of a specific molecule species, only considering solution phase, to a specified quantity. This considers the entire system, and adds or removes molecules as required.
- `enum CMDcode cmdfixmolcountrange(simptr sim, cmdptr cmd, char *line2)`
Fixes the copy number of a specific molecule species, only considering solution phase, to a specified range. This considers the entire system, and adds or removes molecules as required.
- `enum CMDcode cmdfixmolcountonsurf(simptr sim, cmdptr cmd, char *line2)`
Fixes the copy number of surface-bound molecules to a specified value. This considers an entire surface, and adds or removes molecules as required.
- `enum CMDcode cmdfixmolcountonsurf(simptr sim, cmdptr cmd, char *line2)`
Fixes the copy number of surface-bound molecules to a specified range. This considers an entire surface, and adds or removes molecules as required.
- `enum CMDcode cmdfixmolcountincmpt(simptr sim, cmdptr cmd, char *line2)`
Fixes the copy number of molecules in a compartment to a specified value. This adds or removes molecules as required.
- `enum CMDcode cmdfixmolcountrangeincmpt(simptr sim, cmdptr cmd, char *line2)`
Fixes the copy number of molecules in a compartment so that it is within the specified range. This adds or removes molecules as required.
- `enum CMDcode cmdequilmol(simptr sim, cmdptr cmd, char *line2)`
Equilibrates a pair of molecular species, allowing the efficient simulation of rapid reactions. It reads two molecule names from `line2`, followed by a probability value. Then, it looks for all molecules in the live lists with either of the two types and replaces them with the second type using the listed probability or with the first type using 1 the listed probability.
- `enum CMDcode cmdreplacemol(simptr sim, cmdptr cmd, char *line2)`
Replaces one molecule species with another, with some replacement probability. Modified from `cmdreplacecmptmol`.
- `enum CMDcode cmdreplacexyzmol(simptr sim, cmdptr cmd, char *line2)`
Reads the name of a molecule following by a `dim` dimensional point in space from `line2`. Then, it searches the fixed live list for any molecule that is exactly at the designated point. If it encounters one, it is replaced by the listed molecule, and then the live lists are sorted if appropriate. This routine stops searching after one molecule has been found, and so will miss additional molecules that are at the same point.
- `enum CMDcode cmdreplacevolmol(simptr sim, cmdptr cmd, char *line2)`
Replaces one molecule species in a volume with another, with some replacement probability.

5.16 Commands (functions in *smolcmd.c*)

- enum CMDcode cmdreplacecmptmol(simp_{ptr} sim, cmd_{ptr} cmd, char *line2)**
Replaces one molecule species in a compartment with another, with some replacement probability. This is a modification of `cmdreplacevolmol`, and uses a little code from `cmdfixmolcountincmpt`.
- enum CMDcode cmdmodulatemol(simp_{ptr} sim, cmd_{ptr} cmd, char *line2)**
Identical to `cmdequilmol` except that the equilibration probability is not fixed, but is a sinusoidally varying function. After reading two molecule names from `line2`, this routine then reads the cosine wave frequency and phase shift, then calculates the probability using the function $prob = 0.5 * (1.0 - \cos(freq * sim \rightarrow time + shift))$.
- enum CMDcode cmdreact1(simp_{ptr} sim, cmd_{ptr} cmd, char *line2)**
Reads `line2` for the name of a molecule followed by the name of a unimolecular reaction. Then, every one of that type of molecule is caused to undergo the listed reaction, thus replacing each one by reaction products. Molecules are sorted at the end. This might be useful for simulating a pulse of actinic light, for example.
- enum CMDcode cmdsetrateint(simp_{ptr} sim, cmd_{ptr} cmd, char *line2)**
This reads `line2` for the name of a reaction and the new internal rate constant for it. The internal rate constant is set to the new value. Errors can arise from illegal inputs, such as the reaction not being found or a negative requested internal rate constant.
- enum CMDcode cmdshufflemollist(simp_{ptr} sim, cmd_{ptr} cmd, char *line2)**
Shuffles one or more of the live internal molecule lists.
- enum CMDcode cmdshufflereactions(simp_{ptr} sim, cmd_{ptr} cmd, char *line2)**
Shuffles reactions for one or more reactant combinations, for bimolecular reactions. This function is somewhat inefficient in that it shuffles all lists twice if it is called for all reactant pairs, but I decided to keep the inefficiency because improvement, and still allowing for scanning over all of either reactant individually, would lead to slower and messier code.
- enum CMDcode cmdsetsurfcoeff(simp_{ptr} sim, cmd_{ptr} cmd, char *line2)**
Sets the surface interaction rate and then calls `surfacesupdate` to update the probabilities. Zsuzsanna Sukosd wrote this command.
- enum CMDcode cmdsetimestep(simp_{ptr} sim, cmd_{ptr} cmd, char *line2)**
This reads `line2` for the new simulation time step. Nothing else is changed in the simulation, including binding or unbinding radii, so reaction rates may be observed to change.
- enum CMDcode cmdexcludebox(simp_{ptr} sim, cmd_{ptr} cmd, char *line2)**
Allows a region of the simulation volume to be effectively closed off to molecules. The box is defined by its low and high corners, which are read from `line2`. Any molecule, of any type, that entered the box during the last time step, as determined by its `pos` and `posx` structure members, is moved back to its previous position. This is not the correct behavior for a reflective surface, but is efficient and expected to be reasonably accurate for most situations. This routine ought to be replaced with a proper treatment of surfaces in the main program (rather than with interpreter commands), but that's a lot more difficult.
- enum CMDcode cmdexcludesphere(simp_{ptr} sim, cmd_{ptr} cmd, char *line2)**
Like `cmdexcludebox` except that it excludes a sphere rather than a box. The sphere is defined by its center and radius, which are read from `line2`. Any molecule, of any type, that entered

5 C Code: structures and functions

the sphere during the last time step, as determined by its `pos` and `posx` structure members, is moved back to its previous position. This is not the correct behavior for a reflective surface, but is efficient and expected to be reasonably accurate for most situations.

`enum CMDcode cmdincludeecoli(simptr sim, cmdptr cmd, char *line2)`

This is the opposite of the `excludebox` and `excludesphere` commands. Here, molecules are confined to an *E. coli* shape and are put back inside it if they leave. See the user manual for more about it. Unlike the other rejection method commands, this one works even if a molecule was in a forbidden region during the previous time step; in this case, the molecule is moved to the point on the *E. coli* surface that is closest. Because of this difference, this command works reasonably well even if it is not called at every time step.

`enum CMDcode cmdsetreactionratemolcount(simptr sim, cmdptr cmd, char *line2)`

This command sets the reaction rate to be a linear function of one or more molecule counts. I considered letting the user select compartments for the molecules, but that turned out to be too much work. Instead, I suggest that the user use different species names in different regions of space as much as convenient, and use wildcards to collect different species together. That will run faster anyhow.

`enum CMDcode cmdexpandsystem(simptr sim, cmdptr cmd, char *line2)`

This command expands or contracts the entire system, including all molecule positions and surfaces. The work for surfaces is performed by `surftransformpanel`.

`enum CMDcode cmdtranslatecmpt(simptr sim, cmdptr cmd, char *line2)`

This command translates a compartment, including dealing with molecule displacements. All of the work is done by `comparttranslate`.

`enum CMDcode cmddiffusecmpt(simptr sim, cmdptr cmd, char *line2)`

This command diffuses a compartment, including dealing with molecule displacements. All of the work is done by `comparttranslate`. For bounded diffusion, a move that would lead to an illegal location is rejected, and diffusion is attempted again, up to 10 times; if there are 10 failures, then the compartment simply isn't moved at all during that command call.

`enum CMDcode cmdlongrangeforce(simptr sim, cmdptr cmd, char *line2)`

This command simulates a long range force between pairs of molecules. It uses loops over molecule pairs that's largely borrowed from `cmdradialdistribution`. Also, it keeps a list of molecules that have forces on them, which is in `cmd->v1`, used internally as `moleclist`. Each entry in this list includes the molecule serial number (on which the list is sorted), the molecule pointer, and a 4-dimensional vector of doubles. The first three dimensions are for the *x*, *y*, and *z* components of the force on a molecule, and the fourth is 0 for molecules that were never visited during the last round of checks and 1 for molecules that were visited. Unvisited molecules get purged from this list, partly because those molecules might have gone away due to reactions.

The `scanportion` region scans over all molecules that fit the *species*₁ category. Each of these molecules is returned, one at a time, in `mptr`. This `mptr` is added to `moleclist`. The function then scans over all possible *species*₂ molecules that are within `rmax` distance of this molecule and checks to see if they interact. The `duplicate` test investigates whether `mptr2` is an instance of *species*₁ because if it is, then it's already in the `moleclist`, and if it's not then it needs to be put in that list, which is done. If it's put in the list, then that might shift the list by one, in which case *j* becomes incorrect; to address this, *j* is recomputed.

5.16 Commands (functions in *smolcmd.c*)

enum CMDcode cmdtranslatemol(simptr sim, cmdptr cmd, char *line2)

Performs location-dependent drift such as for advection. This translates molecules of the given species and state (each of which are allowed to be “all”) by the displacement that is calculated from the equations, with one equation for each coordinate. These equations are allowed to be functions of the position and/or the time; the position is entered as *x*, *y*, and/or *z*, and the time as *time*. Do not include spaces in the equations and simply enter “0” for an equation if there’s no translation on that axis. For example, in a 3D system, the three equations *y -x 0* cause the molecules to circle the origin in the *x–y* plane and stay stationary along the *z* axis.

enum CMDcode cmdRotateFilamentToY(simptr sim, cmdptr cmd, char *line2)

This is a specialized command for a particular study and not likely to be generally useful. It takes a filament and rotates it about the *x*-axis (translated to the front of the filament) so that the filament’s center of mass lies along the positive *y*-axis. To do so, this first computes the mean position of the nodes (minus any offset from node 0) on the *y* and *z* axes. It then normalizes these values. Then, it rotates the node positions with the rotation equation

$$\begin{bmatrix} y' \\ z' \end{bmatrix} = \begin{bmatrix} m_y & m_z \\ -m_z & m_y \end{bmatrix} \begin{bmatrix} y \\ z \end{bmatrix}$$

Internal functions

void cmdv1free(cmdptr cmd)

Frees array *cmd->v1*.

void cmdv1v2free(cmdptr cmd)

Frees arrays *cmd->v1* and *cmd->v2*.

enum CMDcode conditionalcmdtype(simptr sim, cmdptr cmd, int nparam)

Returns the command type for conditional commands, which are required to return the type of the function that gets called if the condition is true. *cmd* is the conditional command and *nparam* is the number of parameters for the conditional command (e.g. for *cmdifno*, the only parameter is the molecule name, so *nparam* is 1).

int insideecoli(double *pos, double *ofst, double rad, double length)

This is a short utility routine used by the command *cmdincludeecoli*. It returns a 1 if a molecule is inside an *E. coli* shape and a 0 if not. *pos* is the molecule position, *ofst* is the physical location of the cell membrane at the center of the low end of the cell (the cell is assumed to have its long axis along the *x*-axis), *rad* is the cell radius used for both the cylindrical body and the hemispherical ends, and *length* is the total cell length, including both hemispherical ends.

void putinecoli(double *pos, double *ofst, double rad, double length)

This is another short utility routine used by the command *cmdincludeecoli*. It moves a molecule from its initial position in *pos* to the nearest surface of an *E. coli* shape. Parameters are the same as those for *insideecoli*.

int molinpanels(simptr sim, int ll, int m, int s, char pshape)

This function, which might be better off in the *smolsurf.c* code, is used to test if molecule number *m* of live list *ll* is inside any of the *pshape* panels of surface number *s*. Only spheres are allowed currently as panel shapes, because neither rectangles nor triangles can contain

5 C Code: structures and functions

molecules. If `s` is sent in with a value less than 0, this means that all `pshape` panels of all surfaces will be checked.

5.17 Top-level code (functions in `smoldyn.c`)

The top-level source code file, `smoldyn.c`, contains only the `main` function for the stand-alone Smoldyn program. It is declared locally in this source file and so cannot be called from externally, except from the shell. This source code file is not included in `Libsmoldyn`.

```
int main(int argc, char *argv[])
```

`main` is a simple routine that provides an entry point to the program from the shell. It checks the command line arguments, prints a greeting, inputs the configuration file name from the user, and then calls `simInitAndLoad` to load the configuration file and set up all the structures. If all goes well, it calls `smolsimulate` or `smolsimulategl` to run the simulation. At the end, it returns to the shell.

5.18 Libsmoldyn (functions in `libsmoldyn.cpp`)

`Libsmoldyn` is a C/C++ API for the Smoldyn code. It is designed to be redundant with the configuration file input, but accessible from code. This description is far from complete but needs to be finished at some point. See also documentation for it in the `SmoldynManual`.

5.18.1 Error handling

Essentially every user-accessible function returns an error code. Also, functions define their function name in `const char* funcname="function";`. Errors are checked with either the `LCHECK` or `LCHECKNT` macros. These are supposed to differ with `LCHECK` dealing with the error itself and throwing an exception if necessary, and the `LCHECKNT` (no throw) macro not dealing with the error but just copying the message over into global variables. However, in practice, throwing doesn't work anyhow, so the only real difference is that the regular version displays an error message to standard error, whereas the NT version doesn't. Anyhow, if an error is detected, these macros send the function name, the error code, a string error message, and any simulation flags to the `libsmoldyn` functions `smolSetError` or `smolSetErrorNT`.

At `smolSetError` and `smolSetErrorNT`, which are essentially identical, several function inputs are copied over to global variables: the function name to `Liberrorfunction`, the error code to `Liberrorcode`, and the error message to `Liberrorstring`. These items are also displayed to standard error if appropriate.

The global error variables are pretty much only used in one place, which is the `smolGetError` function, which copies their values over to new variables and returns those to the calling code. In addition, each function in the library returns `Liberrorcode`, which it just set if an error occurred.

A few more error handling steps arise if `smolPrepareSimFromFile` or `smolLoadSimFromFile` are used.

In `smolLoadSimFromFile`, this calls the `smolsim.c` function `loadsim`. Here, errors are written by the macro `CHECKS`, which copies the error message to the global variable `ErrorString`. Errors are then handled by `simParseError`, which concatenates the error message from `Parse.ReadFailure` (the file name and file line) with the error message from `ErrorString`. `simParseError` sends the result off to `simLog`, which writes the error to a logfile if one has been created or to standard output otherwise.

5.18 Libsmoldyn (functions in libsmoldyn.cpp)

Here is a typical example of error handling. Starting with stand-alone Smoldyn first, the program starts in `main` in `smoldyn.cpp`, which calls `simInitAndLoad` in `smolsim.c`, which calls `loadsims`, which calls `simreadstring`. In the `simreadstring` function, the input parser encounters an error. It writes an error message to the global variable `ErrorString`, goes to its failure label, and then on to `simParseError`. Here, `Parse_ReadFailure` is called and sent an empty string, and that function returns the string with text that reads “Error reading file in line *number*. \n line: *text of line* \n file: *filename*”; if the input line included substituted text, it also gives the substituted version of the line in the string as well. `Parse_ReadFailure` also closes files and stops the input process. Back in `simParseError`, the input file information is concatenated with the error message to create the complete error string. This is done twice; first to create an error string in the global variable `ErrorLineAndString` and also as the input to `simLog`. The `simLog` function prints the error message to the logging file if one was declared or to standard output if not (unless the ‘s’ runtime flag was selected, in which case the error output is suppressed). Next, `simParseError` returns to `simreadstring`, which returns an error code to `loadsims`, which returns an error code to `simInitAndLoad`, which returns an error code to `main`. Here, the code sees the error message, skips the simulation, and quits with an exit code of 0.

For Libsmoldyn, the flow is similar, but here an entry point is `smolPrepareSimFromFile`. This calls `simInitAndLoad`. As before, an error is encountered, it is displayed to the standard output if the ‘s’ flag isn’t selected (it often should be selected for Libsmoldyn use) and also written to `ErrorLineAndString`. Then, control returns back to `smolPrepareSimFromFile`. This sees the error, so it follows the `LCHECK` macro to head off to `smolSetError`, with the `ErrorLineAndString` contents as the error message. This function prints the error message out a second time, which is a little inelegant, and also copies the message into the library global variable `Liberrorstring`. At this point, control returns to the calling program. If desired, that program can then call `smolGetError` to get the error string again, now in a string variable rather than displayed to the output.

6 Code design

This chapter describes interactions between different portions of the code. The code is, and may always be, in flux. While I try to maintain this section of the documentation, be forewarned that it might not reflect the most recent changes.

6.1 Memory management

The following table, *which is very out of date* shows memory allocation and freeing for the different structures. For both the allocation and freeing columns, the top line shows the function in which the structure is actually allocated or freed, while subsequent lines show the functions that call the preceding functions.

structure	allocation	freeing
moleculestruct	molalloc	molfree
	molexpandlist	molssfree
	molsetmaxmol, molsort	simfree
	simreadstring (max_mol), ?	
	loadsim, ?	
	simInitAndLoad, ?	
moleculesuperstruct	molssalloc	molssfree
	molsetmaxspecies	simfree
	simreadstring (max_species, ..)	
	loadsim	
	simInitAndLoad	
	molsssetgausstable	
	simreadstring (gauss_table_size), setupmols	
	loadsim, simupdate	
	simInitAndLoad; simInitAndLoad, simulatetimestep	
	mollistalloc	
wallstruct	addmollist	
	setupmols, simreadstring (molecule_lists), setupports	
	simupdate, loadsim	
	simInitAndLoad	
	wallalloc	wallfree
	wallsalloc	wallsfree
	walladd	simfree
	simreadstring (boundaries, ..)	
	loadsim	

6 Code design

	simInitAndLoad	
rxnstruct	rxnalloc RxnAddReaction loadsim (reaction, ..) simInitAndLoad	rxnfree rxnssfree simfree
rxnsuperstruct	rxnssalloc RxnAddReaction simreadstring (reaction, ..), loadrxn loadsim, ? simInitAndLoad	rxnssfree simfree
panelstruct	panelsalloc surfreadstring (max_panels) loadsurface simreadstring (start_surface) loadsim simInitAndLoad	panelfree surfacefree surfacessfree simfree
surfacestruct	surfacealloc surfacessalloc simreadstring (max_surface) loadsim simInitAndLoad	surfacefree surfacessfree simfree
surfacesuperstruct	surfacessalloc surfenablesurfaces simreadstring (max_surface) loadsim simInitAndLoad	surfacessfree simfree
boxstruct	boxalloc boxesalloc setupboxes simupdate simInitAndLoad	boxfree boxesfree boxssfree simfree
boxsuperstruct	boxssalloc boxsetsize simreadstring (boxsize, ..), setupboxes loadsim, simupdate simInitAndLoad, simInitAndLoad	boxssfree simfree
compartstruct	compartalloc compartssalloc simreadstring (max_compartment) loadsim	compartfree compartssfree simfree

6.2 Data structure preparation and updating

	<code>simInitAndLoad</code>	
<code>compartsuperstruct</code>	<code>compartssalloc</code> <code>simreadstring (max_compartment)</code> <code>loadsim</code> <code>simInitAndLoad</code>	<code>compartssfree</code> <code>simfree</code>
<code>portstruct</code>	<code>portalloc</code> <code>portssalloc</code> <code>simreadstring (max_port)</code> <code>loadsim</code> <code>simInitAndLoad</code>	<code>portfree</code> <code>portssfree</code> <code>simfree</code>
<code>portsuperstruct</code>	<code>portssalloc</code> <code>simreadstring (max_port)</code> <code>loadsim</code> <code>simInitAndLoad</code>	<code>portssfree</code> <code>simfree</code>
<code>mzrsuperstruct</code>	<code>mzrssalloc</code> <code>mzrssload</code> <code>simreadstring (read_network_rules)</code> <code>loadsim</code> <code>simInitAndLoad</code>	<code>mzrssfree</code> <code>simfree</code>
<code>simstruct</code>	<code>simalloc</code> <code>simInitAndLoad</code>	<code>simfree</code>

6.2 Data structure preparation and updating

The original Smoldyn design was that it read a configuration file, set up internal data structures, ran the simulation, and then quit. Two problems arose. First, the sequence for setting up the internal data structures became increasingly complicated as the program gained features, and second, this overall program flow wasn't always what users wanted. As a result, the code is now more modularized and it includes a "condition" element in each data structure that reports on the overall state of that data structure. These help, but data structure set up and updating are still somewhat complicated, which is the focus of this section.

Note that the condition element and the data structure updating discussed here are designed for relatively infrequent system modifications. That is, they are for changes to the system, performed by the user or by run-time commands. They are not designed for regular simulation operation, such as for updates after molecules diffuse or reactions occur.

The highest level updating function is `simupdate`. This can be called with data structures that have just been loaded in from a configuration file, that have just been created using Libsmoldyn functions, that have been modified using run-time commands, or that are in good running order. It will take care of all necessary updates. `simupdate` calls the other update functions in the order listed below. Because updates cannot be finished until others are started, `simupdate` will sometimes call update functions multiple times, until everything is done.

All update functions return the same types of values. They return 0 for success, 1 if memory

6 Code design

could not be allocated, 2 if a different data structure that needed updating hasn't been adequately updated yet, or larger numbers for other errors.

Structure conditions (SC) are divided into four categories: **SCinit**, **SClists**, **SCparams**, and **SCok**; these are enumerated in the **StructCond** enumerated list. **SCinit** implies that a structure is still being initialized, **SClists** implies that one or more of the lists that comprise a structure need updating, **SCparams** implies that one or more of the structure simulation parameters need updating, and **SCok** implies that the structure is fully updated and ready for use. The items that are initialized for each condition are shown below. This table was updated after changes for version 2.23.

molsupdate

SClists (molsupdatelists)

- gaussian lookup table
- mols->exist (calls rxnisprod and issurfprod)
- creates system molecule lists if none yet
- sets any list lookup values that weren't done yet
- sets molecule list values for molecules in dead list

SCparams (molsupdateparams)

- mols->difstep
- mols->diffuselist

boxesupdate

SClists (boxesupdatelists)

- box superstructure (requires **walls**)
- bptr->indx
- neighbor boxes
- bptr->nwall and bptr->wlist

SCparams (boxesupdateparams)

- bptr->npanel and bptr->panel (requires **surfaces**)
- molecules in boxes (requires **molecules** and assigns mptr->box)

molsort

- resets topl indices
- sorts and compacts live lists (calls boxremovemol and boxaddmol)
- moves molecules from resurrected to reborn lists (calls boxaddmol)
- resets sortl indices

compartupdate

SClists (compartsupdatelists)

- does nothing

SCparams (compartupdateparams)

- finds boxes in compartment (requires **boxes** and **surfaces**)
- finds box volumes and compartment volumes

rxnsupdate

SClists (rxnsupdatelists)

- sets reaction molecule lists (requires **molecule lists**)

SCparams (rxnsupdateparams)

```

sets rates (may require compartments)
sets products
calculates tau values

```

surfupdate

```

SLists (surfupdatelists)
  allocates srfmollist arrays (requires molecule lists)
  sets srfmollist values (requires molecule lists and calls rxnisprod)
SCparams (surfupdateparams)
  sets surface interaction probabilities (requires mols->difc and ->difstep)

```

portsupdate

```

SLists (portsupdatelists)
  sets port molecule lists (uses molecule lists)
SCparams (portsupdateparams)
  does nothing

```

6.3 User input

Text-based input follows the following sequence of functions. This example is for a 4-line input file, of which the first two are parsed by `simreadstring` and the next two by `surfreadstring`.

```

main
  simInitAndLoad
    loadsim
      simreadstring
    loadsim
      simreadstring
    loadsim
      loadsurface
        surfreadstring
      loadsurface
        surfreadstring
      loadsurface
    loadsim
  simInitAndLoad
main

```

Any errors are passed out through the same chain of functions.

6.4 Simulation algorithm sequence

In a sense, the core function of the entire Smoldyn program is `simulatetimestep`, which is in `smolsim.c`. Using the assumption that all data structures are given to this function are in good working order, `simulatetimestep` runs the simulation for one time step, and then cleans up all of the data structures as required, thus leaving them again in good working order. The sequence

6 Code design

of tasks performed here is central to accurate and efficient simulation, as well as data structure maintenance.

As mentioned in the Smoldyn User Manual, the simulation concept is that the system can be observed at a fixed time, then evolves to a new state, can be observed again, and so forth. The details of what happens during the time evolution should not be relevant to the user. Instead, all that should matter is that all aspects of the observable simulated states conform as closely as possible to results that would be found from the ideal and hypothetical continuous-time model system that the simulation is intended to represent. Furthermore, the simulation results should be as independent of the time step as possible, and they should converge to those of the hypothetical model system as the time step is reduced towards zero.

Here, we focus on the details of what happens during a time step. The sequence of simulation algorithms is:

operation	error code	box assignment	molecule sorting	surface sides
_____time = t _____				
observe and manipulate system	er	ok	bad	ok
sort molecule lists	n/a	ok	ok	ok
graphics are drawn	n/a	ok	ok	ok
_____start/end of <code>simulatetimestep</code> _____		ok	ok	ok
diffusion				
molecules diffuse	9	bad	ok	bad
surface collisions				
surface collision interactions	n/a	bad	bad	ok
reactions				
desorption and surface-state transitions	6	bad	bad	ok
assign molecules to boxes	2	ok	bad	ok
0th order reactions	3	ok	bad	bad
1st order reactions	4	ok	bad	bad
2nd order reactions	5	ok	bad	bad
sort molecule lists	6	ok	ok	bad
surface collisions again	n/a	bad	ok	ok
assign again	2	ok	ok	ok
_____time = $t + \Delta t$ _____				

The overall picture is that the simulation sequence is: diffusion, surface interactions, and reactions. The “desorption and surface-state transitions” operation is really a reaction, and so it is debatable whether it should be before or after chemical reactions. I chose to put it where it is because it shares overlap with surface collision interactions, which have to be before reactions.

After commands are run, graphics are displayed to OpenGL if that is enabled. The evolution over a finite time step starts by diffusing all mobile molecules. In the process, some end up across internal surfaces or the external boundary. These are reflected, transmitted, absorbed, or transported as needed. Next, reactions are treated in a semi-synchronous fashion. They are asynchronous in that all zeroth order reactions are simulated first, then unimolecular reactions, and finally bimolecular reactions. With bimolecular reactions, if a molecule is within the binding radii of two different other molecules, then it ends up reacting with only the first one that is checked, which is arbitrary (but not necessarily random). Reactions are synchronous in that reactants are removed from the

system as soon as they react and products are not added into the system until all reactions have been completed. This prevents reactants from reacting twice during a time step and it prevents products from one reaction from reacting again during the same time step. As it is possible for reactions to produce molecules that are across internal surfaces or outside the system walls, those products are then reflected back into the system. At this point, the system has fully evolved by one time step. All molecules are inside the system walls and essentially no pairs of molecules are within their binding radii (the exception is that products of a bimolecular reaction with an unbinding radius might be initially placed within the binding radius of another reactant).

6.5 Wildcards, species groups, and patterns

6.5.1 Patterns

Most of Smoldyn’s treatment of wildcards, species groups, and patterns is in the molecules and reactions portion of the code. The basic idea is that any string that the user enters for a species is treated as a “pattern”. This pattern might represent one species or multiple species.

Patterns are stored in the molecule superstructure. The `PatternData` enumeration, of which there are `PDMAX` options, are for the first `PDMAX` elements of the `mols->patindex` array, called the header. This array, as explained below, lists the species indices that match to text patterns. Each listing requires some information to describe how long the listing is, what’s stored in it, and what it corresponds to, which is included in the header. These header elements should be retrieved using the enumerated values. See below for their description.

Patterns are strings, each of which represents one or more species names. Each string that is used gets recorded as a pattern; however, those that are not used (e.g. species that were generated automatically but never referenced individually by the user) are not recorded here, so there is no certainty that the list here is complete. This list is only updated when it is used.

Patterns are stored in the molecule superstructure. This paragraph describes how they are stored. `maxpattern` slots are allocated for patterns, of which `npattern` are actually used. The actual list of patterns is called `patlist`; these are sorted by alphabetical order. Sorted in parallel are the lists `patindex` which is a list of lists, and `patrname`, which is a list of reaction names. `patrname` is only used if the pattern represents a reaction, and is used to differentiate between multiple different reactions that have identical patterns. In `patindex`, there is one element for each pattern. Each element of `patindex` is a sublist.

The contents of each `patindex` list starts with a header which gives important information about the list. The header occupies the first `PDMAX` list elements. They are:

element	name	meaning
0	<code>PDalloc</code>	allocated sublist size including header values
1	<code>PDnresults</code>	number of stored results
2	<code>PDnspecies</code>	number of species during last sublist update
3	<code>PDmatch</code>	number of “match species” in this pattern
4	<code>PDsubst</code>	number of “replace species” in this pattern
5	<code>PDrule</code>	0 if this is not a “rule” and 1 if it is (only rules generate new species)

`PDalloc` is the total size of the index list. This value is checked when the index list needs to be expanded and is updated exclusively by `molpatternindexalloc`. `PDnresults` is the number of stored results, not including the values in the header. It may be 0 or more. `PDnspecies` is the number of species that existed in the system when the index list was last updated. This value is

6 Code design

initialized to 1, because a minimal simulation that includes molecules includes an empty species. It is increased to `sim->mols->nspecies` at each update. This value can also be set to -1 to indicate that it should not be updated in the future, even if more species are created (appropriate for a pattern that represents only a single species).

The number of spaces used for each stored result is the sum of the match and replace counts. For example, suppose a pattern is `A*\nX*`, which has been found to match and substitute to `AB\nXB`, `AC\nXC`, and `AD\nXD`. If the sublist were allocated with 20 elements, then element 0 would equal 13, element 1 would equal 3, element 2 would equal 6 if `nspecies` was 6 when this sublist was created, element 3 would equal 1 because there's one match word here, element 4 would equal 1 because there's one substitution word here, element 5 would equal 0 if this is not a rule, elements 6 and 7 would be the indices for the "AB" and "XB" species, elements 8 and 9 for "AC" and "XC" species, elements 10 and 11 for "AD" and "XD" species, and elements 12 to 19 would be empty. The list of results is sorted in ascending order.

6.5.2 Control flow and functions - species names

The entry point for most functions that deal with species names that are provided by the user is `molstring2index1`. This function takes a string (typically a species name, a species group name, or a species name with wildcards, optionally followed by the state) and returns an index variable for the species, creating a new index if needed. It also returns the state, returning `MSsoln` if the string did not include a state. For example, adding a species group happens in `moladdspeciesgroup`; the first thing this does is to send the species group name to `molstring2index1` so that it can get back an index variable for the species group. The same function also sends its species name to `molstring2index1` so that it can get back an index variable for the species. `molstring2index1` is also called by `simreadstring` for essentially all input lines that include species, including "dife", "difm", "drift", "surface_drift", "mol", "surface_mol", "compartment_mol", "mol_list", "display_size", "color", and "product_placement". It is also called by a vast number of commands.

Internally, `molstring2index1` sends the string off to `molstring2pattern` where the string is separated into a pattern portion, which might be a species name, a species group name, or a species name with wildcards, and a state for the species. This function can also assemble more complex patterns if used with non-zero mode values, which are for reactions, described below. No wildcard stuff is parsed in `molstring2pattern`. This function simply returns the pattern and the state. Next, `molstring2index1` sends the pattern off to `molpatternindex` which looks to see if the pattern has been expanded before into a list of species and, if so, returns that list of species. If not, it performs any necessary wildcard matching to create a list of species, which it both stores and returns in `index`. At the end, `molstring2index1` returns the index variable.

6.5.3 Reaction rule storage

Reaction rules are stored within the reaction superstructure. A simulation can have up to three reaction superstructures, corresponding to order 0, 1, and 2 reactions, but it's best if all of the rules are stored together because then it's possible to keep them in the order that the user entered them. As a result, Smoldyn first looks for existing rules and if it doesn't find any, then it stores the rules in the lowest order superstructure that is currently defined (this happens at the top of `RuleAddRule`). Any functions that use rules need to look through all of the superstructures to find them.

The rule list, within the reaction superstructure, has allocated size `maxrule`, actual size `nrule` and several lists, which are `rulename`, `rulepattern`, `ruledetails`, and `rulerate`. The `rulename`

list contains strings with the rule names, which are simply the reaction names entered by the user, without any suffix. The `rulepattern` list is a list of patterns which serve as reaction rules; a pattern for a bimolecular reaction has a single-species pattern for the first reactant, which is a species name possibly with wildcard characters, a space, a second single-species pattern, a newline character, and then space-separated product species names. The `ruledetails` list contains arrays of integers, each of which lists the species states cast as integers for the reactants and products in their sequence, then the compartment number if the reaction is restricted to a specific compartment, and then the surface number if the reaction is restricted to a single surface. The `rulerate` list contains the reaction rate entered by the user. The sole difference between reaction rules and reactions (which are not stored here) is that the rules are allowed to generate new species whereas reactions are not. Reaction rules are typically only useful in combination with wildcard characters. These lists are maintained by the `RuleAddRule` function. These rules do not need updating during the simulation. Instead, they are simply stored here and referred to when updating is required. The `ruleonthe-fly` element is initialized to 0. It is then set to 1 if on-the-fly rule generation is desired.

6.5.4 Control flow and functions - reactions

Reactions are initially added to a simulation when the `simreadstring` function encounters a “reaction” or “reaction_rule” word in a configuration file. If that happens, the remainder of the line of input file text gets sent off to `rxnparsereaction` for processing. This function takes care of parsing that is usually done within `simreadstring` but got so complicated that it got separated into its own function. That function reads the reaction name, the names of any compartment or surface that the reaction is restricted to, the reactant names and states, and the product names and states. As it reads the reactants and products, these get appended to `pattern` using the `molstring2pattern` function. This then calls `RxnAddReactionPattern` with the reaction name, the pattern, the reactant and product states, any compartment or surface restrictions, and whether this is a rule. Note that the reaction rate is not read or used so far.

`RxnAddReactionPattern` can be called either when a reaction is first created or during reaction expansion, so it first calls `molpatternindex` without updating to find the current pattern status. If that function does not return an `index` variable, which we’ll assume in this case, that means the pattern is new and an `index` variable needs to be created. The function then calls `molpatternindex` a second time, now with updating, which creates the index variable. The `index` variable includes a list all of the species combinations, given the current species list, that agree with the given pattern. If the reaction does not include wildcards, `index` will list only a single reaction; vice versa, if it does include wildcards, then `index` may list multiple reactions. If the reaction is a rule, then `molpatternindex` will create any species that the reaction rule generates and will add those species to the simulation. Next, `RxnAddReactionPattern` loops over all of the new reactions in the index variable, which is all of them in this case, extracts the species identities, tests to see whether an identical reaction already exists, and calls `RxnAddReaction` to add the reaction if not. If the reaction does exist, then this increases the reaction multiplicity.

`RxnAddReaction` adds a new reaction to the simulation (it also updates an existing reaction if it was only entered with reactants previously). It takes in the reaction name, species, states, and compartment and surface restrictions, creates the appropriate reaction superstructure if needed, creates the requested reaction, adds that reaction to the superstructure, and downgrades the reaction and surface conditions to the `SLists` levels. This triggers reaction molecule list, rate, and product placement computations when updating is next performed. The surface list is here because products with non-zero unbinding radii can cross surfaces, so surface checking for this needs to be turned on.

6 Code design

After all of the reactions for the given pattern have been added, control returns to `rxnparsereaction`, which then reads the reaction rate, if it was entered. It sends this rate off to `RxnSetValue` if the pattern corresponded to only one reaction and to `RxnSetValuePattern` if the pattern corresponded to multiple reactions. In the latter case, `RxnSetValuePattern` simply calls `molpatternindex` without updating to get the index variable, then goes through the specific reactions in it and calls `RxnSetValue` for each of them. Importantly, `RxnSetValuePattern` checks to see if the reaction is being given a rate for the first time, setting the rate to the given value if so, and otherwise it adds the given value to the current reaction rate. This works because this function is only called during setup of multiple reactions or during expansion of rules. Finally, if the reaction is a rule, then `rxnparsereaction` calls `RuleAddRule` to store the rule for later use.

After this, Smoldyn continues on to read more lines from the input file, possibly repeating the same procedure if another reaction is entered.

6.5.5 Control and flow - network generation

Reaction rule expansion occurs through the `RuleExpandRules` function. This function is called if the user uses the “expand_rules” statement, which is parsed by `simreadstring`. It is also called at every time step in the main loop, which is in the `simulatetimestep` function.

The `RuleExpandRules` function first checks to see if expansion is necessary. If so, it goes through the rules list sequentially. For each rule, it calls `molpatternindex` without updating to convert the pattern to an index, it gets the data it needs from the other rule lists, and then it calls `RxnAddReactionPattern` with this information. That function updates the index variable for this pattern and adds any new reactions. After control returns to `RuleExpandRules`, it calls `RxnSetValuePattern` along with the same information and also the starting point for updating, and that function sets the reaction rates for the new patterns, adding the rate to the existing rate.

Within the molecule superstructure, `expand` is a flag for on-the-fly rule-based modeling. It is initialized to 0 and stays that way so long as there have never been any molecules of this species, it is increased to 1 if at least one molecule of this species has been created but it has not yet been used for rule expansion, and is set to either 2 or 3 if it has been used for expansion.

6.6 Units

Units were introduced in Smoldyn version 2.74. The code is surprisingly simple, and is located primarily in `strunits`, in the `string2.c` library. There are a few parsing tasks.

First, input occurs with the `strmathsscanf` function, which is my version of the standard library `sscanf` function. This function describes the units for the value that will be input with text that looks like `%mlg|L`. Here, the `m` character shows that this is a special input that needs to be read by the `strmathsscanf` parser, rather than being sent on to the `sscanf` parser. The subsequent `lg` characters are the standard formatting characters for inputting a double. The vertical bar (pipe) shows that unit information follows, and the `L` shows that the unit should be a length.

The second parsing task occurs if the user enters units with a number. Continuing with the above example, the user might enter `5.3|um`. The pipe again indicates that units follow. This value again needs to be parsed, and then checked that it is consistent with what was expected.

The third parsing task occurs when Smoldyn prints out the model value in the diagnostics. Here, Smoldyn uses `simLog` which reads `%g|L` and then outputs `5.3um`.

6.6.1 Unit formatting

Units start with a pipe character. After that, the code generally uses “L” for length, and “T” for time. Combine multiple units with either a “.” or a “/” character. If units have a power that is not equal to 1, enter it either directly after the unit name or with a carat first. Here are examples:

L.L	L2	L ²
/L	L-1	L ⁻¹
L.L/T	L2/T	L2.T-1
/L/T	L-1.T-1	L ⁻¹ /T

If a number is explicitly unitless, then use a pipe character with nothing following it. For example, the input format string “%mlg|” is used to input a unitless value. If the user includes units, then that is an error.

For output using `simLog`, unit formatting has two additional details. First, it’s allowed to enclose the units in parentheses; for example, “|(L2/T)”. This will display the units in parentheses if there are units and won’t display either parentheses or units if there aren’t units. Second, the units format is allowed to end with exactly one punctuation character, which is kept in the output regardless of whether units are displayed or not, where the punctuation character options are: “, . : ; + - * / < > = !)”.

6.6.2 Unit parsing control flow

Consider inputting a diffusion coefficient using the `difc` statement. This goes through the `difc` input section of `simreadstring`, in `smolsim.c`, which has the line

```
itct=strmathsscanf(line2,"%mlg|L2/T",varnames,varvalues,nvar,&flt1);
```

This line of code sends the user’s input, in `line2` off to `strmathsscanf` for parsing. That function recognizes the “%m” portion of the formatting string and then reads in the user’s expression for this input as a string (variable `expression`). If the user entered a pipe character, to indicate that units follow, then `strmathsscanf` removes it and processes the two halves of the expression string separately; the first half goes off to `strmatheval` for math parsing and evaluation and the second half gets read in as a string (`dimstr`) for unit parsing. Both the “L2/T” portion of the original format string and any units entered by the user get sent off to `strunits`, with option “convert”, for unit parsing.

In `strunits`, the `dimstring`, which is “L2/T” gets parsed first. To do so, it is send back into `strunits` but now with the “parse” option. That option clears the static `dimUnit` array, and then reads the units to find that the power is 2 for the first unit type (length) and -1 for the second unit type (time). It stores these in `dimUnit` so that the array is equal to `{2,-1}`. It also computes the conversion factor between these units and the working units and multiplies that by the user’s value. Returning to the `convert` level of `strunits`, the function copies over the `dimUnit` array to `tempunit` so it won’t be lost and then parses the user’s units, if there are any, in the same manner. Next, it checks unit compatibility and performs conversion. If the user didn’t enter units, then the function needs to know what units are being assumed for this input. To do that, it looks on the `unitStack` to get the latest unit set.

7 Smoldyn modifications

7.0.1 Modifications for version 1.5 (released 7/03)

- Added hierarchical configuration file name support.
- **Zeroreact** assigns the correct box for new molecules.
- The user can choose the level of detail for the bimolecular interactions (just local, nearest neighbor, all neighbor, including periodic, etc.)
- Bimolecular reactions were slow if most boxes are empty. Solution was to go down molecule list rather than box list.
- Absorbing wall probabilities were made correct to yield accurate absorption dynamics at walls.
- Cleaned up and got rid of old commands.
- The current time input was made useful.
- Graphics were improved by adding perspective and better user manipulation.
- Simulation pausing was made possible using graphics and improved without graphics.
- If a command was used with a wrong file name, the command string became corrupted during the final command call. This was fixed by Steve Lay.
- Fixed the neighbor list for bimolecular reactions between mobile and immobile reactants.
- Reactions were made possible around periodic boundaries.
- Molecules were lost sometimes. This bug was fixed: 4 lines before end of **mol**sort: was `while(!live[m])`, is now `while(!live[m]&&nl[11])`.
- Output files now allow the configuration file to be in a different folder as Smoldyn.
- Added an output file root parameter.
- Added the command **replacexyzmol**. Afterwards, the code for the command was sped up considerably.
- Sped up the command **excludebox**.
- Command time reports were fixed for type **b** and **a** commands.
- Added more types of command timing codes.
- Improved accuracy of **unireact** so that it correctly accounts for multiple reactions from one identity.

7 Smoldyn modifications

- Improved product parameter entry and calculation, as well as the output about reaction parameters.
- Added the routine `checkparams` to check that the simulation parameters are reasonable.

7.0.2 Modifications for version 1.51 (released 9/5/03)

- Fixed a minor bug in `doreact` which allowed the molecule superstructure indices to become illegal if not enough molecules were allocated.
- Fixed a minor bug in `cmdreact1` which did not check for errors from `doreact`.
- Added command `molpos`.
- Moved version number from a `printf` statement to a macro, in `smoldyn.c` file.
- Added command `listmols2`, from a file sent to me by Karen Lipkow.
- Fixed a minor bug in `checkparams` that printed warnings for unused reactions.
- In `simulatetimestep` in `smoldyn.c`, the order of operations was `diffuse`, `checkwalls`, and then `assignmolecs`. The latter two were swapped, which should make wall checking more accurate when time steps are used that are so long that rms step lengths are a large fraction of box sizes. The new version is less accurate than before when the simulation accuracy is less than 10, but should be more accurate when it is 10.
- Replaced the `coinrand` call in `unireact`, which determines if a reaction occurred, with `coinrand30` to allow better accuracy with low probabilities. Also changed the relevant check in `checkparams`.
- Improved reactive volume test in `checkparams`.
- Increased `RANDTABLEMAX` from 2047 to 4095.
- Some modifications were made to `random.h`.
- Fixed a major bug in `rxnfree`, regarding the freeing of the `table` elements.

7.0.3 Modifications for version 1.52 (released 10/24/03)

- Changed comments in `rxnparam.h` and `rxnparam.c`, but no changes in code.
- Changed `cmdsavesim` in `smolib2.c` to allow it to compile with `gcc`.
- Added another call to `assignmolecs` in `simulatetimestep` in `smoldyn.c`, after the call to `checkwalls`, to make sure that all molecules are assigned properly before checking reactions. This slows things down some, but should allow slightly longer time steps.
- To the `opengl2.c` file, the `KeyPush` function was modified so now pressing ‘Q’ sets the `GL2PauseState` to 2, to indicate that a program should quit. A few modifications were also made in `smoldyn.c` function `TimerFunction` to make use of this.
- Corrected two significant bugs in the `checkwalls` function in `smolib.c` regarding absorbing walls. First, it didn’t work properly for low side walls. Also, the probability equation was incorrect, which was noticed by Dan Gillespie.

- Fixed a minor bug in `cmdsavesim` in `smollib2.c` file, which caused an output line for `rate_internal` to be displayed for declared but unused reactions.
- Several commented out functions in `loadrxn` were removed because they were obsolete and have been replaced by `product_param`. They were: `p_gem`, `b_rel`, `b_abs`, `offset`, `fixed`, and `irrev`.
- A command superstructure was created, which moved several structure elements out of the simulation structure. No new functionality was created, but the code is cleaner now. New routines are `cmdssalloc` and `cmdssfree`. Updated routines are: `simalloc`, `simfree`, `loads simul`, `setupstructs`, `cmdoutput` (including function declaration), `openoutputfiles` (including function declaration and ending state if an error occurs), `commandpop` (including function declaration), `checkcommand`, `endsimulate`, `savesim`, `main`, and all commands that save data to files.
- Renamed the “test files” folder to “test_files”.

7.0.4 Modifications for version 1.53 (released 2/9/04)

- Cleaned up commands a little more by writing routine `getfptr` in `smollib2.c` and calling it from commands that save data, rather than repeating the code each time.
- All routines that dealt with the command framework were moved to their own library, called SimCommands. This also involved a few function name and argument changes, affecting `smoldyn.c`, `smollib.c`, `smollib.h`, `smollib2.c`, and `smollib2.h`.
- Formatting was cleaned up for structure output routines.
- Swapped drawing of box and molecules, so box is on top. Also increased default box line width to 2 point.
- Computer now beeps when simulation is complete.
- Modified SimCommand library so that each invocation of a command is counted and also changed declaration for `docommand` in `smollib2`. This change was useful for improving the command `listmols2` so it can be run with several independent time counters. Also, wrote command `listmols3`.
- Wrote the new configuration file statement `boxsize`.
- Wrote the new commands `excludesphere` and `includeecoli`.
- Wrote the commands `overwrite` and `incrementfile`, which also involved some changes to the SimCommand library and required the new configuration file statement `output_file_number`.
- Added a new configuration file statement `frame_thickness`.
- When simulation is paused using OpenGL, the simulation time at which it was paused is now displayed to the text window.

7 Smoldyn modifications

7.0.5 Modifications for version 1.54 (released 3/3/04)

- Swapped order of commands and OpenGL drawing so that commands are executed before displaying results. Also wrote section 3.2 of the documentation to discuss this ordering and other timing issues.
- Wrote documentation section 3.3 on surface effects on reaction rates and added the `reactW` set of test files.

7.0.6 Modifications for version 1.55 (released 8/20/04)

- Improved graphics manipulations and added ability to save image as a TIFF file. This is not documented yet.
- Made a few tiny changes in `random.c` and `string2.h` and `.c`.
- The configuration file statement `max_cmd` is now obsolete because the command queue is automatically created and expanded as needed. Also, lots of changes were made to the library file `SimCommand.c` so that there are now two command queues: one is as before and uses floating point times for command execution and the other uses an integer counter for commands that are supposed to happen every, or every *n*'th, iteration.
- Added error strings to commands as well as the macro statement `SCMDCHECK`.

7.0.7 Modifications for version 1.56 (released 1/14/05)

- Made lots of changes in `opengl2.c`.
- `#include` files for `gl.h` and `glut.h` now use brackets rather than quotes.
- Improved graphics significantly.
- Rewrote `TimerFunction` to clarify code.
- Added ability to save TIFF stacks which can be compiled into movies.
- Added `keypress` command.
- Added comments to the code.
- User and programmer parts of documentation were split to separate files.

7.0.8 Modifications for version 1.57 (released 2/17/05)

- Added command `setrateint`.

7.0.9 Modifications for version 1.58 (released 7/22/05)

- Fixed 2-D graphics so they a border is now shown again around the simulation volume.
- Added runtime commands `replacevolmol` and `volumesource`.
- Random number table for diffusion is now shuffled before use, which significantly reduces errors from an imperfect random number generator.
- Added position ranges to `mol` command.

7.0.10 Modifications for version 1.59 (released 8/26/05)

- Random number seed is now stored and is displayed before a simulation starts.

7.0.11 Modifications for version 1.60 (not released, but given to Karen 9/30/05)

- Fixed a small bug in `checkparams`.

7.0.12 Modifications for version 1.70 (released 5/17/06)

- Added reflective, absorbing, and transparent surfaces for 1 to 3 dimensions with panel shapes that can be: rectangle, triangle, and sphere.
- `Geometry.c` and its header `Geometry.h` are new libraries that are used.
- Added background and frame color options.
- Reformatted and significantly updated part 2 of the Smoldyn documentation. Added surface descriptions to part 1 of documentation.
- Changed molecule sorting in `molsort` so that list compacting maintains list order.
- Wrote `reassignmolecs` to replace `assignmolecs`, which should increase efficiency and allow accurate surface treatment.
- Made it possible to load molecule names individually rather than all at once. New configuration file statements are `max_name` and `name`.
- Added pointers to the live molecule lists called `top1` (and renamed `top` to `topd`), which will differentiate old molecules from the new “reborn” ones. This is important for treating surfaces after reactions.

7.0.13 Modifications for version 1.71 (released 12/8/06)

- Added `glutInit` call to `main` function in `smoldyn.c`.
- Changed OpenGL drawing slightly for surfaces, so now 3D surface colors are always the same on the front and back, but can also be semi-transparent, although with OpenGL errors.
- Added command `killmolinsphere`.
- Cleaned up simulation loading some, with minor modifications in `setupstructs`, `loadsimul`, and `setupboxes`, as well as writing of `setdiffusion`. This makes it so that molecule sorting only happens in `molsort`, and it took some unwanted code out of `loadsimul`.
- Added molecule serial numbers to the molecule structure and superstructure.
- Added some elements to command structures so that commands now have storage space.
- Added `RnSort.c` library to project, as well as some new functions in `RnSort.c`.
- Added command `meansqrdisp`.
- Completely rearranged order of functions in `smollib.c` and in documentation part II.

7 *Smoldyn modifications*

- Cleaned up surface code. Fixed rendering of 3-D spheres. Added support for cylinders and hemispheres.
- Added statements: `grid_thickness` and block comments with `/*` and `*/`.
- To `action_front` and similar statements, allowed “all” for molecule name.
- Tried to stop diffusing molecules from leaking across reflective surfaces.

7.0.14 Modifications for version 1.72 (released 2/26/07)

- Finally got reflective surfaces to stop leaking diffusing molecules. This involved many changes in the surface code sections.
- Walls are no longer functional when any surfaces are defined, so new surfaces have to be defined to serve as system boundaries. Also, periodic surfaces are now possible.
- Changed all float data types to doubles throughout `smollib.c`, `smollib2.c`, `Geometry.c`, `smoldyn.c`, and their headers.
- Made it so that bimolecular reactions across surfaces can only happen with transparent surfaces.
- Two dimensional graphics now allow panning and zooming.
- Updated `savesim` command to accomodate surface changes.

7.0.15 Modifications for version 1.73 (released 9/25/07)

- Trivial bug fixed in `loadsurface`, fixed minor bug regarding periodic surfaces.
- Fixed significant bug regarding 2D triangle surfaces.
- Improved `surfaceoutput` so that it only prints first 20 panels of each type for each surface.
- Added new commands: `settimestep`, `beep`, `echo`, `killmolprob`.
- Modified command `meansqrdisp` to read a dimension number. Also made it so it accounts for periodic boundaries.
- Added ‘x’ to possible command stepping options; then improved it 5/25/07.
- Added `wrap` element to molecules and implemented it.
- Lots of trivial changes so that Smoldyn would compile with gcc using “-Wall” flag without warnings.
- Smoldyn can now compile without OpenGL and/or without libtiff.
- Improvements to Makefiles and improvement of compiling advice.

7.0.16 Modifications for version 1.74 (released 10/22/07)

- Coincident surfaces have defined behavior.
- New statement `boundaries` replaces `low_wall` and `high_wall` (but old ones still work).
- New command: `warnescapee`.
- Disk shaped panels are now possible.
- Surface jumps are now possible, which allow for holes in surfaces as well as for better periodic boundaries.
- Significant bug fix in reflections at circular and spherical surfaces.
- If 50 surfaces are encountered by a molecule in one time step, a warning is now printed.
- Cleaned up a lot of surface code.
- Replaced the surface statements `action_front`, `action_back`, and `action_both` with `action`; `color_front`, `color_back`, and `color_both` with `color`; and `polygon_front`, `polygon_back`, and `polygon_both` with `polygon`.

7.0.17 Modifications for version 1.75 (released 11/6/07)

- Added surface-bound molecules. This has involved changes in the molecule superstructure, surface structures, panel structures, and other structures, changes in the graphics to display these states, new input statements (`surface_mol`, `epsilon`, `action`, `neighbor`, changes to `color`, `display_size`, `difc`, `surface_rate`, `surface_rate_internal`, reaction statements, and others), and a lot of new functions in the code. Nearly every function in `smolib.c` required at least some changes to account for this addition. Many additions have not been debugged yet.
- New library dependency: `Sphere.c` and `Sphere.h`.
- Rewrote the graphics code for drawing molecules and part of the code for drawing surfaces.
- Changed panel names, so now they are unique for each panel within a surface, not just each panel within a shape. Also, panel names are available from panels.
- Removed the `assignmolecs` function, which was superseded by `reassignmolecs`.
- Changed several character type structure elements to enums, including `PanelShape`, `PanelFace`, and `SrfAction`.
- Made it so the `read_file` statement can be used within surface or reaction statement blocks. This included adding a short Parse section and structure to the code.
- Renamed `test_files` folder to `examples`. Also, updated the configuration files so that they run with the current syntax.
- Improved error reporting, so now the problematic line of text is displayed.
- Cleaned up reaction loading by adding new functions that allocate reaction structure space for reactants and products.
- Starting adding a user manual section to the documentation.

7 Smoldyn modifications

7.0.18 Modifications for version 1.76 (released 11/7/07)

- Split `smollib.c` and `smollib.h` source code files into `smolload.c`, `smolrun.c`, and their headers. Also, moved some functions from `smoldyn.c` into the libraries.
- Renamed some functions.
- Added and implemented more command-line flags.
- Put most global variables into `sim`, but killed off the event counters.

7.0.19 Modifications for version 1.77 (released 11/18/07)

- Overall, few changes that affect users.
- Made it so surfaces outside of the simulation volume are checked, just like those within the volume. This fixes errors that can occur.
- Rewrote `checkwalls` so that it's more robust for molecules with long rms step lengths and so that its inputs mimic `checksurfaces`.
- Modified `line2nextbox` so that it accounts for boxes that are outside of the system volume.
- Made most files in `smolrun.c` so that they can operate on either complete molecule lists or only those that are in single boxes. This may help for parallelization.

7.0.20 Modifications for version 1.78 (released 11/29/07)

- Minor changes with graphics.
- Changed the `overwrite` and `incrementfile` commands so that they no longer run a second command.
- Added and implemented the `CMDcode` enumerated list in `SimCommand.c`. Also in `SimCommand.c`, added an `iter` parameter to `scmdexecute` and added the functions `scmdcmdtype` and `scmdnextcmdtime`; these may be useful for parallelization.
- Fixed a bug in surface actions.
- New functions: `readmolname`, `molsetdifc`, `molsetdifm`, `molsetdisplaysize`, `molsetcolor`. These clean up code that was elsewhere, allowed more “all” input parameters, and simplify the linking of Smoldyn to other applications. Also, the input formats of some statements (`difc`, `difm`, `surface_mol`, `display_size`, `color`, `permit`, `products`) and a lot of commands were changed to use `readmolname`, which improves them.
- Fixed bugs in lots of commands, which were created by the addition of surface-bound molecules in version 1.75.
- Added function `molcount`, which was used to simplify and improve several runtime commands.
- Rewrote `savesim` command, which involved new functions: `writesim`, `writewalls`, `writesurfaces`, `writereactions`, `writecommands`, and `writemols`.
- Added command timing option ‘j’.

- Added `sname` to surface structure and changed `jump` in panel structure to point to a panel.
- Added reaction reversible parameter 'X'.
- Fixed a major bug in `findreverserxn` that caused it to fail.
- Added checking to `checkrxnparams` for multiple bimolecular reactions with same reactants and for defaults used for the reversible parameter.
- The `permit` element in reactions was changed from size `MSMAX` to size `MSMAX1`, including also implementation in the `bireact` functions and reaction set up functions.
- Changed `addrxnstostruct` parameter list to make it easier to use and to support an improved reaction input format.
- Added command `molcountinbox`.

7.0.21 Modifications for version 1.79 (released 12/6/07)

- Fixed several bugs in `scmdexecute`. Commands no longer execute repeatedly after the simulation is finished.
- (Several of the following corrections were suggested by Kevin Neff; thanks!)
- Added newlines to the ends of header files to avoid compiler warnings.
- Fixed types of some stand-in functions for compiling without OpenGL.
- In `opengl2.c`, `WriteTIFF` function, changed `uint32` to `unsigned int`. Also, changed inclusion of "tiffio.h" to `<tiffio.h>`.
- Changed the Makefile so that it is easier to configure for no OpenGL and/or no Libtiff. Also, included linked library option (thanks to Upi) and an option for converting line terminating characters from Mac to Linux.
- Added basic compartments.
- Added commands `molcountincmpt` and `molcountonsurf`.
- Fixed a bug in the `molcount` function.
- The `savesim` command now works again, at least mostly.
- Replaced all `%lf` format specifiers in `printf` statements with `%g` and `%lf` in `scanf` to `%lg`.
- Changed the program license from "permission is granted for non-commercial use and modification of the code" to LGPL.

7 Smoldyn modifications

7.0.22 Modifications for version 1.80 (released 12/22/07)

- Changed the exiting code some so ‘Q’ quits the program.
- Commands that don’t work now print errors only once and aren’t repeated.
- Modified `setrates`, `setproducts`, and `setupsurfaces` so that the time step can be modified during the simulation. Also modified command `setimestep` to work properly.
- Command `savesim` now does not save the output file name in the list of `output_files`, which means that the saved result can be run without overwriting it.
- Implemented `rxn->permit` element in `setrates` and also fixed a bug in `setrates`. A slight change in `setrates` is that `rate2` was not overwritten if it already had a positive value; now it is overwritten.
- Added and implemented `Simsetrandseed` function.
- Wrote part of `checksurfaceparams` function.
- Finished basic compartment implementation. `setupcomparts` was written, as well as new functions `boxrandpos` and `compartadddbox`.
- Added basic molecule porting for supporting MOOSE, with a new “port” surface action. Currently, this is more of a hack than a proper addition. New functions include `molgetexport` and `molputimport`, as well as `mollistnum`, `moldummyporter`, and `getnextmol`. Also, there is an `nlist` element in molecule superstructures, the number of live lists was expanded from 2 to 3, and there is a `list` element in molecules. `molsort` required some changes, as did various other functions.

7.0.23 Modifications for version 1.81 (released 1/22/08)

- Changed developing environments from Macintosh Codewarrior to Macintosh gcc with XCode as an editor.
- Restructured code from `smolload.c`, `smolrun.c`, and `smolib2.c` libraries, to `smolmolec.c`, `smolsurface.c`, `smolboxes.c`, `smolgraphics.c` and others. All headers are now in `smoldyn.h`.
- Rearranged and tidied up all code.
- Rearranged `else if` order in `loadsim` function.
- Replaced the fixed number of live lists (2 up to version 1.79, 3 for version 1.80) with a variable number. Also, changed the molecule superstructure significantly to implement these multiple lists. Most molecule functions required changes, as did some reaction functions, and many others.
- Added molecule porting properly instead of as a hack, which involved new structures, new functions, and a new source code file.
- Added a lot of new molecule functions, as well as some surface and reaction functions.
- Added statements `gauss_table_size`, `molecule_lists`, and `mol_list`.

- Added and implemented `prob` and `tau` elements to the reaction structure.
- Fixed many minor bugs throughout the code.
- Cleaned up various minor issues in diagnostics output.
- Added simulation event counters to `sim`, which involved minor changes in several top-level functions including `checksurfaces`, `checkwalls`, `zeroreact`, `unireact`, and `bireact`.
- Added and implemented event counters to the simulation structure.
- Lots of work went into part I of the documentation, along with many new and improved example files.

7.0.24 Modifications for version 1.82 (released 2/28/08)

- Changed `readmolname` so now a name without a state implies `MSsoln` rather than `MSall`, as it was before.
- Fixed `bireact` and `unireact` so that now molecule identities are checked before reactions are performed to make sure that they didn't just get killed off in a prior reaction.
- Reactions were overhauled with all new configuration file commands, new data structures, a lot of new functions, new diagnostics output, changes in every function in `smolreact.c`, improved example configuration files, and updated text in the user manual. A bug was fixed in which reaction permissions did not always correspond to those that were entered. Another bug was fixed in which reaction products frequently escaped the simulation volume.

7.0.25 Modifications for version 1.83 (released 3/14/08)

- Fixed minor bug with allosteric reactions.
- Changed order 1 reaction probabilities so that multiple reaction pathways now simulate with the correct rates (example file is `unireactn.txt`).
- Added command `setrandseed`.
- Actual surface action rates are now displayed.
- Verified that surface sticking rates are correct.
- Fixed a minor bug in `line2nextbox` function.
- Added drift to molecular motions.

7.0.26 Modifications for version 1.84 (released 4/11/08)

- Fixed a very minor bug in `line2nextbox` function.
- Added `-o` command line option
- Changed command `meansqrdisp` so that it now outputs $\langle r^4 \rangle$ as well as $\langle r^2 \rangle$.
- Rewrote command for changing the time step, which also included some new set time step functions and changes in a few setup functions.

7 Smoldyn modifications

- Added an enumeration for Smoldyn structures (enum SmolStruct), as well as text input and output functions to support them.
- Made a minor change to `reassignmolecs`, so the function doesn't bother running if there is only 1 box.
- Added commands `killmoloutsidesystem`, `diagnostics`, `setgraphics`, and `setgraphic_iter`.
- Fixed a bug in `fixpt2panel` regarding cylinders with their fronts inside.
- Cross-compiled Smoldyn for Windows using mingw compiler.
- Added order 0 and order 1 compartment reactions.

7.0.27 Modifications for version 1.85 (released 6/3/08)

- Changed command execution timing for 'x' type commands so that they do not execute any more often than the simulation time step.
- Added command `molcountspace`.
- Fixed some minor bugs in `srfcalspace`.
- Replaced the use of the system-supplied random number generator with the SIMD-oriented Fast Mersenne Twister (SFMT).
- Fixed a bug in which molecules leaked out of the corners of sticky boxes.
- Fixed a bug that was created in version 1.84 in which reversible reactions led to leaky surfaces.
- Fixed a trivial bug in which the wall list wasn't freed upon program termination.
- Fixed minor bugs in circle and arc drawing code.
- Ran Valgrind on Smoldyn with several configuration files and found no memory errors in the Smoldyn code (although several are reported for OpenGL).
- Fixed a bug in `findreverserxn` for reactions with 3 or more products.
- Fixed a bug in `rxnsetproduct` regarding type `RPpagemax` parameters.
- Cleaned up set up functions some, with new documentation.
- Fixed a bug in zeroth order compartment reactions.
- Parser code was separated from molecules source file, and `smoldyn.h` header file, and put into `parse.h` and `parse.c`. Also, this code got a lot of improvements, including support for macro substitution. New statements include `define`, `define_global`, `undefine`, `ifdefine`, `ifundefine`, `else`, and `endif`.
- Renamed allosteric reactions to conformational spread reactions. This included changes in `smoldyn.h`, `smolreact.c`, and the documentation.

7.0.28 Modifications for version 1.86 (released 11/17/08)

- Small progress on implementing accurate adsorption algorithms.
- Added display size checking for molecules.
- Fixed a bug that arose if both `max_names` and `names` were used.
- Fixed a bug in `molcountspace` command.
- Fixed a bug in `define` statements
- Changed makefiles to not assume availability of SSE2 processors
- Changed `queue.c` source file so that the `voidcomp.c` source file is no longer needed; this should enable compiling on 64-bit computers without problems.
- Fixed a bug in which molecules that were created only through commands did not interact with surfaces.

7.0.29 Modifications for version 1.87 (released 12/7/08)

- Fixed a minor bug in line parsing.
- Vastly improved the `wrl2smol` utility program.

7.0.30 Modifications for version 1.88 (released 1/16/09)

- Added accurate adsorption, desorption, and partial transmission algorithms.
- Added commands `fixmolcount` and `fixmolcountonsurf`, `molcounthead`.
- Fixed minor bug regarding 256+ character config. files lines.
- Smoldyn now recognizes Macintosh format config. files.
- Added `SmolCrowd` utility program to the distribution.
- Added areas to surface outputs.
- Changed config. file molecule names from “names” to “species”.
- Simplified config. file for `start_reaction`, `start_surface`, `start_compartment`, and `start_port`.
- Added command timing options A, B, &, I, E, and N. Also, fixed a minor bug in integer command queue timing (these were changed from ints to long long ints).
- Added compartment definitions that logically combine other compartments.

7 Smoldyn modifications

7.0.31 Modifications for version 1.89 (released 2/11/09)

- Wrote documentation for “simulation settings” and started documentation for “network generation”.
- Added some code for libMoleculizer linking. This is very preliminary.
- Added surface-specific reactions, and also made compartment-specific reactions work for bimolecular reactions.
- Completed support for adsorption, desorption, and partial transmission. This required many bug fixes and lots of work on `SurfaceParam.c`. All of these functions were tested and seem to work well. Also, reasonably thorough quantitative tests were run for all algorithms and all appear excellent.
- Lots of work on code modularization. This involved some data structure modifications (moved the `nspecies`, `maxspecies`, and `spnames` elements from the `simstruct` to the molecule superstructure), lots of cleaning up in `smolmolec.c`, and a lot of changes in `loadsim`. One result is that many config. file statements that could only be entered once before can now be entered multiple times, and each time overwrites the last one. This work is needed for a command line user interface and for Moleculizer linking.
- Progress towards making the setup functions able to run more than once, to allow for mid-simulation updates.
- Removed `math2.h` dependency of `rxnparam.c` file, and also improved documentation for it.
- Removed the `confspread` flag from the reaction structure since it was redundant with `rparamt` being equal to `RPconfspread`.
- Modified the `molcountspace` command slightly so that its spatial domains no longer include the endpoints; i.e. to (low, high), from [low, high].
- Changed surface molecule lists so that surface-bound reaction products are now checked for their positions after reactions and surface-bound actions (e.g. desorption) is handled better and more efficiently.
- Added “bounce” reaction product placement option. Molecules can now have excluded volume.

7.0.32 Modifications for version 2.00 (released 2/17/09)

- Fixed a trivial bug so that the `time.start` statement now sets both the start time and the current time.
- Updated the version number to 2.00 to reflect the fact that all aspects of simulations with surfaces are now complete, as are a tremendous number of other improvements.

7.0.33 Modifications for version 2.01 (released 3/3/09)

- Fixed a bug with compartment volume calculations, and improved compartment setup. Renamed `compartaddbox` to `compartupdatebox`, and largely rewrote it.

- Fixed a bug in `SurfaceParam.c` regarding rate calculation for reversible transmission.
- Added command `setsurfcoeff`, which Zsuzsanna Sukosd wrote.
- Separated `loadsims`, `loadsurface`, `loadcompart`, and `loadport` into two functions each, where the new functions are `simreadstring`, `surfreadstring`, `compartreadstring`, and `portreadstring`. Also added new statements to make the use of statement blocks optional within configuration files.
- Added command `set`.
- Made it possible for panel neighbors to be on other surfaces. As part of this, wrote function `surfchangeneighbors`, along with a few other functions.
- Substantially revised and improved surface diffusion. Checked all 2-D and 3-D panel shapes for basic functionality.

7.0.34 Modifications for version 2.02 (released 5/5/09)

- Changed `molcount` function so that it now works for unsorted molecule lists as well as for sorted molecule lists. This will affect most molecule counting commands.
- Some code modularization so that most statements in `simreadstring` now work through functions rather than directly. More work is needed for `surfreadstring`, in particular. Also, decreased dependence of modules on each other.
- Added and implemented a `condition` element to each superstructure, which will allow for improved modularity. This will also allow changes during run-time, such as new species.
- Fixed bugs in which reactions between solution-phase and surface-bound molecules didn't always happen.
- Added command `molcountincmpt`.
- Added command `meansqrdisp2`.

7.0.35 Modifications for version 2.03 (released 5/22/09)

- Changed build system from a simple makefile to the GNU Autoconf/Automake system. This works for many systems but not all.
- Added parallel operation using pthreads. This seems to work well, but doesn't actually lead to much speed up.
- Added support for libMoleculizer, which includes a lot of new code.

7.0.36 Modifications for version 2.04 (released 6/27/09)

- Lots of improvements to the same changes that were made for version 2.03. Now, all of these features appear to be stable and the build system appears to work for most Mac or Linux systems. However, some documentation is still lacking.
- Added boundary absorption for effective unbounded diffusion.
- Improved `SurfaceParam.c` some to remove all dependencies and to make it run a little faster.

7 *Smoldyn modifications*

7.0.37 Modifications for version 2.05 (released 7/23/09)

- Improved Smoldyn build system so that it now works for most Mac and Linux platforms.
- Some updates to the documentation, particularly in the compiling section.
- The Smoldyn release now includes the libmolecularizer documentation.
- Changed Smoldyn license from LGPL to GPL because of realization that large parts of the code could easily be used in proprietary software, which would hurt development of the Smoldyn project.

7.0.38 Modifications for version 2.06 (released 11/6/09)

- Lots of code cleanup, including formatting changes, removal of unused variables, adding function declarations, etc.
- Subversion site cleanup, with removal of unimportant files.
- Added jump behaviors for peripheral surface-bound molecules.
- Improved conversions between strings and enumerated types for surfaces.
- Added a graphics time delay for simulations in pause mode or completed simulations.
- Added optional species conversion at surfaces.
- Fixed a bug in molecule existence checking.
- Tidied up execution time reporting.
- Fixed several minor bugs in parameter checking.
- Fixed a bug in which the maximum number of species was incremented twice for empty molecules.
- Got compilation for Windows to work again, using a hand-written makefile (no libmolecularizer support).
- New release shell script.

7.0.39 Modifications for version 2.07 (released 11/17/09)

- Trivial updates to `surface_mol` and `compartment_mol` (to support 0 molecules)
- Fixed a bug in `molexpandlist`.
- Fixed a bug in `setupmols`, regarding `fixedlist` and `diffusinglist`.
- Fixed a minor bug in `rxnsetproducts`.
- Fixed several bugs in `savesim` command and added `writemolecules` and `molpos2string`.
- Fixed a major bug in diffusion of surface-bound molecules between panels.
- Fixed a bug in `Geometry.c`, in finding closest point on ring.

- Fixed a bug in reaction product placement, which shouldn't affect function but should speed code up.
- Fixed a bug in execution time reporting, and added time output to diagnostics.
- Fixed a minor bug with `neighbor_dist` statement.

7.0.40 Modifications for version 2.08 (unofficially released 11/20/09)

- Fixed a bug for reactions between 2 surface-bound molecules, where the destination panel was sometimes wrong.
- Cleaned up release files some.

7.0.41 Modifications for version 2.09 (released 1/6/10)

- Nathan fixed “make dist” build function.
- Fixed a small bug in `wrl2smol`, in which it didn't print out the input file name.
- Fixed a bug in `unbounded_emitter` surfaces which caused Smoldyn crashes.
- Fixed a minor bug in `surfstring2dm` for “polygon none” statement.
- Added command `molcountincmpt2`.
- Added statement and data structure portions for surface shininess, surface edge stippling, and graphics lights.
- Spun off the graphics stuff from `simstruct` and put into new graphics superstructure.
- Substantial work on graphics, including new `opengl_better` graphics drawing option.
- Added function `readsurfacename`.
- Added command `movesurfacemol`. It had a bug during initial release, which was fixed for re-release (1/12/10).
- Fixed two minor bugs in `opengl2.c` that only applied to compiling without `opengl` support. Re-released 1/20/10.

7.0.42 Modifications for version 2.10 (released 3/24/10)

- Lots of work on `libmoleculizer`. Mostly rewrote `smolmoleculizer.c` file, and redid most of the Smoldyn moleculizer data structures. More work is needed though.
- Fixed a few minor bugs.
- Fixed a bug in which cylinders weren't always put into boxes.
- Moved text from the `INSTALL_NOTES` file into the manual and killed the file.

7.0.43 Modifications for version 2.11 (released 5/4/10)

- Nathan added `-lglut` flag to standard configure for building.

7 Smoldyn modifications

7.0.44 Modifications for version 2.12 (released 6/10/10)

- Changed panel input so that a panel can be input multiple times and the new data will overwrite the old.
- Added commands `replacecmptmol` and `molcountincmpts`.
- Changed `molchangeident` to allow it to kill molecules.
- Changed `readmolname` to allow it to read empty molecules, and also updated the code where this function was called. Now, many commands allow “empty” as a molecule species.

7.0.45 Modifications for version 2.13 (released 7/15/10)

- Fixed reaction output slightly, so that activation-limited reaction rate is displayed and binding radius is not displayed for order 1 reactions.
- Various bug fixes in `smolmoleculizer.c`.
- Changed all `molcount` commands to be able to handle `nspecies` changing during the simulation.
- Improved `molsetexist` function, and set existence for `lmzr` species.
- Tidied `libmoleculizer-dependent` commands `speciesstreamcounthead` and `speciesstreamcount`.
- Substantial `libmoleculizer` work. Nathan made generated species have complex-form names rather than mangled names and he improved Python error reporting. I improved diffusion coefficients for generated species, including several new Smoldyn statements. More work is needed though.

7.0.46 Modifications for version 2.14 (released 7/18/10)

- Fixed `#ifdef` portion of `queue.h` so that it will compile on Windows correctly.
- Added default product placement parameter for continuation reactions.

7.0.47 Modifications for version 2.15 (released 7/20/10)

- Fixed bug that I just introduced for version 2.14 with default product placement parameter for continuation reactions.
- Fixed bug that didn’t allow `libmoleculizer` to run.

7.0.48 Modifications for version 2.16 (released 9/24/10)

- Added statement `expand_network`.
- Lots of work on `libmoleculizer` documenting and examples.
- Fixed yet another bug in `rxnsetrate` regarding automatic setting of product placement.
- Fixed a bug in which desorbed molecules could escape being checked for surface crossings.

7.0.49 Modifications for version 2.17 (released 11/19/10)

- Added compile flag `-lGLU` to `configure.ac` file, which seems to be necessary for Ubuntu systems.
- In `SurfaceParam.c`, `surfaceprob` function, `SPArevTrans` algorithm section, replaced `exp(2*c1*c1)*erfcd(SQRT2*c1)` with `experfcD(SQRT2*c1)` to expand the input domain.

7.0.50 Modifications for version 2.18 (released 1/6/11)

- Fixed a small bug in `molchangeident` function. Before, panel data was retained for molecules that desorbed to `MSbsoln`, whereas it should not have been.

7.0.51 Modifications for version 2.19 (released 2/11/11)

- Converted programmer's documentation from Word to LaTeX.
- Changed build configuring so that Libmoleculizer is now disabled by default, but can be enabled using the flag `--enable-libmoleculizer`. Before, it was enabled by default, but could be disabled.
- Cleaned up and documented pthreads code that Nathan wrote in the `smoldyn.h`, `smolsim.c`, and `smolsurf.c` files. More work is still needed for other files.
- Cleaned up `simulatetimestep` function, for slight efficiency improvements and better code.
- Renamed unthreaded functions from Nathan's names, which included a “_unitary” portion of the name, back to what I had before.
- Moved code for putting surface-bound molecules back on their surfaces from `checksurfaces` to `diffuse`.
- Enabled surface collisions for surface-bound molecules. This involved adding the surface action data structure, redesigning surface rate inputs and outputs, and redesigning surface action and rate indices. It also involved lots of new functions in `smolsurface.c`.
- Wrote documentation section on the simulation algorithm sequence.
- Improved surface and surface superstructure allocation so that they can be called more than once and the `max_surfaces` statement is no longer required.
- Modularized code in `smolsurface.c`, so that all statements recognized by `srfreadstring` now call functions rather than assigning directly to data structures.
- Swapped positions of `PFnone` and `PFboth` in the enumeration definition (from 3 and 2, respectively, to 2 and 3, respectively). I don't think this will cause problems, but I'm not certain.
- Added color words for all color inputs.
- Edited `SimCommand.c` to enable dynamic allocation for output files.
- Added and enabled `append_file` statement.
- Added `ifincmpt` and `killmolincmpt` commands.
- Added “-define” option for command line.

7 Smoldyn modifications

7.0.52 Modifications for version 2.20 (released 3/4/11)

- Fixed bug which causes crashes if species change at surface was to “empty”.
- Fixed a major bug, which caused Smoldyn to crash if it was compiled without threading support (this included the Windows version).
- Threading is now disabled by default.
- Fixed a few define statement issues. Trailing whitespace after replacement text is no longer part of the replacement. Multiple replacements no longer occur in the same piece of text (e.g. the AB of ABC is replaced with DE, making DEC, and then the EC is replaced with FG to yield DFG). Finally, long keys now take priority over short keys when both are able to match the same text.
- Added `display_define` statement.

7.0.53 Modifications for version 2.21 (released 3/11/11)

- Fixed bugs in `Geometry.c` for nearest triangle column and nearest triangle points.
- Trivial change to `opengl2.c`, so that it now displays the file name as the window name.
- Fixed rare bug in which `line2nextbox` couldn't find a next box due to roundoff error.
- Enabled diffusion for surface-bound molecules from spherical or similar panels to neighboring panels that happens when a collision occurs (new code in `dosurfinteract`).
- Increased the number of panel neighbors that can be entered on one line from 32 to 128.
- Improved but did not fix problem in `checksurfaces` where a molecule interacts with surface A and when it is placed on the correct side of surface A it is inadvertently placed on the wrong side of surface B due to round-off error.
- Improved SmolCrowd user interface.

7.0.54 Modifications for version 2.22 (released 3/22/11)

- Fixed a trivial mistake in `opengl2.c` that I created in version 2.21 that arose when compiling without `opengl`.
- Worked on shared library compilation and hand-coded makefiles. Added `libsmoldyn.c` and `libsmoldyn.h` to the Smoldyn source and added the `libsmoldyn` directory to the Subversion site and distribution.

7.0.55 Modifications for version 2.23 (released 6/24/11)

- Wrote most of `libsmoldyn.h` and `libsmoldyn.c`.
- Improved modularization and dynamic memory allocation for compartments.
- Improved modularization and dynamic memory allocation for ports.
- Added commands `shufflemollist` and `shufflereactions`.

- Fixed some minor bugs regarding jump panel display and error checking.
- Improved robustness and conciseness of molecule parameter display.
- Added display of time and molecule counts to graphics window.
- Cleaned up conditional compiling switching in `opengl2.c` and `smolgraphics.c`, so that each function only has one opening line, which is the same for all configure options.
- Added default product placement type for bimolecular reactions with two products to `RPpagemaxw`.
- Added commands `ifflag`, `ifprob`, and `setflag`.
- Removed the need for `max_species` and `max_mol` statements by implementing automatic memory allocation for both.
- Added a `simupdate` call from `simulatetimestep`, which will take care of data structure changes that commands cause.
- Fixed a minor bug in `parse.c`, `Parse_ReadLine`, for define statements that have no replacement text.
- Renamed, edited, and wrote several functions so that each module now has a function for updates, and this updating function now calls separate functions for updating at the lists level and at the parameters level. `setupmols` became `molsupdate`, `molcalcpars` became `molsupdateparams`, `setupboxes` became `boxesupdate`, `setupcompart` became `compartupdate`, `setuprxns` became `rxnsupdate`, `rxnsetmollist` became `rxnsupdatelists`, `rxnsetimestep` became `rxnsupdateparams`, `setupsurfaces` became `surfupdate`, `surfsetimestep` became `surfupdateparams`, and `setupports` became `portsupdate`.
- Checked entire code to make sure that data structure condition elements are modified correctly upon system changes. With these changes, all system changes, performed either before or during simulations, should trigger the necessary updates and thus take care of themselves.
- Edited `Parse_ReadLine` so that it now checks for lines that are too long and returns errors. Comment lines are allowed to be any length.
- Improved `rxnsetproduct` so that it now sets `unbindrad` and `prdpos` to 0, rather than assuming that they already were 0. Also, it copes better with conformational spread and other reaction types.
- Added checks for order 0 reactions with surface-bound products and no defined surface. Bug reported by Christine Hoyer.
- Added command `listmolsmpt`.
- Fixed a bug in which reactions between surface-bound molecules and solution-phase molecules could result in products that went to the wrong side of the surface. This involved changes to the panel assignment of the reaction position, and the reaction position, in `doreact`.

7 Smoldyn modifications

- Added a margin to each panel, so that diffusing surface-bound molecules that need to get moved back to panels that they diffused off of (or onto neighboring panels) no longer get moved to the exact panel edge, but to distance margin inside of the panel edge. Adding this involved adding a margin element to the surface superstructure, adding a margin statement, adding the `surfsetmargin` function, and adding margin parameters to `movept2panel` and `movemol2closepanel`.

7.0.56 Modifications for version 2.24 (released 7/27/11)

- Modularized graphics input, so all assignments are now made in `smolgraphics` functions rather than directly in `simreadstring`.
- Added basic graphics support to `libsmoldyn`.
- Edited `unireact` so that it now tests for reaction probability first, and also it tests for compartments and surfaces independently. This version should be slightly faster than prior ones.
- `Libsmoldyn` now compiles as both a static and a dynamic library with the GNU autotools. This is mostly from additions to `Makefile.am` in the `source/Smoldyn` directory.
- Lots of edits to `configure.ac`. This included some cleaning up, adding `-enable-libsmoldyn` option, and removing search for pre-installed `Libmoleculizer`. I also added an option for `-enable-swig`, although that will probably need modification soon.
- Added and implemented condition element in graphics superstructure.
- Moved most of the contents of `smolsimulategl` function into new graphics updating functions.
- Moved top level OpenGL functions out of `smoldyn.c` and into `smolgraphics.c`.
- Removed function declarations from `smoldyn.h` and put them in `smoldynfuncs.h`. This hides them from `Libsmoldyn` users.
- Removed library dependencies from `smoldyn.h`.
- `sim-icmds` is now a `void*` instead of a `cmdssptr`. This requires type casting in the code, but was required to remove a `Simcommand.h` dependency from `smoldyn.h`.
- Removed `GLfloat` data types from the graphics superstructure to remove an OpenGL dependency from `smoldyn.h`.
- Fixed all OpenGL data types in `smolgraphics.c` so that they will work even if, for example, `GLdouble` differs from `double`. I have not fixed data types in `opengl2.c` yet.
- Status note: both `Smoldyn` and `Libsmoldyn` compile well with `AutoTools`. For Mac, this includes all configure options, including with `Libmoleculizer` enabled. However, `Libsmoldyn` doesn't seem to actually work when it calls `Libmoleculizer`. With `MinGW`, both `Smoldyn` and `Libsmoldyn` seem to compile well and install to the windows directory, although not with `Libmoleculizer`. The `Libsmoldyn` build with `MinGW` is suffixed with `.a`; I don't know if it's supposed to have a `.dll`, nor how to create such a thing.
- Added `molcountspecies` and `mollistsize` commands.

- Wrote `release.sh` shell script for releasing, including building pre-compiled versions for the Mac and windows directories. Also wrote `install.sh` shell script for Mac version.
- Added `simversionnumber` and `smolGetVersion` functions.
- Got SWIG working for Python, although more work is needed in AutoTools stuff for SWIG code.

7.0.57 Modifications for version 2.25 (released 9/26/11)

- In `configure.ac` file, line 278, swapped sequence of a couple of lines. For version 2.24, the 3 lines starting with `"# This is probably redundant"` were first and then the 2 lines starting with `OPENGL_CFLAGS`. I swapped this sequence. See e-mail 9/13 with Pascal Bochet regarding compiling on Ubuntu. Apparently, this change did not fix his problem, but instead he added `"LIBS=-lglut -lGLU"` to his `./configure` line and that worked.
- Added functions `surfexpandmaxspecies` and `rxnexpandmaxspecies`, and also edited `molenablemols` and `rxnssalloc`. This should fix bugs that arose when users added species without allocating memory using `max_species`.

7.0.58 Modifications for version 2.26 (released 3/2/12)

- Added some support for species name entry using enhanced wildcards. This included new matching functions in `string2.c`, a logic expansion function in `string2.c`, the function `molwildcardname`, additions to many parameter setting functions, and additions to `molcount`. Support is fairly limited so far, and needs to be added to lots of commands and reactions.
- Added `simfuncfree` function.
- Libmoleculizer is close to working fully, including returning correct reaction rates. Membrane-bound states still pose problems. Also, checking has still been minimal.
- Fixed bug that made it impossible for excluded volume reactions to occur across periodic boundaries. In the process, also slightly improved reaction location determination for reactions that occur across periodic boundaries.
- Fixed a minor bug with `define` statement in which recursive defines sometimes got overlooked.

7.0.59 Modifications for version 2.27 (released 7/26/12)

- Merging in VCell changes.
- VCell people are now accessing Smoldyn code from `hedgehog.fhcrc.org` server. They used my user name on 4/10/12 but should be using their own user names from 4/11/12 onwards.
- Changed Smoldyn to Cmake. This conversion still needs substantial work. It needs Libmoleculizer, it needs to compile SmolCrowd and `wrl2smol`, it needs to work with MinGW, and it needs to create statically linked binaries that will work elsewhere.
- Changed all for loops over enumerated types to include explicit type conversions.

7 Smoldyn modifications

- Major work on error handling. Moved CHECK and CHECKS from individual files to smoldynfuncs.h. Changed their definitions some and added CHECKMEM and CHECKBUG macros. The basic error handling design is changed. Now, functions typically catch errors using one of the CHECK macros. This writes an error string to the global variable ErrorString and sends control to the local failure label. At this point, functions clean up as needed and call simLog to report the error. simLog may throw an exception to higher up or may return. On the latter case, the function with the error passes an error code up to its calling function. The calling function usually catches this error with CHECK but does not report it with simLog because it has already been reported. As a result of this change, functions no longer pass strings with error messages, but instead messages always transmit via the global ErrorString function. The exception is for library functions, such as Parse.c.
- Replaced all printf functions in the core code with calls to simLog.
- Added triangle area function to geometry.c and replaced Ye Li's function with a more numerically stable one.
- Improved initial window placement in opengl2.c (Jim Schaff's change).
- Changed char* to const char* for C++ compatibility in: boxsetsize and lots of other functions.
- Added a **strict** parameter to the **panelside** function to fix bugs that arose from coincident panels that face in opposite directions.
- Changed molecular desorption from surfaces from a simple change in molecule state and position, to a killing of the original molecule and a creation of a new molecule with the correct new parameters. This causes the desorbed molecule to be checked for further surface crossings with the other reborn molecules, which prevents it from leaking out of the system in case it desorbs across a surface. Also, changed surface checking for reborn molecules so that all reborn molecules are checked and not just those that can undergo reactions.
- When existing molecules are replaced by new ones, which occurs now in desorption and also in bounce reactions and conformational spread reactions, they now keep their serial numbers and their posoffset vectors.
- Added support for molecular drift that is relative to the local panel orientation. This included the **surfdrift** part of the molecule superstructure and all of its setting, allocating, freeing, and functioning.
- For bounce type reactions, added a default behavior where the unbinding radius isn't fixed but is the binding radius plus the amount of overlap between the two molecules, when they were collided.
- Added a precision to the numerical output for commands.
- SCMDCHECK was changed to accept a format and message string, rather than just a static message.
- VCell group added lots of stuff to compartments. In particular, it can now use volume sampling to determine what fraction of each box is in which compartment, rather than testing random points. This is vastly faster. compartmentIdentifierpair name is compartment name and pixel is the compartment id. As part of this, VolumeSamples is struct in

smoldyn.h. num[3] is size of grid on x, y, z. size[3] is the width of each volume element on each dimension. origin[3] is the grid origin. volsamples is array with each unsigned char equal to the compartment ID for the center point of that volume element. ncmptidpair is number of possible compartments (maximum ID number + 1). This functionality is not available in stand-alone Smoldyn. getcompartmentid is new function. posincompartment is rewritten. Lots of new stuff is on top. Uses volumesamples method. compartupdatebox is rewritten, and now called compartupdatebox_volumesample. VCell group added zlib dependency to unpack volume samples data. Also added fromHex function and loadHighResVolumeSample function. They made lots of changes in compartupdateparams and added compartupdateparams_volumesample.

- Added some error catching stuff to main.
- Added a separate VCell main
- Added a little stuff to sim struct.
- In retrospect, this release was okay, but the build system was far from adequate. It did not cross-compile for Windows and the compiled version did not work on Mac OS 10.5.

7.0.60 Modifications for version 2.28 (released 8/28/12)

- Fixed a bug in `morebireact`, in `smolreact.c`, where molecule positions were moved to account for periodic boundary wrapping, but the `posoffset` element was not updated.
- Added `listmols4` command.
- Improved the build system. Now, zlib is not a standard dependency.
- `product_placement` can now be entered multiple times.
- Commands no longer output using `fprintf`, but now use `scmdfprintf`. This enables output using user-chosen precision. It will also be useful for output from simulations to `libsmoldyn`, but that hasn't been added yet. Because of this change, there is a whitespace change in essentially all text output files. As a result, the regression tests show that all files differ from version 2.27 to version 2.28. However, I went through the files and verified that there are no other changes.

7.0.61 Modifications for version 2.29 (released 4/10/13)

- Fixed a minor bug in `molismobile`, in `smolmol.c`, which caused bugs for some species with surface drifts.
- Surrounded most VCell code with conditional compiling statements, so it's no longer part of standard Smoldyn builds. The reason was primarily so `Libsmoldyn` could be C compatible.
- Various minor `Libsmoldyn` fixes (`smolAddReaction` was missing a check on product states, changed default throwing threshold and debug mode behaviors, and added call to `checksimparams` to `smolDisplaySim`).
- Improved `CMakeLists.txt` file for `Libsmoldyn` compiling.

7 *Smoldyn modifications*

- Fixed CMakeLists.txt file so that it now installs Python libraries for Libmoleculizer.
- Added several new colors to the color words that can be used, in `smolgraphics.c`.
- Added command called `executiontime`.
- Renamed `rand_seed` statement to `random_seed`, but left it backward compatible.
- Added extern "C" stuff to `libsmoldyn.h`.

7.0.62 Modifications for version 2.30 (released 8/21/13)

- Fixed a bug in `Geo_Cyl2Rect`, in the Geometry library. It was causing leaking surfaces.
- Added function `checksurfaces1mol`, for checking surfaces after a molecule hits a port. This also required minor changes in `molchangeident`, `molkill`, and `surfinteract`.
- Fixed a minor bug in `surfaddpanel` that didn't allow panels to be redefined.
- Changed the termination text because some users didn't understand how to quit properly.
- Added `meansqrdisp3` command.
- Added `ifchange` command.
- Added `residencetime` command.
- Changed license from GPL to LGPL to promote Smoldyn's use in other software. In particular, VCell can't use Smoldyn under the GPL license. Also, I think that NSF and other funding agencies prefer LGPL.

7.0.63 Modifications for version 2.31 (released 9/9/13)

- Added command called `setreactionratemolcount`.
- Added a line to `rxnsetproduct` that enables product setting even if the reverse reaction has rate 0.
- Modified command `meansqrdisp3` so that it weights diffusion coefficients based on molecule lifetimes and also so it works with all states of a species.

7.0.64 Modifications for version 2.32 (released 8/29/14)

- Overhaul of wildcard support. This included work on the wildcard match and substitute functions in `string2.c`, addition of pattern data structure components in the `sim- $\mu$ mols` superstructure, addition of some pattern handling functions including `molstring2pattern`, `molpatternindex`, `molstring2index1`, and memory allocation functions for patterns and pattern lists, a rewrite of the `molcount` function, and implementation of wildcards in `simreadstring`. This task is not complete yet; see below in the wish/ to do list.

- Added lattice support by merging nsv branch into trunk. The following functions are new and some of these still require documentation: smolcmd.c: cmdwriteVTK, cmdprintLattice, edits to cmdmolcount, cmdmolcountspace, cmdsaveSim; libsmoldyn.c: smolAddLattice, smolGetLatticeIndex, smolGetLatticeIndexNT, smolGetLatticeName, smolAddLatticeMolecules, smolAddLatticePort, smolAddLatticeSpecies, smolAddLatticeReaction; smolport.c: portgetmols2, portputmols2; all of smollattice.c; smolsurface.c: checkSurfaces1mol; smolgraphics.c: RenderMolecs.
- Documentation was partly updated but both doc1 and doc2 still require many updates.
- Substantial editing of CMakeLists.txt file and release scripts.
- Added BioNetGen to Smoldyn utility programs. They still need work, but they're included in the distribution now.
- Removed Libmoleculizer, pthreads code, and the last vestiges of the AutoTools build system. None of these things ever worked quite as intended, so they were removed to clean up the code.
- Improved "bounce" reactions so that they now enable simulations of cluster formation.
- Wrote and added SmoldynQuickGuide to the Smoldyn documentation collection.
- Fixed a minor bug that recorded periodic boundaries in boxes even when surfaces were defined. This led to surface leaking.
- Fixed a minor bug in which the "keypress" command would cause crashes if the simulation was run without OpenGL support.

7.0.65 Modifications for version 2.33 (released 10/9/14)

- Fixed a minor but important bug in changemolident.
- Made input names UK English compatible.
- Fixed typos in the QuickGuide.
- Updated all example files, so they no longer use reaction blocks and they don't use static memory allocation with max_ statements.
- Fixed a bug in which desorbed molecules could cross surfaces without detection.
- Added cmdtrackmol function.

7.0.66 Modifications for version 2.34 (released 1/8/15)

- Nearly complete BNG conversion code.
- Added lattice code to default (pre-compiled) distribution.
- Fixed several small bugs.
- Changed lattice code so that molecules from continuous space to lattice space only get placed into first subvolume. This leads to more accurate diffusion results.

7.0.67 Modifications for version 2.35 (released 4/23/15)

- Fixed substantial bugs in `rxnsetrate` and `rxnncalcrate`, in `smolreact.c`. The bug was that unimolecular reactions that had multiple reaction channels (e.g. $A \rightarrow B$ and $A \rightarrow C$) computed the conditional probability (the probability of reaction given that the prior reactions did not happen) for the second and subsequent reactions incorrectly. This was only a problem with relatively high reaction probabilities.

7.0.68 Modifications for version 2.36 (released 6/9/15)

- Fixed a bug in `doreact` in which bounce reactions did not work when both reactants had exactly the same locations.
- Changed the default graphic iteration value from 1 to 20.
- All output commands now flush the output buffer after the command is done.

7.0.69 Modifications for version 2.37 (released 10/7/15)

- Fixed a bug in `parse.c` in which global definitions weren't being passed to upstream files.
- Added a check at the `cmd` statement to make sure that simulation times have been entered.
- Fixed a bug in which surface-bound molecules diffused onto new panels incorrectly. This required quite a lot of work including many new functions in `Geometry.c` (all of the exit functions and several new triangle functions), a new function in `Sphere.c` (`Sph_RotateVectWithNormals3D`), expansion of the panel `points` elements to include edge normals, and lots of work in `smolsurface.c`. This work included: a complete rewrite of `movemol2closepanel`, additions to `surfaddpanel`, `movept2panel`, `closestpanelpt`, and the new `lineexitpanel` and `paneledgenormal` functions. I think all bugs are fixed but if not, see debugging advice in `movemol2closepanel`. One result of this fix is that surface-bound molecules that diffuse onto a panel that has a jump action no longer jump when that happens. That feature was added for Hugo's work, but I'm removing it at this point because I think there are better ways of accomplishing the same thing.
- Added `reaction_serialnum` statement and implemented it. This involved a small amount of code, mostly in `smolreact.c`.
- Added `reaction_log` and `reaction_log_off` statements and implemented them. This mostly involved a fairly small amount of code in `smolreact.c`. Also, I started a new library file called `List.c` for this. This addition, and the `reaction_serialnum` addition, should make it much easier to do particle tracking simulations.

7.0.70 Modifications for version 2.38 (released 10/22/15)

- Added line to `cmake` file to enable building for Macs with OS 10.5 or above.
- Changed `boundaries` statement to allow 'x', 'y' and 'z' inputs for dimension values. Same change to `low_wall`, `high_wall`, `panel rectangle`, and some commands.

- Added `updategraphics` command.
- Added `quit_at_end` statement.
- Improved `product_placement` so that it now checks product states and it allows the user to specify products with “`product_n`”.
- Fixed minor bug in `parse.c` that required all lines to end with `n`, which didn’t work if there was an end of file within the line.
- Fixed minor bug in `simParseError` that didn’t return an error if a file wasn’t found.
- Changed `sim-lcmds` from a `void*` to a `cmdssptr`. I have no idea why it was a `void*` before.
- Improved reaction logging output by changing from `fprintf` to `scmdfprintf`.
- Edited `scmdfprintf` to allow for `NULL cmds` input.
- Edited `closestsurfacept` to allow for a box input.
- Fixed bug in `doreact` which occurred when the reactants were bound to two different panels, and the reaction occurred in yet a third panel. Also, substantially cleaned up `doreact`, `RxnSetRevparam`, and `rxnsetproduct` to make code clearer, more consistent, and generally better, and to improve documentation. Now, `confspread` and bounce reactions can have any number of products. Also, bounce reactions now measure the vector from the reaction position rather than from the center between the two reactants (typically the same, but not for curved surfaces). I suspect that bounce reactions are now as good as they can be made, within the limitations of the basic algorithm (e.g. only one bounce per time step).
- Fixed `doreact` so that reaction products that are surface-bound get moved to the surface right away.
- Changed `parse.c` so that Smoldyn can read Mac and, hopefully, Windows files without them needing to be converted to Linux format first.
- Added `systemcenter` function to `smolwalls.c`.
- Added `expandsystem` command and `surftransformpanel`, as first steps towards having moving surfaces. They work ok (at least the parts that I tested), but the lack of panel distortion with anisotropic expansion means that molecules can escape across surfaces in some cases.
- Added `fixmolcountrange`, `fixmolcountrangeonsurf`, and `fixmolcountrangeincmpt` commands.
- Added `intersurface` reactions, including the `reaction_intersurface` statement, the `RxnSetIntersurfaceRules` function, and implementation in `morebireact` and `doreact`.
- Fixed a minor bug in `smollattice.c` which arose when surface panel definitions were changed in version 2.37.
- Fixed bugs in `doreact` that caused surface-bound molecules in bounce type reactions to jump between surface panels. This was caused by incorrect setting of the new panel value.

7.0.71 Modifications for version 2.39 (released 1/15/16)

- Fixed a bug in `Geo_NearestTrianglePt2`, which led to incorrect answers for points that were outside of the triangle. This fixed a bug that was detected in `S7_surfaces/surfacediffuse/simple3a.txt`.
- Added `cmdmolcountspaceradial` function.
- Made several very minor modifications requested by VCell. Mostly, changed `#if OPTION_VCELL` to `#ifdef OPTION_VCELL`, but also changed use of `NAN` and `INFINITY`.
- Added a check to `surfupdatelists` so that the `SMLdiffuse` flag is not set if all molecule-surface actions are “no” or “transmit” and there are no other details of interest. In other words, this prevents surface checking if there are no surface actions to check for.
- Modified `reassignmolecs`, from `smolboxes.c`, to make it much faster. Before, if a molecule needed to move boxes, this searched through the current box’s molecule list for the molecule, which it then removed. Now, the same routine is used if only reborn molecules are assigned but otherwise this clears all box molecule lists and assigns them from scratch. This is vastly faster if boxes have lots of molecules in them.
- Fixed minor bugs in `meansqrdisp`, `meansqrdisp2`, and `meansqrdisp3` which did not allow computations for all axes.

7.0.72 Modifications for version 2.40 (released 3/22/16)

- This is somewhat of a beta release, since a lot of things have changed and not all are documented or fully tested.
- BioNetGen support now works.
- Improved install for Mac slightly and Windows substantially.
- Added support for wildcards and formulas in essentially all statements and commands.
- Added “csv” output format and implemented it.
- Added `^` operator to `strmatheval` function.
- Edited `molpatternindex` and added `moladdspeciesgroup` function. Also implemented this with the new `species_group` statement.
- Added species groups.
- Added wildcard reaction expansion, including several new parts to the data structures.
- Removed `readmolname` and `molwildcardname` functions.
- Added `molgeneratespecies` function.
- Modified `graphicsenablegraphics` trivially so that graphics are not enabled by default.
- Improved graphics so surfaces in 3D now allow combinations of polygon values.

- Added `molcountspacepolarangle` command.
- Fixed several bugs in surface-bound molecule diffusion.
- Started updating the user's manual, but there's much more to be done.

7.0.73 Modifications for version 2.41 (released 4/8/16)

- Fixed a trivial bug that caused crashes.
- Added a "Getting Started" chapter to the User's Manual and reformatted the whole manual.
- Added molecule-surface interactions for new species to the BioNetGen code.
- Fixed a bug in `molkill` calls in commands which prevented molecules from being killed.
- Added an improved bounce reaction algorithm, for ballistic reflection.
- Improved the Windows installation script. It might not work yet, but it's probably closer.

7.0.74 Modifications for version 2.42 (released 4/29/16)

- Added command for radial distribution function, also including a `boxscan` function.
- Fixed bugs in excluded volume reactions, so they now work with periodic boundaries and have been thoroughly validated.
- Modified all `listmols` commands so that they now output molecule serial numbers. This causes all of the regression tests to report a different behavior, but that's because of the expanded output, not different simulation results.

7.0.75 Modifications for version 2.43 (released 5/18/16)

- Renamed `Smoldyn_doc1` to `SmoldynUsersManual`, and split `Smoldyn_doc2` to `SmoldynCodeDoc` and `LibsmoldynManual`.
- Fixed excessive include files for `Libsmoldyn`. Also fixed some bugs in `Libsmoldyn`. It now works for the test files.
- Added `gaussiansource` and `molcountspace2d` commands.

7.0.76 Modifications for version 2.44 (released 6/27/16)

- Added reversible reaction definitions. In the process, moved reaction parsing to the new function `rxnparsereaction` and added function `molreversepattern`.
- Fixed `release.sh` slightly so that `freeglut.dll` is included in the windows release.

7 *Smoldyn modifications*

7.0.77 Modifications for version 2.45 (released 7/15/16)

- Fixed bugs in `equilmol` and `modulatemol` where they returned the wrong probabilities. These were caused by the change to using `molscan` in version 2.40.
- Fixed minor bug in `molreversepattern` that causing compiling errors on Red Hat Linux.
- Renumbered User Manual chapters to re-align them with the example file directories.
- Fixed a bug in `strmatheval` that read 2^{-3} as a binary - sign when it should be a unary - sign.
- Fixed a bug in `nsv_print` in which it could only print to a fixed size buffer. It now allocates the buffer as needed.
- Added some checks to the lattice code to ensure that lattices have an integer number of compartments on each axis, that ports are at compartment boundaries, and that there is at least one compartment on either side of the port.

7.0.78 Modifications for version 2.46 (released 7/30/16)

- Added section `S94_archive` to the examples files. It currently has files for my 2016 Bioinformatics paper.
- Added `multiplicity` element to the reaction structure and implemented it. As part of this, added a “`reaction.multiplicity`” statement to the input format.
- Major overhaul of `molpatternindex`. As part of this, it now has 2 levels of updating.
- Overhaul of `RxnTestExist`.
- On-the-fly wildcard expansion is now truly on-the-fly so that it expands around individual species as they are populated rather than expanding one entire level at a time. Added several files to the regression tests that all agree with the BioNetGen expansion, making them validated models.

7.0.79 Modifications for version 2.47 (released 8/30/16)

- Added a new rules section to the code. This includes a rules portion of the data structure, the `smolrules.c` file, and functions in that file. In the process, I removed rules from the reaction superstructure. I also replaced the `RxnAddRule` and `RxnExpandRules` functions with `RuleAddRule` and `RuleExpandRules` functions. This is a cleaner design and should accomodate diffusion and other rule types.
- Added rules for molecule properties. These include for diffusion coefficients, diffusion matrices, drifts, surface drifts, molecule lists, display sizes, colors, surface actions, surface rates, and internal surface rates.
- Added a reaction template to the rule structure so now reaction modification statements, such as `product_placement`, can write to this template. The relevent portions of the template are then copied over when the rule is expanded.

- Expanded `readrxnname` so that it now returns reactions with the given name, rules with the given name, or reactions that start with the given name. Also, implemented this so that reactions entered with wildcards, whether rules or not, can be modified using reaction modification statements.
- Added `molcountspecieslist` command.
- Made the `text_display` statement functional with rules.
- Fixed bugs in `strEnhWildcardMatchAndSub` so that it now works with empty pattern or destination strings and also does one to many, many to one, zero to many, many to zero, many to many, etc.
- Updated a lot of commands so that if they request a species name and it doesn't exist, but there are reaction rules, then this isn't an error. Instead, an implicit assumption is made that the species will exist later on.
- Fixed bugs in reaction logging.
- Fixed a compiling bug in `smolmolec.c` in which `const char*` was being converted to `char*` in a `strchr` function.

7.0.80 Modifications for version 2.48 (released 11/16/16)

- Fixed a minor bug in `strmathsscanf`.
- Added several example files in archive directory for Bioinformatics and MMB papers.
- Fixed bug in `latticeaddmols`.
- Added requirement that 'dim' has to be entered before reactions because not doing this causes crashes. It's likely required before other things, but I didn't modify others.

7.0.81 Modifications for version 2.49 (released 2/17/17)

- Fixed bugs in `molcount` and `molcountspace` commands that arose if lattice species didn't align with Smoldyn species.
- Made it an error to call the `reaction_log` statement multiple times with the same reaction and different filenames. Previously, this was allowed and new filenames overwrote the old ones, but this behavior was confusing because it seemed as though there should be multiple output files.
- Wrote `SmolEmulate`, including 2D emulation code. This is not integrated into Smoldyn yet and is not ready for general release.
- Added `cmdtranslatecmpt` and `cmddiffusecmpt` commands. In the process, wrote several new functions for surface, panel, and compartment moving. Also, surfaces now have a list of their membrane-bound molecules.
- Added function parsing to `string2.c`, so Smoldyn now recognizes `sqrt`, `sin`, `cosh`, etc.

7 *Smoldyn modifications*

7.0.82 Modifications for version 2.50 (released 2/27/17)

- Added infrastructure for Smoldyn functions. This involved the command `cmdevaluate`, the functions `fnmolcount`, `molscanfn`, `loadsmolfuctions`, and others. Also renamed `molscan` to `molscancmd`.
- Added `oldpoint` and `oldfront` elements to surface panels. Now, when a panel is moved, the old location (`oldpos`) is stored in these old elements and the new ones are updated. Then, molecule adjustment is determined using the correct combination of old and new locations. Molecules still cross surfaces incorrectly if they get squeezed between a moving surface and a static surface (or two moving surfaces), but this is unavoidable. Several functions now include an `oldpos` input to see if the function should use the old or new panel positions.
- Added `if` command.
- Added `molcountonsurf` function.
- Added a `moleculetouch` element that records each time the molecules were touched. It lets functions not recompute if the touch value is the same as it was during a prior call.

7.0.83 Modifications for version 2.51 (released 3/15/17)

- Fixed a bug in which surface-bound molecules could leak through another surface. The problem was that the surface-bound molecule code was designed with relatively large panels in mind, and is generally exact for those panels, when they are planar. However, small panels are more complicated. As a result, I added the `pnlx` element to the molecule structure to keep track of the prior panel. Also, I changed the reflection algorithm in `dosurfinteract` so now the back-up plan in the case that reflection does not place the molecule on the correct side of the reflecting panel is to put the molecule back where it started from. The prior back-up plan was to fix the molecule to the correct side of the reflecting surface, but this made the system extremely sensitive to any future molecule adjustments. I also added the `fixpt2panelnocross` function which fixes a point to a panel but doesn't cross any other panels in the process. The final result works well, but it is still conceivable to have molecule leaks in rare situations.... See Modifications for version 2.74.

7.0.84 Modifications for version 2.52 (released 5/16/17)

- Martin Robinson fixed a minor bug in `smolcmd.c` that had to do with the lattice code.

7.0.85 Modifications for version 2.53 (released 5/25/17)

- Fixed minor bugs in `smolcmd.c` that had to do with the lattice code. Part of the version 2.49 modifications was that I “fixed bugs in `molcount` and `molcountspace` commands that arose if lattice species didn't align with Smoldyn species.” However, I now think the code was correct originally (at version 2.48) and that these changes for version 2.49 were incorrect. Martin effectively reverted one of them in `molcountspace` for version 2.52. I just finished that reversion and also reverted the one in `molcount`. Now, that portion of the code is identical to how it was in 2.48 and it seems to run well. I still don't know what prompted me to make these changes in the first place.

- Documented some of the `nsvc.cpp` functions.
- In `strloadmathfunctions`, which is in `string2.c`, the code now assigns function pointers to variables before sending them off to `strevalfunction`. This provides the context that C++ compilers need to determine which address to choose for the functions because their names are overloaded.

7.0.86 Modifications for version 2.54 (released 8/27/17)

- Added a line to `molpatternindex` that sets the `PDrule` element of the `index` variable if update is requested, `index` already existed, and the element wasn't already set. The only thing this does is to prevent a bug warning in the output.
- Fixed two small bugs in `lineXpanel` for sphere and cylinder crossing checking. These failed if the two points were the same because then the geometry function returned NaN and then that failed the test. The solution was to move to negative logic so that NaN would have the correct result. These failure arose because I changed to `fixpt2panelnocross` and that checks for lines crossing panels, often with the same beginning and ending points. Added `surfacedrift2` and `3` to the regression files because these were the ones that failed.

7.0.87 Modifications for version 2.55 (released 7/16/18)

- Added a 'time' option to the `random_seed` statement for setting the value to the current time.
- Added math functionality for command timing.
- Trivial improvement to wall output dimensions from numbers to letters.
- Wrote `bindingradiusprob` and `numrxnrateprob` functions in `rxnparam.c`.
- The MinGW build system that I've been using for years (`/opt/local/bin/i386-mingw32-gcc`) quit working because I did a Mac OS upgrade last summer which somehow killed it (the upgrade was after Smoldyn 2.54 was released). After lots of effort, I found a MinGW version that was precompiled for Mac (an older OS, but that doesn't seem to matter), which I installed. It seems to work well. It's at `/usr/local/gcc-4.8.0-qt-4.8.4-for-mingw32/win32-gcc/bin/i586-mingw32-gcc`. I updated the `Toochain-mingw32` file for this change.
- The Windows install wasn't working because the BioNetGen directory was called `bin` rather than `BioNetGen`. I think that's fixed by now.
- Added several lines of code to the `adsorb` section of `dosurfinteract` so that surface-bound molecules that adsorb onto a new surface end up on the correct side of the original surface, as well as on the correct side of the new surface. In particular, molecules starting in the up or down states end up equally distributed above and below the original surface.
- Largely rewrote `molsetlistlookup` so that it now sets the table for both `MSsoln` and `MSbsoln`. This change also required expanding the `listlookup` element slightly in the memory allocation function to allow for the `MSbsoln` option. The `listlookup` values are always set to the same lists for both solution phase options. As another part of this change, `molismobile` now allows for a `MSbsoln` input.

7.0.88 Modifications for version 2.56 (released 9/18/18)

- Fixed a bug that arose when the system was expanded and then some molecules were killed. System expansion triggered `boxupdateparams` which then removed molecules from boxes but didn't clear the `mptr_box` value and then `boxremovemol` crashed due to the molecule not being in the box. Fix involved allowing for molecule not in box in `boxremovemol` and not adding dead molecules to boxes in `boxupdateparams`. In the process, I discovered that molecule assignment may be done more often than necessary.
- Removed a call to `molsort` from `cmdgaussiansource`. I'm fairly certain it was unnecessary, and tests showed that it still worked without it.
- Lots of struggles with MinGW. The version at `/usr/local/gcc-4.8.0-qt-4.8.4-for-mingw32/win32-gcc/bin/i586-mingw32-gcc` didn't work. I got it with MacPorts but that didn't work either. Yet more versions also didn't work. I finally killed off all of them and tried fresh. Solved MinGW compiling but not with LibTiff, which remains to be solved. See the MinGW cross-compiling documentation section.
- The MinGW compiler is more strict and found several minor bugs. Fixed some of them, but there are still possible string overwrite and truncation issues that it complains about.
- Implemented most of lambda-rho algorithm. As part of this, I did a lot of work on the `SmolEmulate` utility, computed reaction rate lookup tables for non-unit reaction probabilities, debugged the `numrxnrateprob` and `bindingradiusprob` functions, and then updated all calls to `numrxnrate` and `bindingradius` to these probabilistic functions. These updated functions call the original functions in the usual case that the reaction probability is 1. The functions in `rxnparam.c` are generally good now, but still not quite perfect.
- Another part of lambda-rho, added the input file statement "reaction_chi" and implemented it. This addition seems to work reasonably well, but Smoldyn isn't always able to align the simulated and requested reaction rates yet. In particular, it has trouble with the `pgem`, `pgemmax`, and `pgemmaxw` reverse parameters.

7.0.89 Modifications for version 2.57 (released 3/21/19)

- Moved Smoldyn source from the Fred Hutchinson Center's "hedgehog" subversion server to github, where it is called "Smoldyn_Official".
- Fixed minor bugs in `cmdwriteVTK` that caused compiling errors when compiling with `OPTION_VTK` turned on.
- Added support for scientific notation in string math parsing.
- Added option for species change upon surface action. Example is `s7_surfaces/ReflectToNewSpecies.txt`.
- Major change with serial numbers to enable single molecule tracking where two molecules can bind together and retain both serial numbers and then dissociate and have their original serial numbers. Changed molecule serial numbers to `unsigned long long`, but left most other serial number values as `long int`.

7.0.90 Modifications for version 2.58 (released 3/29/19)

- Fixed minor bug in `string2.c` that caused loss of numerical precision when using `strmathsscanf`.
- Changed the serial number coding for single-molecule tracking some, which also had some effects on the user's input format.
- Fixed important bug in `molscan cmd` which caused it to scan over some molecules twice in some cases. For example, if called from `cmdkillmolprob`, this would test molecules twice to see if they should be killed, leading to incorrect kill probabilities. The problem was that some commands change the `sortl` elements, making them bad things to assume aren't changing.
- Added `replacemol` command.
- Re-released 4/2/19 with following changes:
 - In `CMakeLists.txt`, I changed the minimum deployment for Mac from OS 10.5 to 10.9. This removed a compiling warning about `libstdc++` being discontinued.
 - Copied `libstdc++-6.dll` from `/opt/local/x86_64-w64-mingw32` to `Smoldyn-official/windows`. Also, edited the `windows install.bat` file so that this dll is copied when Smoldyn is installed. This will hopefully fix an error that Kevin Chen told me about.

7.0.91 Modifications for version 2.59 (released 5/15/19)

- Trivial changes to the code to remove compiler errors and warnings.
- Major effort to compile for Windows. I failed at getting cross-compiling working and actually made it worse by replacing a cross-compiler that worked for non-graphics Smoldyn to one that didn't work at all. Instead, I now compile by hand on a Windows computer. This still doesn't support LibTiff or NSV but at least it does work.

7.0.92 Modifications for version 2.60 (released 6/7/19)

- Further improvements on Windows support. This release now builds `SmolCrowd` and `wr12smol` on a Windows computer. It also uses static linking so that dll files aren't needed. I also improved the Windows install script.
- Modified the `BNG2_path` statement so that it now reads the entire line of text after the statement word rather than just the first word in the line, thus allowing for spaces in file names.
- Fixed a trivial bug for the "stipple" statement, which was trying to read hex code with the math parser.

7.0.93 Modifications for version 2.61 (released 7/22/19)

- Rewrote `boxscansphere`, which is called by the commands `radialdistribution` and `radialdistribution2`. This function only worked for systems that had their low corner of space at (0, 0, 0) and that used periodic boundaries. Now, it works for all systems.
- Fixed a very minor bug in `Geo_CylisXaabb3`, in which the code didn't check for a zero denominator, although I don't think this actually affected results at all.
- When boxes determine what panels are in them, the edge boxes used to extend to $\pm\text{MAX_DBL}$. However, this caused problems in some geometry functions, so I defined a new global variable called `VERYLARGE`, which I set to 10^{20} , which is now effectively the edge of simulated space, at least for determining what panels are in what boxes.

7.0.94 Modifications for version 2.62 (released 10/14/20)

- Extended the `pointsource`, `volumesource`, and `gaussiansource` commands so that they now allow non-integer numbers of molecules to be requested. If a non-integer value is requested, these functions choose a Poisson-distributed random number of molecules.
- Windows version is now compiled with MSVC and includes support for hybrid simulation with NSV package and saving images with LibTiff.
- Added Python support, which was mostly Dilawar Singh's work. This includes the `pybind11` library and new Python code. The pre-compiled Mac version now comes with a Python wheel. The Python wheel can also be downloaded from PyPI with `pip`. This also included substantial editing of the Libsmoldyn documentation.
- Substantial CMakeLists.txt clean-up, along with updates to the build system documentation.
- Cross-compiling from Mac to Windows with MinGW works again, but not with graphics, so this was used for the `SmolCrowd` and `wrl2smol` utilities, but not for `Smoldyn`.
- Dilawar added several things that I haven't addressed yet: Travis continuous integration, Doxygen, readthedocs website compatibility, and several scripts.

7.0.95 Modifications for version 2.63 (released 12/15/20)

- Documented Dilawar's work on Python bindings in code documentation and Libsmoldyn manual.
- Dilawar added test for `SMOLDYN_NO_PROMPT` environment variable, which is essentially the same as setting the `sim->quitatend` flag.
- I made lots of minor fixes in the Python bindings, including fixing the Python example files. Dilawar made more fixes and also added many more Python example files.
- Dilawar added `smolSetMoleculeColor`, `smolSetMoleculeSize`, `smolSetIntersurfaceRules` to libsmoldyn.
- Lots of changes by Dilawar to scripts, including particularly for documentation. He also fixed some issues in the CMake build system.

- Dilawar added Python callback functionality. This includes a few minor changes in `smoldyn.h`, including new code in the `sim` structure, and a short bit of code in `simulatetimestep`, in `smolsim.c`; this code is in C++, so the file got renamed to `smolsim.cpp`.
- Dilawar made some minor changes in `libsmoldyn.c` `smolGetPanelIndex` and `smolGetPanelIndexNT` for searching for a panel in a panel list. As part of the same changes, he added a C++ array to `smoldyn.h`. These changed both files from C to C++. I don't understand these updates.
- Added data tables to the command superstructure, so that output data can be stored internally rather than only in external files. As part of this, wrote double-double list types in `Lists.c`, added several functions to support data arrays in `SimCommand.c`, edited all observation commands to work with the new function calls and also to enable data storage where possible, added statement "output_data" and command "printdata", edited reaction logging in `smolreact.c`, added Libsmoldyn functions `smolAddOutputData` and `smolGetOutputData`.
- Fixed bugs in `cmddiffusecmpt`.
- Added `maxmollist` element to surface structure because memory management here was sloppy, which led to errors.
- Fixed some bugs with `reassignmolecs` function. First of all, it's now called as the system is set up, before the simulation starts. Also, it now runs even if there is only 1 box, which is a little less efficient in rare cases, but more reliable.
- Tried to get `freeglut` to work, without success.
- Added Chapter 15 to user's manual on Python bindings.

7.0.96 Modifications for version 2.64 (released 3/25/21)

- Dilawar rewrote the documentation conversions, in the docs directory.
- Added `warnescapecmpt` command.
- Fixed a minor bug in which reversible reactions that took place on surfaces led to an initial unbinding separation that didn't address surface crossing, which led to products leaking across the surface.
- Converted the User's manual from Word to LaTeX and also renamed it from `SmoldynUsersManual` to `SmoldynManual`.
- Merged the Libsmoldyn user's manual into the regular user's manual.
- Had problems with static `libtiff` for Mac, so downloaded a new version and put the compiled static library in `Smoldyn/Mac`.
- Minor change in `CMakeLists.txt` so that the version number could be input from the command line.

7 *Smoldyn modifications*

7.0.97 Modifications for version 2.65 (released 5/18/21)

- Dilawar moved some functionality from `smoldyn.py` to `module.cpp`, in order to support loading simulations from files using Python.
- Fixed several bugs that arose when displaying diagnostics with partially set up structures.
- Dilawar added a `changelog.md` file to the package.
- Changed `SimCommand.c` and `.h`, so that commands are now added to a list of templates rather than directly to the command queues. Also added a condition element to the superstructure, telling whether commands have been added to the queues or not. This required a major overhaul of command running code.
- Moved a few lines of code about commands from `endsimulate` to `smolsimulate`.
- Dilawar made major changes to the build system, so that it's more modular now.
- Added long range forces, with `cmdlongrangeforce` and several additions to `List.c` and `List.h`.
- Made some minor changes to `SimCommand` so that its output now pays attention to simulation flags.
- Fixed some minor bugs in `molsetlistlookup` and `molupdatelists`, which set the wrong list number for empty molecules. This only affected partially set up data structures.
- Edited `molsort` so that it can cope with empty molecules in the resurrected list.
- Fixed `smolReadConfigString` so that it now works again, for both Libsmoldyn and the Python low-level interface.
- Added `smolSetSimFlags` and `smolRunCommand` functions to Libsmoldyn and also to the Python low-level interface.
- Fixed minor issues with `smolRunTimeStep` and commands, which were causing commands to not run after the first time step.

7.0.98 Modifications for version 2.66 (released 6/18/21)

- Substantial work on filaments. Some example files work, but several still don't.
- Dilawar fixed documentation formatting issues for the `readthedocs` site.
- Various additions for compatibility with Biosimulators.
- Minor additions to Python API: `Simulation.counts` and `getError`.
- Added "-s" command line flag for silent operation.
- Improved error reporting for Libsmoldyn `smolPrepareSimFromFile` and `smolLoadSimFromFile` functions.
- Added a test for Perl availability in `bngrunBNGL`, in `smolbng.c`.
- Added `smolGetSpecies` function to libsmoldyn. I haven't tested it yet.

7.0.99 Modifications for version 2.67 (released 11/2/21)

- Dilawar fixed a typo in the filament code.
- Fixed the `longrangeforce` command, where an index value could be off by one.
- Jonathan Karr and Dilawar updated the Biosimulators links.

7.0.100 Modifications for version 2.68 (released 7/18/22)

- Added `listmolssurf` command.

7.0.101 Modifications for version 2.69 (released 10/17/22)

- Fixed bug in which commands added mid-simulation were not added to the command queue.
- Fixed bug in the Python interface, in which commands with an '@' timing type weren't accepting an "on" time.

7.0.102 Modifications for version 2.70 (released 10/24/22)

- Fixed bug, or maybe just a lack of feature availability, in libsmoldyn and Python interface in which reaction probabilities could not be set for bimolecular reactions.

7.0.103 Modifications for version 2.71 (released 2/5/23)

- Fixed bug in which graphics were being called from Python API even when they were turned off. This arose from `gl2glutInit` being called in `Simulation.cpp` `Simulation::runSim` for all cases; it's now called in `smolgraphics.c` `smolsimulategl` instead.
- Made it an error to enter command timing with either non-integer values for integer type commands, or to enter timing steps that are less than or equal to zero.
- Improved command timing so that all commands that execute at a particular time are performed with integer commands first, then float commands (unchanged from before) and also, within these lists, in the order that they are listed in the configuration file (new functionality).
- Added `translatemol` command. This involved several other changes. Added `molmovemol` function and also called this function from `longrangeforce` rather than moving molecules locally (in general, very few functions, including especially those in the commands, should be allowed to move molecules by themselves). Added `checkwallsimol` for wall checking for a single molecule at a time. Moved `molchangeident` to a different region of the file. In `string2.c`, added a `const` declaration to both `strevalfunction` and `strmathevalint` and cleaned up the code slightly. Also, edited `strmatheval` so that the `expression` input isn't changed by this function and is now a `const`.

7.0.104 Modifications for version 2.72 (released 8/17/23)

- Made some minor modifications in `smolgraphics.c`. Tiff files were saving using the latest screen update, which occurs asynchronously, rather than at the requested time, which is now fixed. To do this, I added a `swapbuffers` input to `RenderSim` and also now draw the output without swapping buffers before recording tiff files. I also removed redundant system drawings by testing if it was just drawn, and then not bothering to redraw if unnecessary.
- Changed `molchangeident` to include a `crsspt` input parameter. This is used to rescale molecule motions in cases where the diffusion coefficient changes at a surface crossing.
- Improved `surfaceoutput` so that it displays more details about molecule-surface interactions, particularly with verbose output.
- Several minor clean-up things and bug fixes in `surfupdateparams`.
- Edited `dosurfinteract` so that molecule diffusion coefficients change at the moment that molecule identities change, which basically means that the correct parameters get passed to the updated version of `molchangeident`.
- Lots of changes to clean up error handling and command line input. In `smolsim.c`, rewrote `simSetLogging` and made some small edits in `simLog` and the top of `loadsim`. Added a `logfile` input to the non-VCell version of `simInitAndLoad`. Edited `smolSetLogging` in `libsmoldyn.cpp`. There's also a small edit in `python/module.cpp`, `init_and_run` to account for the change to `simInitAndLoad`.
- Added command line options with double-hyphens. Previously, the only support for this was with “-define”. Now, all options can be entered with a double-hyphen version. Also, options with single hyphens can be entered all at once or one at a time. This required lots of edits to `main` in `smoldyn.cpp`.
- Added command line input for logging file with `--logfile`.
- In `python/module.cpp`, removed the `overwrite` parameter from the `runSimUntil` python function because it had no purpose and was causing compiling warnings.
- In `NextSubVolume/NextSubvolumeMethod.cpp`, commented out the `beta` computation in the `ReactionList::recalculate_propensities` function because it wasn't being used and was causing compiling warnings.
- Wrote function `surfacetransmit` in `SurfaceParam.c`. This computes transmission probabilities from transmission coefficients or vice versa, while accounting for diffusion coefficient changes.
- Edited `surfupdateparams` to support species changing during surface interactions.
- Extended `srreverseaction` to return the action type.
- Edited several functions to allow for surface interaction rates of `-1`, which indicates the maximum possible rate. These functions included `surfreadstring`, `surfaceoutput`, `srcompareaction`, `writesurfaces`, and `surfsetrate`. Another function is `surfaceprob` in `SurfaceParam.c`.

- Overhauled `srfcalsrate` and `srfcalsprob` so that they allow for maximum possible rates as inputs and so that they account for surface actions when determining reversible rates and probabilities. They also account for species changes and diffusion coefficient changes at surfaces and use the new `surfacetranmit` function to compute reversible transmission parameters.
- Fixed a tiny bug in `surfreadstring` in which the “rate_rule” statement didn’t actually call the appropriate setting function.
- Changed `surfupdateparams` so that reverse probabilities are now recorded after all molecule probabilities are computed, to allow for species changes at surfaces.
- Increased standard string lengths. In `string2.c`, increased `STRCHAR` from 256 to 512 and `STRCHARLONG` from 4000 to 4096. Edited `parse.c` so that all strings are now 4096 characters long rather than 256 characters. This fixes bugs that arose from very long input lines (or at least pushes the problem threshold a little farther away). I tried to allow arbitrarily long input lines, but it would have been very difficult. This editing is far from perfect; I tried to allow 4096 characters for error strings, input strings, and filenames, and 512 characters for statement and object names. However, I was not particularly careful, so some may be set incorrectly. However, by increasing all sizes, I’m hoping that bugs won’t be a problem.
- Fixed a bug in `Geo_SphereReflectSphere` that caused problems with excluded volume reactions, using the “bounce -2” setting. This function implements the reflection method that I described in my 2017 paper. Previously, it was designed under the assumption that all collisions occur at the front of a molecule, meaning the direction that it’s traveling toward. However, collisions can also occur at the back, which was leading to incorrect vector reflection. This is now fixed and has been validated some, but still needs additional validation.
- Merged an automated pull request that might fix dependabot alerts about out of date scipy.
- Removed `public std::unary_function<type, type>` instances from `tree_iterator.hpp` (part of NSV code) because this is deprecated.

7.0.105 Modifications for version 2.73 (released 2/27/24)

- Investigated diffusion coefficients in crowded systems. I now believe that the overlap method is better than the reflection method. This is discussed in the reactions chapter of the Smoldyn manual. I changed the default behavior to overlap method.
- Added a line to the CMake file. In the Python module section, edited a line to read `find_package(Python3 3.10 EXACT COMPONENTS Interpreter Development REQUIRED)`, where the new portion is the Python 3.10 requirement. This prevents an error of `invalid command 'bdist_wheel'` because it now uses the same Python version for everything. This was reverted afterward. The below `PYTHON_VERSION` update fixes this problem.
- Edited `Smoldyn/source/pybind11/tests/requirements.txt` to simplify the file and hopefully stop warnings from the gitHub dependabot.

7 Smoldyn modifications

- Added `PYTHON_VERSION` as an option to the build, with new code in the master `CMakeLists.txt` file. Using this fixes the problem of building for one Python version and running a different one.
- Fixed bug in `cmdsetreactionratemolcount`. It wasn't loading the species name correctly, leading to a non-functional command. I'm not sure why it wasn't caught earlier, but it wasn't.

7.0.106 Modifications for version 2.74 (released 1/22/25)

- Changed Mac install and uninstall scripts so that individual commands use `sudo`, but the whole script does not. This is important because `sudo` should not be used with `pip`.
- Lots of work on filaments, including changes to data structures, input syntax, examples, and most functions in `smofilament.c`. Removed the biology enumeration because this is better handled with a standard library approach. Removed the `frontseg` element, which allowed filament segments to be stored in their array in a shifted position rather than up against the front end; this worked but had minimal benefit and added substantial complication. Filament dynamics now basically work in both 2D and 3D. Improved filament graphics.
- Improved code robustness when encountering input file parsing errors.
- Added `crossVVD` to `Rn.c`.
- In `Sphere.c`: removed clamping from `Sph_Xyz2Xyz`, added `Sph_Eax2Xyz`, `Sph_DcmxCart`, `Sph_DcmtxCart`, converted the documentation from Word to LaTeX (`Sphere.tex`), updated the documentation, and checked lots of equations using `docs/libSteve/SupportingFiles/-Sphere2.nb`.
- Smoldyn now supports units for everything except filaments. This is mostly covered in `string2.c`, `strunits` function and also, to lesser extend in `strmathsscanf`. It's also in `simLog`. Every input line that takes in a floating point number, with only a couple of exceptions, now support units as part of that input, and will convert as needed. Output lines that print floating point numbers often display units, but I didn't work as hard there (e.g. reaction rate units depend on reaction order and system dimensionality, so they're a pain to deal with). This has been documented, but could be documented more. The only test file so far is `S2_config/units.txt`.
- Fixed a few minor issues with lighting in `smolgraphics.c`, including mostly documentation.
- Added command `cmdprintFilament`, primarily for debugging purposes.
- Fixed a bug in the reflect portion of the `dosurfinteract` function. This runs the reflection with `surfacereflect`, moves the crossing point to the correct side of the panel, and then changes panels for surface-bound molecules if needed. For up to version 2.50, molecules were then automatically moved to the correct panel side, just in case they didn't reflect correctly. I changed this in version 2.51 (see the notes for that release), so that the code started checking whether molecules ended up on the correct side and then moved any failed molecules back to their starting points. This addressed problems that arose with the `examples/S95_regressions/surfacediffuse.txt` file, which uses a large number of tiny surface panels. An undesirable side effect of this change occurred for molecules inside spheres or cylinders (e.g. `dumbell.txt`),

where a molecule that started inside the surface could cross the surface and be reflected but still end up outside due to a long diffusion step. The check caught the result and put the molecule back where it started from, thereby preventing any leaking but also leading to incorrect diffusion simulation. If the check hadn't existed, the molecule would have reflected a second time, which would have been the correct simulation method and still prevented leaks. However, simulation errors were minor enough that they weren't caught. In 1/2025, Fabian Ormersbach found a problem which I determined arose from molecules inside 2D cylinders simultaneously diffusing out of the cylinder and across a periodic boundary, and then the check here put these molecules back in their starting points but the `posoffset` parameter was not reset. The solution was to realize that the check is only needed for surface-bound molecules; the function also displays a warning when molecules are returned to prior locations. This modification changes the regression results for several files (`dumbbell.txt`, `surf2.txt`, `surf3.txt`) but everything appears to simulate correctly. Added regression file `surfacesmsd.txt`.

7.0.107 Modifications for version 2.75 (released 7/30/25)

- The functions `checksurfaces`, `checksurfaces1mol`, and `dosurfinteract` all have checks to make sure that molecules are being placed correctly, putting them back to their prior location if not. Downgraded these errors to warnings, increased thresholds from 50 to 100 failures, and added debugging output during verbose simulations. Found that the `checksurfaces` warning can trigger if a molecule starts just inside a concave surface and bounces right along the inside of the surface; this triggers a warning and the molecule is moved back, but is actually simulated correctly.
- Improved error messages and the manual for unbounded emitters. A user declared an emitter at the corner of a box of surfaces, and it was unclear why it didn't work correctly. It turned out that the unbounded emitter was supposed to be inside the box, not at the corner.
- Cleaned up filament code, including particularly the Euler, RK2, and RK4 integration functions.
- Renamed “monomer codes” in filaments to “sequence code”.
- Restructured the `filament` data structure so that all working stuff was moved into the `filamentwork` data structure.
- Wrote EulerMat integration method, which performs Euler integration with a force matrix. It's only here as a step toward implicit integration.
- Wrote implicit integration code.
- Added node mobility to the filament data structure and a new `node_mobility` statement.
- Added command `cmdRotateFilamentToY`.
- Added and also renamed several filament example files.

7.0.108 Modifications for version 2.76 (not released yet)

- Added test for coincident surface panels. This is a new test in `checksurfaceparams` and a new function `sr coincidentpanels`, both in `smolsurface.c`. This only looks for panels that are identical to each other, including with a rotated point sequence or front/back flipped. It does not detect overlap in only one region of a panel, such as a pair of partially coincident rectangles.
- Fixed a minor bug in `strunits`, which created errors if units were used and commands included numbers; the problem was that default units only apply while a configuration file is being read, so there are no default units mid-simulation when a command is being parsed. As a result, Smoldyn couldn't convert the command values. Now, Smoldyn still can't convert values in commands, but this doesn't create an error.
- Implemented use of `strmatherror` for error reporting. Before, math errors were generated but never read. I created the `CHECKM` macro, which behaves the same as the `CHECKS` macro, except that it also checks for math errors and appends math error strings to the error outputs. It should be used after every call to `strmathsscanf`. I didn't create a new macro within `smolcmd.c`, but those calls to `strmathsscanf` also check for math errors now.
- The following items were written by Junjie Sun using Claude. I proofread them, added them to the main branch, and added edits.
- Fixed a minor bug in `filAlloc`. Fix: `filAlloc` never reassigned `fil->branchspots` on branch-array growth.
- Add `printFilaments` command: dump all filament node coordinates.

8 The wish/ to do list

8.1 Bugs and issues to fix

- The `comparttranslate` function isn't written well for molecules that get squeezed. What's needed is to write a second version of `checksurfaces1mol` that explicitly accounts for surface motion, rather than to just use this function, which is written for molecule motion, in the surface's moving reference frame.
- Surfaces now have lists of molecules adsorbed to them. The code would be faster if these lists were used as appropriate, such as for checking the actions of surface-bound molecules.
- Smoldyn shouldn't terminate if pattern or species name is over-length during on-the-fly expansion, but ignore it and issue a warning.
- Should expand length from `STRCHAR` to `STRCHARLONG` for patterns because they are for reactions too.
- The binding radius if time step is 0 output doesn't account for surface-bound species.
- The unbinding radius if time step is 0 output doesn't always work (it sometimes displays -2). This may have been fixed for version 2.13.
- If `time_now` is not start time, then all commands from start time to current time get run. They shouldn't.
- Reaction molecule states should be improved. At present, `rxnXsurface` doesn't allow, for example, `fsoln` molecules to reaction with back molecules, but this is over-ridden in `bireact`. It would be better to use the reaction `permit` array more effectively, and also to allow users to enter relative reactions, such as same side, opposite side, surface-bound, etc.
- Excluded volume reactions do not work with periodic boundaries. Fixed for version 2.26. They also don't work with the `meansqrdisp` commands.
- Need to finish writing `checksurfaceparams` function. In particular: check probabilities, if panels are wholly outside volume, and if jump panels source and destination panels are compatible.
- 3-D box cross testing, done in `Geometry.c`, isn't correct for all shapes (triangles and disks in particular, maybe others).
- Surface rendering still has some problems.
- If there are multiple reaction channels for a first order reaction and if the rate of one is set as a rate while the rate of a later one is set with the probability, and then the simulation time step is then changed by a command, then the reaction probability of the later one does not properly account for the changed conditional probability of the first occurring. The solution

8 The wish/ to do list

is to add a data structure element called **probreq**, for probability requested, and to then calculate the simulation probability from this, so that the original user-entered information is never overwritten.

8.2 Desired features

8.2.1 Libsmoldyn

- Test libsmoldyn.
- Modularize commands.
- Libsmoldyn needs functions such as: `smolGetMolCount`, `smolGetMolCountCmpt`, `smolGetMolCountPort`, `smolGetMolCountSurf`, `smolGetMolPositions`, `smolGetNumSpecies`, etc. Also, matched functions `smolGetMolecules` and `smolSetMolecules`.
- Prefix all pointer inputs with **const** if the memory that they point to won't be changed.
- Finish swigging for Python, and then swig for R and Octave.

8.2.2 Code acceleration

- There are lots of little tweaks that could be made to speed up the code. For example:
- Replace most if...else constructs in the core simulation code with switch
- Replace `mol->via` with a local vector.
- Check for system dimensionality from 3 to 2 to 1, rather than in ascending order.
- Pre-compute and store the panel margin data with the panels rather than re-computing frequently.
- Unroll loops when possible, often creating separate code for 1D, 2D, and 3D.
- Try to replace division with multiplication and get rid of square roots and trig. functions, as possible.

8.2.3 Electrostatics

- Arnd Pralle wants Smoldyn to simulate electrostatic interactions. At first, I was thinking that this would be a big pain. But then I realized that it's actually just a very minor tweak of the bounce type reactions. Instead of the product distance being some fixed distance or the collision distance plus the overlap distance, the product distance should be a function of the overlap. For attractive interactions, if the overlap is small, then the molecules get sucked in slightly, and if the overlap is large then they get sucked in more. An issue to deal with is that there will likely be multiple reactions listed for the same pair of molecules, and the right ones need to be executed at the right times. For example, there might be electrostatic interaction at relatively long distances but also chemical reactions at short distances.

8.2.4 BioNetGen

- Need to finish groups stuff. More testing would be good, too.

8.2.5 Math

- Add functions that have parentheses, such as $\sin(x)$ and $\text{atan2}(x, y)$. Also add a `molcount` function.

8.2.6 Commands

- Commands need to be overhauled. They need to be modularized, they need to support wildcards, and they need to be able to export data to Libsmoldyn.
- To commands, add an argument to the functions for data export by observation commands. I'm thinking that this should be a data structure that includes a header, `maxcolumns`, `ncolumns`, `maxrows`, `nrows`, and a data table. Also create a command adding function that Libsmoldyn can use to add commands to the system, and also get data back from commands.
- Simple spatially separated macromolecular complexes. Karen wants to be able to model Tar-CheA dimers, plus their interactions with their neighbors. My idea is to add a command which creates a "child" CheA for each Tar molecule, which are the "parents", and this CheA would have a fixed offset position from the Tar. When this command is called later on, every child is moved to the fixed offset value, thus causing the child molecules to always track their parents. This motion would ignore surfaces.
- It should be possible to change single reaction rates during a simulation. It's possible at present to change the internal value for the reaction rate, but not more user-friendly values. I think this is now possible in version 2.23 but I haven't verified it yet.
- It would be nice if commands could communicate with each other. An idea for this is to establish a bulletin board within the command superstructure, on which commands could post and read memos. More generally, this could be expanded into an entire programming language if desired, although it would take some thought on how to do it in the best way.
- A new runtime command for more versatile text output. Rather than having a pile of specialized output commands, it would be nice to have something akin to a `print` statement, where any of a wide variety of simulation variables could be printed with a user-defined format.

8.2.7 Distribution

- Sourceforge. The account has been set up but needs code, etc. Then, I should e-mail everyone who might be interested to invite them to join the mailing lists.
- The Calibayes project might be a good place to integrate Smoldyn into, to enable parameter fitting.
- Need to improve Windows distribution. Write a batch file for installing. Note that the download location appears to be `%USERPROFILE%\Downloads\smoldyn-2.xx-windows.zip`. It should be moved to `%PROGRAMFILES%`. Instructions: user needs to right-click the downloaded zip file, which will normally be in `C:\Users\your_name\Downloads`, and then select "Extract without confirmation" or something else that extracts the zipped file into the current directory. Next, run the `install.bat` file.

8 The wish/ to do list

Next step: Try to release again and test on Windows. To install on Windows, use notepad to edit the install file, which opens a gui, then edit the destination directory to something that I don't need admin privileges for. Mac stuff seems good now. Trying to get Windows good in all ways next, including with BioNetGen.

For Windows, FROMDIR should be the current working directory. Also the DESTDIR thing doesn't work because it isn't being substituted correctly.

- Smoldyn distribution should be made much easier. For example, distribute pre-compiled Mac software, and add Smoldyn to MacPorts, Fink, etc. code databases. This has been done to some extent.

8.2.8 Core Smoldyn

- Add a surface panel shape called a holey sphere. This would be a sphere but with as many holes as desired. It could serve as a neuron junction, the end of a dividing cell, etc. Another idea is to add a new surface feature for holes. These could be added to spheres, cylinders, triangles, etc. The simplest hole is a spherical hole; it says that the normal surface rules apply unless the molecule is within this sphere, in which case the surface effectively isn't there. This would be easy to define, easy to implement, and efficient, but would be a challenge for graphics.
- Movable internal surfaces. Lots of people have asked for this collection of new features. At a minimum, it should be possible to move surfaces using commands. Note that these moving commands should not use absolute coordinates but relative ones; this is so a repeated command would cause continuous motion. Eventually, surfaces should also move with diffusion coefficients, in response to specific molecules, etc. Also, deformable surfaces. A tremendous amount of work could be done here and would be very useful.
- Adaptive time steps.
- Check bimolecular reaction rates for surface-bound molecules.

8.2.9 Graphics and I/O

- Off-screen rendering would be very helpful. Two options seem most promising. (1) Off-screen rendering with OpenGL. This requires a frame buffer object, which looks relatively straightforward but still has some challenges. (2) Dump the data to a text file that can be post-processed by VMD (visualization of molecular dynamics). For this, John Stone (VMD author) e-mailed me 3/31/11 to suggest that I use the XYZ file format for the molecules plus a VMD script for the surfaces.
- A minor but helpful feature would be the possibility of graphing molecule concentrations at the bottom of the graphics window.
- In `loadsim`, `output_files`, a file name that is identical to the config. file name should result in a warning and ask user if it's desired. (There is already a file overwrite warning, so this isn't really needed; it's also a bit harder to add than it seemed it might be.)
- Keypress 'i' should enter interactive mode in which the user can type commands into the standard input and they will be executed right away.

- There should be statements and Libsmoldyn functions for removing things from Smoldyn. Like removing surfaces, surface panels, reactions, molecules, etc. Species removal would be harder.

8.2.10 Code improvements

- First order reactions would be more efficient with an event queue rather than a probabilistic likelihood at each time step.
- I think I see two ways to improve performance in `checksurfaces`: (1) instead of `for(bp1=pos2box...)`, do `for(bp1=mptr->box, ...)`. (2) End this same box loop when `crossmin<2`.
- Speed for 3D systems can be improved by always checking for 3D first, rather than last.
- Various cleanups would be nice for reactions: (1) `doreact` could be streamlined slightly for order 2 reactions by precomputing x , (3) `rxnss->table` symmetry is performed by having separate identical sides, but one side could just as easily point to the same data as does the other side, (5) allostery may need improvement, (6) derive theory for non-one reaction probabilities, (7) implement unimolecular equilibrium constant reactions in which the user just enters the reactants, the equilibrium constant, and the time constant (which may be 0) and Smoldyn calculates transition probabilities.

8.2.11 Major additions

- Fibers (such as DNA, actin, microtubules, MinD, FtsZ, etc.), fiber-bound molecules, etc. Also, membrane-bound polymers would be nice.
- Intrinsic molecular parameters. In this idea, the user enters fundamental intrinsic molecular parameters, and then Smoldyn calculates the model parameters from them. For example, the user enters molecular weight, and the occupied membrane area for surface-bound molecules, and Smoldyn calculates the diffusion coefficient from those things (see my diffusion coefficient rule-of-thumb in the MMB paper and the Saffman-Delbruck equation in PNAS 1975). Also, the user enters intrinsic reaction rates or activation energies, and ideal (short time step) binding radii, and Smoldyn calculates reaction rate constants from them. Ideally, the user could also specify viscosities for different surfaces and different compartments, and Smoldyn would adjust diffusion coefficients and binding radii automatically. Christine (Le Novere group) has already written some code that may be relevant to this idea. Having this feature would save the user from creating an explosion of molecular species for the various molecular environments. However, it might slow the program down substantially to check the environment for every potential dynamic aspect.
- Inclusion of continuous concentrations for chemical species that are abundant. Ideally, these concentrations should be updated with ODEs, PDEs, spatial- or non-spatial Langevin dynamics, or spatial- or non-spatial Gillespie algorithm, according to the user's choice.
- Derive theory for bimolecular reaction rates with probabilistic reactions.

Bibliography

- [1] Steven S Andrews. Methods for modeling cytoskeletal and dna filaments. *Physical biology*, 11(1):011001, 2014.